

CSCI-351

Data communication and Networks

Lecture 8: Network Layer (Putting the Net in Internet)

Announcement

2

- Project 1 Milestone due

Quiz

3


Recap

4

Wireless – Modulation


5

- From Analog to Digital
 - AM (Amplitude Modulation)
 - FM (Frequency Modulation)
 - ASK (Amplitude Shift Key)
- QAM (Quadrature Amplitude Modulation)
 - 802.11ac 256-QAM



8-QAM


6



The Danger of Loops

7

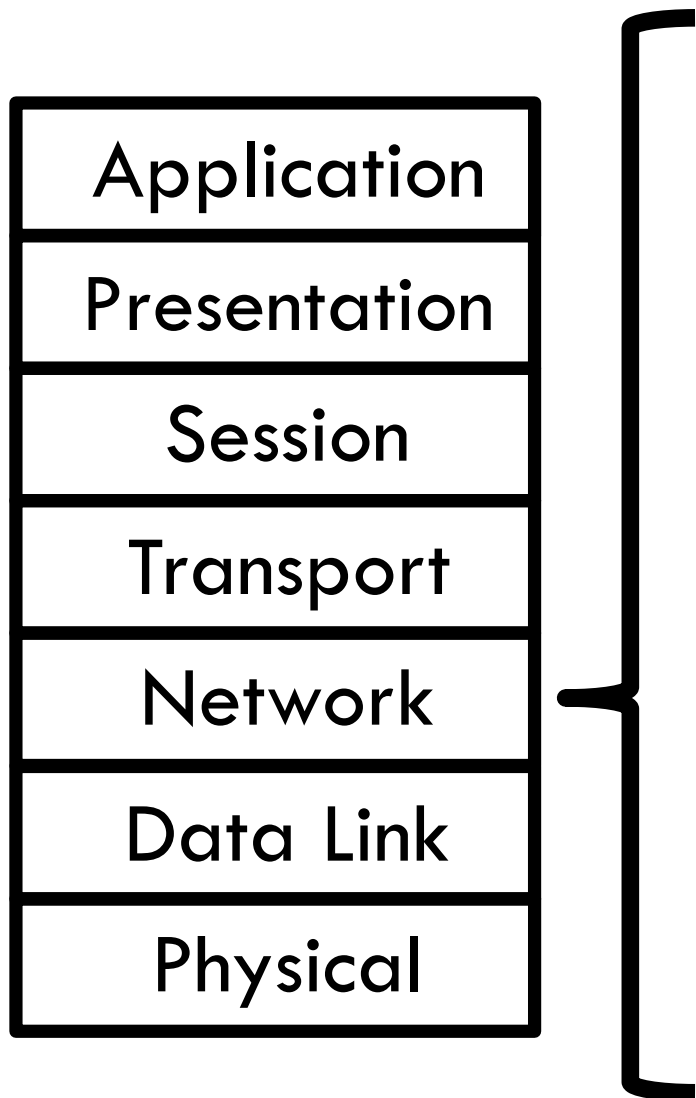
- ❑ $\langle \text{Src}=\text{AA}, \text{Dest}=\text{DD} \rangle$
- ❑ This continues to infinity
 - ▣ How do we stop this?
- ❑ Remove loops from the topology
 - ▣ Without physically unplugging cables
- ❑ 802.1 uses an algorithm to build and maintain a spanning tree for routing



8

Network Layer

9




- Function:
 - ▣ Route packets end-to-end on a network, through multiple hops
- Key challenge:
 - ▣ How to represent addresses
 - ▣ How to route packets
 - Scalability
 - Convergence

Routers, Revisited


10

- How to connect multiple LANs?
- LANs may be incompatible
 - ▣ Ethernet, Wifi, etc...
- Connected networks form an internetwork
 - ▣ The Internet is the best known example



Structure of the Internet

11



- Ad-hoc interconnection of networks
 - ▣ No organized topology
 - ▣ Vastly different technologies, link capacities
- Packets travel end-to-end by hopping through networks
 - ▣ Routers “peer” (connect) different networks
 - ▣ Different packets may take different routes

Internetworking Issues

12

- Naming / Addressing
 - ▣ How do you designate hosts?
- Routing
 - ▣ Must be scalable (i.e. a switched Internet won't work)
- Service Model
 - ▣ Internet Service Model
Best-effort (i.e. things may break)
 - ▣ What gets sent?
 - ▣ How fast will it go?
Store-and-forward datagram network
 - ▣ What happens if there are failures?
 - ▣ Must deal with heterogeneity
 - Remember, every network is different

13 Outline

Addressing

-  Class-based

-  CIDR

IPv4 Protocol Details

-  Packed Header

-  Fragmentation

IPv6

Possible Addressing Schemes

14

□ Flat

- ▣ e.g. each host is identified by a 48-bit MAC address
- ▣ Router needs an entry for every host in the world
 - Too big
 - Too hard to maintain (hosts come and go all the time)
 - Too slow (more later)

□ Hierarchy

- ▣ Addresses broken down into segments
- ▣ Each segment has a different level of specificity

Example: Telephone Numbers

15

1-585-475-3234



R·I·T

Updates are Local




RIT GCCIS RIT GCCIS
Room 4567 Room 1234

Very Specific

Binary Hierarchy Example

16



IP Addressing

17

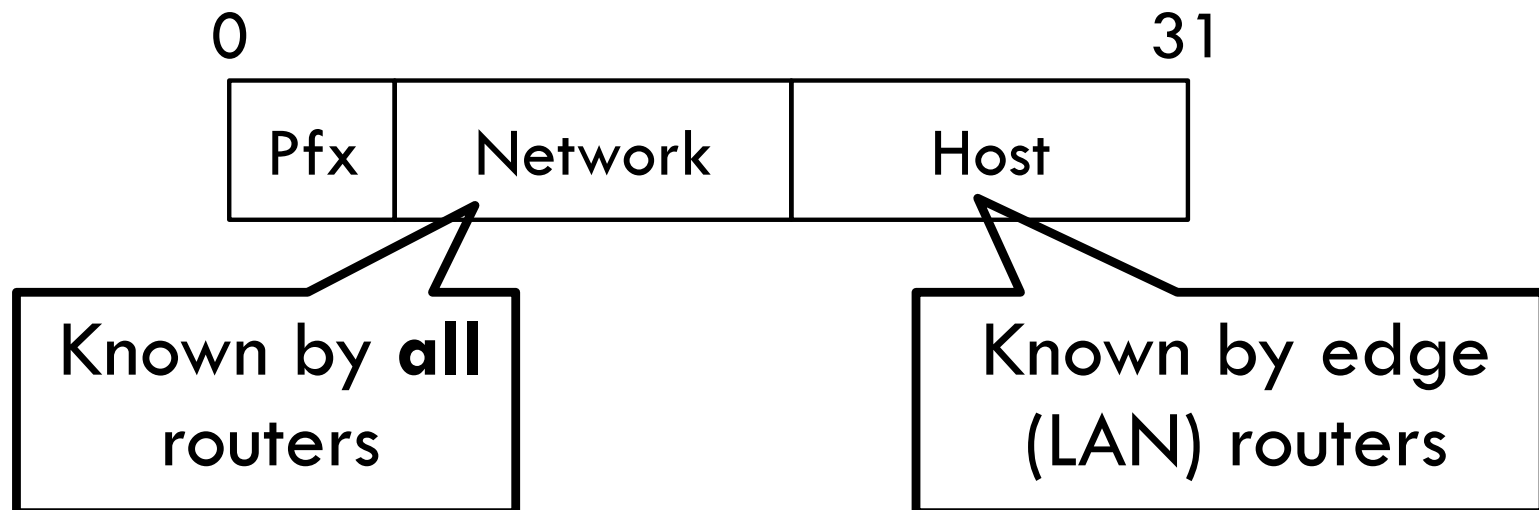
- IPv4: 32-bit addresses
 - ▣ Usually written in dotted notation, e.g. 192.168.21.76
 - ▣ Each number is a byte
 - ▣ Stored in Big Endian order

	0	8	16	24	31
Decimal	192	168	21	76	
Hex	C0	A8	15	4C	
Binary	11000000	10101000	00010101	01001100	

IP Addressing and Forwarding

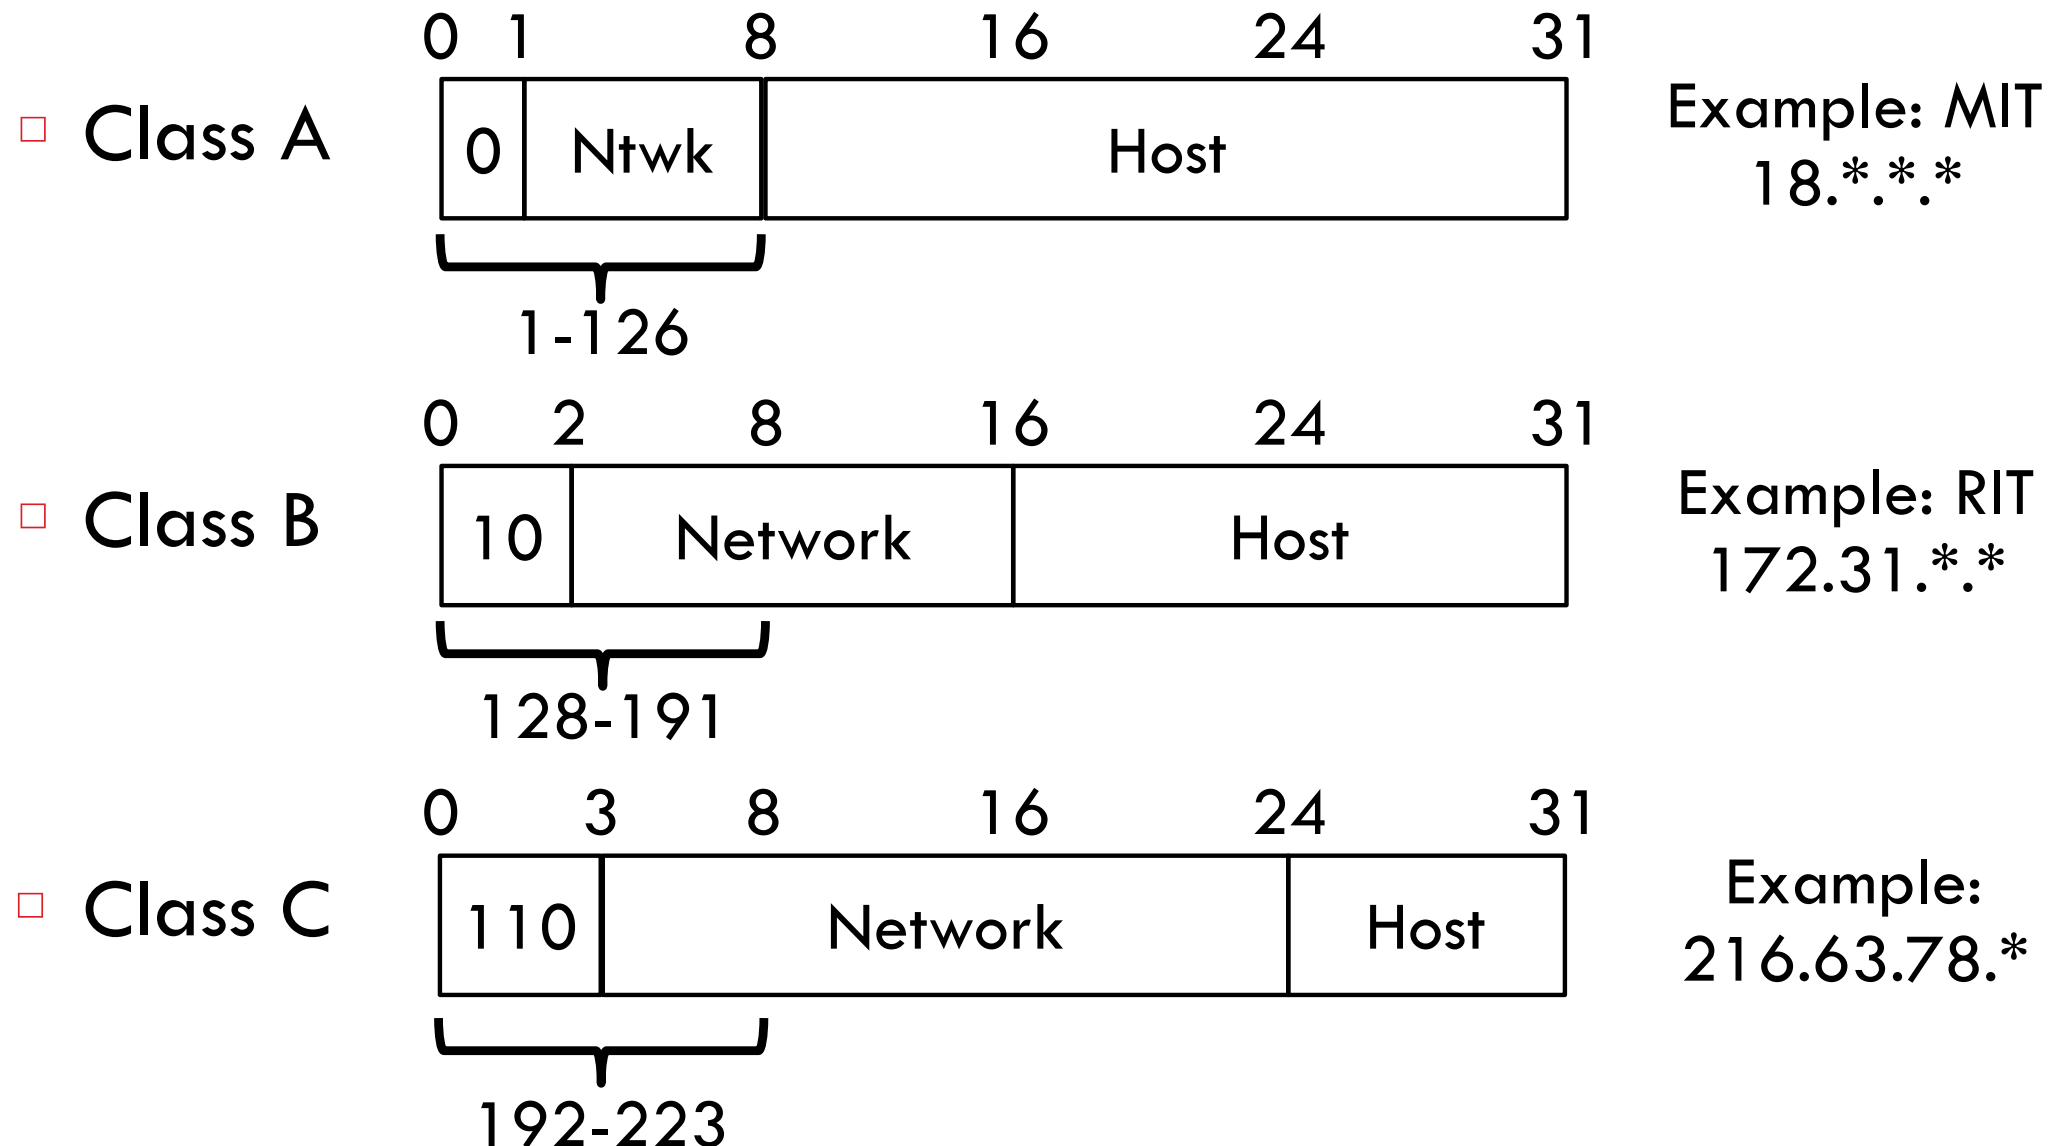
18

- Routing Table Requirements
 - ▣ For every possible IP, give the next hop
 - ▣ But for 32-bit addresses, 2^{32} possibilities!
 - ▣ **Too slow:** 48GE ports and 4x10GE needs 176Gbps bandwidth
DRAM: ~1-6 Gbps; **TCAM** is fast, but 400x cost of DRAM
- Hierarchical address scheme
 - ▣ Separate the address into a network and a host



Classes of IP Addresses

19



How Do You Get IPs?

20

- IP address ranges controlled by IANA

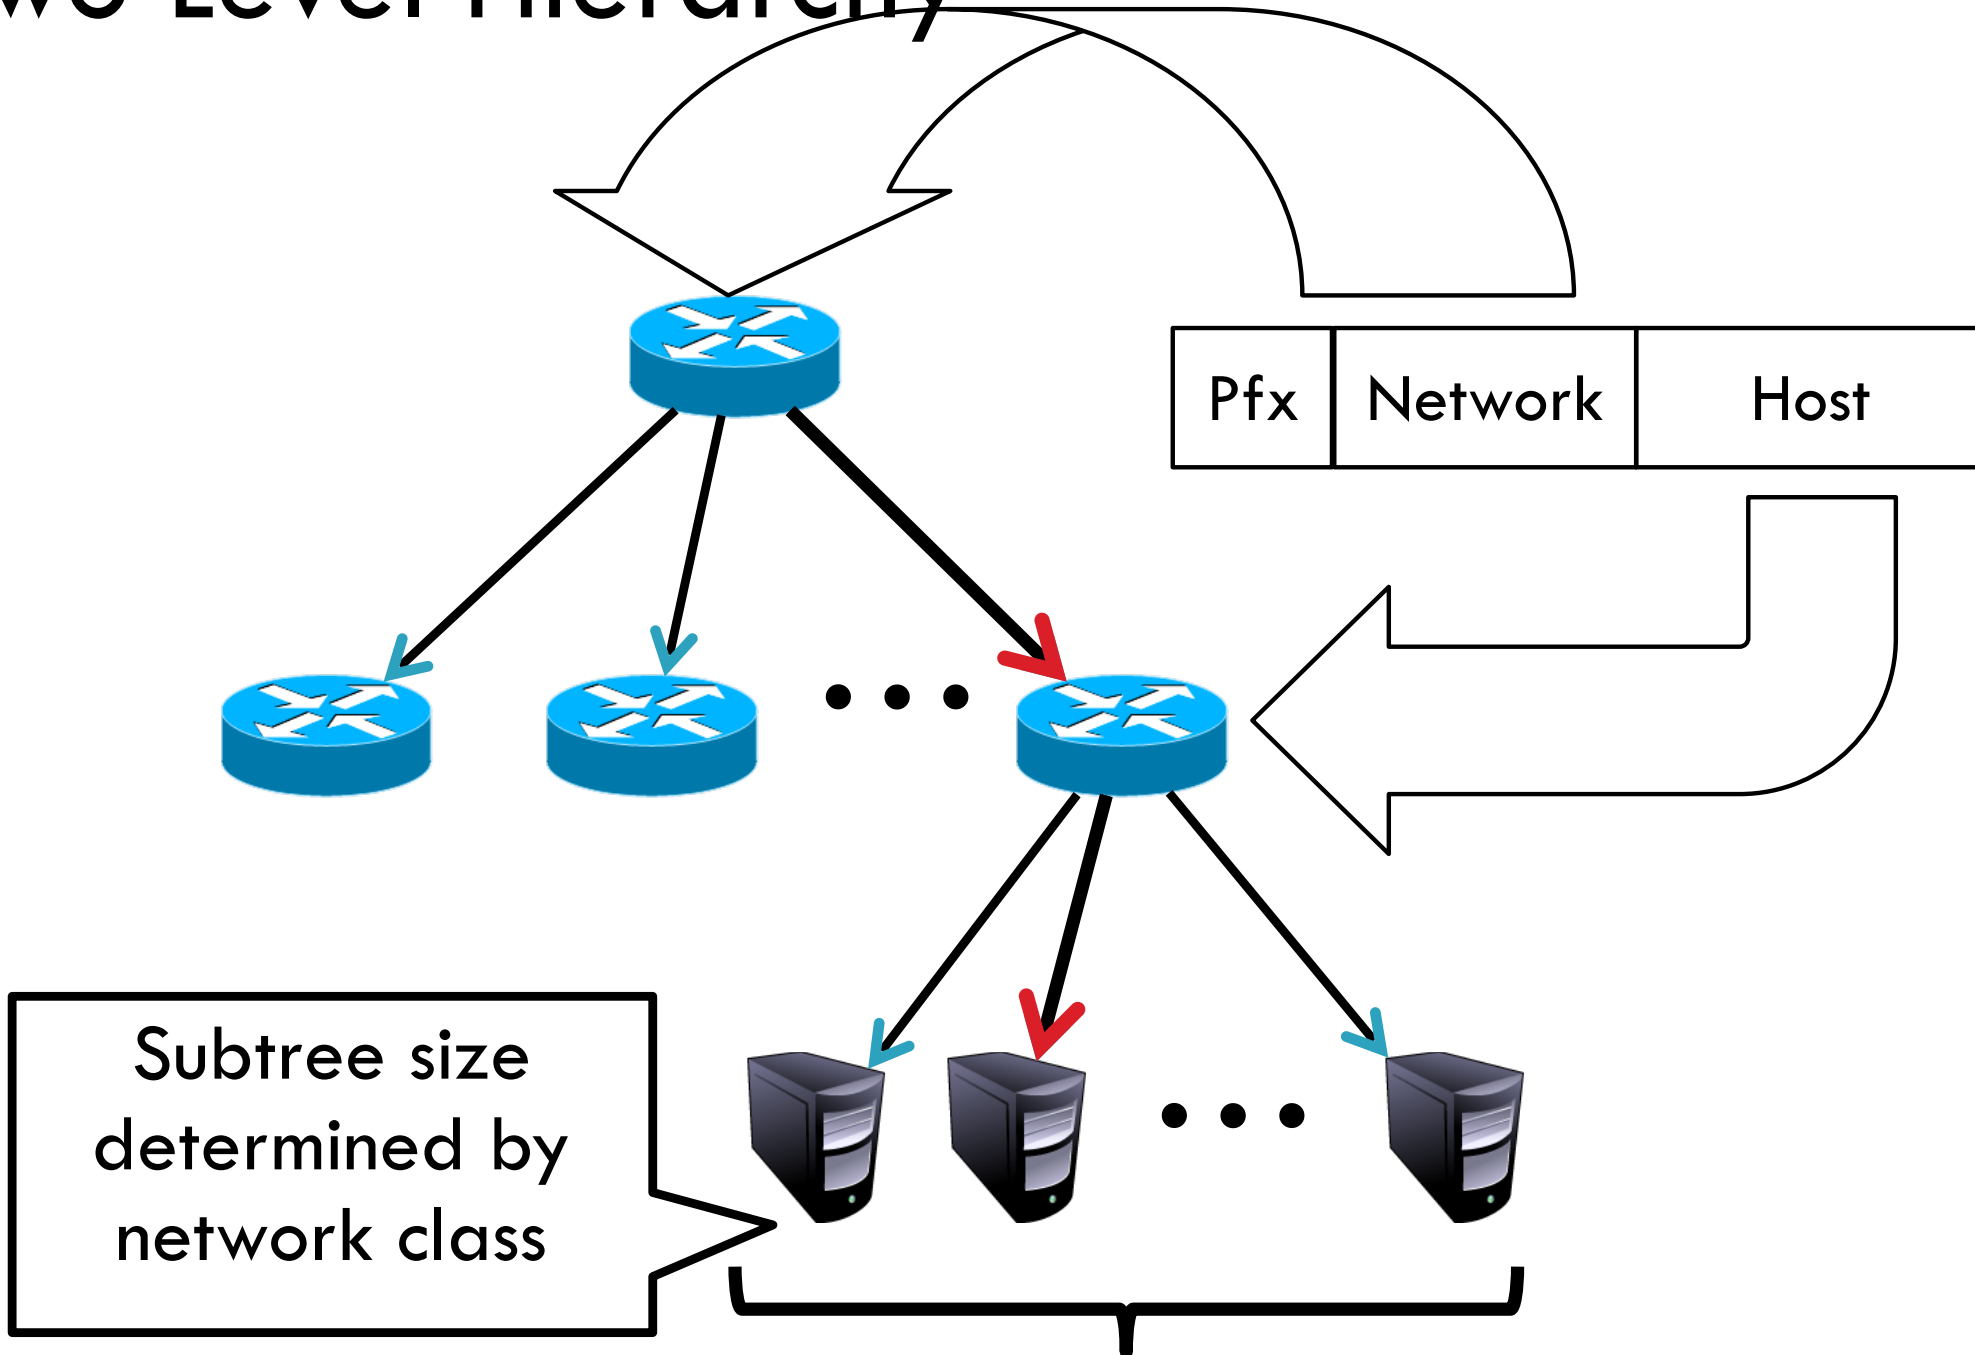


Internet Assigned Numbers Authority

- Internet Assigned Number Authority
- Roots go back to 1972, ARPANET, UCLA
- Today, part of ICANN
- IANA grants IPs to regional authorities
 - ARIN (American Registry of Internet Numbers) may grant you a range of IPs
 - You may then advertise routes to your new IP range
 - There are now secondary markets, auctions, ...

Two Level Hierarchy

21



Class Sizes

22

Way too big

Class	Prefix Bits	Network Bits	Number of Classes	Hosts per Class
A	1	7	$2^7 - 2 = 126$ (0 and 127 are reserved)	$2^{24} - 2 = 16,777,214$ (All 0 and all 1 are reserved)
B	2	14	$2^{14} = 16,398$	$2^{16} - 2 = 65,534$ (All 0 and all 1 are reserved)
C	3	21	$2^{21} = 2,097,512$	$2^8 - 2 = 254$ (All 0 and all 1 are reserved)
			Total: 2,114,036	

Too many
network IDs

Too small to be
useful

Subnets

23

- Problem: need to break up large A and B classes
- Solution: add another layer to the hierarchy
 - ▣ From the outside, appears to be a single network
 - Only 1 entry in routing tables
 - ▣ Internally, manage multiple subnetworks
 - Split the address range using a subnet mask

Pfx	Ntwk	Subnet	Host
-----	------	--------	------

Subnet Mask: 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0

Subnet Example

24

□ Extract network:


IP Address:	10110101	11011101	01010100	01110010
Subnet Mask: &	11111111	11111111	11000000	00000000
Result:	10110101	11011101	01000000	00000000

□ Extract host:

IP Address:	10110101	11011101	01010100	01110010
Subnet Mask: &	~(11111111	11111111	11000000	00000000)
Result:	00000000	00000000	00010100	01110010

N-Level Subnet Hierarchy

25



Example Routing Table

26

Address Pattern	Subnet Mask	Destination Router
0.0.0.0	0.0.0.0	Router 4
18.0.0.0	255.0.0.0	Router 2
128.42.0.0	255.255.0.0	Router 3
128.42.128.0	255.255.128.0	Router 5
128.42.222.0	255.255.255.0	Router 1

- Question: 128.42.222.198 matches four rows
 - ▣ Which router do we forward to?
- Longest prefix matching
 - ▣ Use the row with the longest number of 1's in the mask
 - ▣ This is the most specific match

Subnetting Revisited

27

- Question: does subnetting solve all the problems of class-based routing?

NO

- Classes are still too coarse
 - ▣ Class A can be subnetted, but only 126 available
 - ▣ What if MIT does not use all IPs?
 - ▣ Class C is too small
 - ▣ Class B is nice, but there are only 16,398 available
- Routing tables are still too big
 - ▣ 2.1 million entries per router

Classless Inter Domain Routing

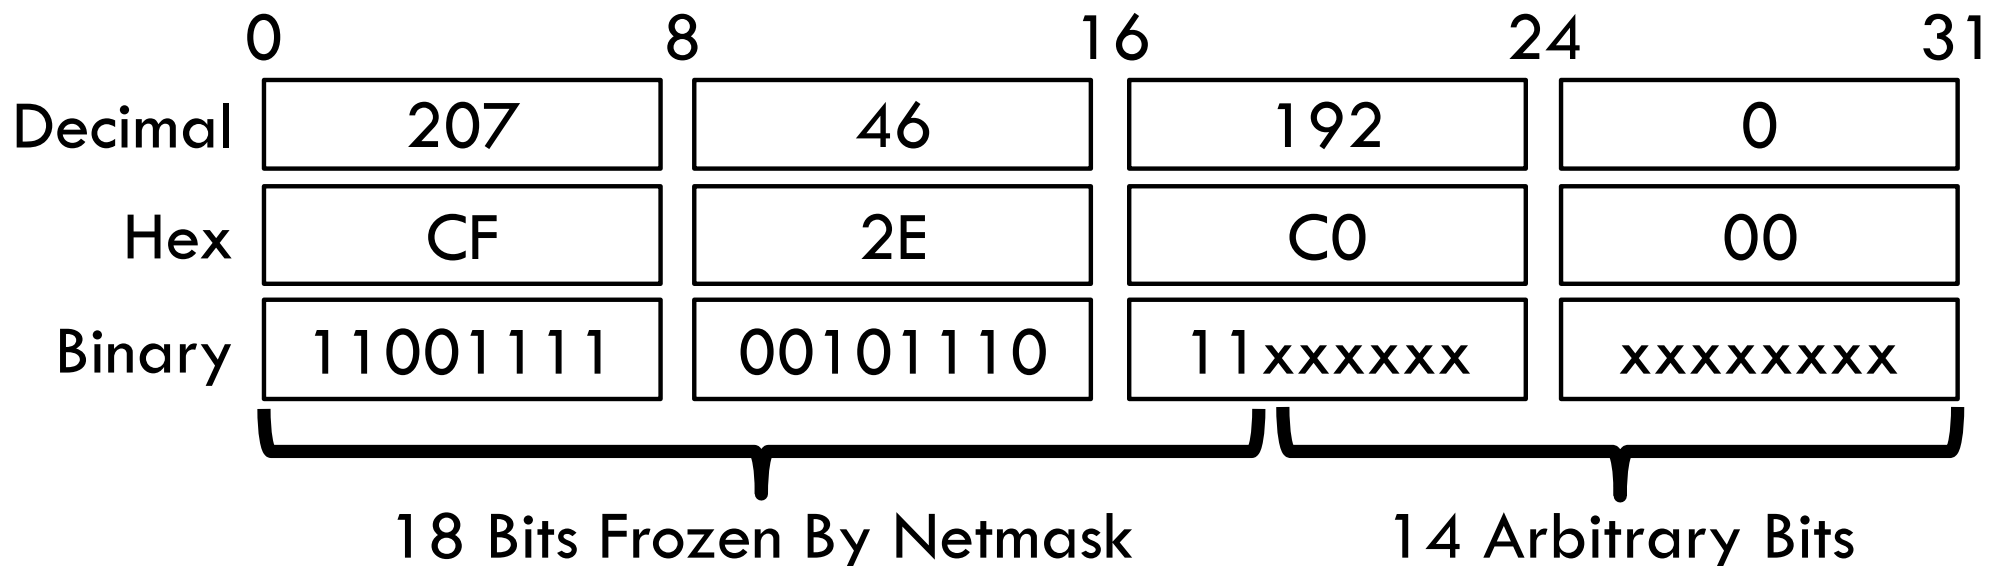
28

- CIDR, pronounced 'cider'
- Key ideas:
 - ▣ Get rid of IP classes
 - ▣ Use bitmasks for all levels of routing
 - ▣ **Aggregation** to minimize FIB (forwarding information base)
- Arbitrary split between network and host
 - ▣ Specified as a bitmask or prefix length
 - ▣ Example: Rochester Institute of Technology
 - 172.31.0.0 with netmask 255.255.0.0
 - 172.31.0.0 / 16

Aggregation with CIDR

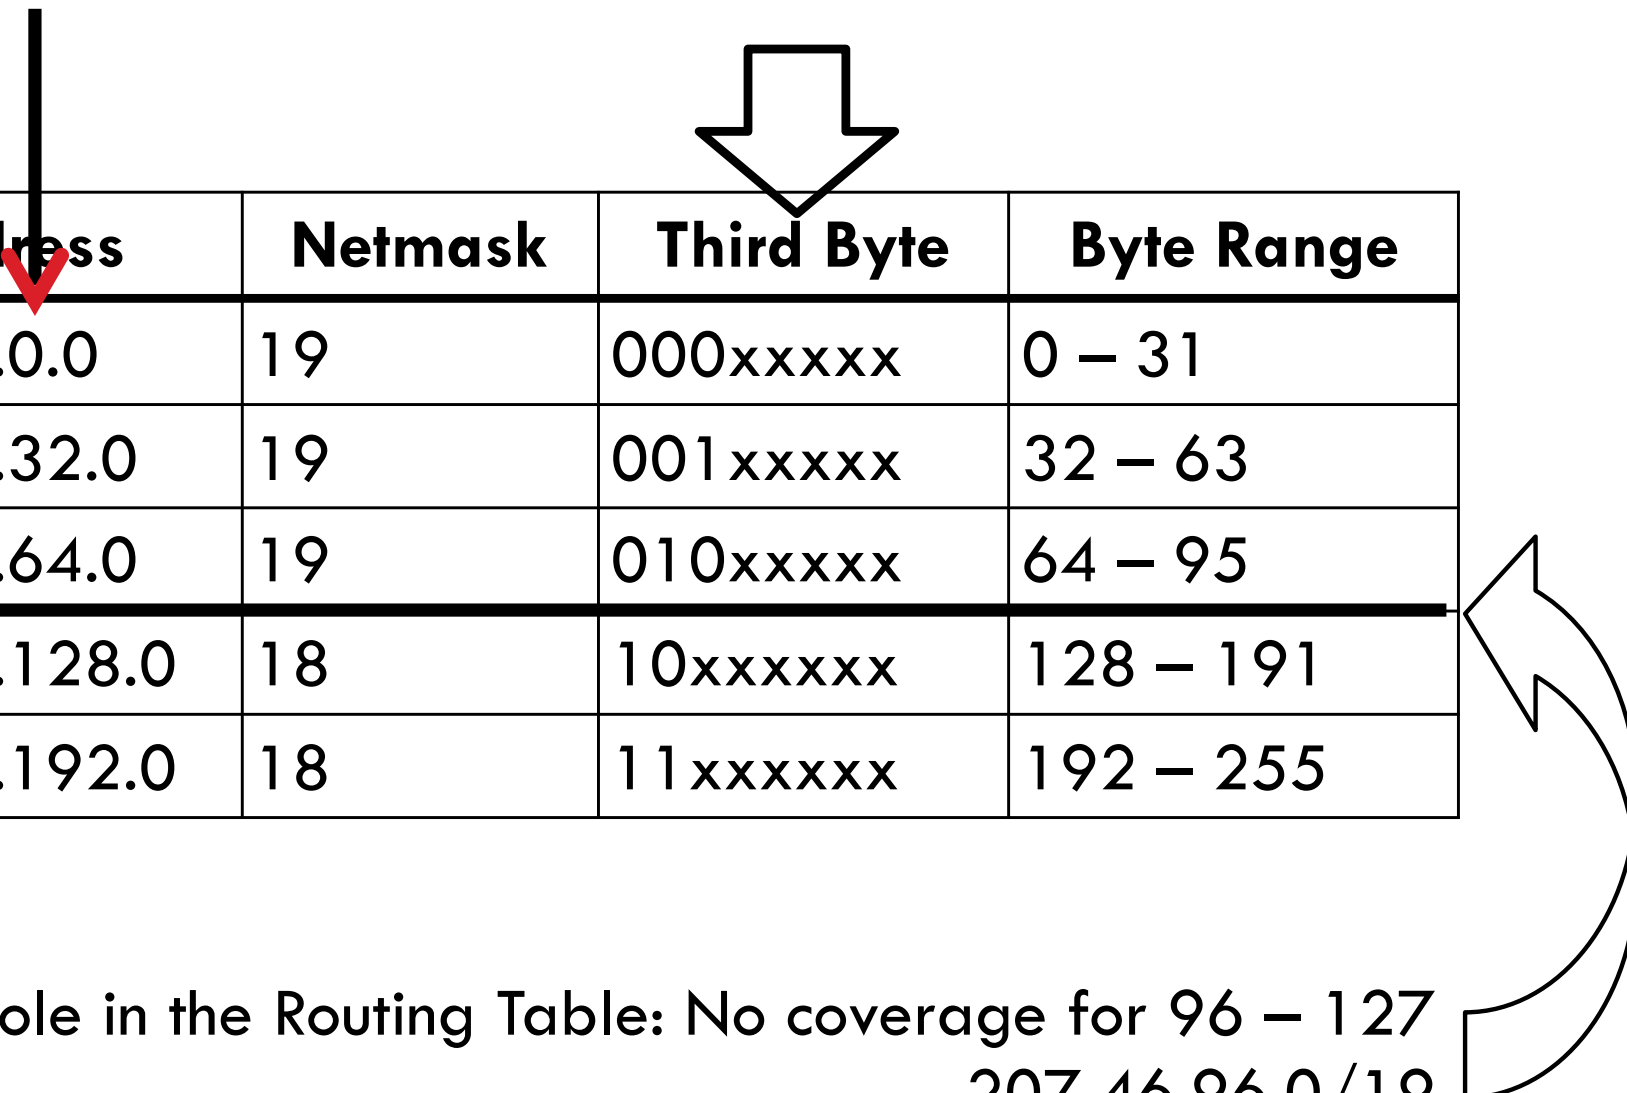
29

- Original use: aggregating class C ranges
- One organization given contiguous class C ranges
 - ▣ Example: Microsoft, 207.46.192.* – 207.46.255.*
 - ▣ Represents $2^6 = 64$ class C ranges
 - ▣ Specified as CIDR address 207.46.192.0/18



Example CIDR Routing Table

30




Address	Netmask	Third Byte	Byte Range
207.46.0.0	19	000xxxxx	0 – 31
207.46.32.0	19	001xxxxx	32 – 63
207.46.64.0	19	010xxxxx	64 – 95
207.46.128.0	18	10xxxxxx	128 – 191
207.46.192.0	18	11xxxxxx	192 – 255

Hole in the Routing Table: No coverage for 96 – 127
207.46.96.0/19

Size of CIDR Routing Tables

31



- From www.cidr-report.org
- CIDR has kept IP routing table sizes in check
 - ▣ Currently ~450,000 entries for a complete IP routing table
 - ▣ Only required by backbone routers

Takeaways

32

- Hierarchical addressing is critical for scalability
 - ▣ Not all routers need all information
 - ▣ Limited number of routers need to know about changes
- Non-uniform hierarchy useful for heterogeneous networks
 - ▣ Class-based addressing is too coarse
 - ▣ CIDR improves scalability and granularity
- Implementation challenges
 - ▣ Longest prefix matching is more difficult than schemes with no ambiguity

Outline



Addressing



Class-based



CIDR



IPv4 Protocol Details



Packed Header



Fragmentation



IPv6

IP Datagrams

34

- IP Datagrams are like a letter
 - ▣ Totally self-contained
 - ▣ Include all necessary addressing information
 - ▣ No advanced setup of connections or circuits

0	4	8	12	16	19	24	31
Version		HLen		DSCP/ECN		Datagram Length	
Identifier				Flags	Offset		
TTL		Protocol		Checksum			
Source IP Address							
Destination IP Address							
Options (if any, usually not)							
Data							

IP Header Fields: Word 1

35

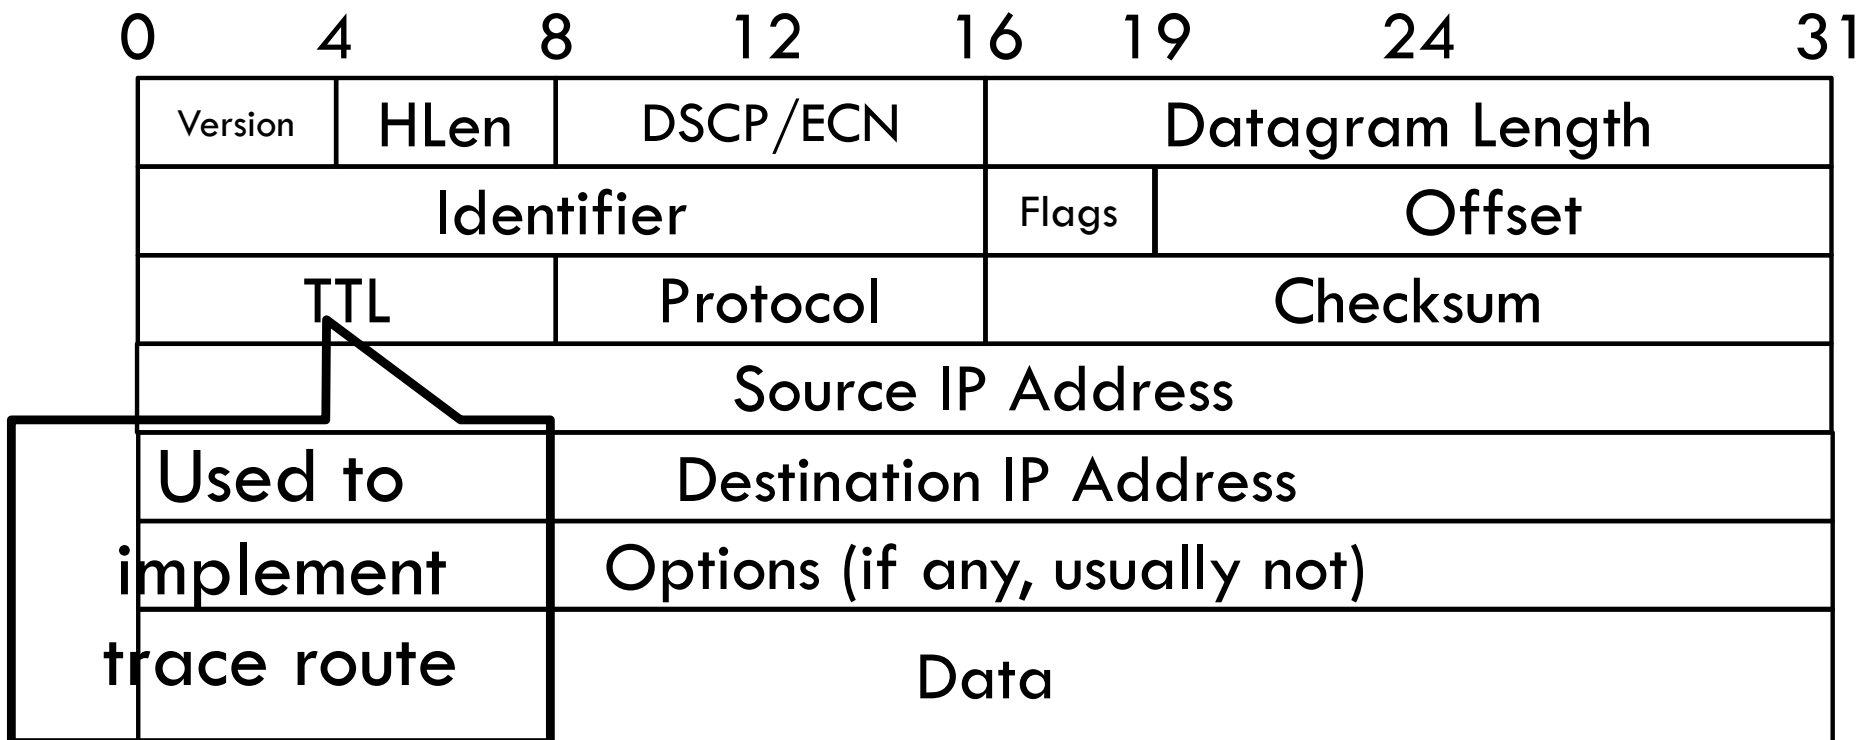
- Version: 4 for IPv4
- Header Length: Number of 32-bit words (usually 5)
- Type of Service: Priority information (unused)
- Datagram Length: Length of header + data in bytes

0	4	8	12	16	19	24	31
Version	HLen	DSCP/ECN		Datagram Length			
Identifier				Flags	Offset		
TTL		Protocol		Checksum			
Source IP Address				Destination IP Address			
Options (if any, usually not)				Data			

IP Header Fields: Word 3

36

- Time to Live: decremented by each router
 - ▣ Used to kill looping packets
- Protocol: ID of encapsulated protocol
 - ▣ 6 = TCP, 17 = UDP
- Checksum



IP Header Fields: Word 4 and 5

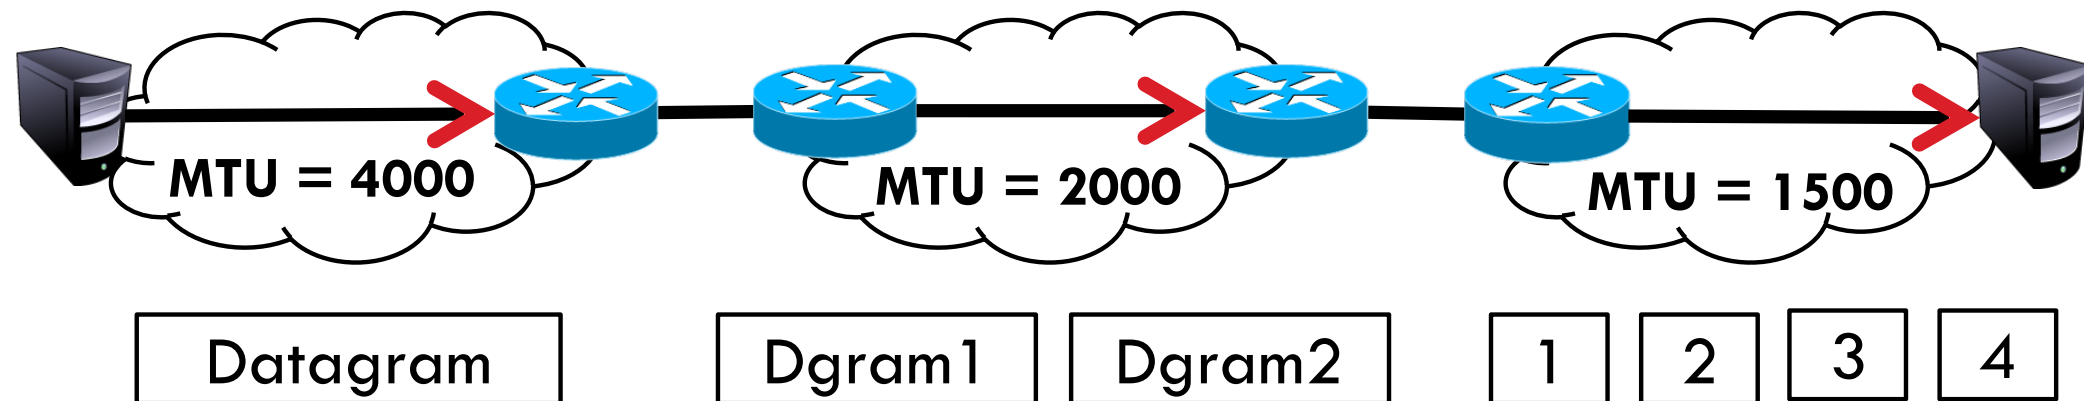
37

- Source and destination address
 - ▣ In theory, must be globally unique
 - ▣ In practice, this is often violated

0	4	8	12	16	19	24	31
Version		HLen		DSCP/ECN		Datagram Length	
Identifier				Flags	Offset		
TTL		Protocol		Checksum			
Source IP Address							
Destination IP Address							
Options (if any, usually not)							
Data							

Problem: Fragmentation

38



- Problem: each network has its own MTU
 - ▣ DARPA principles: networks allowed to be heterogeneous
 - ▣ Minimum MTU may not be known for a given path
- IP Solution: fragmentation
 - ▣ Split datagrams into pieces when MTU is reduced
 - ▣ Reassemble original datagram at the receiver

IP Header Fields: Word 2


39

- ❑ Identifier: a unique number for the original datagram
- ❑ Flags: M flag, i.e. this is the last fragment
- ❑ Offset: byte position of the first byte in the fragment
 - ▣ Divided by 8

0	4	8	12	16	19	24	31
Version		HLen		TOS		Datagram Length	
Identifier				Flags	Offset		
TTL		Protocol		Checksum			
Source IP Address							
Destination IP Address							
Options (if any, usually not)							
Data							

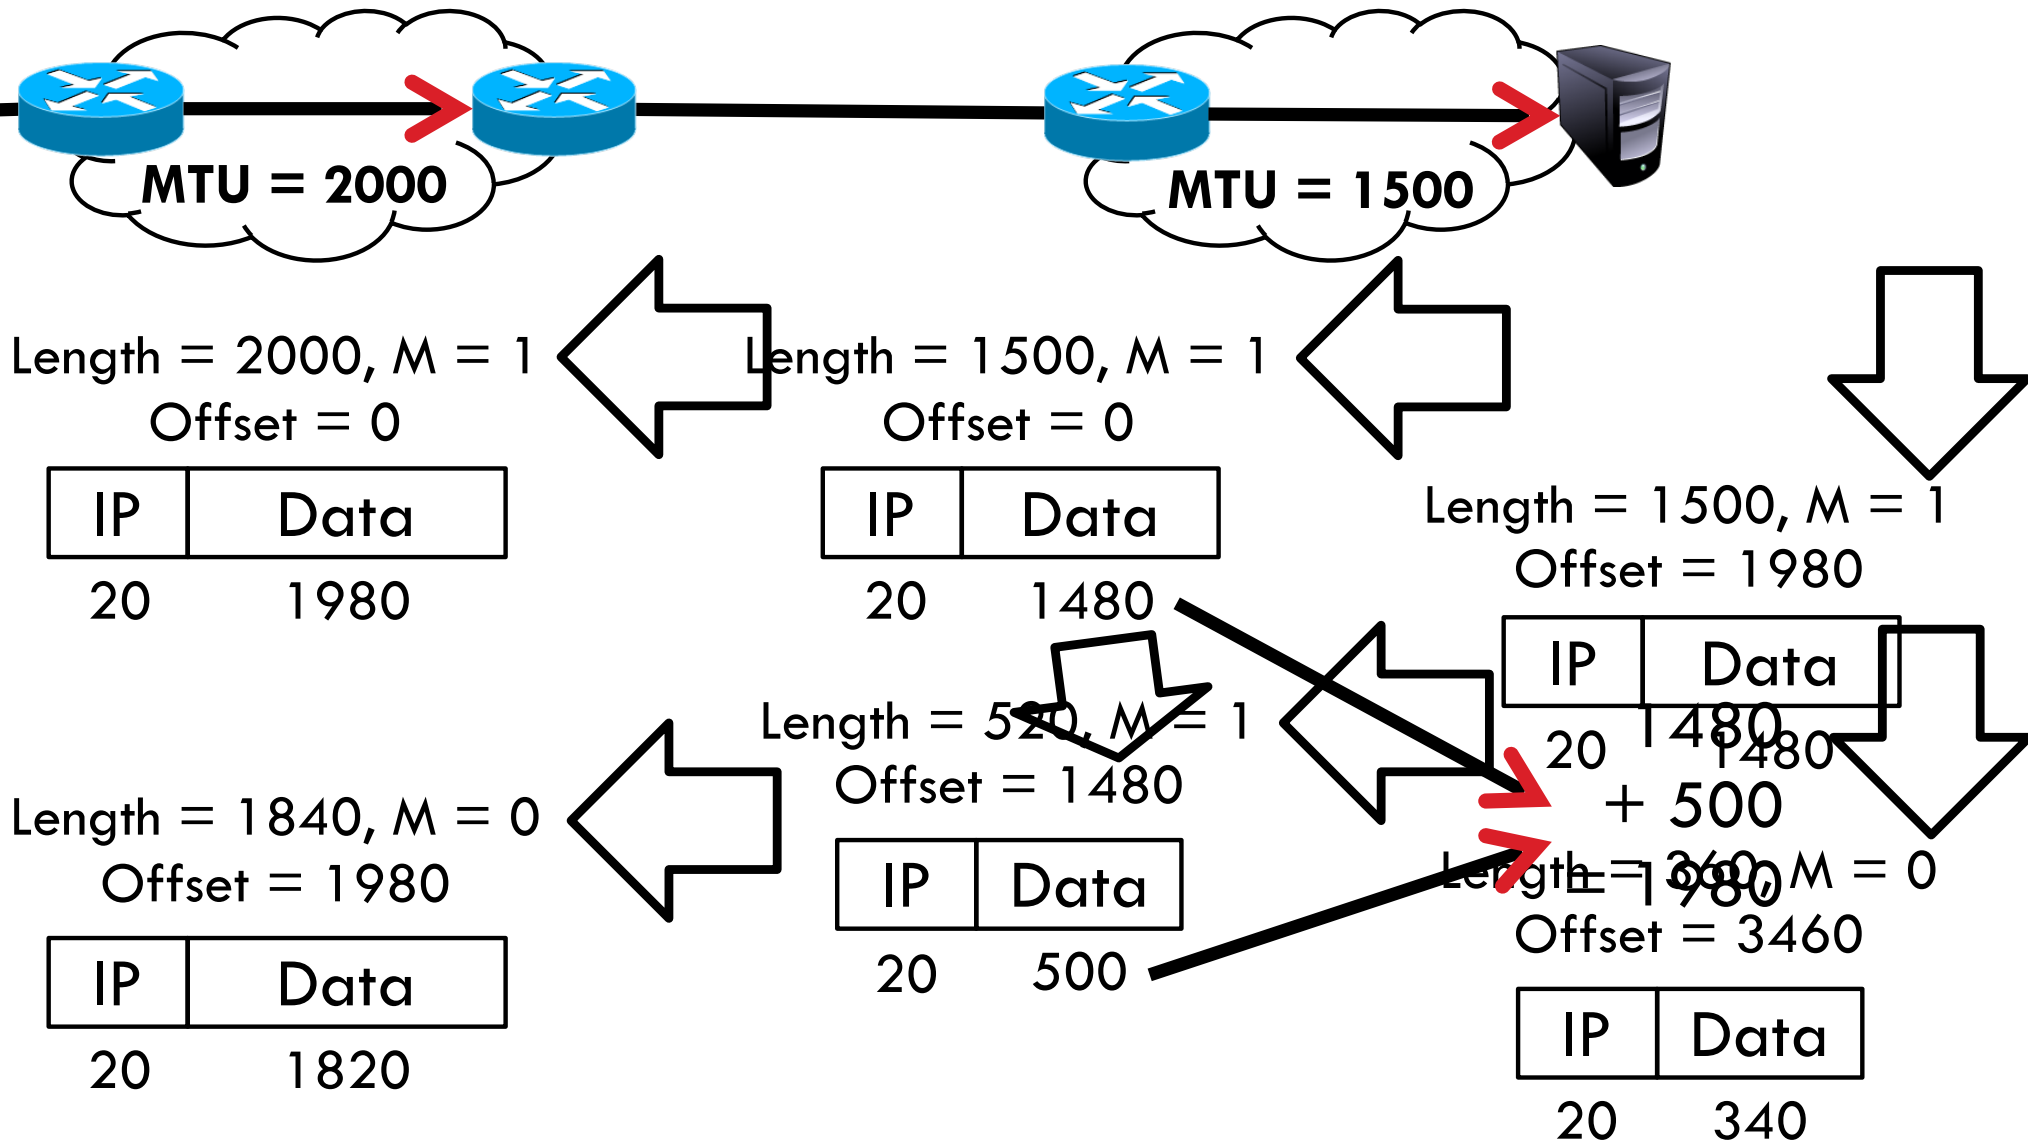
Fragmentation Example

40



Fragmentation Example

41



IP Fragment Reassembly

42

Length = 1500, $M = 1$, Offset = 0

IP	Data
20	1480

Length = 520, $M = 1$, Offset = 1480

IP	Data
20	500

Length = 1500, $M = 1$, Offset = 1980

IP	Data
20	1480

Length = 360, $M = 0$, Offset = 3460

IP	Data
20	340

- Performed at destination
- $M = 0$ fragment gives us total data size
 - ▣ $360 - 20 + 3460 = 3800$
- Challenges:
 - ▣ Out-of-order fragments
 - ▣ Duplicate fragments
 - ▣ Missing fragments
- Basically, memory management nightmare

Fragmentation Concepts

43

- Highlights many key Internet characteristics
 - ▣ Decentralized and heterogeneous
 - Each network may choose its own MTU
 - ▣ Connectionless datagram protocol
 - Each fragment contains full routing information
 - Fragments can travel independently, on different paths
 - ▣ Best effort network
 - Routers/receiver may silently drop fragments
 - No requirement to alert the sender
 - ▣ Most work is done at the endpoints
 - i.e. reassembly

Fragmentation in Reality

44

- Fragmentation is expensive
 - ▣ Memory and CPU overhead for datagram reconstruction
 - ▣ Want to avoid fragmentation if possible
- MTU discovery protocol
 - ▣ Send a packet with “don’t fragment” bit set
 - ▣ Keep decreasing message length until one arrives
 - ▣ May get “can’t fragment” error from a router, which will explicitly state the supported MTU
- Router handling of fragments
 - ▣ Fast, specialized hardware handles the common case
 - ▣ Dedicated, general purpose CPU just for handling fragments

45 Outline

Addressing

-  Class-based

-  CIDR

IPv4 Protocol Details

-  Packed Header

-  Fragmentation

IPv6

The IPv4 Address Space Crisis

47

- Problem: the IPv4 address space is too small
 - ▣ $2^{32} = 4,294,967,296$ possible addresses
 - ▣ Less than one IP per person
- Parts of the world have already run out of addresses
 - ▣ IANA assigned the last /8 block of addresses in 2011

Region	Regional Internet Registry (RIR)	Exhaustion Date
Asia/Pacific	APNIC	April 19, 2011
Europe/Middle East	RIPE	September 14, 2012
North America	ARIN	13 Jan 2015 (Projected)
South America	LACNIC	13 Jan 2015 (Projected)
Africa	AFRINIC	17 Jan 2022(Projected)

IPv6

48

- IPv6, first introduced in 1998(!)
 - ▣ 128-bit addresses
 - ▣ $4.8 * 10^{28}$ addresses per person
- Address format
 - ▣ 8 groups of 16-bit values, separated by ':'
 - ▣ Leading zeroes in each group may be omitted
 - ▣ Groups of zeroes can be omitted using '::'

2001:0db8:0000:0000:0000:ff00:0042:8329

2001:0db8:0:0:0:ff00:42:8329

2001:0db8::ff00:42:8329

IPv6 Trivia

49

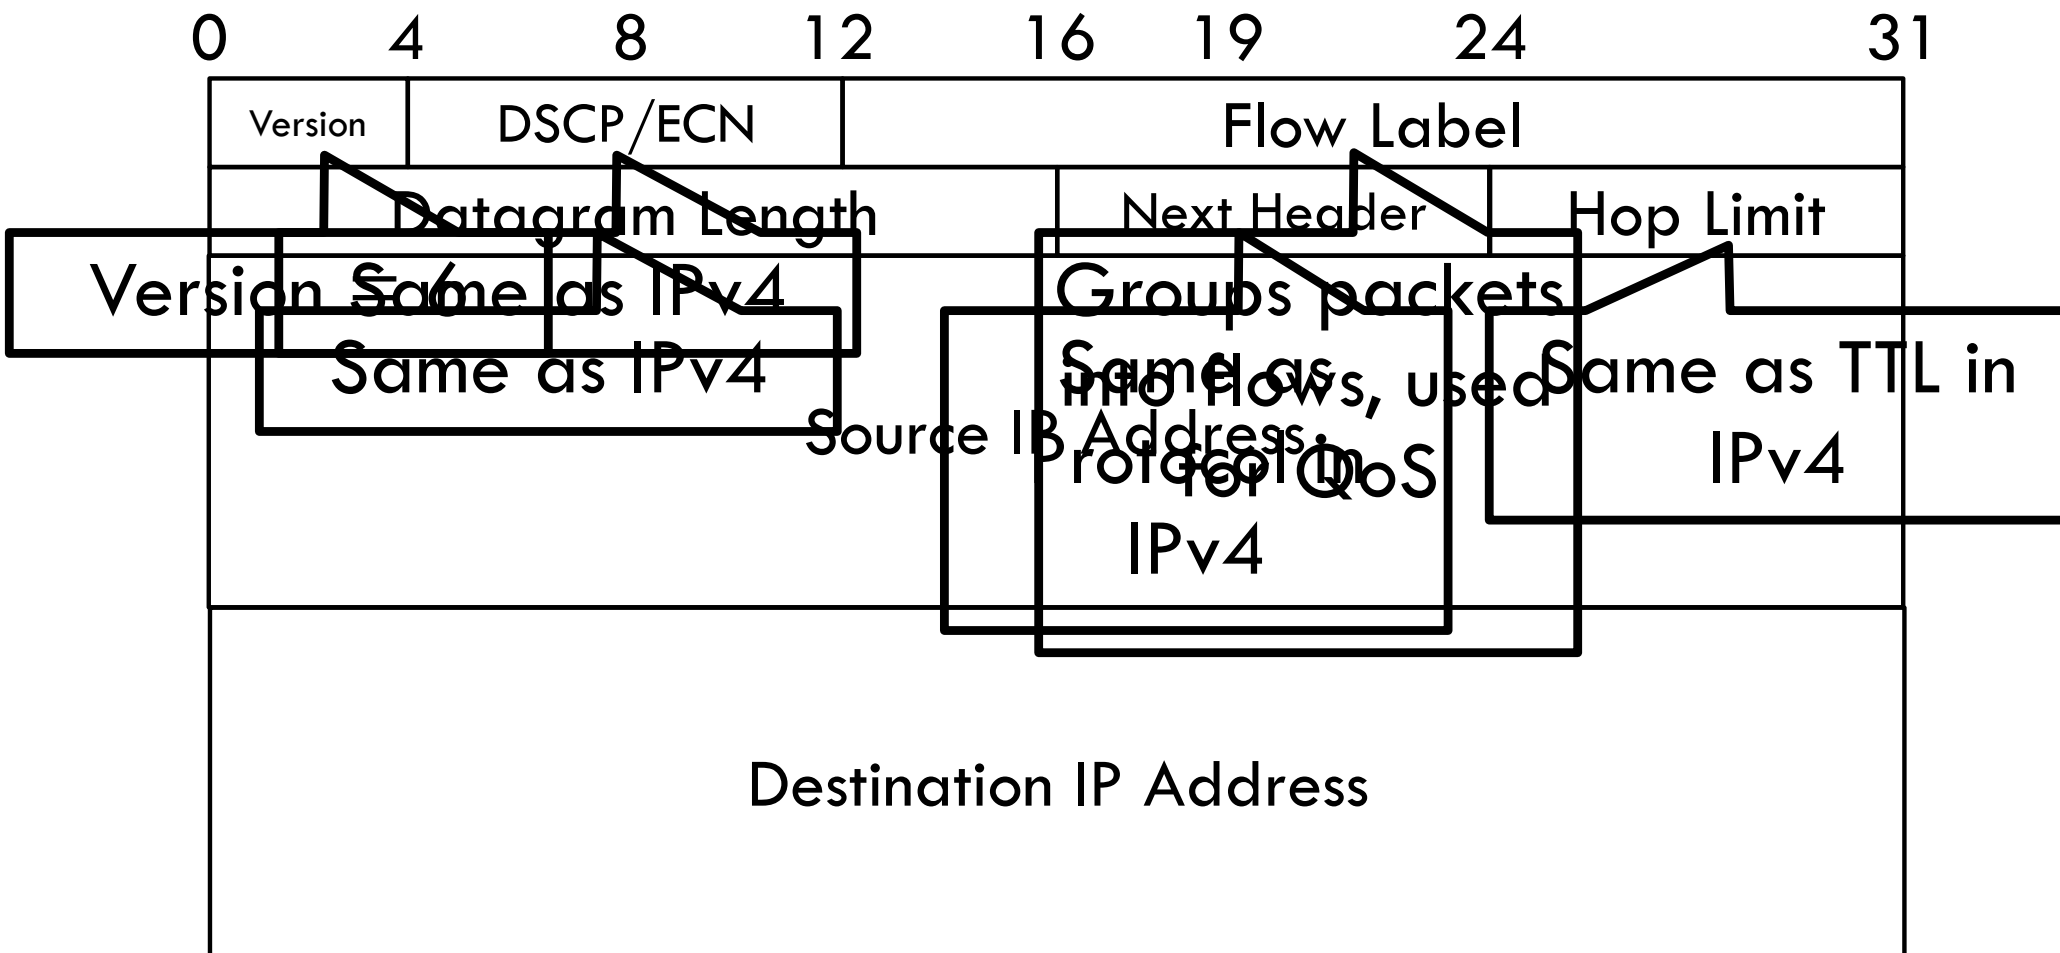
- Who knows the IP for localhost?
 - ▣ 127.0.0.1

- What is localhost in IPv6?
 - ▣ ::1

IPv6 Header

50

- Double the size of IPv4 (320 bits vs. 160 bits)



Differences from IPv4 Header

51

- Several header fields are missing in IPv6
 - ▣ Header length – rolled into Next Header field
 - ▣ Checksum – was useless, so why keep it
 - ▣ Identifier, Flags, Offset
 - IPv6 routers do not support fragmentation
 - Hosts are expected to use path MTU discovery
- Reflects changing Internet priorities
 - ▣ Today's networks are more homogeneous
 - ▣ Instead, routing cost and complexity dominate
- No security vulnerabilities due to IP fragments

Performance Improvements

52

- No checksums to verify
- No need for routers to handle fragmentation
- Simplified routing table design
 - ▣ Address space is huge
 - ▣ No need for CIDR (but need for aggregation)
 - ▣ Standard subnet size is 2^{64} addresses
- Simplified auto-configuration
 - ▣ Neighbor Discovery Protocol
 - ▣ Used by hosts to determine network ID
 - ▣ Host ID can be random!

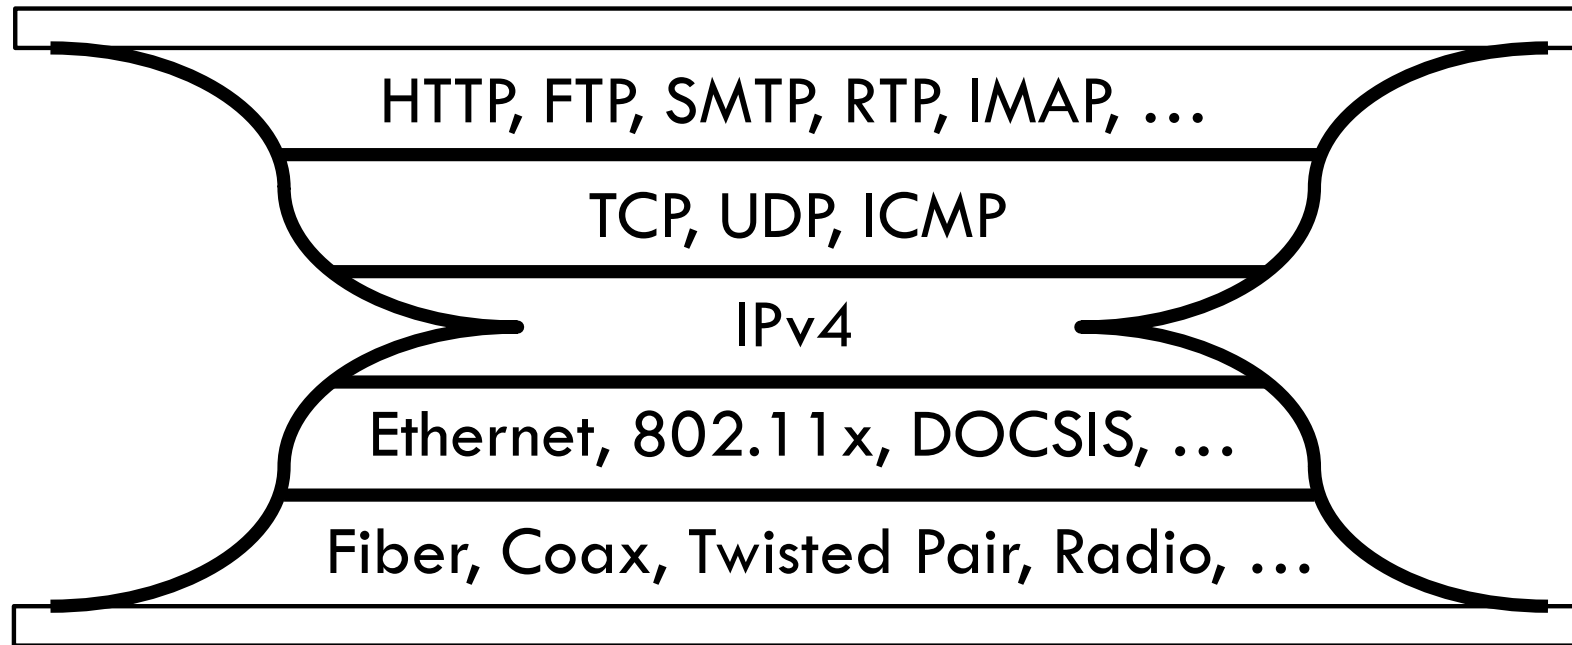
Additional IPv6 Features

53

- Source Routing
 - ▣ Host specifies the route to wants packet to take
- Mobile IP
 - ▣ Hosts can take their IP with them to other networks
 - ▣ Use source routing to direct packets
- Privacy Extensions
 - ▣ Randomly generate host identifiers
 - ▣ Make it difficult to associate one IP to a host
- Jumbograms
 - ▣ Support for 4Gb datagrams

Deployment Challenges

54




- Switching to IPv6 is a whole-Internet upgrade
 - ▣ All routers, all hosts
 - ▣ ICMPv6, DHCPv6, DNSv6
- 2013: 0.94% of Google traffic was IPv6, 2.5% today

Transitioning to IPv6

55

- How do we ease the transition from IPv4 to IPv6?
 - ▣ Today, most network edges are IPv6 ready
 - Windows/OSX/iOS/Android all support IPv6
 - Your wireless access point probably supports IPv6
 - ▣ The Internet core is hard to upgrade
 - ▣ ... but a IPv4 core cannot route IPv6 traffic



Transition Technologies


56

- How do you route IPv6 packets over an IPv4 Internet?
- Transition Technologies
 - ▣ Use tunnels to encapsulate and route IPv6 packets over the IPv4 Internet
 - ▣ Several different implementations
 - 6to4
 - IPv6 Rapid Deployment (6rd)
 - Teredo
 - ... etc.

6to4 Basics

57


- Problem: you've been assigned an IPv4 address, but you want an IPv6 address
 - ▣ Your ISP can't or won't give you an IPv6 address
 - ▣ You can't just arbitrarily choose an IPv6 address
- Solution: construct a 6to4 address
 - ▣ 6to4 addresses always start with 2002::
 - ▣ Embed the 32-bit IPv4 inside the 128-bit IPv6 address



Routing from 6to4 to 6to4

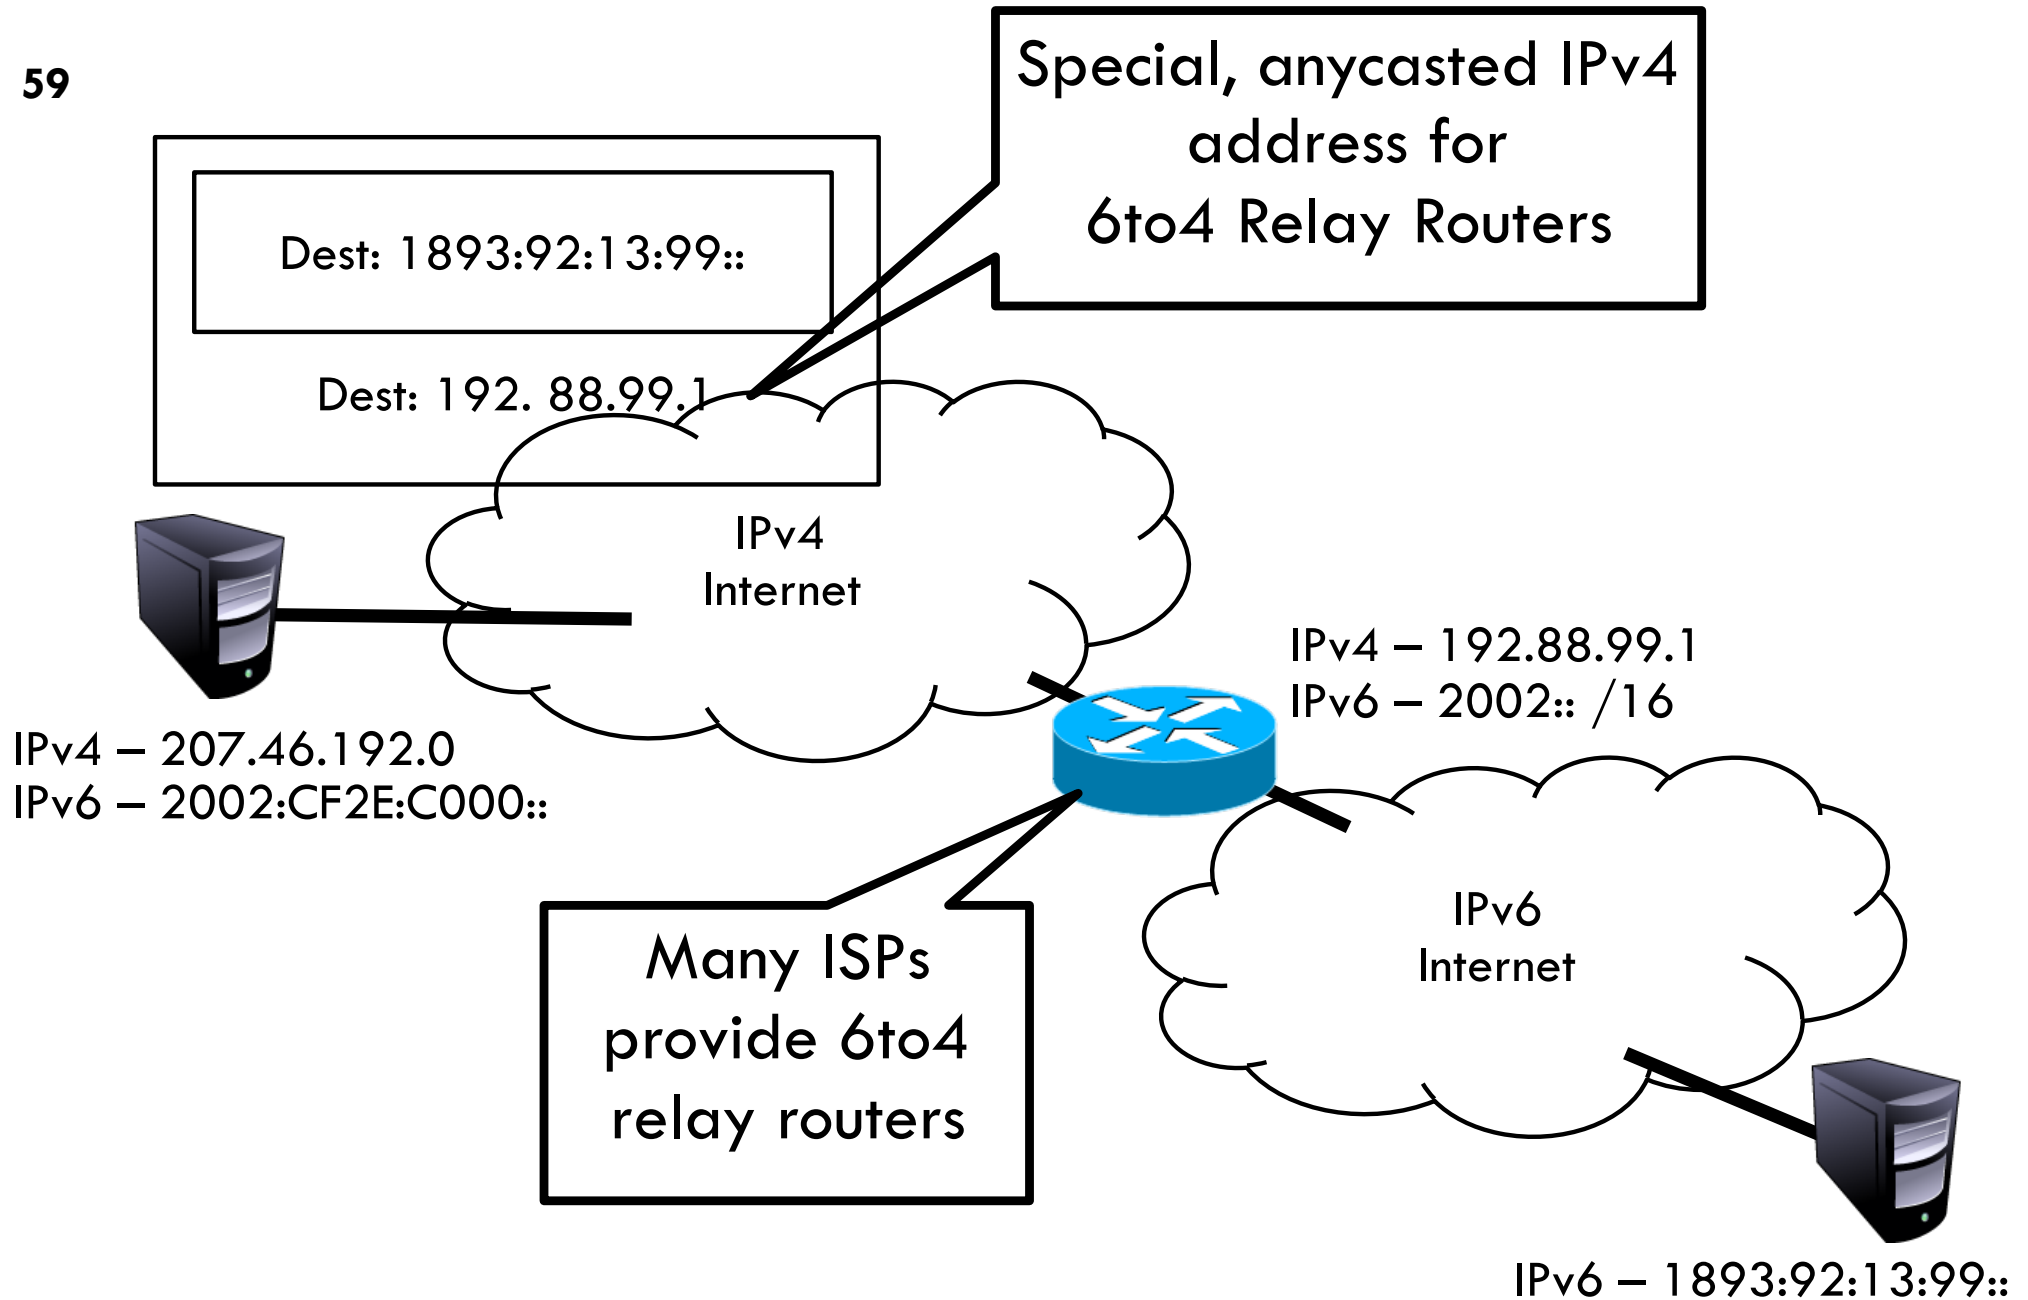
58

- How does a host using 6to4 send a packet to another host using 6to4?



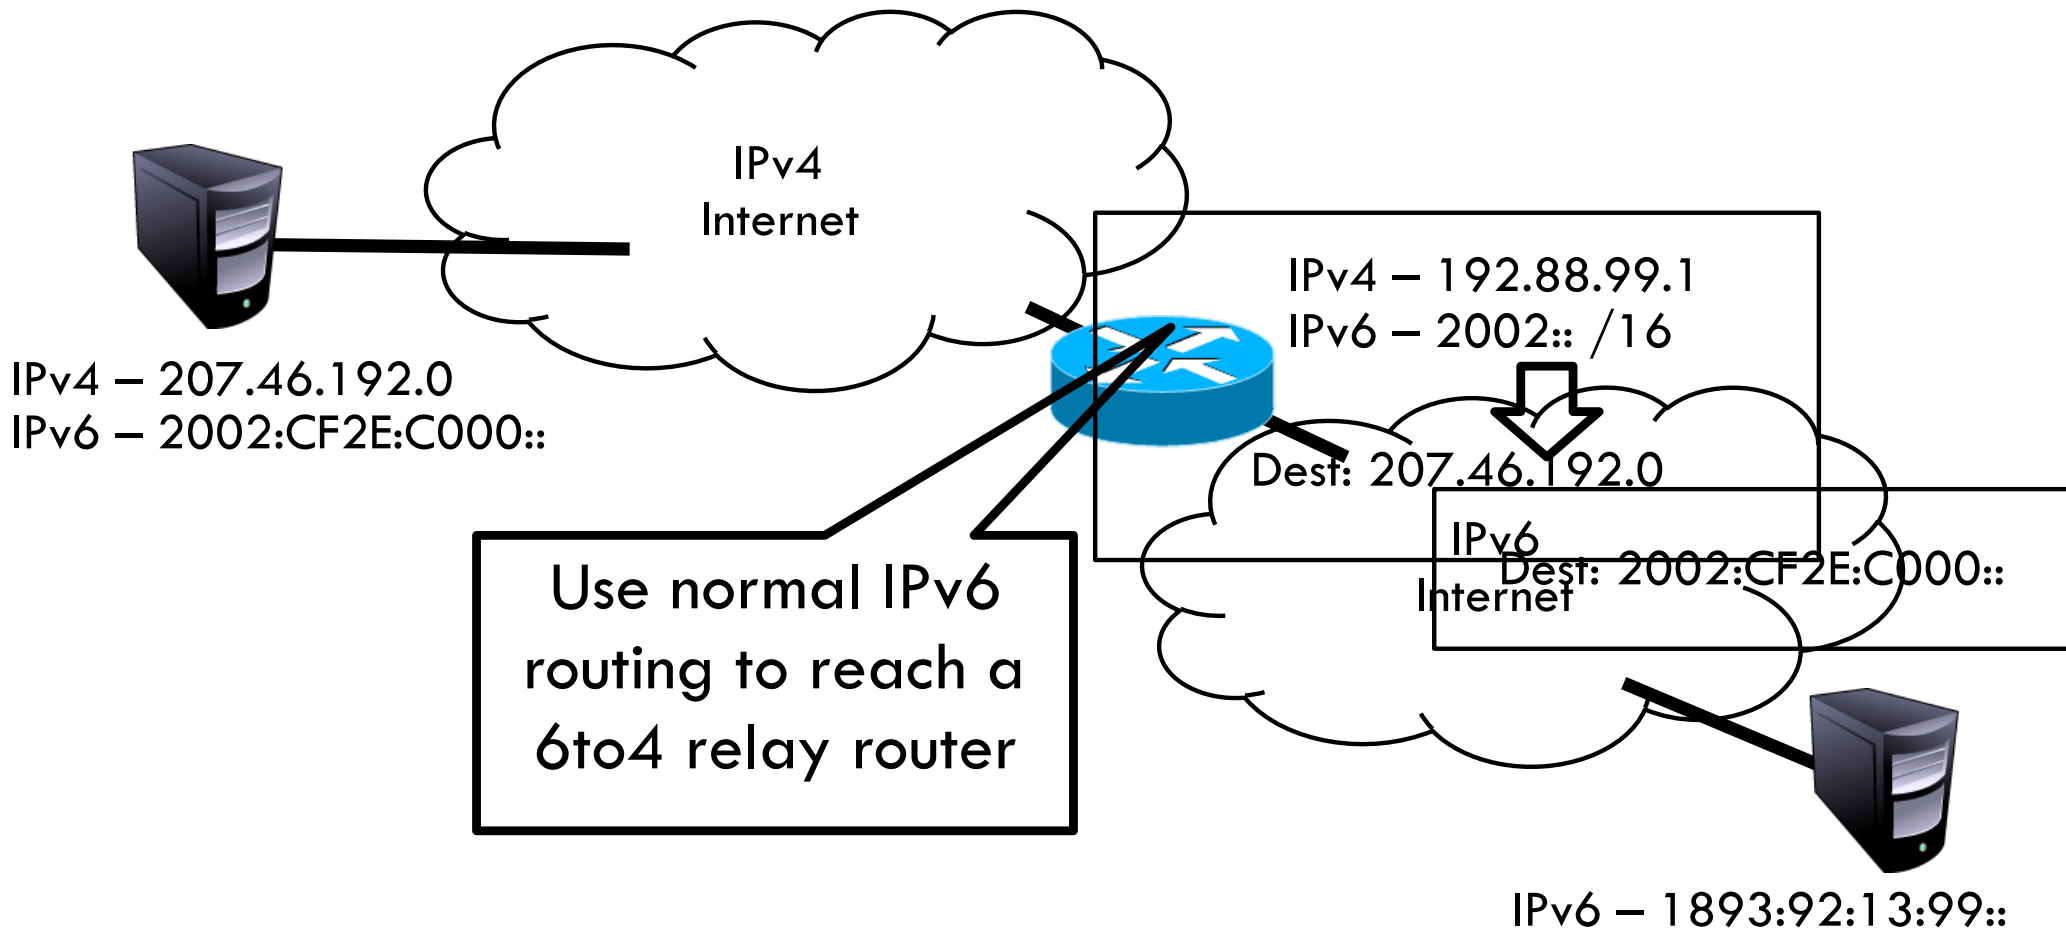
Routing from 6to4 to Native IPv6

59



Routing from Native IPv6 to 6to4

60



Problems with 6to4

61

- Uniformity
 - ▣ Not all ISPs have deployed 6to4 relays
- Quality of service
 - ▣ Third-party 6to4 relays are available
 - ▣ ...but, they may be overloaded or unreliable
- Reachability
 - ▣ 6to4 doesn't work if you are behind a NAT
- Possible solutions
 - ▣ IPv6 Rapid Deployment (6rd)
 - Each ISP sets up relays for its customers
 - Does not leverage the 2002:: address space
 - ▣ Teredo
 - Tunnels IPv6 packets through UDP/IPv4 tunnels
 - Can tunnel through NATs, but requires special relays

Consequences of IPv6

62

- Beware unintended consequences of IPv6
- Example: IP blacklists
 - ▣ Currently, blacklists track IPs of spammers/bots
 - ▣ Few IPv4 addresses mean list sizes are reasonable
 - ▣ Hard for spammers/bots to acquire new IPs
- Blacklists will not work with IPv6
 - ▣ Address space is enormous
 - ▣ Acquiring new IP addresses is trivial