

CSCI-351

Data communication and Networks

Lecture 8: Network Layer (Putting the Net in Internet)

Recap

2

Network Layer

3




- Function:
 - ▣ Route packets end-to-end on a network, through multiple hops
- Key challenge:
 - ▣ How to represent addresses
 - ▣ How to route packets
 - Scalability
 - Convergence

Routers, Revisited


4

- How to connect multiple LANs?
- LANs may be incompatible
 - ▣ Ethernet, Wifi, etc...
- Connected networks form an **internetwork**
 - ▣ The Internet is the best known example



Structure of the Internet

5



- Ad-hoc interconnection of networks
 - ▣ No organized topology
 - ▣ Vastly different technologies, link capacities
- Packets travel end-to-end by hopping through networks
 - ▣ Routers “peer” (connect) different networks
 - ▣ Different packets may take different routes

Internetworking Issues

6

- Naming / Addressing

- How do you designate hosts?

- Routing

- M

- Ser

- W

- H

- What happens if there are failures?

- Must deal with **heterogeneity**

- Remember, every network is different

Internet Service Model

- Best-effort (i.e. things may break)
- Store-and-forward datagram network

- ❑ Addressing
 - ❑ Class-based
 - ❑ CIDR
- ❑ IPv4 Protocol Details
 - ❑ Packed Header
 - ❑ Fragmentation
- ❑ IPv6

Possible Addressing Schemes

8

□ Flat

- ▣ e.g. each host is identified by a 48-bit MAC address
- ▣ Router needs an entry for every host in the world
 - Too big
 - Too hard to maintain (hosts come and go all the time)
 - Too slow (more later)

□ Hierarchy

- ▣ Addresses broken down into segments
- ▣ Each segment has a different level of specificity

Example: Telephone Numbers

9

1-585-475-3278



Very General

R·I·T

Updates are Local



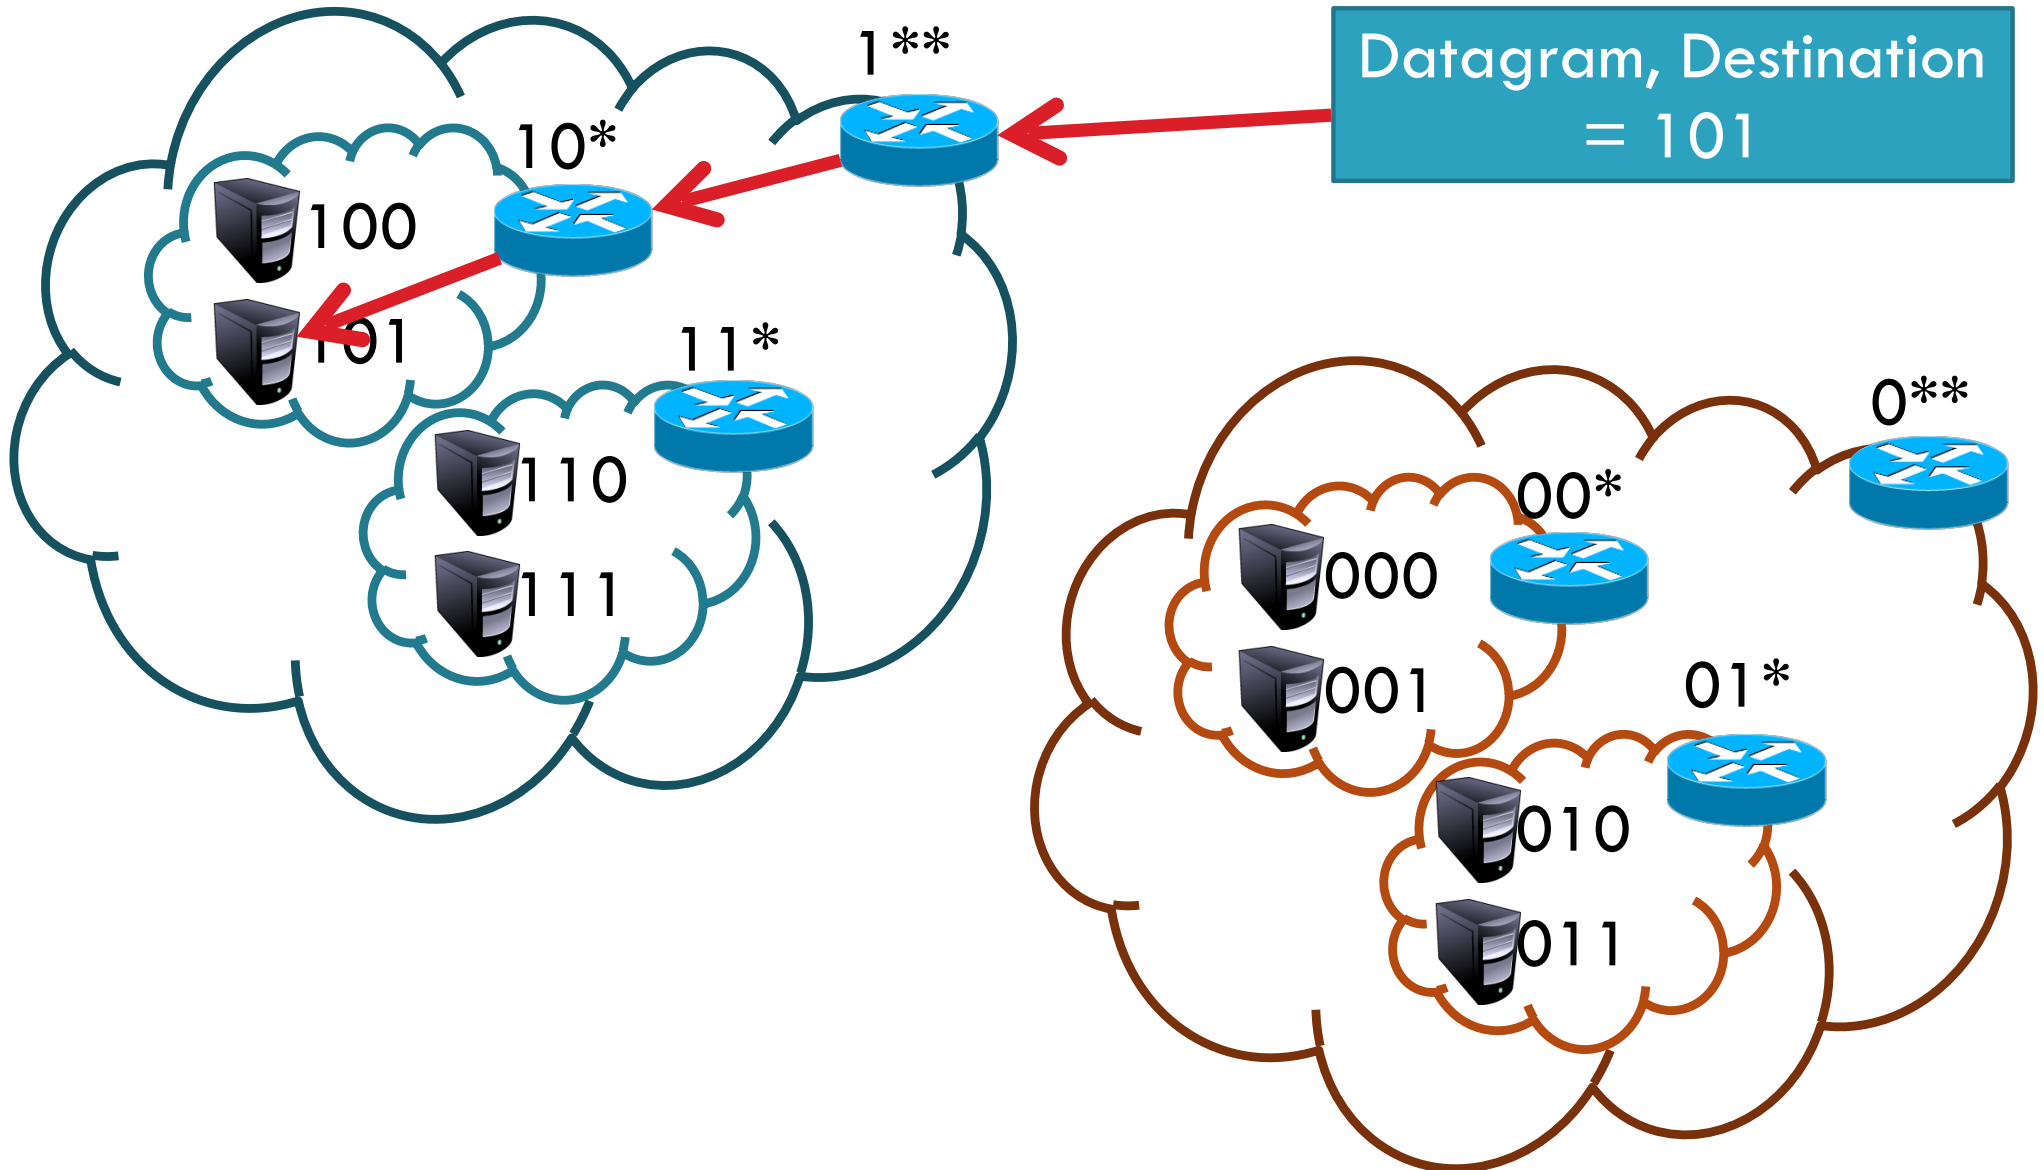
RIT GCCIS RIT GCCIS
Room 4567 Room 1234



Very Specific

Binary Hierarchy Example

10



IP Addressing

11


- IPv4: 32-bit addresses
 - ▣ Usually written in dotted notation, e.g. 192.168.21.76
 - ▣ Each number is a byte
 - ▣ Stored in Big Endian order

	0	8	16	24	31
Decimal	192	168	21	76	
Hex	C0	A8	15	4C	
Binary	11000000	10101000	00010101	01001100	

IP Addressing and Forwarding

12

- Routing Table Requirements
 - ▣ For every possible IP, give the next hop
 - ▣ But for 32-bit addresses, 2^{32} possibilities!
 - ▣ **Too slow:** 48GE ports and 4x10GE needs 176Gbps bandwidth
DRAM: ~1-6 Gbps; **TCAM** is fast, but 400x cost of DRAM
- Hierarchical address scheme
 - ▣ Separate the address into a network and a host



Recap

13


IP address

14

- IPv4: 32-bit addresses
 - Too big to fit into a routing table (2^{32})
 - Let's make a hierarchy!
 - Class-based IP address
 - Some address spaces are wasted
 - Still too many hosts in one network
 - Let's make another hierarchy!
 - Subnet Mask
 - Classless Routing

Classes of IP Addresses

15



How Do You Get IPs?

16

- IP address ranges controlled by IANA

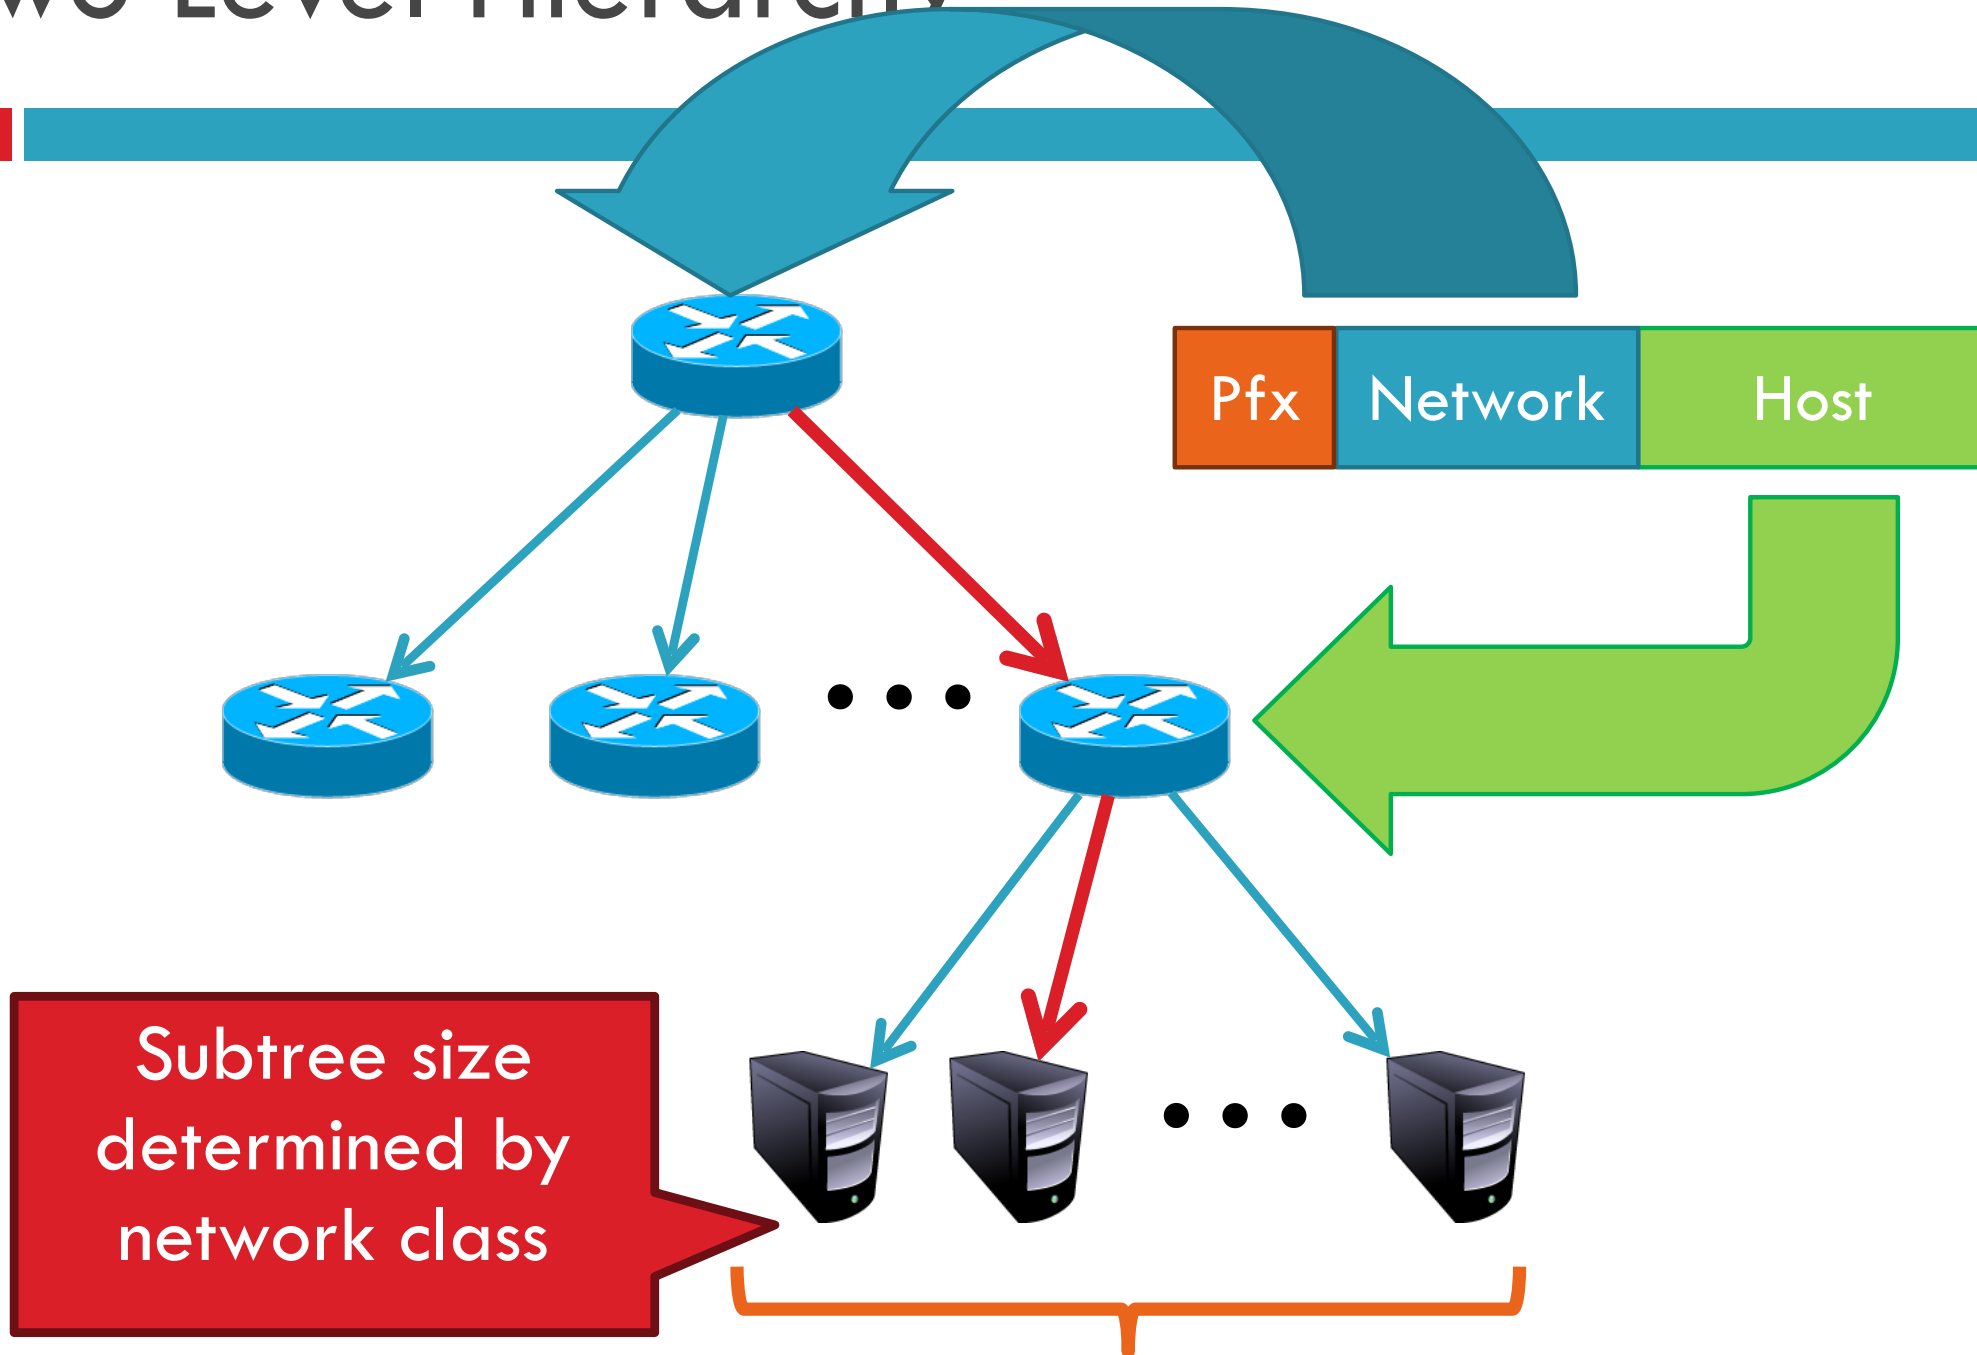


Internet Assigned Numbers Authority

- Internet Assigned Number Authority
 - Roots go back to 1972, ARPANET, UCLA
 - Today, part of ICANN
- IANA grants IPs to regional authorities
 - ARIN (American Registry of Internet Numbers) may grant you a range of IPs
 - You may then advertise routes to your new IP range
 - There are now secondary markets, auctions, ...

Two Level Hierarchy

17



Class Sizes

18

Way too big

Class	Prefix Bits	Network Bits	Number of Classes	Hosts per Class
A	1	7	$2^7 - 2 = 126$ (0 and 127 are reserved)	$2^{24} - 2 = 16,777,214$ (All 0 and all 1 are reserved)
B	2	14	$2^{14} = 16,398$	$2^{16} - 2 = 65,534$ (All 0 and all 1 are reserved)
C	3	21	$2^{21} = 2,097,512$	$2^8 - 2 = 254$ (All 0 and all 1 are reserved)
			Total: 2,114,036	

Too many
network IDs

Too small to be
useful

Subnets

19

- Problem: need to break up large A and B classes
- Solution: add another layer to the hierarchy
 - ▣ From the outside, appears to be a single network
 - Only 1 entry in routing tables
 - ▣ Internally, manage multiple subnetworks
 - Split the address range using a **subnet mask**



Subnet Mask:

Subnet Example

20

□ Extract network:

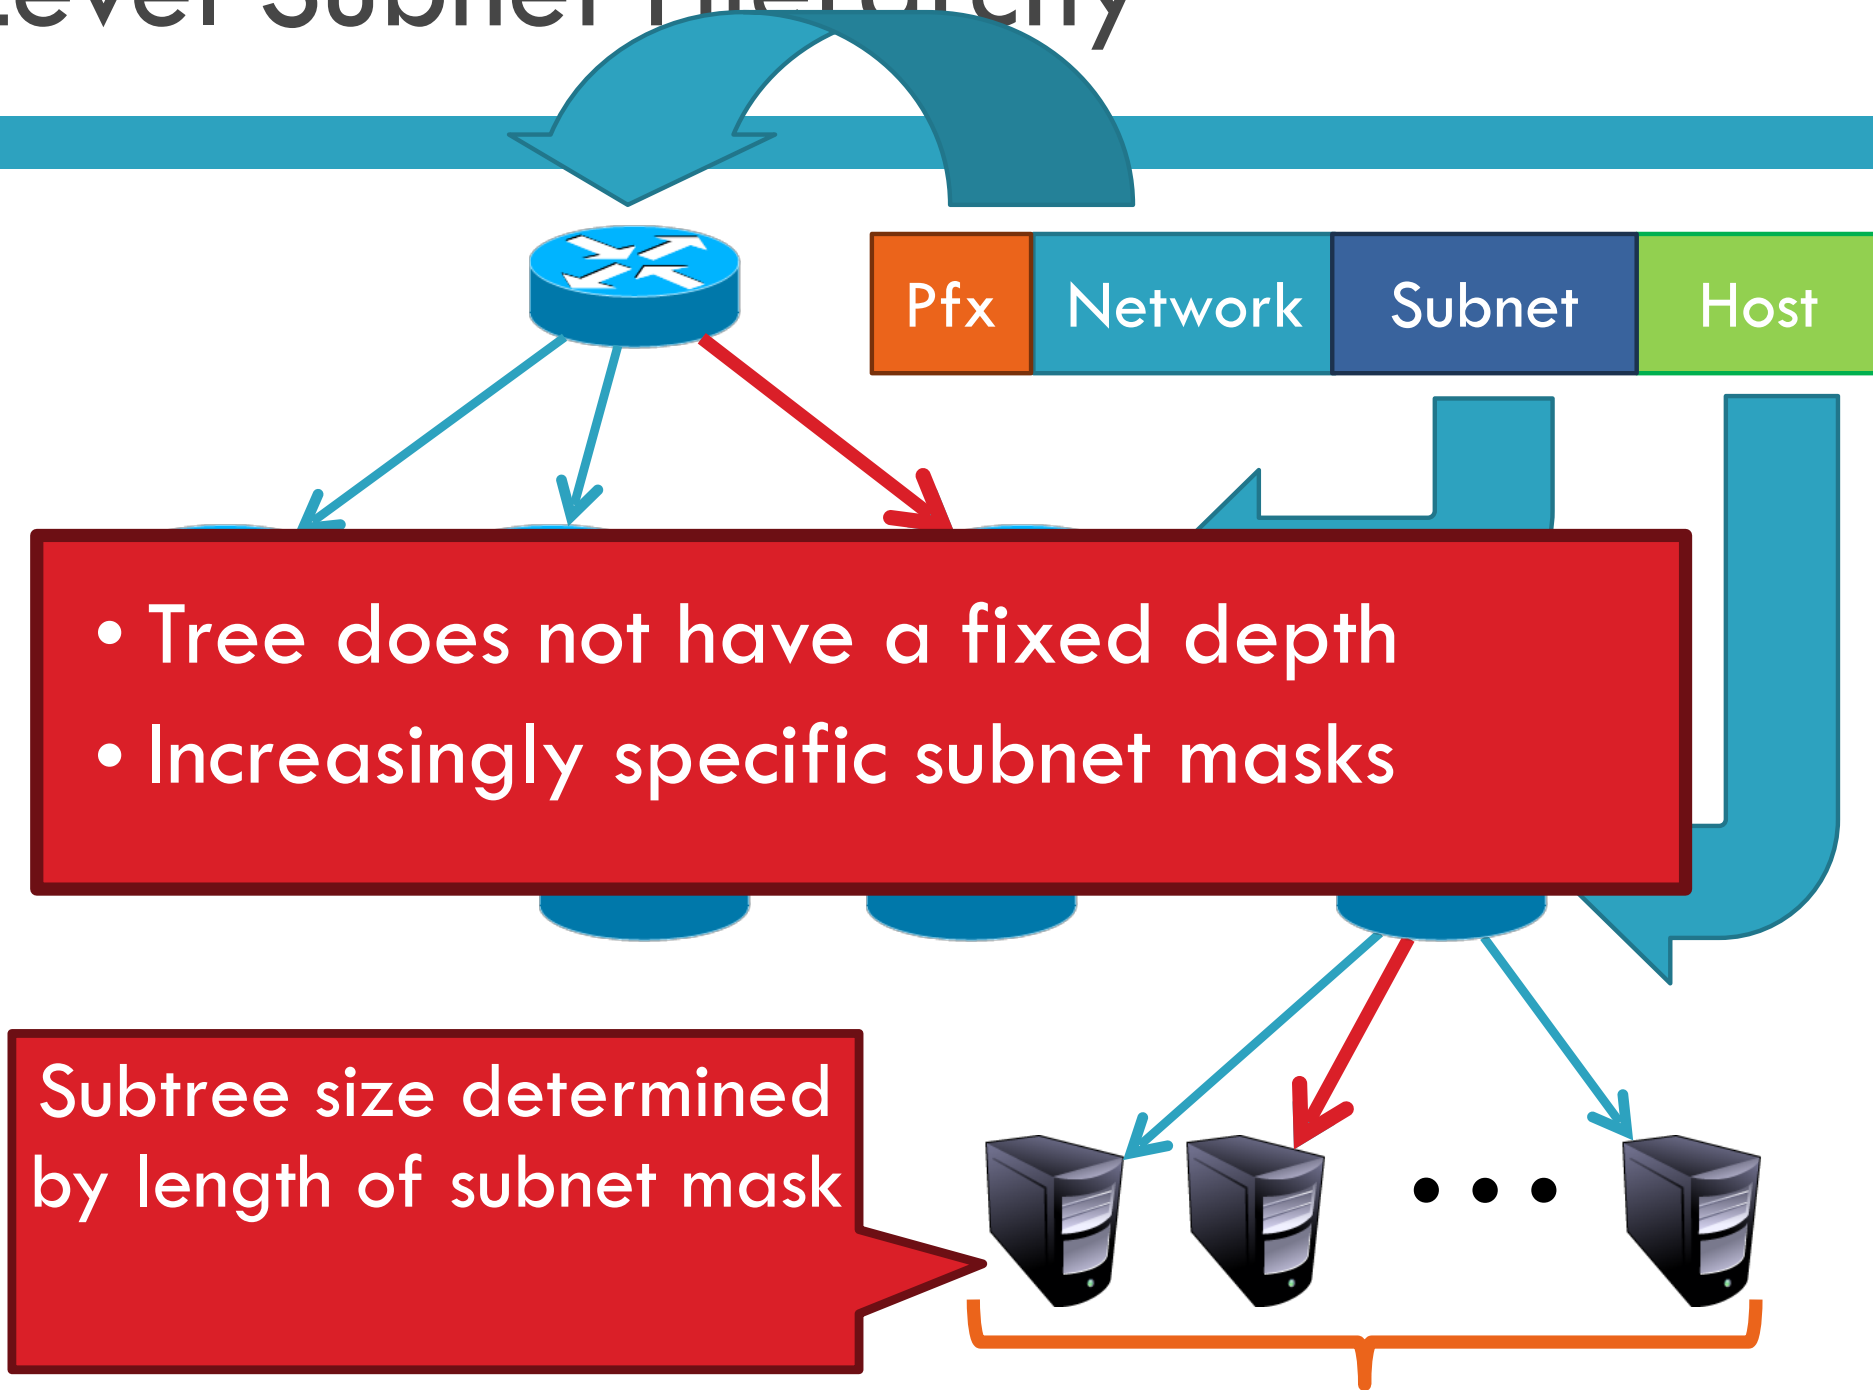
IP Address:	10110101 11011101 01010100 01110010
Subnet Mask:	& 11111111 11111111 11000000 00000000
Result:	10110101 11011101 01000000 00000000

□ Extract host:

IP Address:	10110101 11011101 01010100 01110010
Subnet Mask:	& ~(11111111 11111111 11000000 00000000)
Result:	00000000 00000000 00010100 01110010

N-Level Subnet Hierarchy

21



Example Routing Table

22

Address Pattern	Subnet Mask	Destination Router
0.0.0.0	0.0.0.0	Router 4
18.0.0.0	255.0.0.0	Router 2
128.42.0.0	255.255.0.0	Router 3
128.42.128.0	255.255.128.0	Router 5
128.42.222.0	255.255.255.0	Router 1

- Question: 128.42.222.198 matches four rows
 - ▣ Which router do we forward to?
- Longest prefix matching
 - ▣ Use the row with the longest number of 1's in the mask
 - ▣ This is the **most specific match**

Subnetting Revisited

23

- Question: does subnetting solve all the problems of class-based routing?

NO

- Classes are still too coarse
 - ▣ Class A can be subnetted, but only 126 available
 - ▣ What if MIT does not use all IPs?
 - ▣ Class C is too small
 - ▣ Class B is nice, but there are only 16,398 available
- Routing tables are still too big
 - ▣ 2.1 million entries per router

Classless Inter Domain Routing

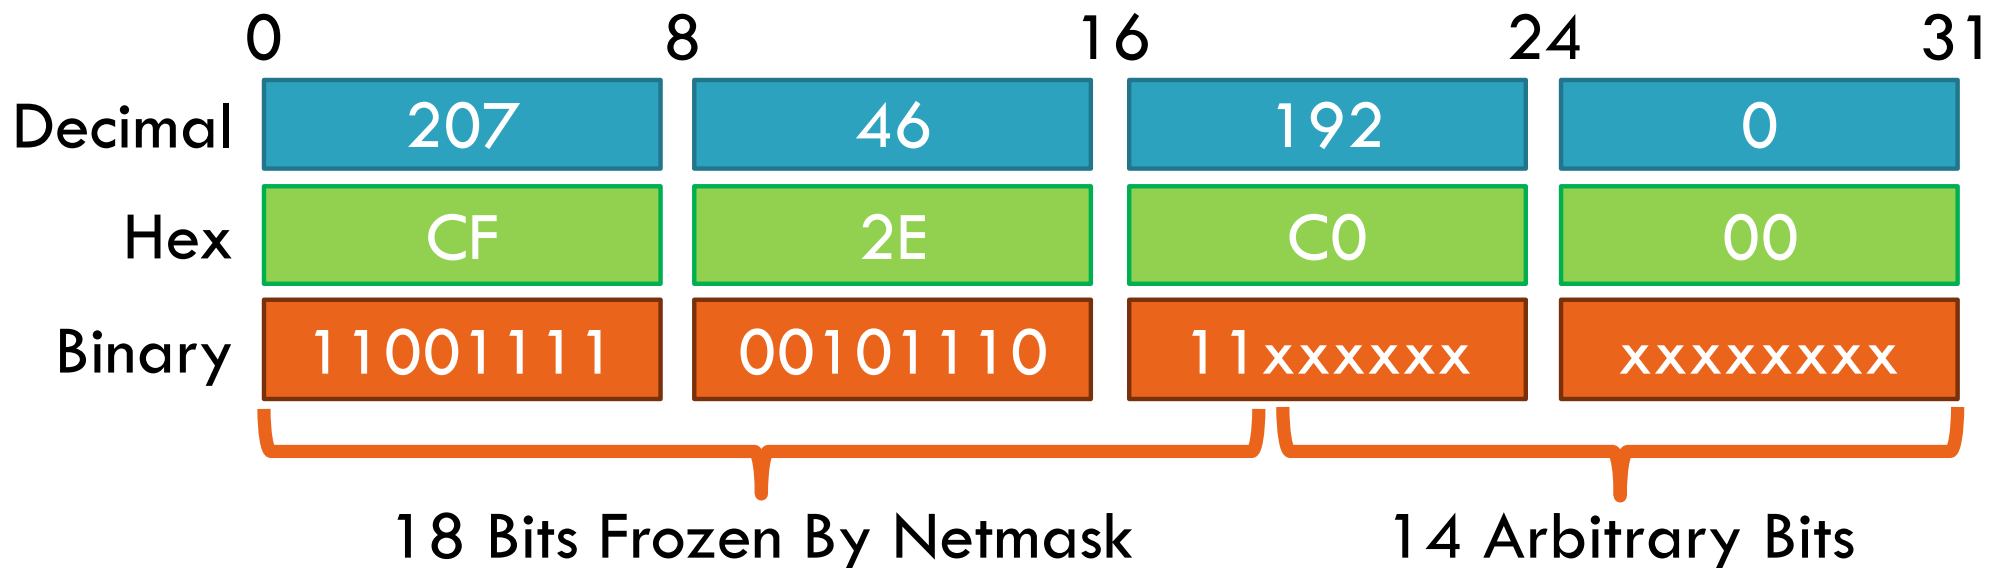
24

- CIDR, pronounced 'cider'
- Key ideas:
 - ▣ Get rid of IP classes
 - ▣ Use bitmasks for all levels of routing
 - ▣ **Aggregation** to minimize FIB (forwarding information base)
- Arbitrary split between network and host
 - ▣ Specified as a bitmask or prefix length
 - ▣ Example: Rochester Institute of Technology
 - 172.31.0.0 with **netmask** 255.255.0.0
 - 172.31.0.0 / 16

Aggregation with CIDR

25



- Original use: aggregating class C ranges
- One organization given contiguous class C ranges
 - ▣ Example: Microsoft, 207.46.192.* – 207.46.255.*
 - ▣ Represents $2^6 = 64$ class C ranges
 - ▣ Specified as CIDR address 207.46.192.0/18




Example CIDR Routing

Insert an routing entry to cover all addresses under 207.46.0.0/16

26



Address	Netmask	Third Byte	Byte Range
207.46.0.0	19	000xxxxx	0 – 31
207.46.32.0	19	001xxxxx	32 – 63
207.46.64.0	19	010xxxxx	64 – 95
207.46.128.0	18	10xxxxxx	128 – 191
207.46.192.0	18	11xxxxxx	192 – 255



Hole in the Routing Table: No coverage for 96 – 127
207.46.96.0/19

Example CIDR Routing

Insert an routing entry to cover all addresses under 207.46.0.0/16

27

Address	Bit-represented Address	
207.46.0.0/19	11001111.00101110.00000000.00000000/19	
207.46.32.0/19	11001111.00101110.00100000.00000000/19	
207.46.64.0/19	11001111.00101110.01000000.00000000/19	
207.46.128.0/18	11001111.00101110.10000000.00000000/18	
207.46.192.0/18	11001111.00101110.11000000.00000000/18	

Example CIDR Routing


Insert an routing entry to cover all addresses under 207.46.0.0/16

28

Address	Bit-represented Address	
207.46.0.0/19	11001111.00101110.00000000.00000000/19	
207.46.32.0/19	11001111.00101110.00100000.00000000/19	
207.46.64.0/19	11001111.00101110.01000000.00000000/19	
207.46.96.0/19	11001111.00101110.01100000.00000000/19	
207.46.128.0/18	11001111.00101110.10000000.00000000/18	
207.46.192.0/18	11001111.00101110.11000000.00000000/18	

Size of CIDR Routing Tables

29



- From www.cidr-report.org
- CIDR has kept IP routing table sizes in check
 - ▣ Currently ~450,000 entries for a complete IP routing table
 - ▣ Only required by backbone routers

Takeaways

30

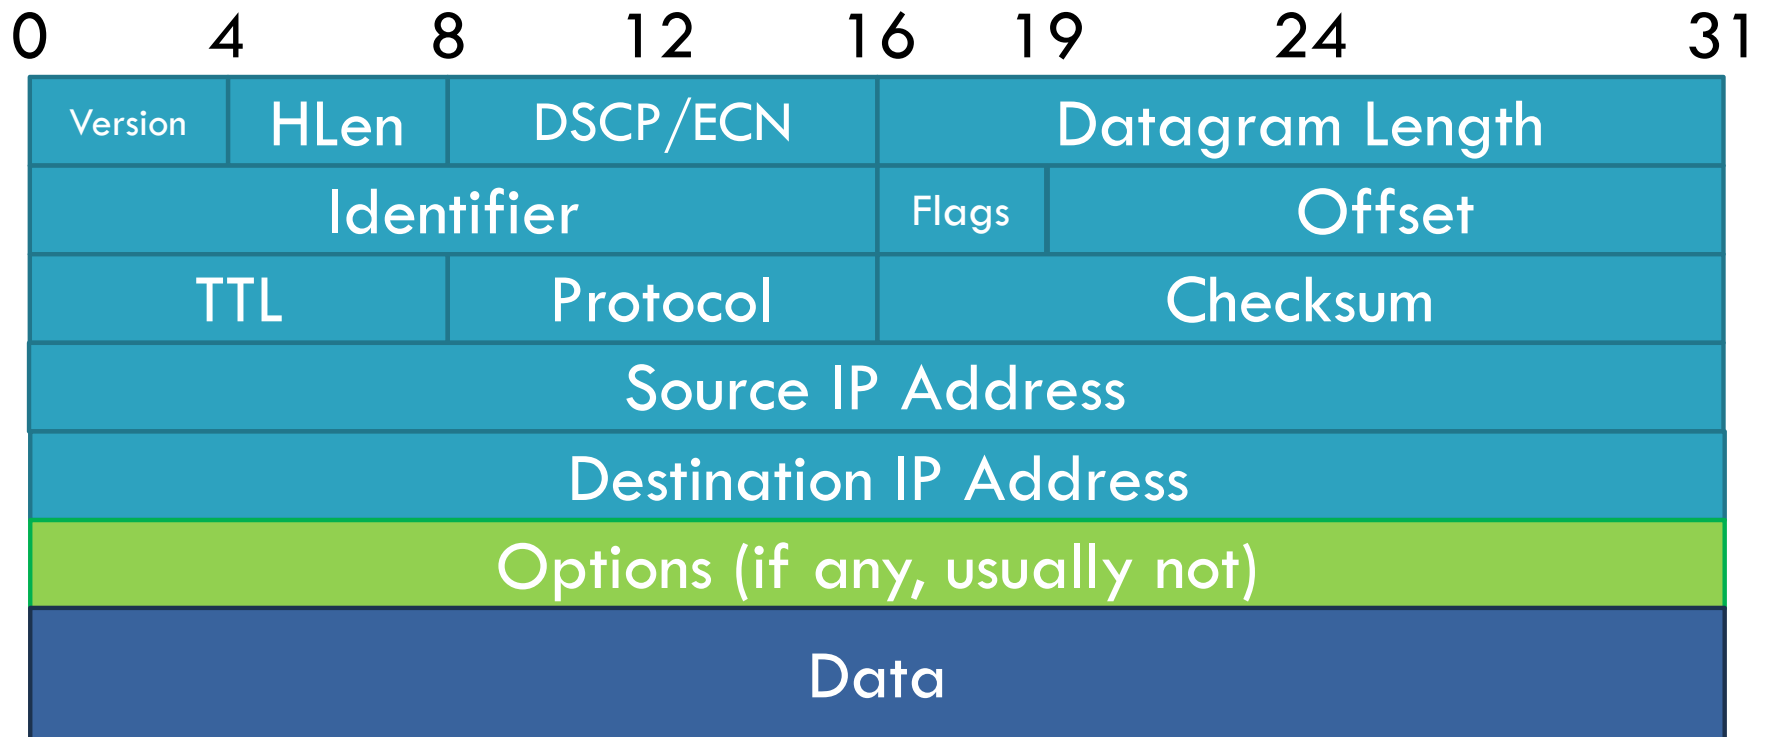
- Hierarchical addressing is critical for scalability
 - ▣ Not all routers need all information
 - ▣ Limited number of routers need to know about changes
- Non-uniform hierarchy useful for heterogeneous networks
 - ▣ Class-based addressing is too coarse
 - ▣ CIDR improves scalability and granularity
- Implementation challenges
 - ▣ Longest prefix matching is more difficult than schemes with no ambiguity

- ❑ Addressing
 - ❑ Class-based
 - ❑ CIDR
- ❑ IPv4 Protocol Details
 - ❑ Packed Header
 - ❑ Fragmentation
- ❑ IPv6

IP Datagrams

32

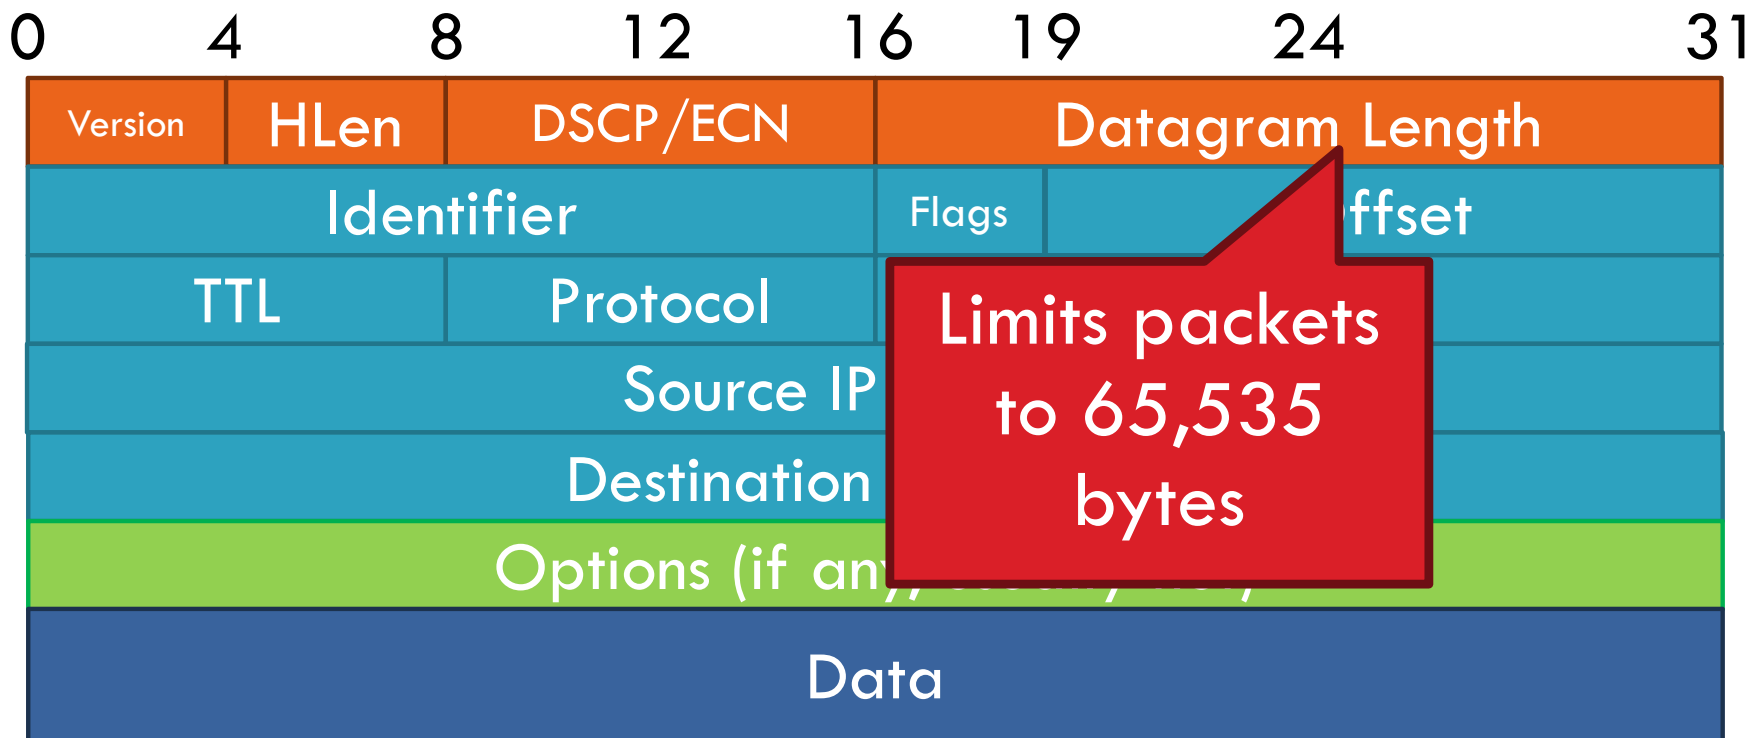
- IP Datagrams are like a letter
 - ▣ Totally self-contained
 - ▣ Include all necessary addressing information
 - ▣ No advanced setup of connections or circuits



IP Header Fields: Word 1

33

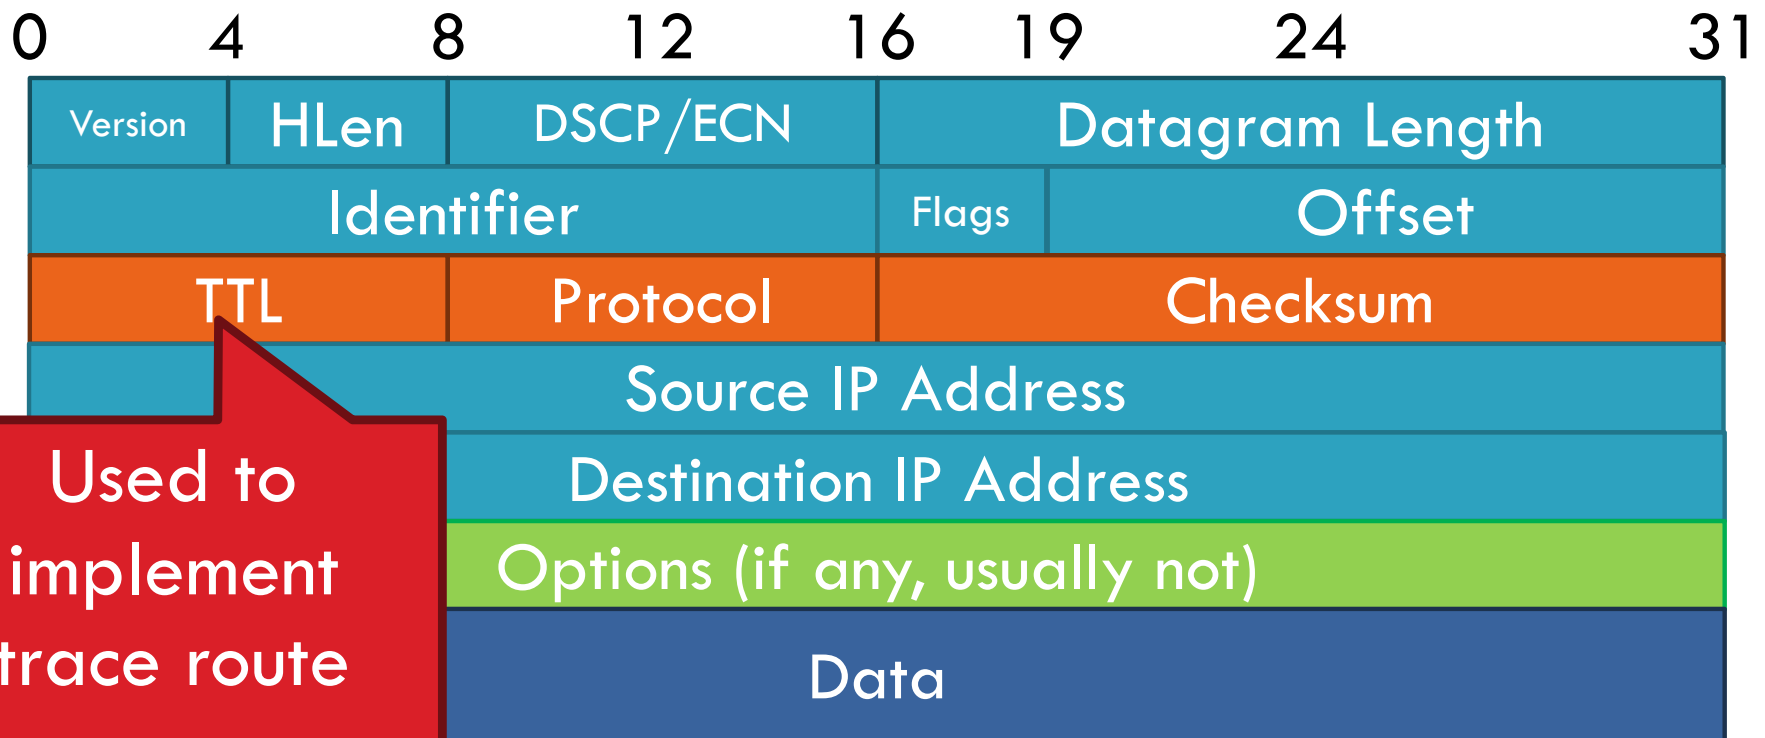
- Version: 4 for IPv4
- Header Length: Number of 32-bit words (usually 5)
- Type of Service: Priority information (unused)
- Datagram Length: Length of header + data in bytes



IP Header Fields: Word 3

34

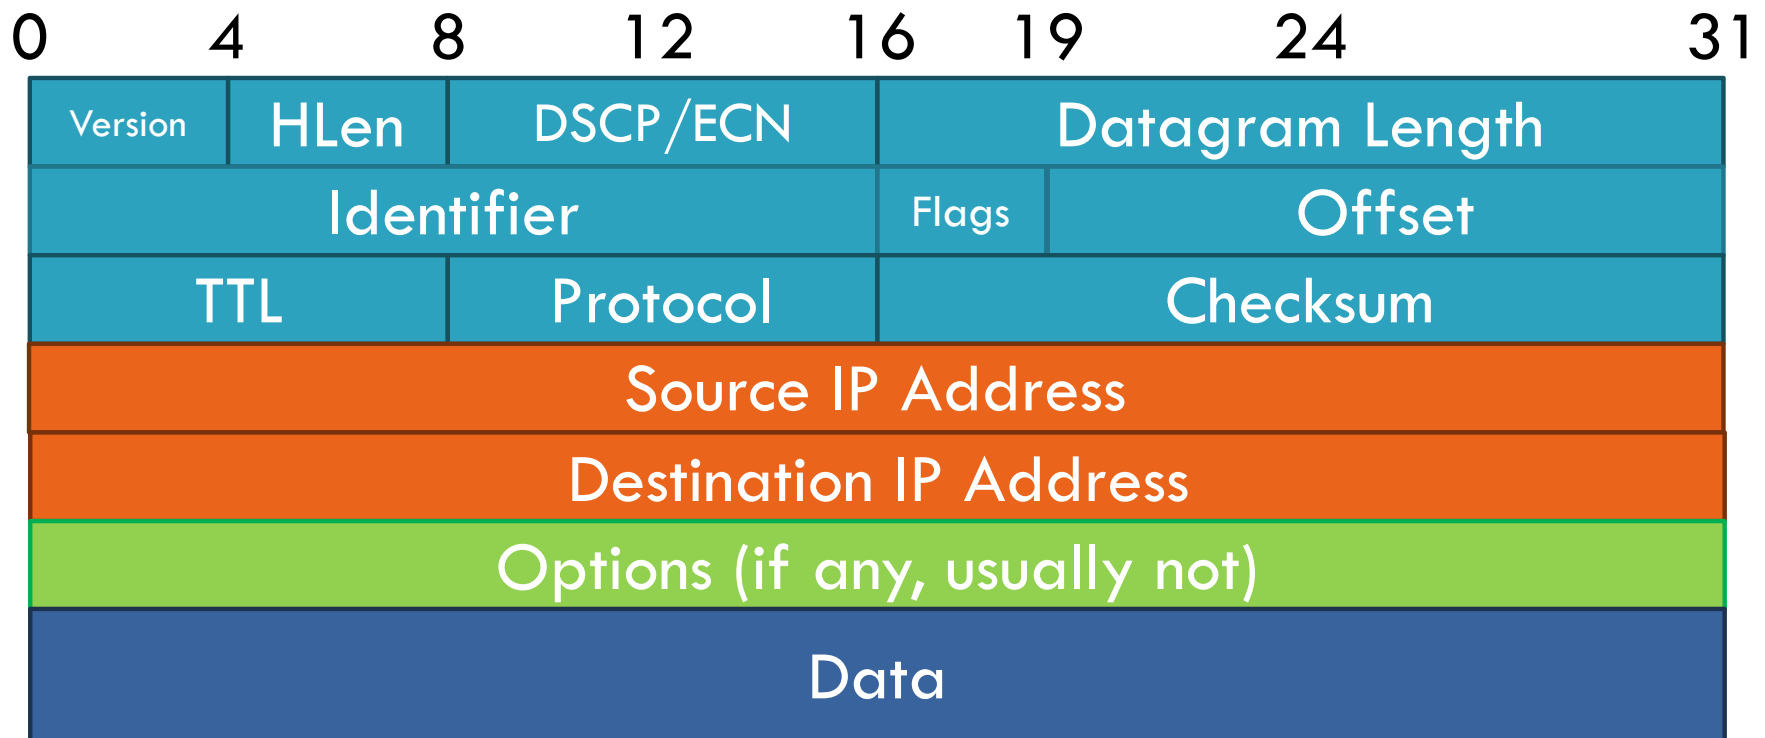
- Time to Live: decremented by each router
 - ▣ Used to kill looping packets
- Protocol: ID of encapsulated protocol
 - ▣ 6 = TCP, 17 = UDP
- Checksum



IP Header Fields: Word 4 and 5

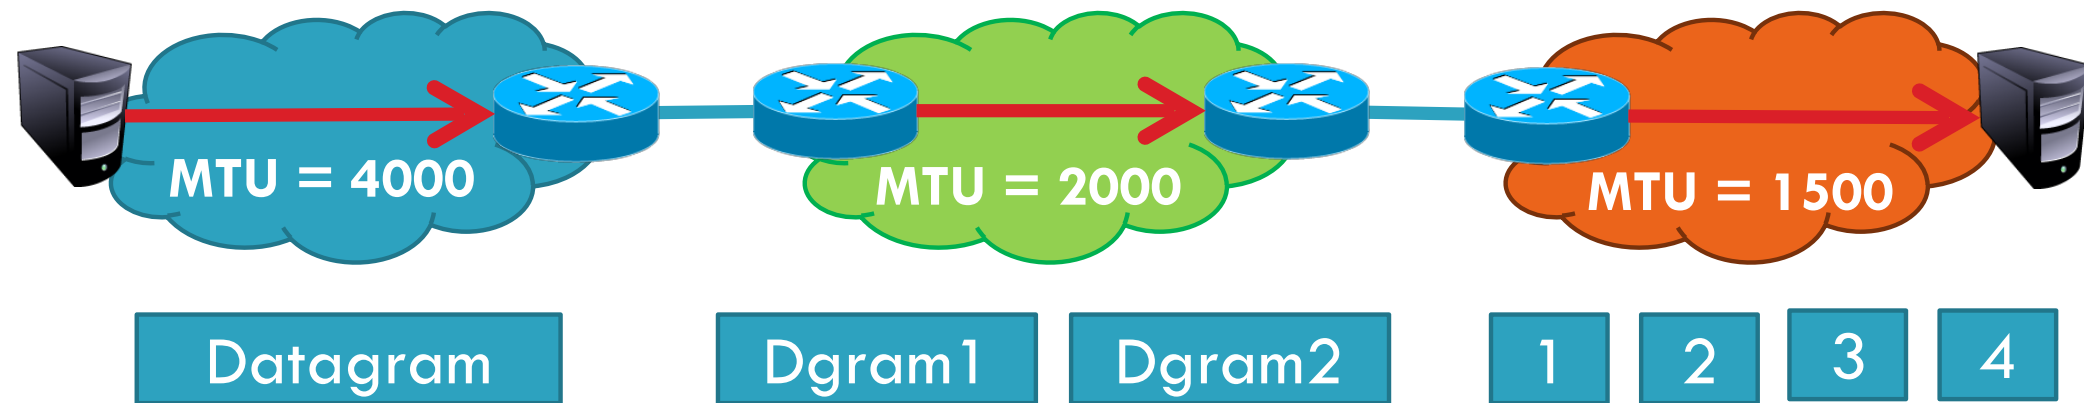
35

- Source and destination address
 - ▣ In theory, must be globally unique
 - ▣ In practice, this is often violated



Problem: Fragmentation

36

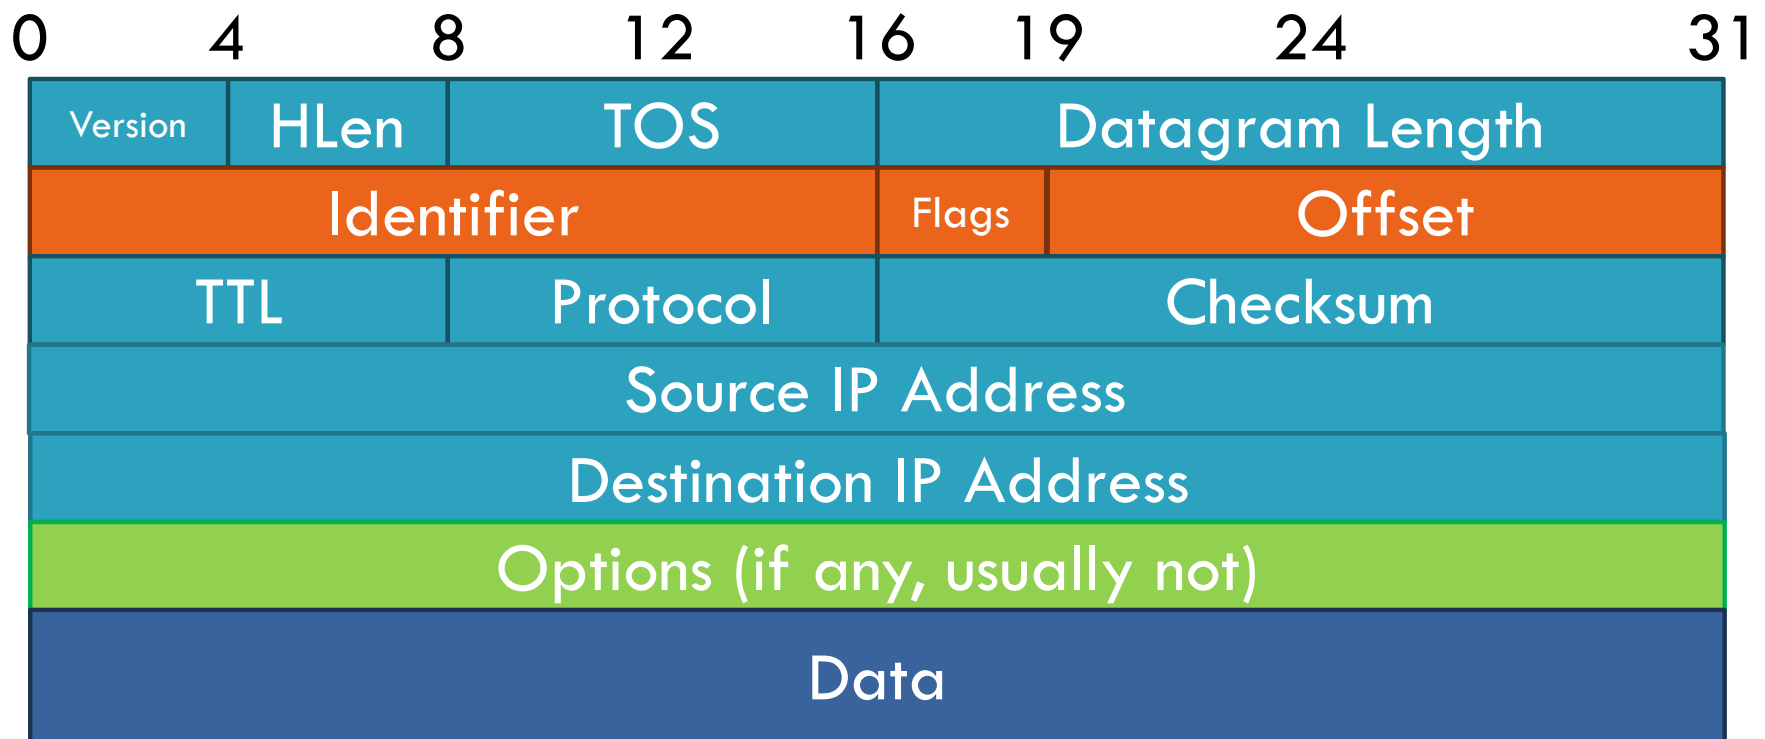


- Problem: each network has its own MTU
 - ▣ DARPA principles: networks allowed to be heterogeneous
 - ▣ Minimum MTU may not be known for a given path
- IP Solution: fragmentation
 - ▣ Split datagrams into pieces when MTU is reduced
 - ▣ Reassemble original datagram at the receiver

IP Header Fields: Word 2

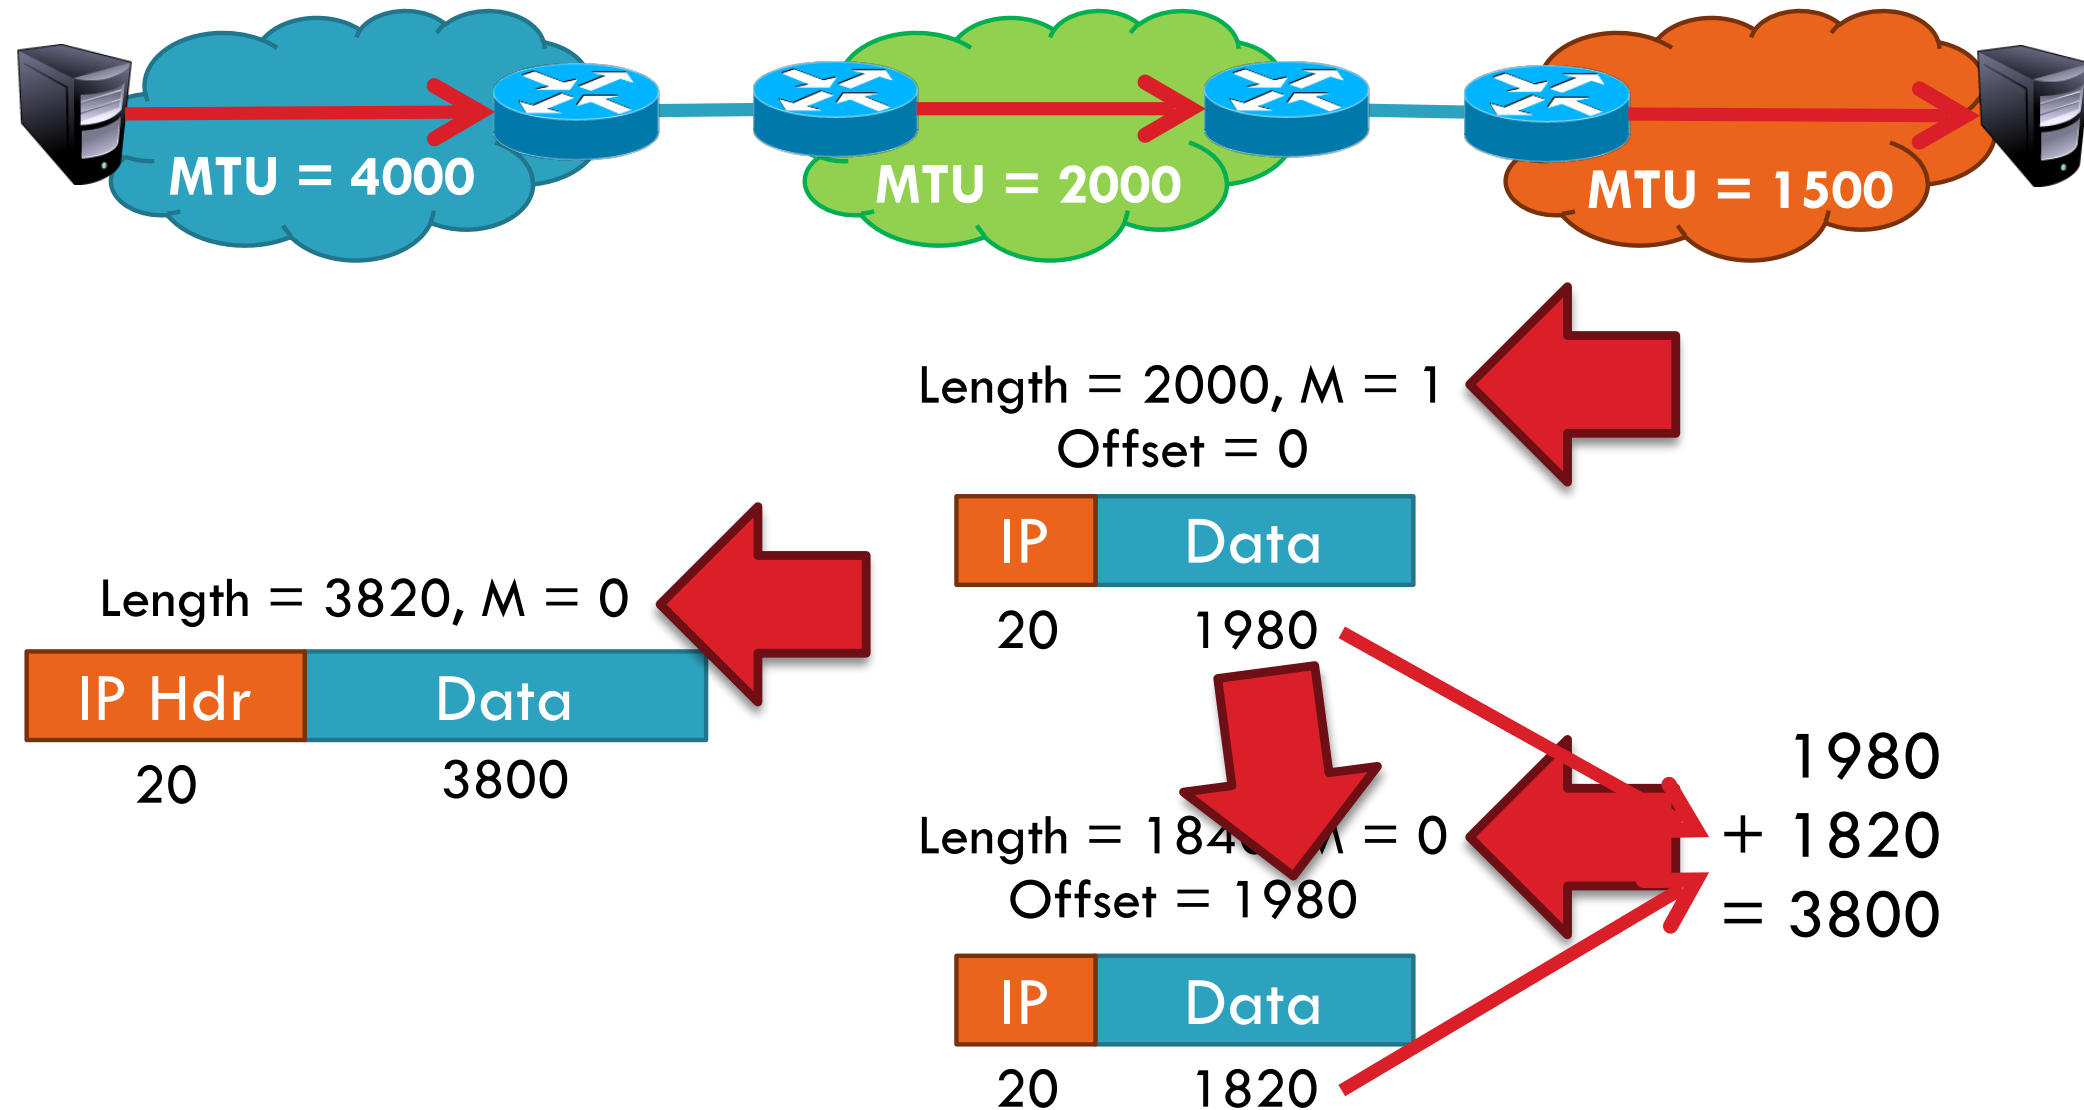
37

- Identifier: a unique number for the original datagram
- Flags: M flag (3 bits)
- Offset: byte position of the first byte in the fragment
 - ▣ Divided by 8



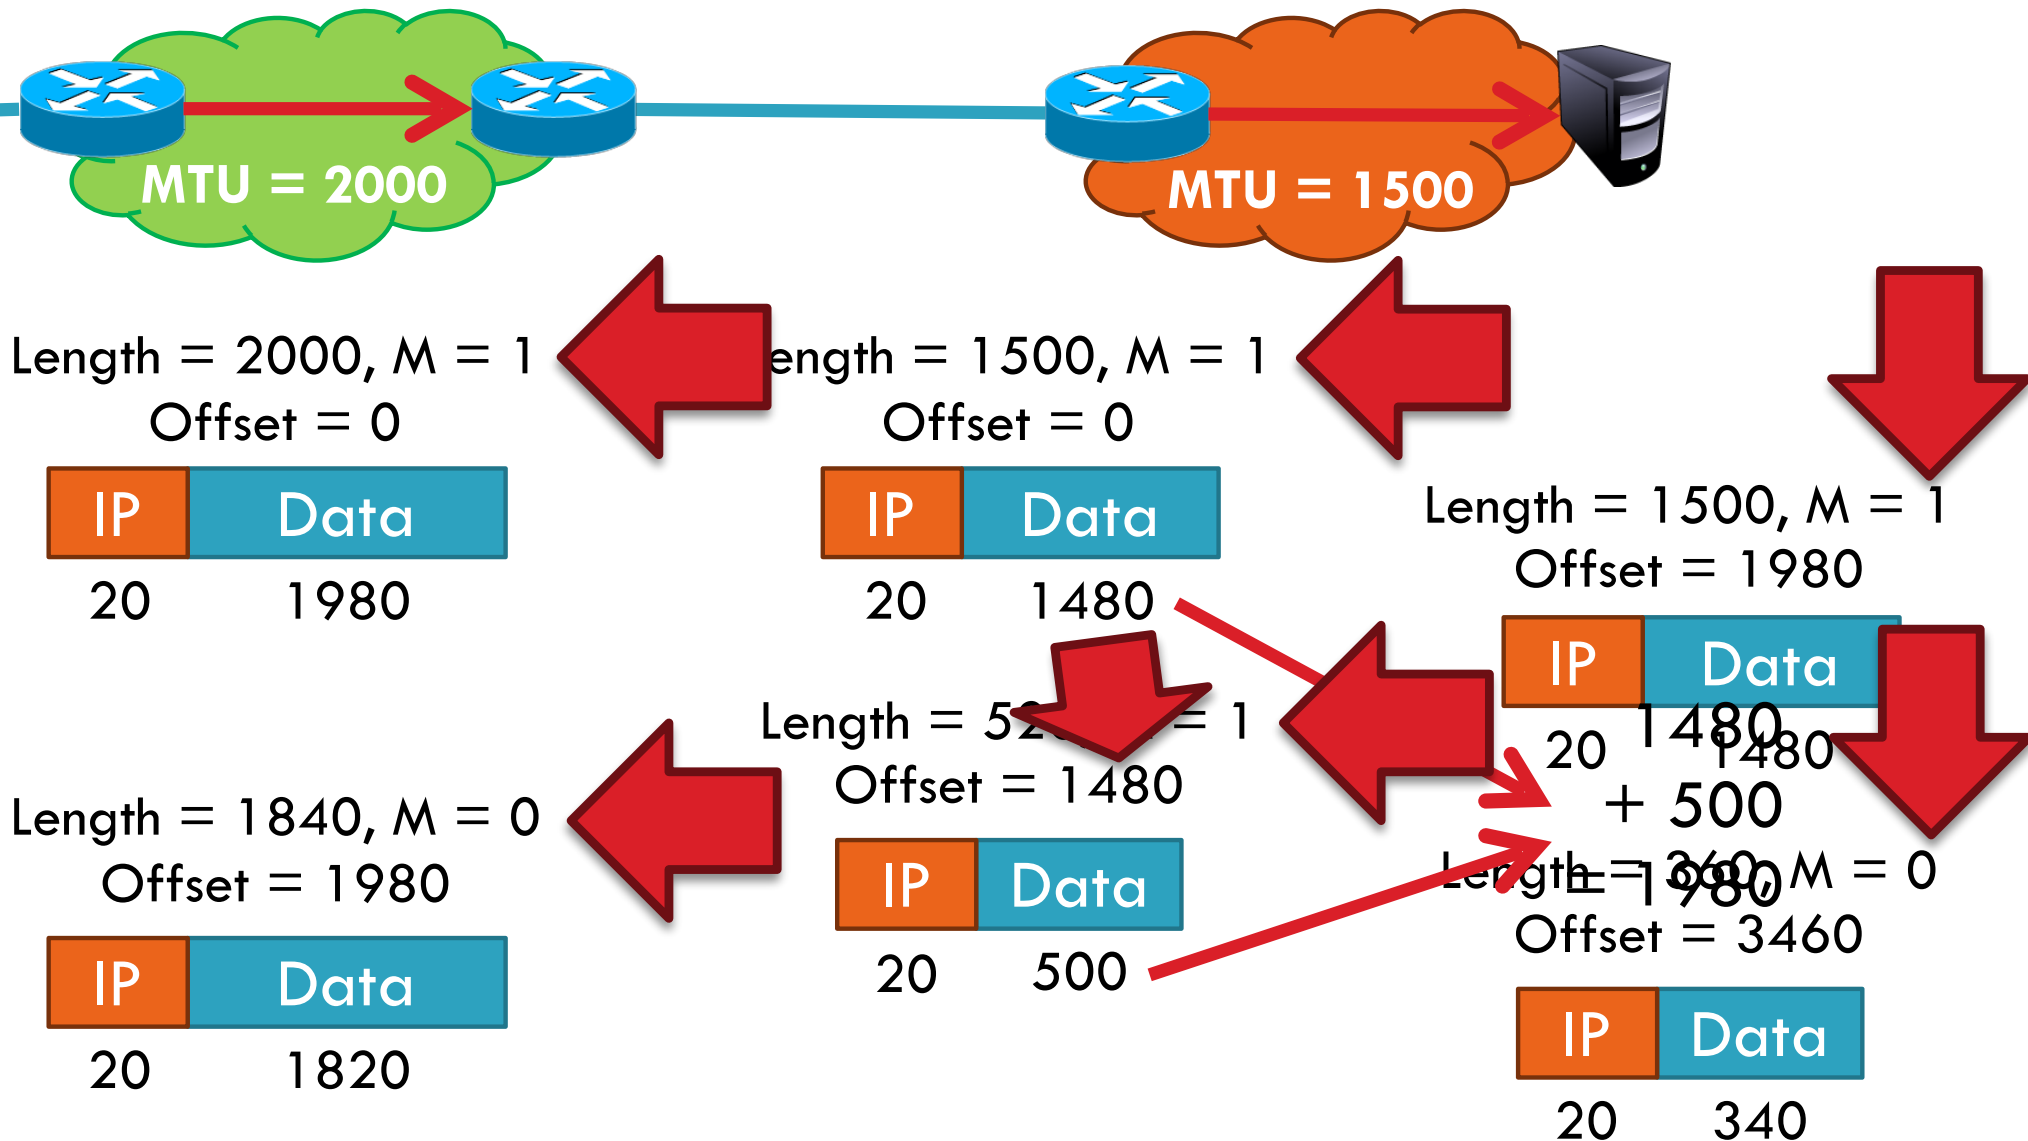
Fragmentation Example

38



Fragmentation Example

39




IP Fragment Reassembly

40


Length = 1500, $M = 1$, Offset = 0




Length = 520, $M = 1$, Offset = 1480



Length = 1500, $M = 1$, Offset = 1980



Length = 360, $M = 0$, Offset = 3460



- Performed at destination
- $M = 0$ fragment gives us total data size
 - ▣ $360 - 20 + 3460 = 3800$
- Challenges:
 - ▣ Out-of-order fragments
 - ▣ Duplicate fragments
 - ▣ Missing fragments
- Basically, memory management nightmare

Fragmentation Concepts

41

- Highlights many key Internet characteristics
 - ▣ Decentralized and heterogeneous
 - Each network may choose its own MTU
 - ▣ Connectionless datagram protocol
 - Each fragment contains full routing information
 - Fragments can travel independently, on different paths
 - ▣ Best effort network
 - Routers/receiver may silently drop fragments
 - No requirement to alert the sender
 - ▣ Most work is done at the endpoints
 - i.e. reassembly

Fragmentation in Reality

42

- Fragmentation is expensive
 - ▣ Memory and CPU overhead for datagram reconstruction
 - ▣ Want to avoid fragmentation if possible
- MTU discovery protocol
 - ▣ Send a packet with “don’t fragment” bit set
 - ▣ Keep decreasing message length until one arrives
 - ▣ May get “can’t fragment” error from a router, which will explicitly state the supported MTU
- Router handling of fragments
 - ▣ Fast, specialized hardware handles the common case
 - ▣ Dedicated, general purpose CPU just for handling fragments
- Examples
 - ▣ Censorship?

- ❑ Addressing
 - ❑ Class-based
 - ❑ CIDR
- ❑ IPv4 Protocol Details
 - ❑ Packed Header
 - ❑ Fragmentation
- ❑ IPv6

The IPv4 Address Space Crisis

44

- Problem: the IPv4 address space is too small
 - ▣ $2^{32} = 4,294,967,296$ possible addresses
 - ▣ Less than one IP per person
- Parts of the world have already run out of addresses
 - ▣ IANA assigned the last /8 block of addresses in 2011

Region	Regional Internet Registry (RIR)	Exhaustion Date
Asia/Pacific	APNIC	April 19, 2011
Europe/Middle East	RIPE	September 14, 2012
North America	ARIN	September 24, 2015
South America	LACNIC	June 10, 2014
Africa	AFRINIC	2019 (Projected)

IPv6

45

- IPv6, first introduced in 1998(!)
 - ▣ 128-bit addresses
 - ▣ $4.8 * 10^{28}$ addresses per person
- Address format
 - ▣ 8 groups of 16-bit values, separated by ':'
 - ▣ Leading zeroes in each group may be omitted
 - ▣ Groups of zeroes can be omitted using '::'

2001:0db8:0000:0000:0000:ff00:0042:8329

2001:0db8:0:0:0:ff00:42:8329

2001:0db8::ff00:42:8329

IPv6 Trivia

46

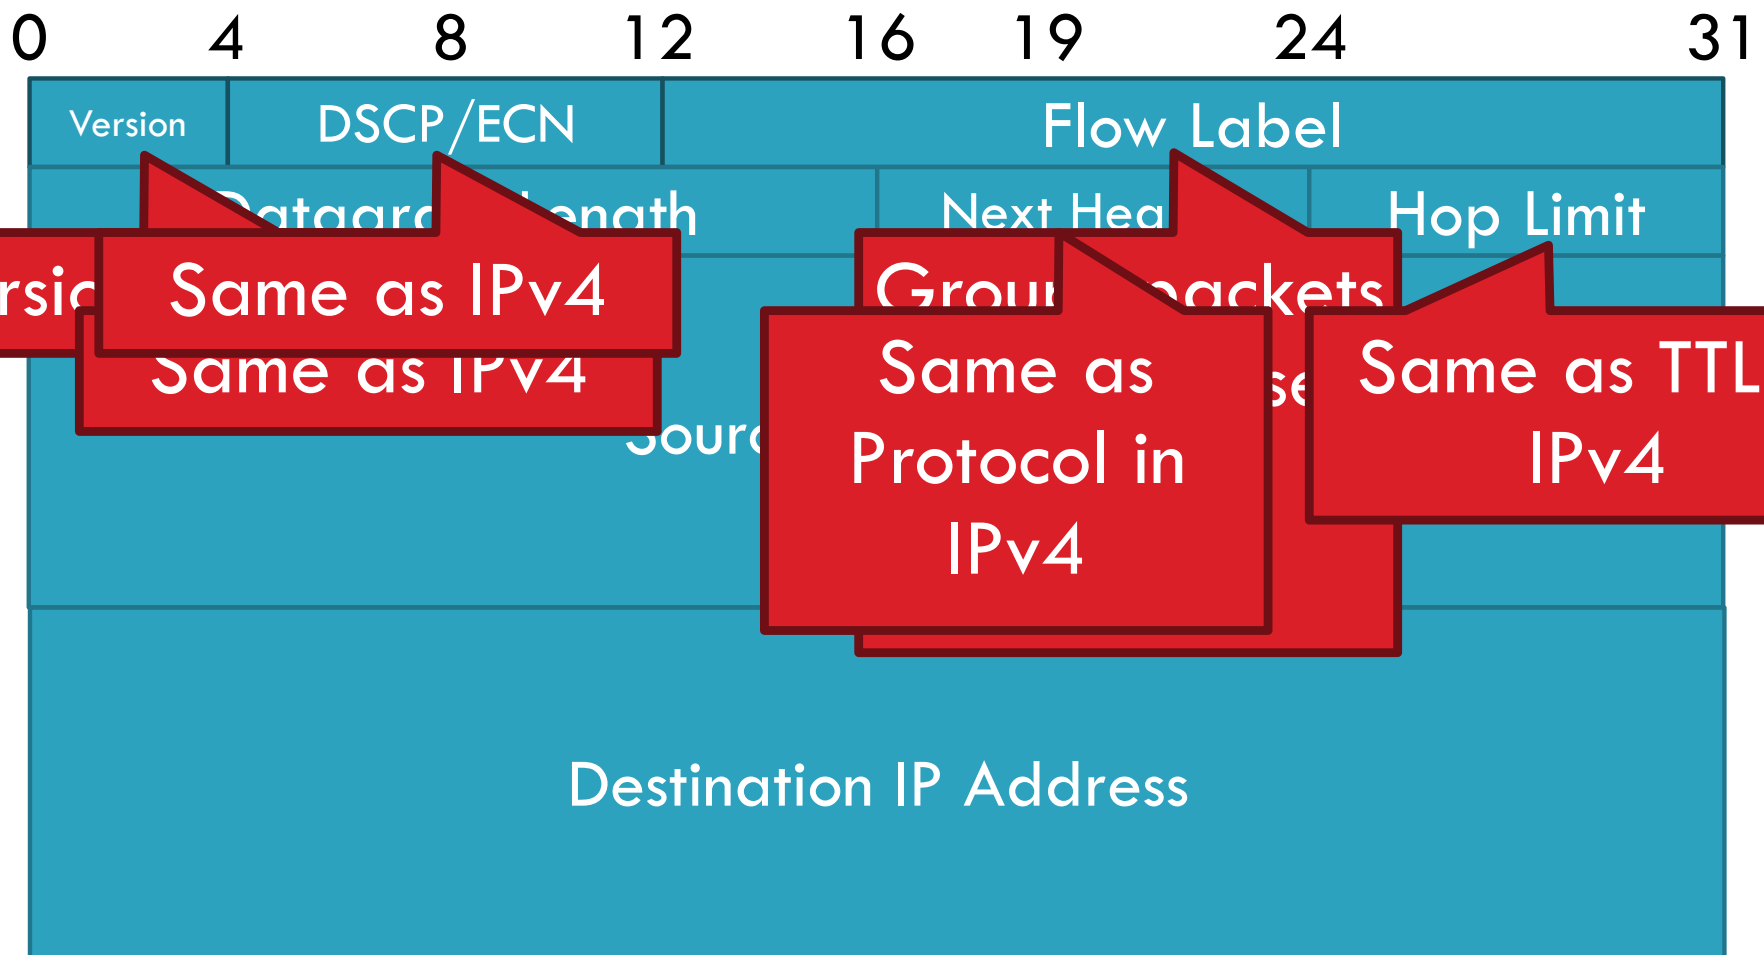
- Who knows the IP for localhost?
 - ▣ 127.0.0.1

- What is localhost in IPv6?
 - ▣ ::1

IPv6 Header

47

- Double the size of IPv4 (320 bits vs. 160 bits)



Differences from IPv4 Header

48

- Several header fields are missing in IPv6
 - ▣ Header length – rolled into Next Header field
 - ▣ Checksum – was useless, so why keep it
 - ▣ Identifier, Flags, Offset
 - IPv6 routers do not support fragmentation
 - Hosts are expected to use path MTU discovery
- Reflects changing Internet priorities
 - ▣ Today's networks are more homogeneous
 - ▣ Instead, routing cost and complexity dominate
- No security vulnerabilities due to IP fragments

Performance Improvements

49

- No checksums to verify
- No need for routers to handle fragmentation
- Simplified routing table design
 - ▣ Address space is huge
 - ▣ No need for CIDR (but need for aggregation)
 - ▣ Standard subnet size is 2^{64} addresses
- Simplified auto-configuration
 - ▣ Neighbor Discovery Protocol
 - ▣ Used by hosts to determine network ID
 - ▣ Host ID can be random!

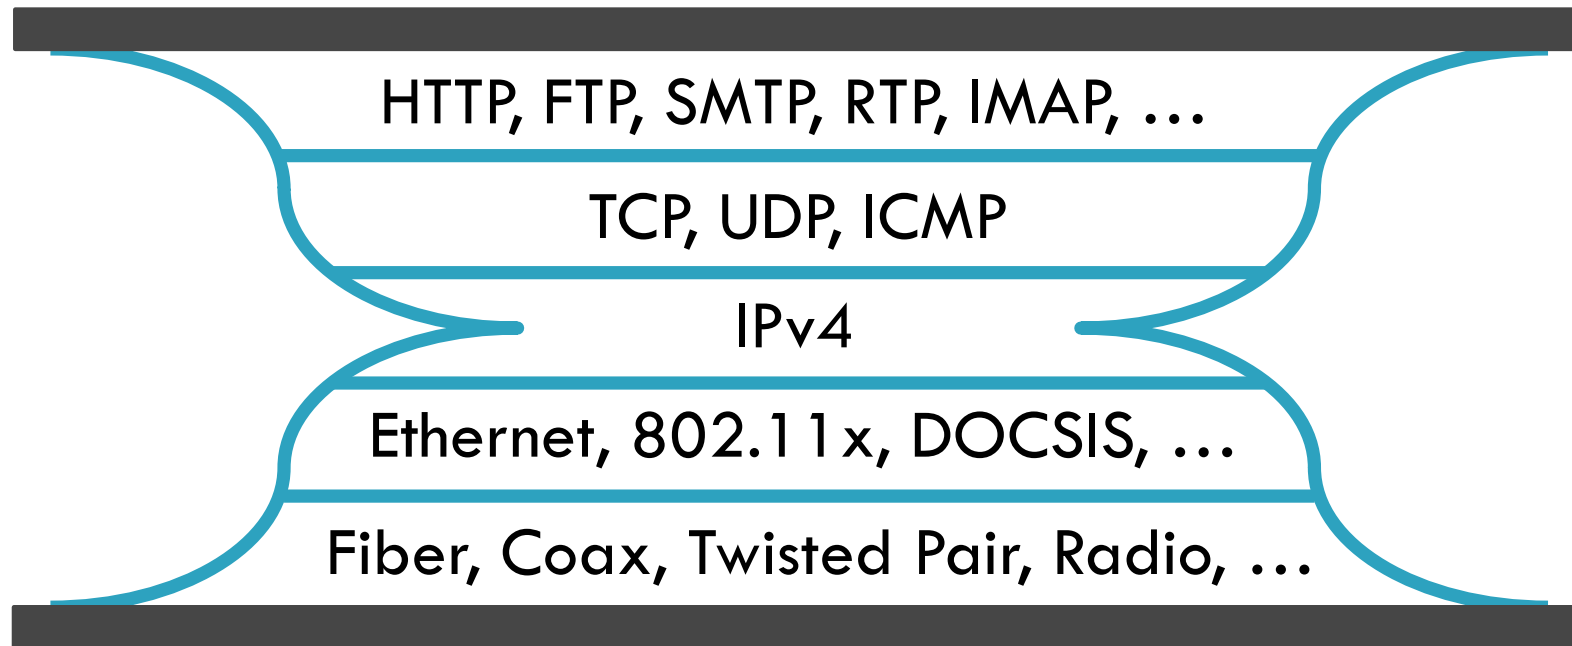
Additional IPv6 Features

50

- Source Routing
 - ▣ Host specifies the route to wants packet to take
- Mobile IP
 - ▣ Hosts can take their IP with them to other networks
 - ▣ Use source routing to direct packets
- Privacy Extensions
 - ▣ Randomly generate host identifiers
 - ▣ Make it difficult to associate one IP to a host
- Jumbograms
 - ▣ Support for 4Gb datagrams

Deployment Challenges

51



- Switching to IPv6 is a whole-Internet upgrade
 - ▣ All routers, all hosts
 - ▣ ICMPv6, DHCPv6, DNSv6
- 2013: 0.94% of Google traffic was IPv6


<https://www.google.com/intl/en/ipv6/statistics.html>

52

IPv6 Adoption

We are continuously measuring the availability of IPv6 connectivity among Google users. The graph shows the percentage of users that access Google over IPv6.

Native: 9.90% 6to4/Teredo: 0.01% Total IPv6: 9.91% | Dec 12, 2015




Server-side reachability

53

Percentage of Alexa Top 1000 websites currently reachable over IPv6


Measurements every hour from AS35425



Transitioning to IPv6

54

- How do we ease the transition from IPv4 to IPv6?
 - ▣ Today, most network edges are IPv6 ready
 - Windows/OSX/iOS/Android all support IPv6
 - Your wireless access point probably supports IPv6
 - ▣ The Internet core is hard to upgrade
 - ▣ ... but a IPv4 core cannot route IPv6 traffic



Transition Technologies


55

- How do you route IPv6 packets over an IPv4 Internet?
- Transition Technologies
 - ▣ Use **tunnels** to **encapsulate** and route IPv6 packets over the IPv4 Internet
 - ▣ Several different implementations
 - 6to4
 - IPv6 Rapid Deployment (6rd)
 - Teredo
 - ... etc.

6to4 Basics

56


- Problem: you've been assigned an IPv4 address, but you want an IPv6 address
 - ▣ Your ISP can't or won't give you an IPv6 address
 - ▣ You can't just arbitrarily choose an IPv6 address
- Solution: construct a 6to4 address
 - ▣ 6to4 addresses always start with 2002::
 - ▣ Embed the 32-bit IPv4 inside the 128-bit IPv6 address



Routing from 6to4 to 6to4

57

- How does a host using 6to4 send a packet to another host using 6to4?



Routing from 6to4 to Native IPv6

58

Special, anycasted IPv4
address for
6to4 Relay Routers

Dest: 1893:92:13:99::

Dest: 192.88.99.1

IPv4
Internet

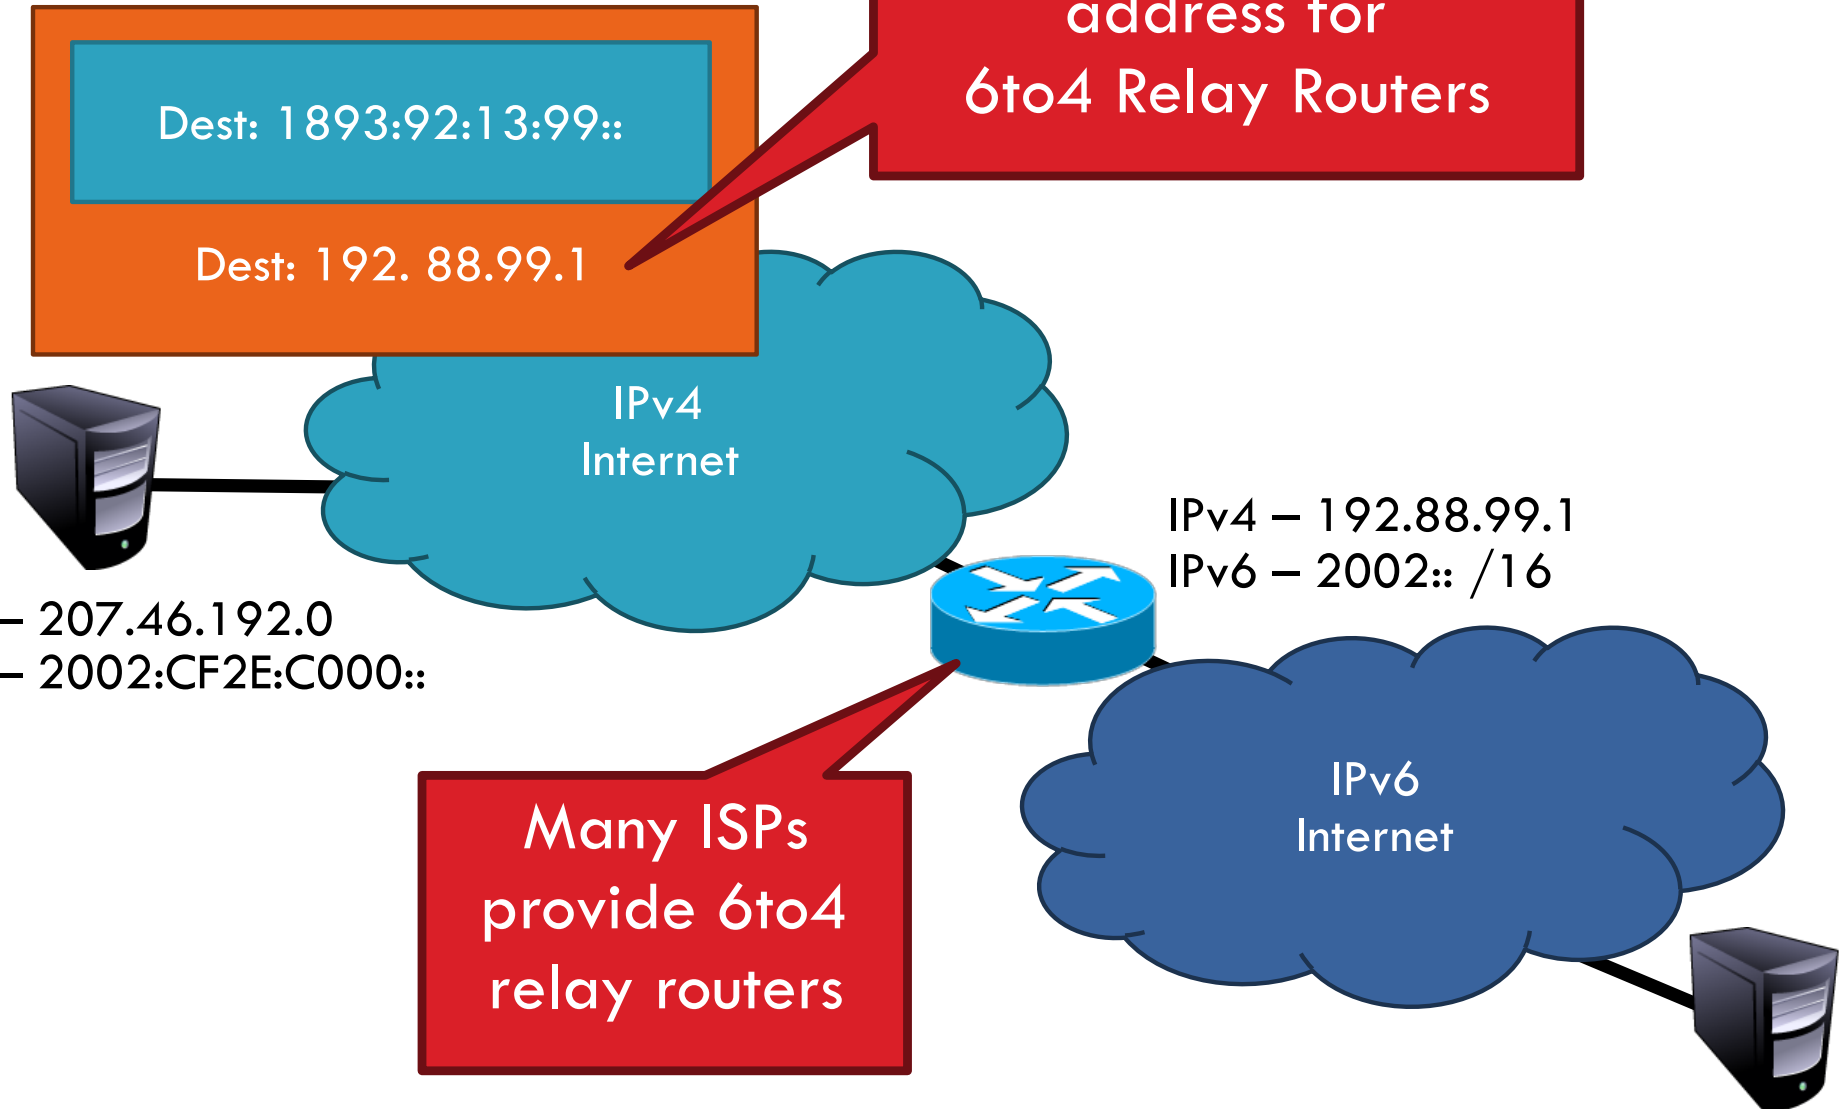
IPv4 – 192.88.99.1
IPv6 – 2002:: /16

IPv6
Internet

Many ISPs
provide 6to4
relay routers

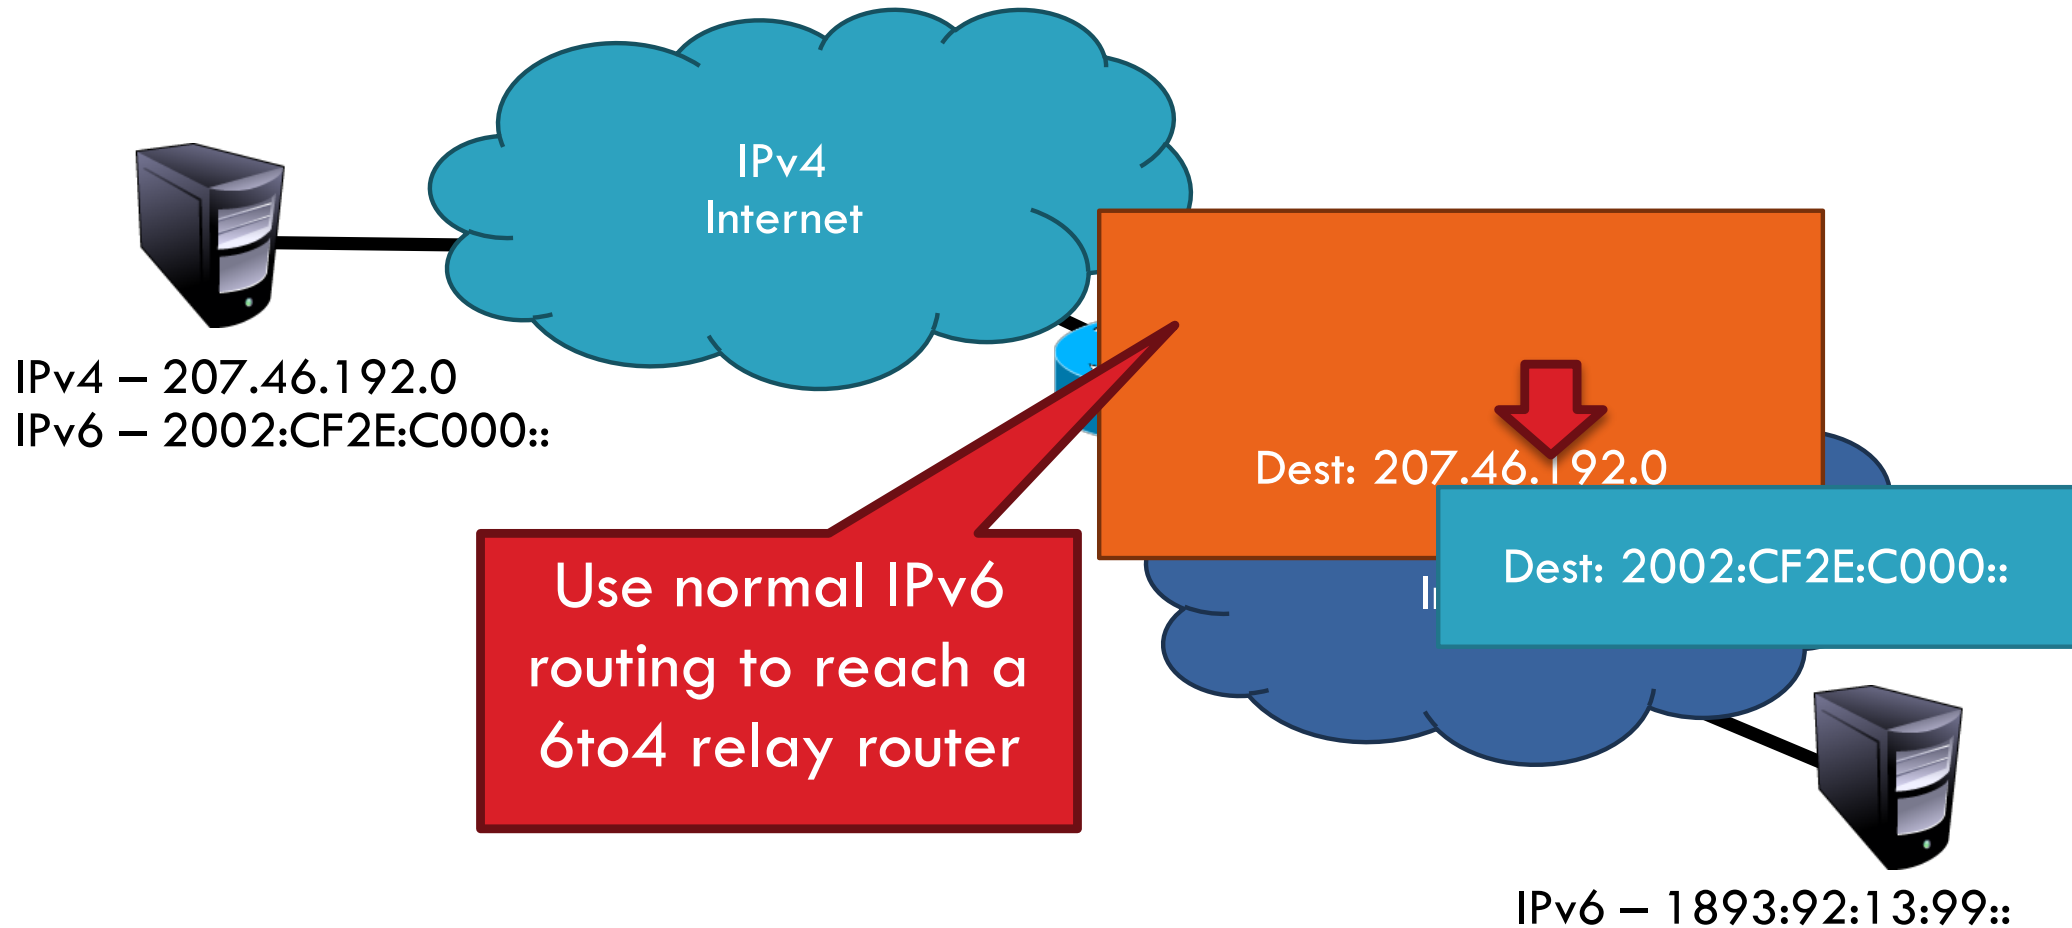
IPv6 – 1893:92:13:99::

IPv4 – 207.46.192.0
IPv6 – 2002:CF2E:C000::



Routing from Native IPv6 to 6to4

59



Problems with 6to4

60

- Uniformity
 - ▣ Not all ISPs have deployed 6to4 relays
- Quality of service
 - ▣ Third-party 6to4 relays are available
 - ▣ ...but, they may be overloaded or unreliable
- Reachability
 - ▣ 6to4 doesn't work if you are behind a NAT
- Possible solutions
 - ▣ IPv6 Rapid Deployment (6rd)
 - Each ISP sets up relays for its customers
 - Does not leverage the 2002:: address space
 - ▣ Teredo
 - Tunnels IPv6 packets through UDP/IPv4 tunnels
 - Can tunnel through NATs, but requires special relays

Consequences of IPv6

61

- Beware unintended consequences of IPv6
- Example: IP blacklists
 - ▣ Currently, blacklists track IPs of spammers/bots
 - ▣ Few IPv4 addresses mean list sizes are reasonable
 - ▣ Hard for spammers/bots to acquire new IPs
- Blacklists will not work with IPv6
 - ▣ Address space is enormous
 - ▣ Acquiring new IP addresses is trivial