

CSCI-351

Data communication and Networks

Lecture 17: Peer-to-Peer System and BitTorrent

(I swear I only use it for Linux ISOs)

2 Outline

- ❑ Peer-to-Peer Overview
- ❑ Example: Bittorrent

Traditional Internet Services Model

3

□ Client-server

- ▣ Many clients, 1 (or more) server(s)
- ▣ Web servers, DNS, file downloads, video streaming

□ Problems

- ▣ Scalability: how many users can a server support?
 - What happens when user traffic overload servers?
 - Limited resources (bandwidth, CPU, storage)
- ▣ Reliability: if # of servers is small, what happens when they break, fail, get disconnected, are mismanaged by humans?
- ▣ Efficiency: if your users are spread across the entire globe, how do you make sure you answer their requests quickly?

The Alternative: Peer-to-Peer

4

- A simple idea
 - ▣ Users bring their own resources to the table
 - ▣ A cooperative model: clients = peers = servers
- The benefits
 - ▣ Scalability: # of “servers” grows with users
 - BYOR: bring your own resources (storage, CPU, B/W)
 - ▣ Reliability: load spread across many peers
 - Probability of them all failing is **very** low...
 - ▣ Efficiency: peers are distributed
 - Peers can try and get service from nearby peers

The Peer-to-Peer Challenge

5

- What are the key components for leveraging P2P?
 - ▣ Communication: how do peers talk to each other
 - ▣ Service/data location: how do peers know who to talk to
- New reliability challenges
 - ▣ Network reachability, i.e. dealing with NATs
 - ▣ Dealing with churn, i.e. short peer uptimes
- What about security?
 - ▣ Malicious peers and cheating
 - ▣ The Sybil attack

Centralized Approach

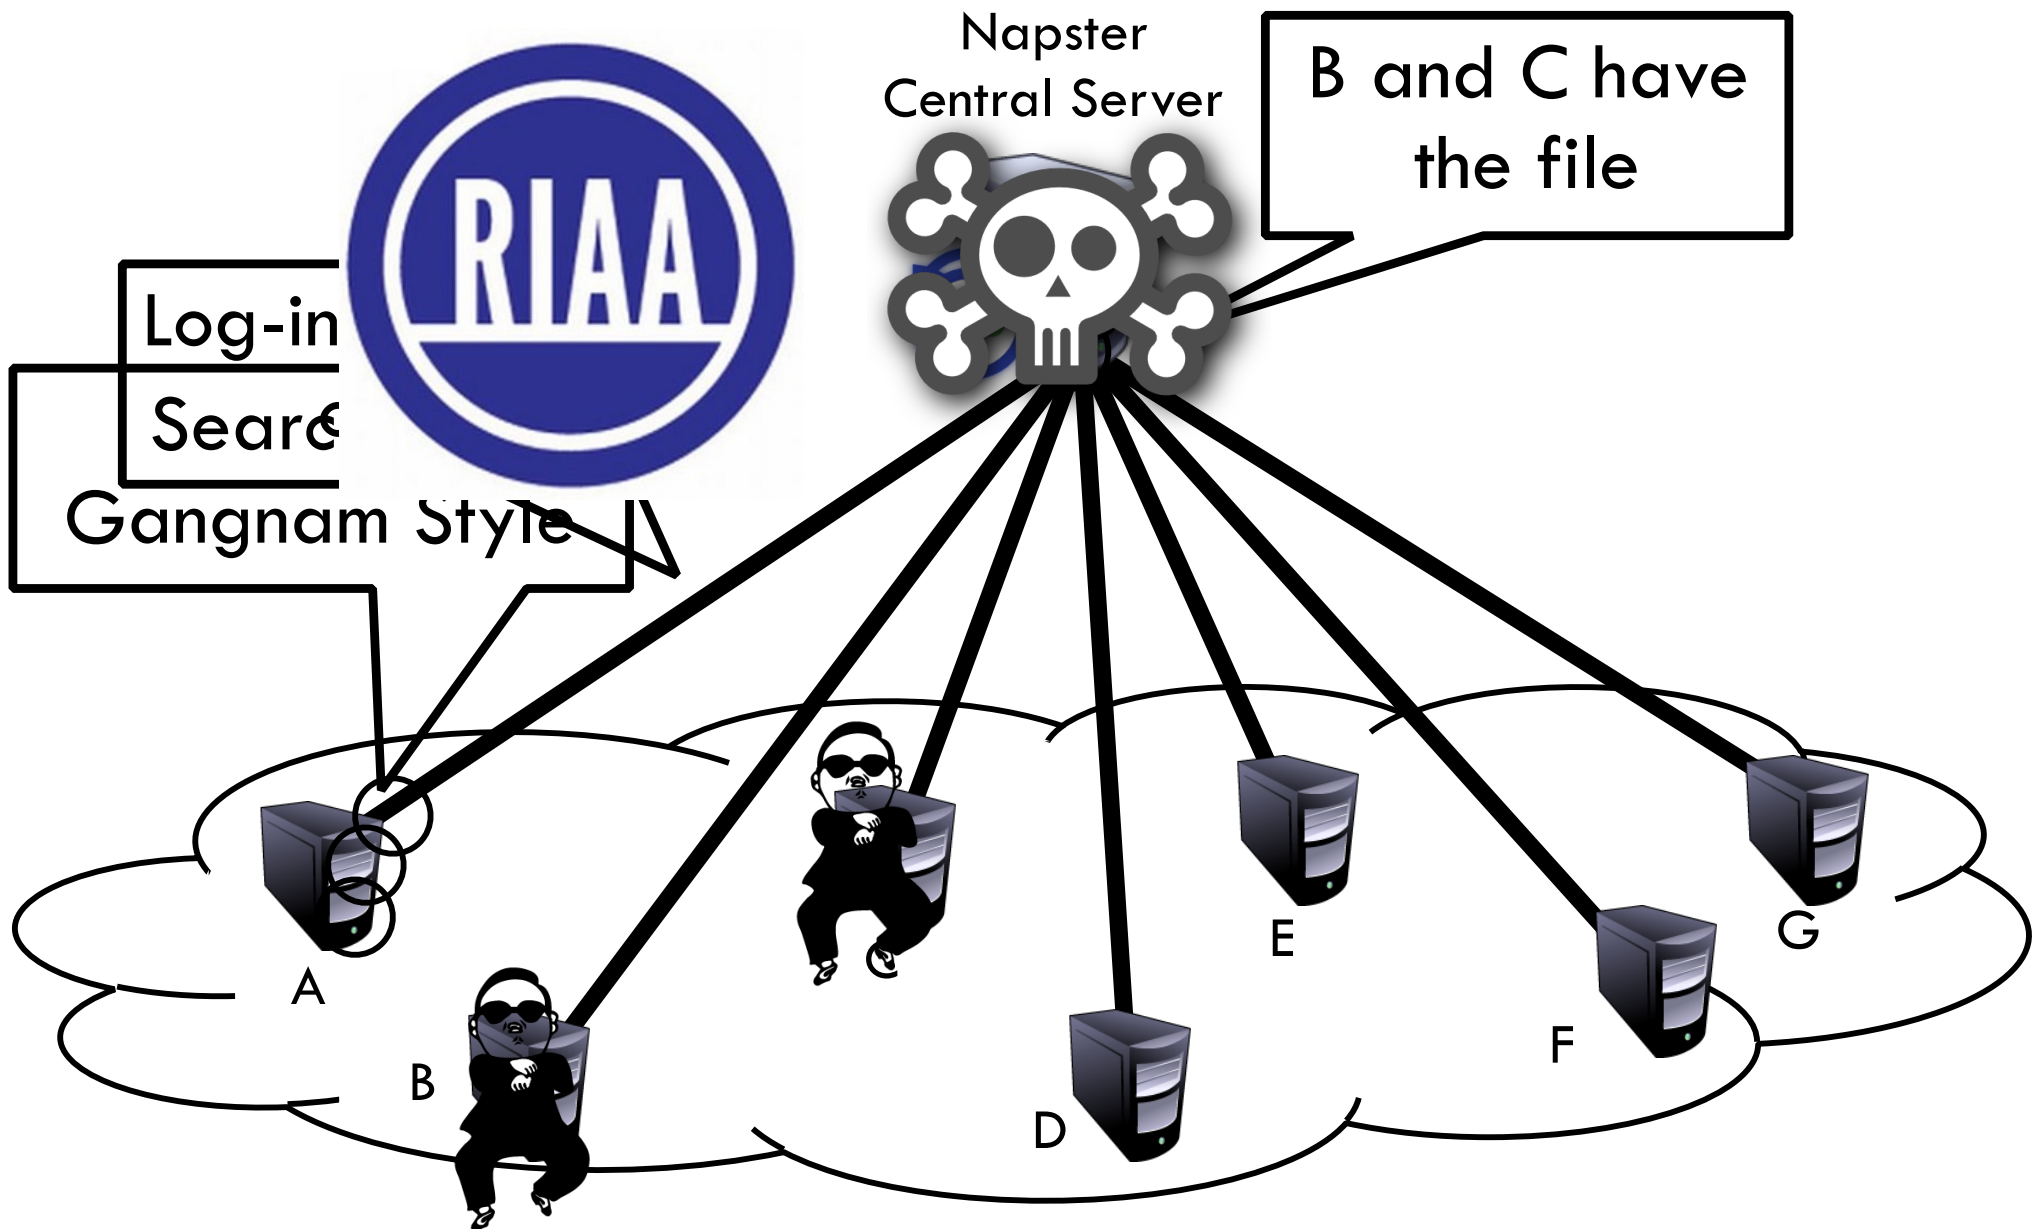
6

- The original: Napster
 - ▣ 1999-2001
 - ▣ Shawn Fanning, Sean Parker
 - ▣ Specialized in MP3s (but not for long)
- Centralized index server(s)
 - ▣ Supported all queries
- What caused its downfall?
 - ▣ Not scalable
 - ▣ Centralization of liability



Napster Architecture

7



What is BitTorrent




8

- Designed for fast, efficient content distribution
 - ▣ Ideal for large files, e.g. movies, DVDs, ISOs, etc.
 - ▣ Uses P2P file swarming
- Not a full fledged P2P system
 - ▣ Does not support searching for files
 - ▣ File swarms must be located out-of-band
 - ▣ Trackers acts a centralized swarm coordinators
 - Fully P2P, trackerless torrents are now possible
- Insanely popular
 - ▣ 35-70% of all Internet traffic in early 2010

BitTorrent Overview

9





.torrent File

10

- Contains all meta-data related to a torrent
 - ▣ File name(s), sizes
 - ▣ Torrent hash: hash of the whole file
 - ▣ URL of tracker(s)
- BitTorrent breaks files into pieces
 - ▣ 64 KB – 1 MB per piece
 - ▣ .torrent contains the size and SHA-1 hash of each piece
- Basically, a .torrent tells you
 - ▣ Everything about a given file
 - ▣ Where to go to start downloading

Torrent Sites

11



- Just standard web servers
 - ▣ Allow users to upload .torrent files
 - ▣ Search, ratings, comments, etc.
- Some also host trackers
- Many famous ones
 - ▣ Mostly because they host illegal content
- Legitimate .torrents
 - ▣ Linux distribution
 - ▣ World of Warcraft patches

Torrent Trackers

Tracker



12

- Really, just a highly specialized webserver
 - ▣ BitTorrent protocol is built on top of HTTP
- Keeps a database of swarms
 - ▣ Swarms identified by torrent hash
 - ▣ State of each peer in each swarm
 - IP address, port, peer ID, TTL
 - Status: leeching or seeding
 - Optional: upload/download stats (to track fairness)
 - ▣ Returns a random list of peers to new leechers

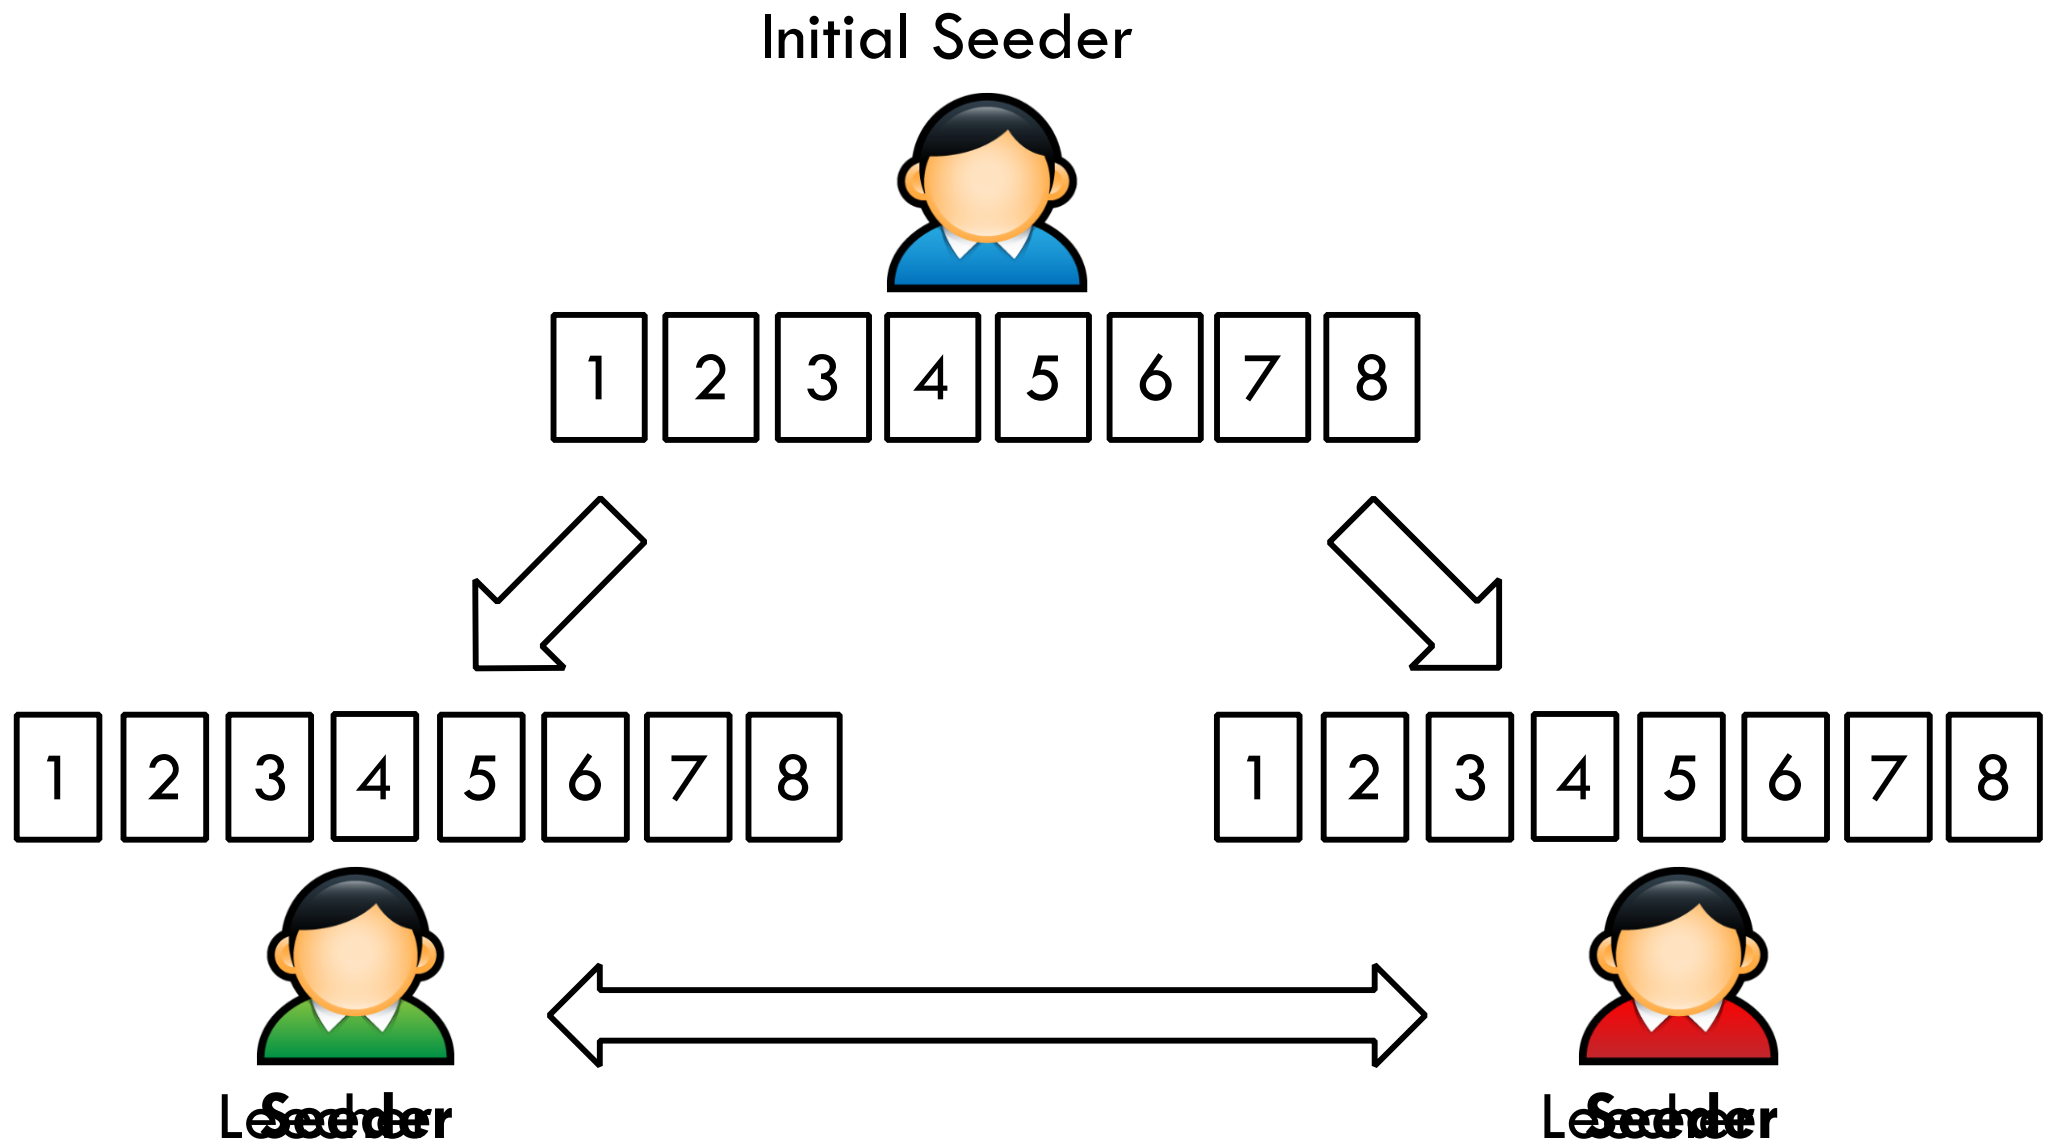
Peer Selection

13

- Tracker provides each client with a list of peers
 - ▣ Which peers are best?
 - Fastest bandwidth
- Option 1: learn dynamically
 - ▣ Try downloading from many peers
 - ▣ Keep only the best peers
 - ▣ Strategy used by BitTorrent
- Option 2: use external information
 - ▣ E.g. Some torrent clients prefer peers in the same ISP

Sharing Pieces

14




The Beauty of BitTorrent

15

- More leechers = more replicas of pieces
- More replicas = faster downloads
 - ▣ Multiple, redundant sources for each piece
- Even while downloading, leechers take load off the seed(s)
 - ▣ Great for content distribution
 - ▣ Cost is shared among the swarm

Typical Swarm Behavior

16



Sub-Pieces and Pipelining

17

- Each piece is broken into sub-pieces
 - ▣ ~16 KB in size
- TCP Pipelining
 - ▣ For performance, you want long lived TCP connections (to get out of slow start)
 - ▣ Peers generally request 5 sub-pieces at a time
 - ▣ When one finished, immediately request another
 - ▣ Don't start a new piece until previous is complete
 - Prioritizes complete pieces
 - Only complete pieces can be shared with other peers

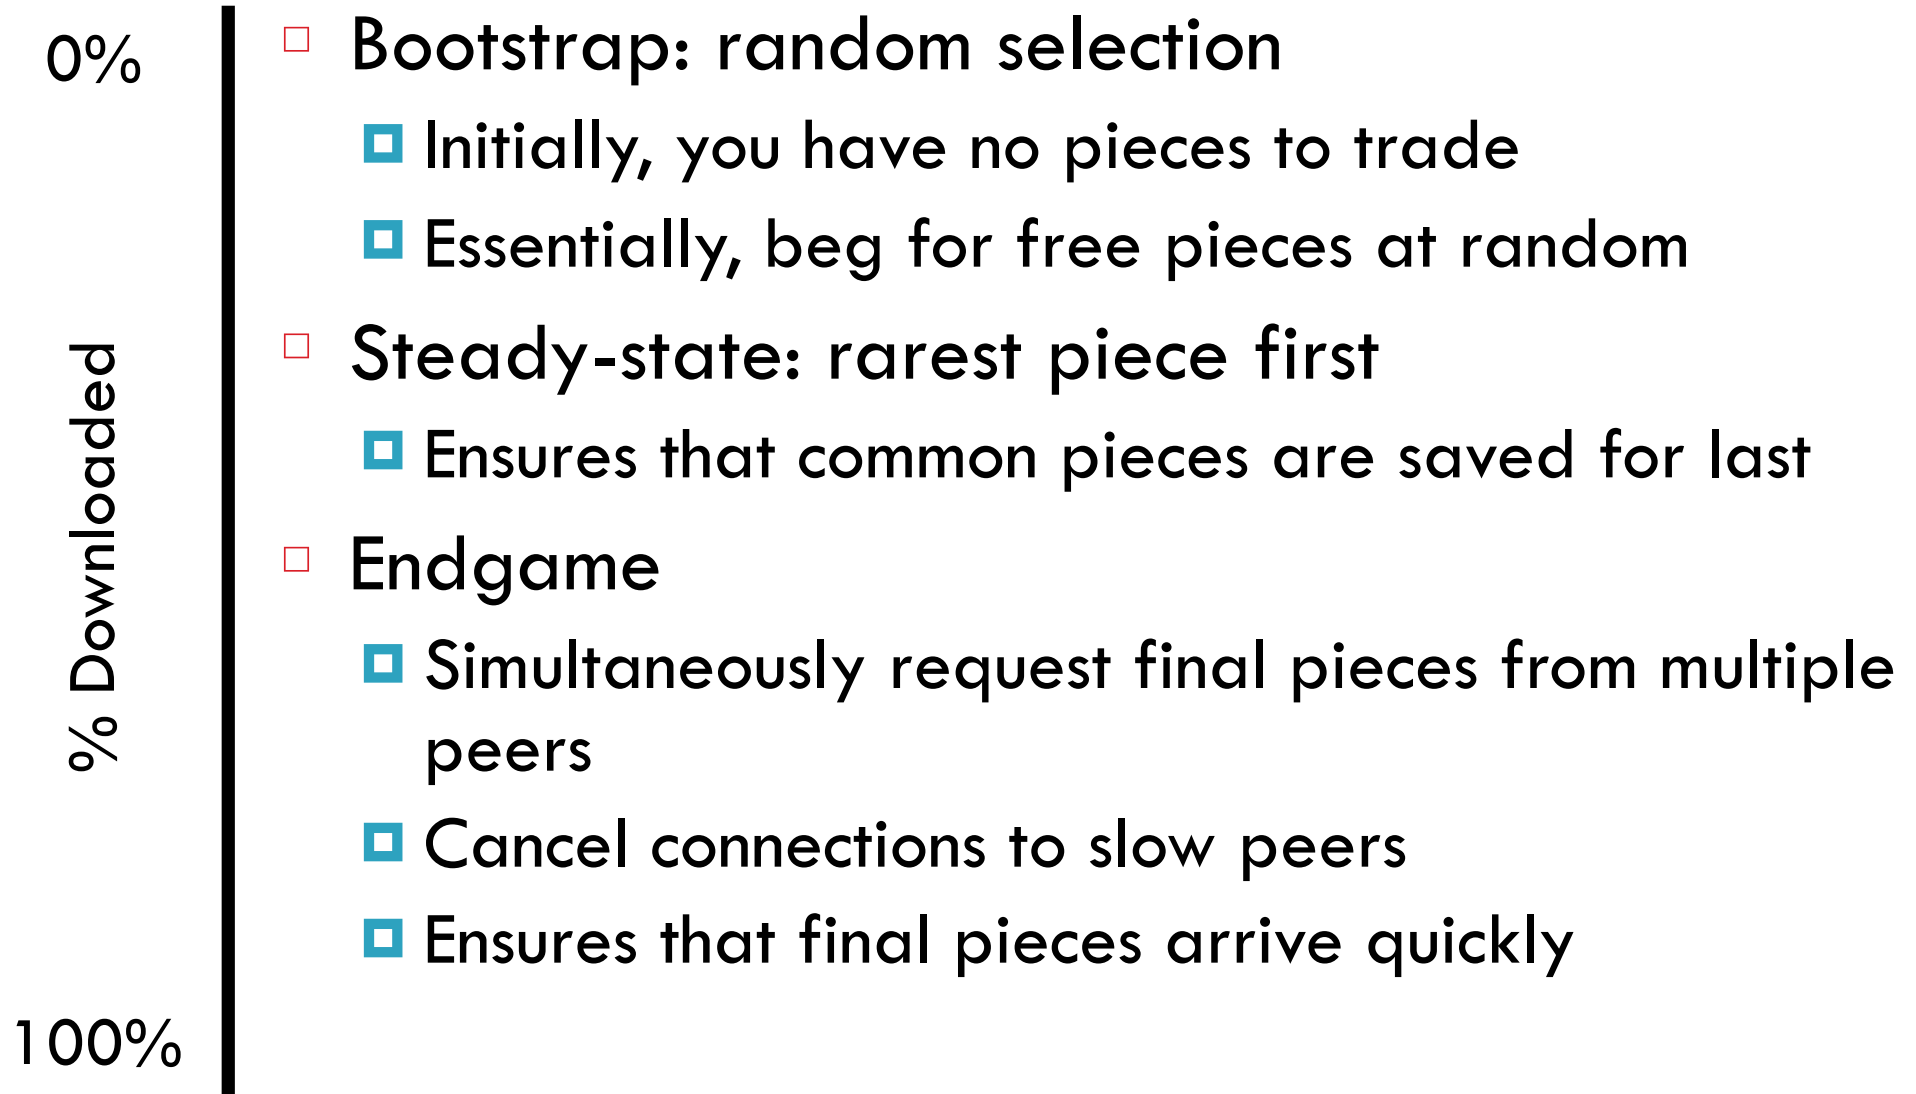
Piece Selection

18

- Piece download order is critical
 - ▣ Worst-case scenario: all leeches have identical pieces
 - Nobody can share anything :(
 - ▣ Worst-case scenario: the initial seed disappears
 - If a piece is missing from the swarm, the torrent is broken
- What is the best strategy for selecting pieces?
 - ▣ Trick question
 - ▣ It depends on how many pieces you already have

Download Phases

19



Upload and Download Control

20

- How does each peer decide who to trade with?
- Incentive mechanism
 - ▣ Based on tit-for-tat, game theory
 - ▣ “If you give a piece to me, I’ll give a piece to you”
 - ▣ “If you screw me over, you get nothing”
 - ▣ Two mechanisms: choking and optimistic unchoke

A Bit of Game Theory

21

- Iterated prisoner's dilemma
- Very simple game, two players, multiple rounds
 - ▣ Both players agree: +2 points each
 - ▣ One player defects: +5 for defector, +0 to other
 - ▣ Both players defect: +0 for each
- Maps well to trading pieces in BitTorrent
 - ▣ Both peers trade, they both get useful data
 - ▣ If both peers do nothing, they both get nothing
 - ▣ If one peer defects, he gets a free piece, other peer gets nothing
- What is the best strategy for this game?

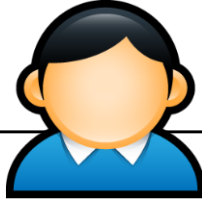

Tit-for-Tat

22

- Best general strategy for iterated prisoner's dilemma
- Meaning: “Equivalent Retaliation”

Rules

1. Initially: cooperate
2. If opponent cooperates, cooperate next round
3. If opponent defects, defect next round

Round			Points
1	Cooperate	Cooperate	+2 / +2
2	Cooperate	Defect	+0 / +5
3	Defect	Cooperate	+5 / +0
4	Cooperate	Cooperate	+2 / +2
5	Cooperate	Defect	+0 / +5
6	Defect	Defect	+0 / +0
7	Defect	Cooperate	+5 / +0
		Totals:	+14 / +14

Choking

23

- Choke is a temporary refusal to upload
 - ▣ Tit-for-tat: choke free riders
 - ▣ Cap the number of simultaneous uploads
 - Too many connections congests your network
 - ▣ Periodically unchoke to test the network connection
 - Choked peer might have better bandwidth


Optimistic Unchoke

24

- Each peer has one optimistic unchoke slot
 - ▣ Uploads to one random peer
 - ▣ Peer rotates every 30 seconds
- Reasons for optimistic unchoke
 - ▣ Help to bootstrap peers without pieces
 - ▣ Discover new peers with fast connections

BitTorrent Protocol Fundamentals

25



- BitTorrent divides time into rounds
 - ▣ Each round, decide who to upload to/download from
 - ▣ Rounds are typically 30 seconds
- Each connection to a peer is controlled by four states
 - ▣ Interested / uninterested – do I want a piece from you?
 - ▣ Choked / unchoked – am I currently downloading from you?
- Connections are bidirectional
 - ▣ You decide interest/choking on each peer
 - ▣ Each peer decides interest/choking on you

Connection States

26

Error states.

Connection should
Download control
be closed

- ❑ d – interested and choked
- ❑ D – interested and unchoked
- ❑ K – uninterested and unchoked
- ❑ S – snubbed (no data received in 60 seconds)
- ❑ F – piece(s) failed to hash

Upload control

- ❑ u – interested and choked
- ❑ U – interested and unchoked
- ❑ O – optimistic unchoke
- ❑ ? – uninterested and unchoked

Connection information

- ❑ I – incoming connection
- ❑ E/e – Using protocol encryption

Most peers are d or D. No
need to connect with
uninteresting peers.

General Trackers Peers Pieces Files Speed Logger					
IP	Client	Flags	%	Down S...	Up Speed
bl20-87-69.dsl...	µTorrent 3.2.3	ud IXP	8.6		0.3 kB/s
0545651f.skyb...	Vuze 5.0.0.0	D IXP	100.0	3.6 kB/s	
14-202-18-1.st...	µTorrent Mac	d IXP	100.0		
S010600265ac...	µTorrent 2.0.4	d IXeP	100.0		
S0106586d8f3...	BitTorrent 7.0.	d IX	100.0		
S010624ab81...	Transmission 2.	d IXEP	35.6		
c-24-130-191-...	µTorrent 3.3	d IXe	100.0		
27-33-0-184.t...	µTorrent 2.2.1	d IXP	100.0		
em36-244-251...	BitTorrent 7.8.	d XP	100.0		
41.78.77.178 [...]	BitTorrent 7.8	ud IHXP	47		0.4 kB/s

More on this
later...

- ❑ h – used UDP hole punching
- ❑ P – connection uses µTP

How was this peer located?

- ❑ H – DHT (distributed hash table)
- ❑ L – local peer discovery (multicast)
- ❑ X – peer exchange

Upload-Only Mode

27

- Once a peer completes a torrent, it becomes a seed
 - ▣ No downloads, no tit-for-tat
 - ▣ Who to upload to first?
- BitTorrent policy
 - ▣ Upload to the fastest known peer
 - ▣ Why?
 - ▣ Faster uploads = more available pieces
 - ▣ More available pieces helps the swarm

BitTorrent and TCP

28

- BitTorrent accounts for 35-70% of all Internet traffic
- Thus, BitTorrent's behavior impacts everyone
- BitTorrent's use of TCP causes problems
 - ▣ Long lived, BitTorrent TCP flows are “elephants”
 - Ramp up past slow start, dominate router queues
 - ▣ Many applications are “mice,” get trampled by elephants
 - Short lived flows (e.g. HTTP traffic)
 - Delay sensitive apps (i.e. VoIP, SSH, online games)
- Have you ever tried using SSH while using BitTorrent?

Conclusions

29

- BitTorrent is an extremely efficient tool for content distribution
 - ▣ Strong incentive system based on game theory
 - ▣ Most popular file sharing client since 2001
 - ▣ More active users than YouTube and Facebook combined
- However, BitTorrent is a large system with many different mechanisms
 - ▣ Ample room to modify the client, alter behavior