

# CSCI-351

## Data communication and Networks

### **Lecture 13: HTTP**

# HTTP Basics

2

- HTTP layered over bidirectional byte stream
- Interaction
  - Client sends **request to server**, followed by **response from server** to client
  - Requests/responses are encoded in text
- Stateless
  - Server maintains no information about past client requests

# HTTP Request

3

**GET /foo/bar.html HTTP/1.1**

## □ Request line

### □ Method

- GET – return URI
- HEAD – return headers only of GET response
- POST – send data to the server (forms, etc.)
- ...

### □ URL (relative)

- E.g., /index.html

### □ HTTP version

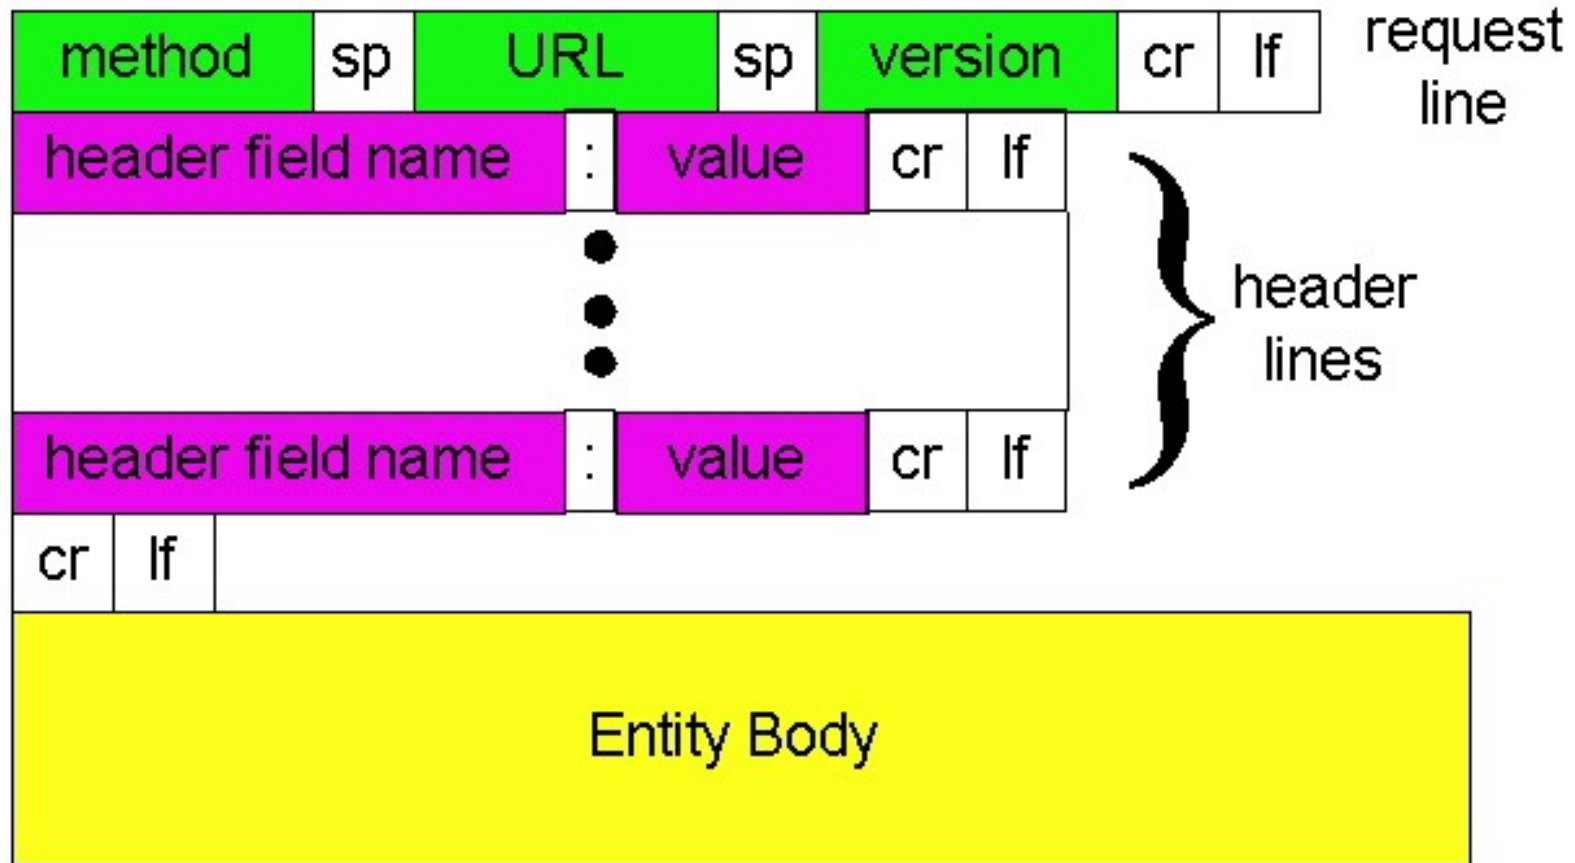
# HTTP Request

4

- Request headers (each ended with CRLF)
  - Acceptable document types/encodings
  - From – user email
  - If-Modified-Since
  - Referrer – what caused this page to be requested
  - User-Agent – client software
  - Cookie - previously stored information
  - Content-Length - Size of data (only on POST)
- Blank-line (CRLF)
- Body

# HTTP Request (visual)

5



# HTTP Request Example

6

**GET /blah.html?foo=bar HTTP/1.1**

**Accept: \*/\***

**Accept-Language: en-us**

**Accept-Encoding: gzip, deflate**

**User-Agent: Mozilla/4.0 (compatible; MSIE 5.5; Windows NT 5.0)**

**Host: [www.intel-iris.net](http://www.intel-iris.net)**

**Connection: Keep-Alive**

# HTTP Response

7

## □ Status-line

- HTTP version

- 3 digit response code

  - 1XX – informational

  - 2XX – success

    - 200 OK

  - 3XX – redirection

    - 301 Moved Permanently

    - 303 Moved Temporarily

    - 304 Not Modified

  - 4XX – client error

    - 404 Not Found

  - 5XX – server error

    - 505 HTTP Version Not Supported

- Reason phrase

# HTTP Response (cont.)

8

## ☐ Headers

- ☐ Location – for redirection
- ☐ Server – server software
- ☐ WWW-Authenticate – request for authentication
- ☐ Allow – list of methods supported (get, head, etc)
- ☐ Content-Encoding – E.g x-gzip
- ☐ Content-Length
- ☐ Content-Type
- ☐ Expires
- ☐ Last-Modified

## ☐ Blank-line

## ☐ Body



# HTTP Response Example

9

**HTTP/1.1 200 OK**

**Date: Tue, 27 Mar 2018 03:49:38 GMT**

**Server: Apache/1.3.14 (Unix)**

**Last-Modified: Mon, 29 Jan 2001 17:54:18 GMT**

**ETag: "7a11f-10ed-3a75ae4a"**

**Accept-Ranges: bytes**

**Content-Length: 4333**

**Keep-Alive: timeout=15, max=100**

**Connection: Keep-Alive**

**Content-Type: text/html**

**Cache-Control: private**

**...DATA...**

# Web pages

10

- Multiple (typically small) objects per page
  - E.g., each image, JS, CSS, etc downloaded separately
- Single page can have 100s of HTTP transactions!
- File sizes
  - Heavy-tailed
  - Most transfers/objects very small
- Problem: Browser can't render complete page until all downloaded

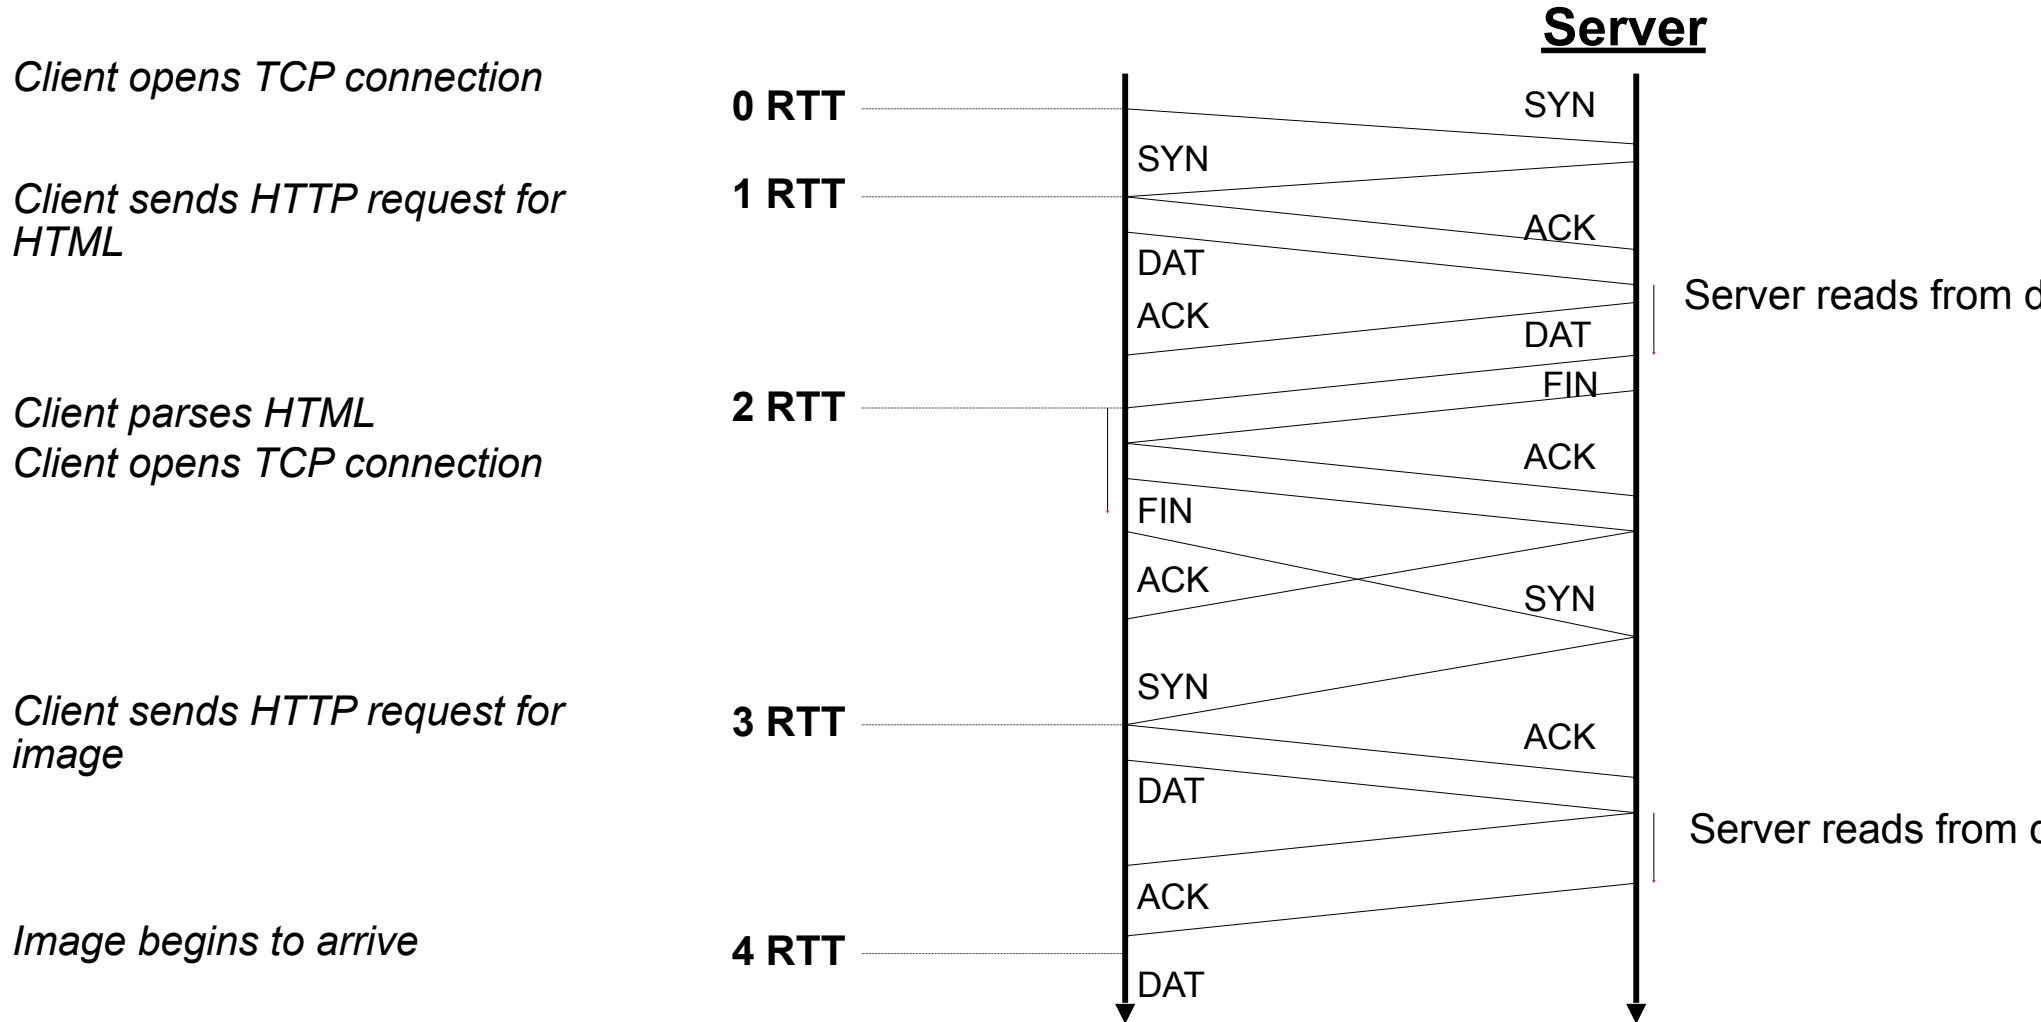
# HTTP 0.9/1.0

11

- One request/response per TCP connection
  - Simple to implement
  
- Disadvantages
  - Multiple connection setups → three-way handshake each time
    - Several extra round trips added to transfer
  - Multiple slow starts

# Single Transfer, One Image

12



# More Problems

13

- ❑ Short transfers are hard on TCP
  - ❑ Stuck in slow start
  - ❑ Loss recovery is poor when windows are small
  - ❑ SYN/ACK overhead is highest
  
- ❑ Lots of extra connections
  - ❑ Increases server state/processing
  
- ❑ Server also forced to keep TIME\_WAIT connection state
  - ❑ Why must server keep these?
  - ❑ Tends to be an order of magnitude greater than # of active connections, why?

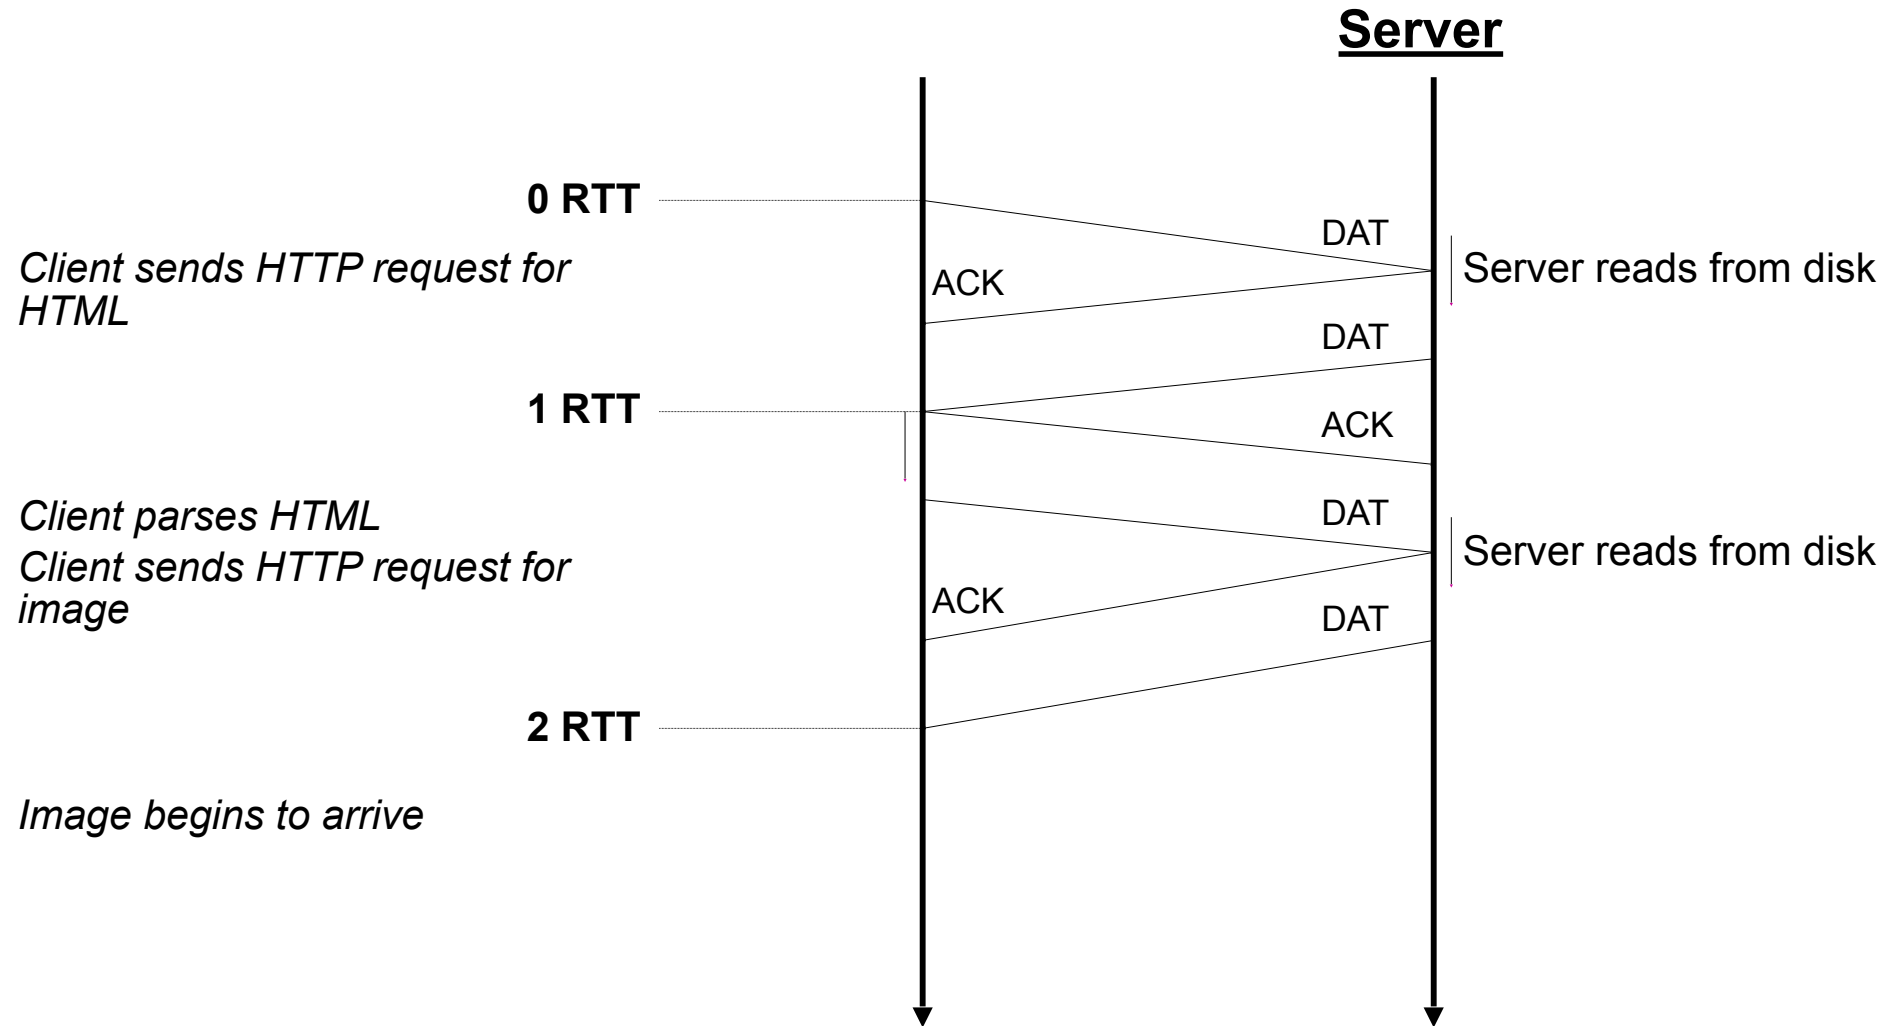
# Persistent Connections

14

- Multiplex multiple transfers onto one TCP connection
  
- Client keeps connection open
  - Can send another request after the first completes
  - Must announce intention via a header
    - Connection: keepalive
  - Server must say how long response is, so client knows when done
    - Content-Length: XXX

# Persistent Connection Example

15



# HTTP Caching

16

- Clients often cache documents
  - Challenge: update of documents
  - If-Modified-Since requests to check
    - HTTP 0.9/1.0 used just date
    - HTTP 1.1 has an opaque “etag” (could be a file signature, etc.) as well
  
- When/how often should the original be checked for changes?
  - Check every time?
  - Check each session? Day? Etc?
  - Use Expires header
    - If no Expires, often use Last-Modified as estimate



# Example Cache Check Request

17

**GET / HTTP/1.1**

**Accept: \*/\***

**Accept-Language: en-us**

**Accept-Encoding: gzip, deflate**

**If-Modified-Since: Mon, 29 Jan 2018 17:54:18 GMT**

**If-None-Match: "7a11f-10ed-3a75ae4a"**

**User-Agent: Mozilla/4.0 (compatible)**

**Host: www.intel-iris.net**

**Connection: Keep-Alive**

# Example Cache Check Response

18

**HTTP/1.1 304 Not Modified**

**Date: Tue, 27 Mar 2018 03:50:51 GMT**

**Server: Apache/1.3.14 (Unix)**

**Connection: Keep-Alive**

**Keep-Alive: timeout=15, max=100**

**ETag: "7a11f-10ed-3a75ae4a"**

# Content in today's Internet

19

- Most flows are HTTP
  - Web is at least 52% of traffic
  - Median object size is 2.7K, average is 85K (as of 2007)
  
- HTTP uses TCP, so it will
  - Be ACK clocked
  - For Web, likely never leave slow start
  
- Is the Internet designed for this common case?
  - Why?