

# CSCI-351

# Data communication and Networks

**Lecture 17: HTTPS**

**Warning: This may be hard to understand. Do not lose  
yourself during the class and keep asking questions**



# SSL/TLS

- Application-layer protocol for confidentiality, integrity, and authentication between clients and servers
  - Introduced by Netscape in 1995 as the Secure Sockets Layer (SSL)
  - Designed to encapsulate HTTP, hence HTTPS
- Transport Layer Security
  - Defined in an RFC in 1999
  - Supersedes SSL: SSL is known to be insecure and should not be used
- Sits between transport and application layers
  - Thus, applications must be TLS-aware

# Goals of TLS

- Confidentiality and integrity: use BofA's public key to negotiate a session key; encrypt all traffic
- Authentication: BofA's cert can be validating by checking Verisign's signature



*https://www.bankofamerica.com*

- Contains BofA's public key
- Signed by Verisign



Bank of America



# Let's Talk about Certificates

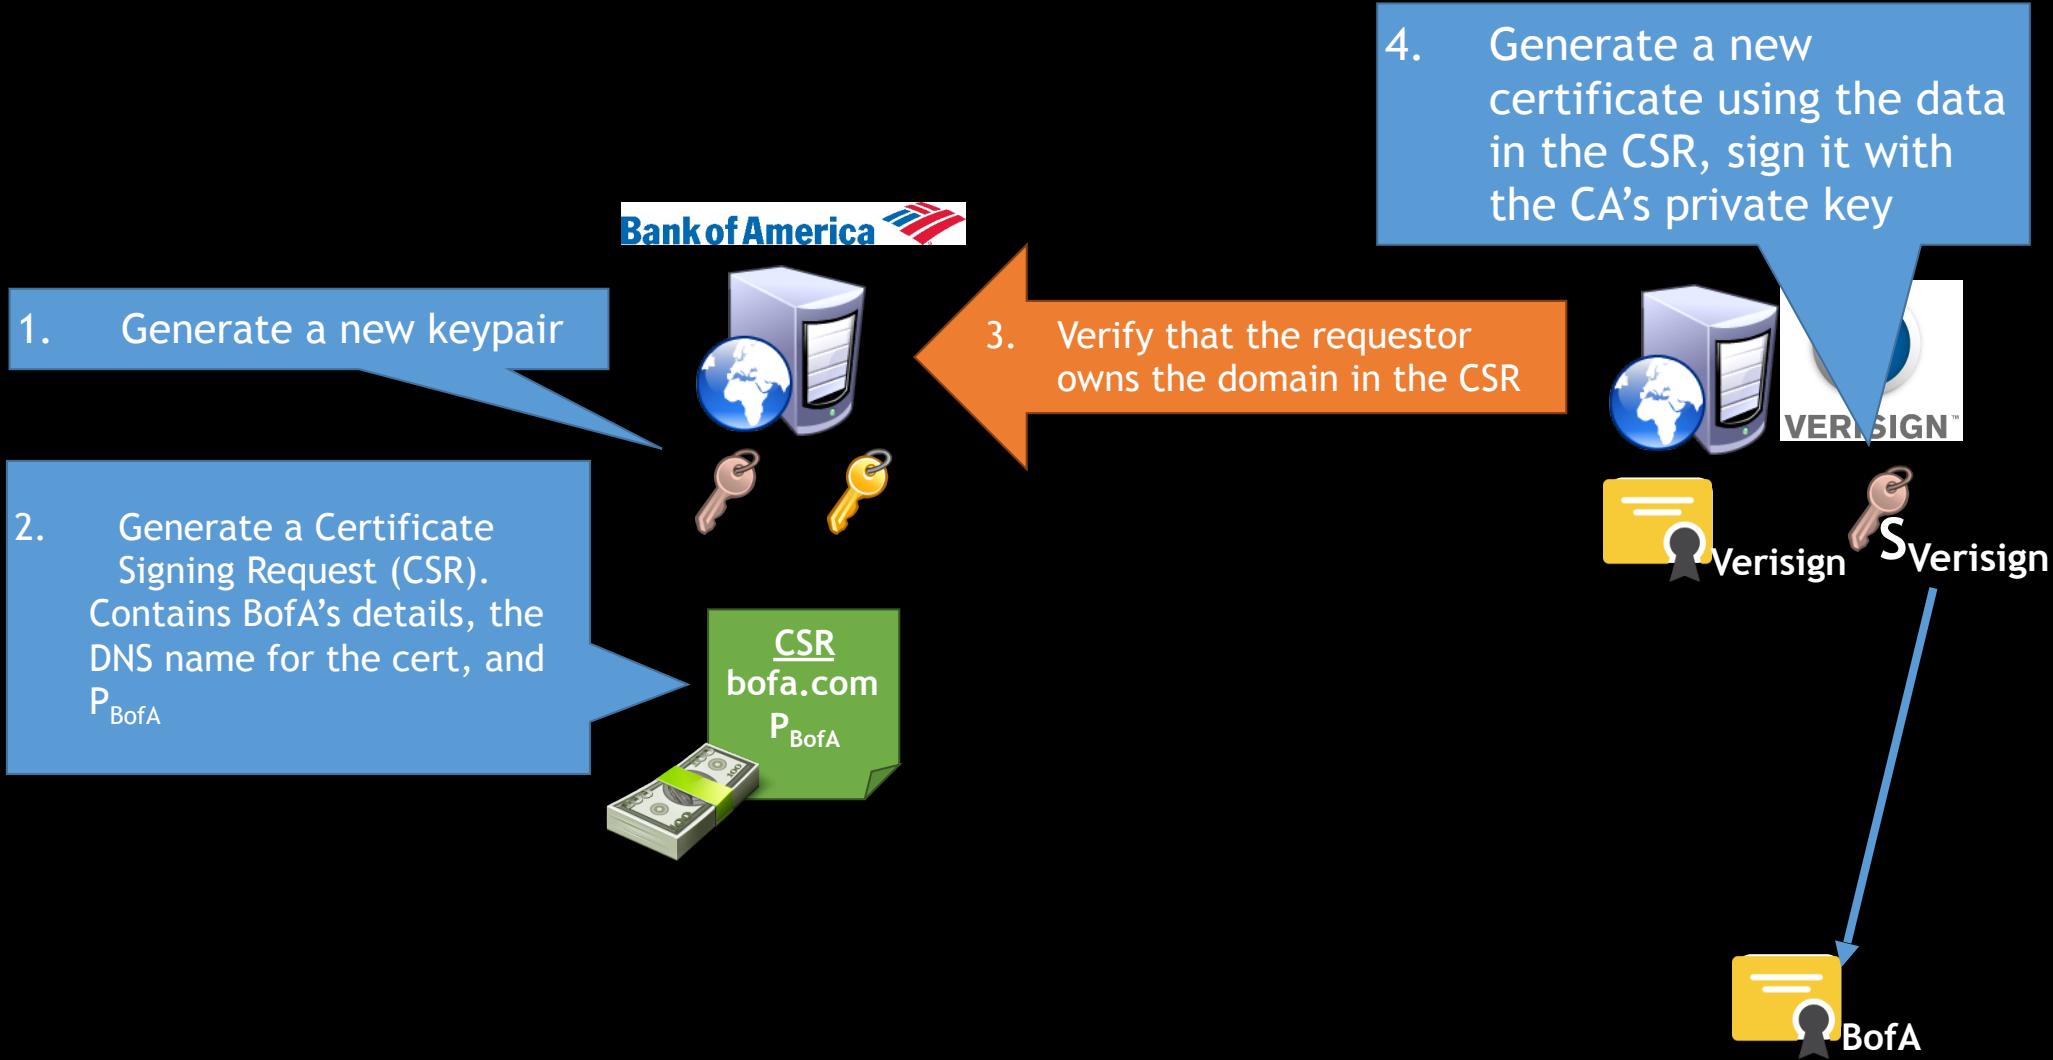


- Suppose you start a new website and you want TLS encryption
  - You need a certificate. How do you get one?
- Option 1: generate a certificate yourself
  - Use *openssl* to generate a new asymmetric keypair
  - Use *openssl* to generate a certificate that includes your new public key
- Problem?
  - Your new cert is *self-signed*, i.e. not signed by a trusted CA
  - Browsers cannot authenticate your cert to a trusted root CA
  - Users will be shown a scary security warning when they visit your site

# Certificate Authorities

- Certificate Authorities (CAs) are the roots of trust in the TLS PKI
  - Symantec, Verisign, Thawte, Geotrust, Comodo, GlobalSign, Go Daddy, Digicert, Entrust, and hundreds of others
  - Issue signed certs on behalf of third parties
- How do you become a CA?
  - I. Create a self-signed root certificate
  2. Get all the major browser vendors to include your cert with their software
  3. Keep your private key secret at all costs
- What is the key responsibility of being a CA?
  - Any CA can issue a cert for any domain!
  - The only thing that stops me from buying a cert for *google.com* is a manual verification process

# Acquiring a Certificate



# X.509 Certificate (Part I)

Certificate:

Data:

Version: 3 (0x2)

Serial Number:

0c:00:93:10:d2:06:db:e3:37:55:35:80:11:8d:dc:87

Issuer: who generated this cert? (usually a CA)

Signature Algorithm: sha256WithRSAEncryption

Issuer: C=US, O=DigiCert Inc, OU=www.digicert.com, CN=DigiCert SHA2 Extended

Validation Server CA

Validity

Not Before: Apr 8 00:00:00 2014 GMT

Not After :Apr 12 12:00:00 2016 GMT

Certificates expire

Used for revocation

Subject: businessCategory=Private Organization/1.3.6.1.4.1.311.60.2.1.3=US/1.3.6.1.4.1.311.60.2.1.2=Delaware/serialNumber=5157550/street=548 4th Street/postalCode=94107, C=US, ST=California, L=San Francisco, O=GitHub, Inc., CN=github.com

Subject Public Key Info:

Public Key Algorithm: rsaEncryption

Public-Key: (2048 bit)

Modulus:

00:b1:d4:dc:3c:af:fd:f3:4e:ed:c1:67:ad:e6:cb:

GitHub's public key

- Subject: who owns this cert?
- This is GitHub's certificate
- Must be served from [github.com](https://github.com)

# X.509 Certificate (Part 2)

X509v3 extensions:

X509v3 Subject Alternative Name:

DNS:github.com, DNS:www.github.com

Additional DNS names that  
may serve this cert

X509v3 CRL Distribution Points:

Full Name:

URL:<http://crl3.digicert.com/sha2-ev-server-g1.crl>

If this cert is revoked, its serial will be in the lists at these URLs

Full Name:

URL:<http://crl4.digicert.com/sha2-ev-server-g1.crl>

X509v3 Certificate Policies:

Policy: 2.16.840.1.114412.2.1

CPS: <https://www.digicert.com/CPS>

Authority Information Access:

OCSP - URI:<http://ocsp.digicert.com>

This cert's revocation status  
may also be checked via OSCP

# TLS Connection Establishment



# Quick question

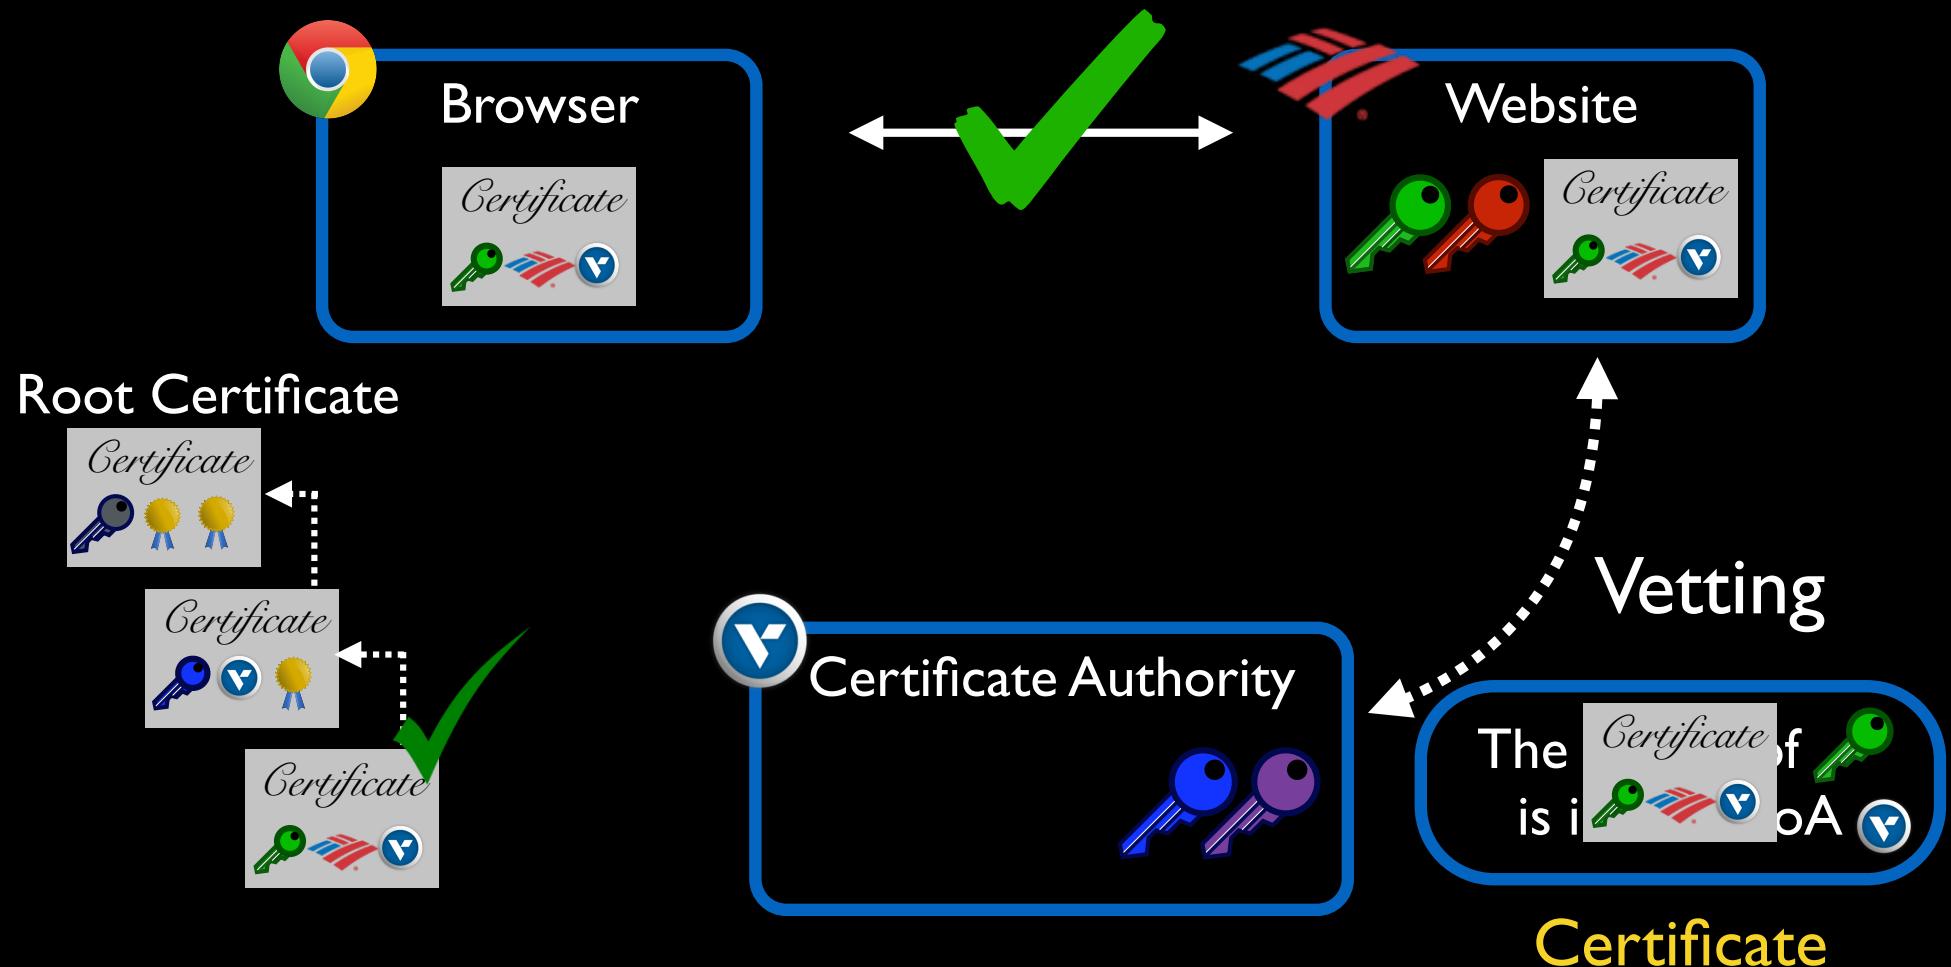
- TLS is based on the Transport Layer
  - The layer below domain name service (DNS)
- All message after TLS handshake encrypted
- If one server (with IP address) serves one domain name, it will be trivial
  - What about the server serving multiple domains (virtual hosting?)
- SNI, DNS, ESNI, DNS-over-TLS, and so on.

# TLS Authentication

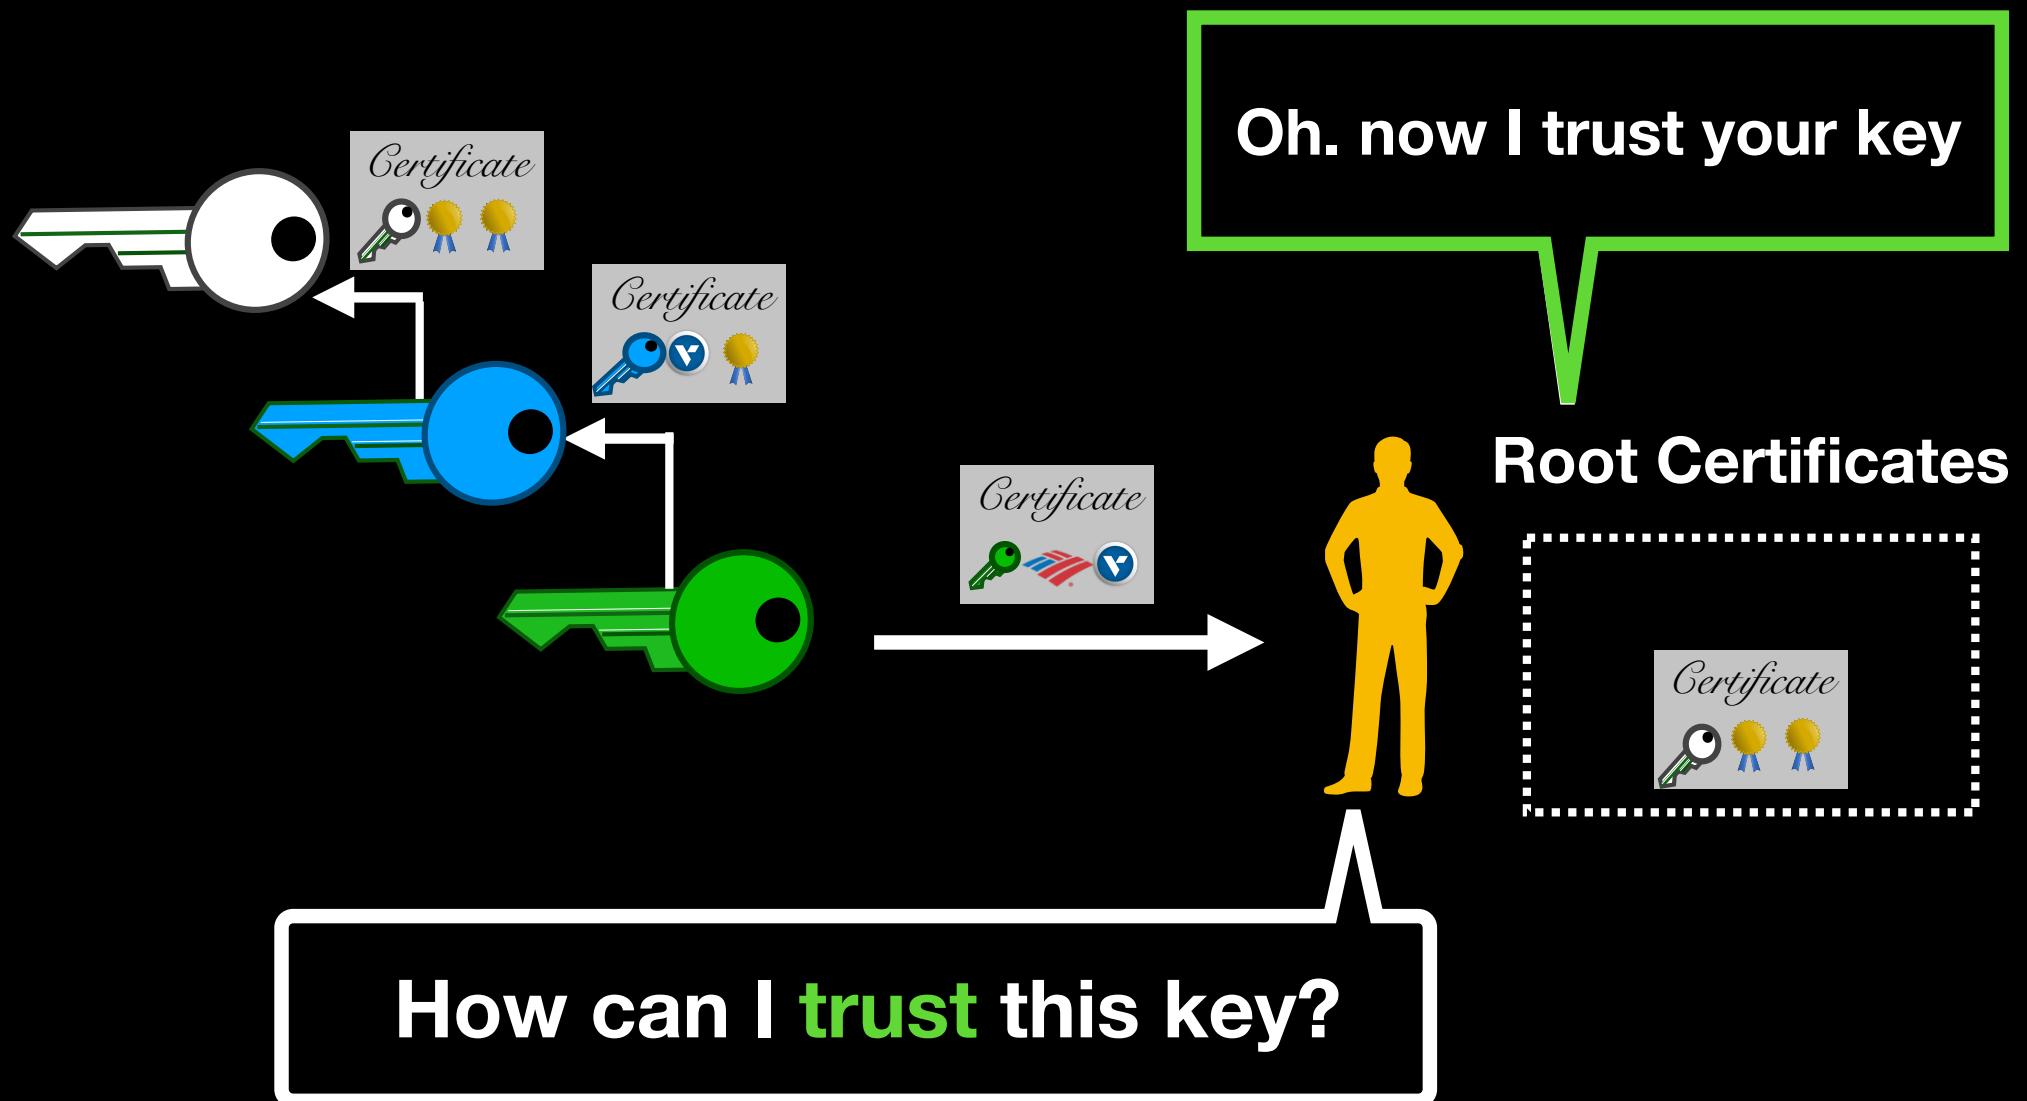
- During the TLS handshake, the client receives a **certificate chain**
  - Chain contains the server's cert, as well as the certs of the signing CA(s)
- The client must **validate** the certificate chain to establish trust
  - i.e. is this chain authentic, correct, cryptographically sound, etc.
- Client-side validation checks
  - Does the server's DNS name match the common name in the cert?
    - E.g. *example.com* cannot serve a cert with common name *google.com*
  - Are any certs in the chain expired?
  - Is the CA's signature cryptographically valid?

# How HTTPS Works

How can users truly know with whom they are communicating?



# HTTPS: Hierarchical PKI



# X.509 Format

Version: 3 (0x2)

Serial Number:

0e:77:76:8a:5d:07:f0:e5:79:59:ca:2a:9d:50:82:b5

Signature Algorithm: sha1WithRSAEncryption

Issuer: C=US, O=DigiCert Inc, OU=www.digicert.com,  
CN=DigiCert High Assurance EV CA-1

Validity

Not Before: May 27 00:00:00 2011 GMT

Not After : Jul 29 12:00:00 2013 GMT

Subject: C=US, ST=California, L=San Francisco,  
O=GitHub, Inc., CN=github.com

Subject Public Key Info:

Public Key Algorithm: rsaEncryption

Public-Key: (2048 bit)

Modulus:

00:ed:d3:89:c3:5d:70:72:09:f3:33:4f:1a:72:74:  
d9:b6:5a:95:50:bb:68:61:9f:f7:fb:1f:19:e1:da:

# X.509 Format

- Real world examples

# CA Trustworthiness (I)

- A CA is essentially a trusted third party
  - Certificate signatures are attestations of authenticity for the server and (optionally) the client
  - Remember: trust is bad and should be minimized!
- If a CA mistakenly (or purposefully) signs a certificate for a domain and provides it to a malicious principal, TLS can be subverted
- Not only must we trust root CAs, but also intermediate CAs that have been delegated signing authority

# CA Trustworthiness (2)

- Clearly, the CA secret key must be protected at all costs
  - Possession of the CA secret key grants adversaries the ability to sign any domain
  - **Attractive target for adversaries**
- Signatures should only be issued after verifying the identity of the requester
  - Also known as domain validation
  - Should be easy, right?

# CA Failures

Issued to: Microsoft Corporation

Issued by: VeriSign Commercial Software Publishers CA

Valid from 1/29/2001 to 1/30/2002

Serial number is 1B51 90F7 3724 399C 9254 CD42 4637 996A

Issued to: Microsoft Corporation

Issued by: VeriSign Commercial Software Publishers CA

Valid from 1/30/2001 to 1/31/2002

Serial number is 750E 40FF 97F0 47ED F556 C708 4EB1 ABFD

- In 2001, VeriSign issued two executable signing certificates to someone claiming to be from Microsoft
  - Could be used to issue untrusted software updates

# Comodo

## Independent Iranian hacker claims responsibility for Comodo hack

Posts claiming to be from an Iranian hacker responsible for the Comodo hack ...

by Peter Bright - Mar 28 2011, 11:15am EDT

65

1. Hello
- 2.
3. I'm writing this to the world, so you'll know more about me..
- 4.
5. At first I want to give some points, so you'll be sure I'm the hacker:
- 6.
7. I hacked Comodo from InstantSSL.it, their CEO's e-mail address [mfpenco@mfpenco.com](mailto:mfpenco@mfpenco.com)
8. Their Comodo username/password was: user: gtadmin password: [trimmed]
9. Their DB name was: globaltrust and instantsslcms

The alleged hacker's claim of responsibility on pastebin.com

The hack that resulted in [Comodo creating certificates](#) for popular e-mail providers including Google Gmail, Yahoo Mail, and Microsoft Hotmail has been claimed as the work of an independent Iranian patriot. A [post](#) made to data sharing site pastebin.com by a person going by the handle "comodohacker" claimed responsibility for the hack and described details of the attack. A second [post](#) provided source code apparently reverse-engineered as one of the parts of the attack.

# Diginotar

## Another fraudulent certificate raises the same old questions about certificate authorities

For the second time this year, Iranian hackers have created a fraudulent ...

by Peter Bright - Aug 29 2011, 11:12pm EDT

42

Earlier this year, an [Iranian hacker](#) broke into servers belonging to a reseller for certificate authority Comodo and issued himself a range of certificates for sites including Gmail, Hotmail, and Yahoo! Mail. With these certificates, he could eavesdrop on users of those mail providers, even if they use SSL to protect their mail sessions.

It's happened again. This time, Dutch certificate authority DigiNotar has issued a fraudulent certificate for google.com and all subdomains. As before, Gmail appears to be the target. The perpetrator also appears to be Iranian, with [reports](#) that the certificate has been used in the wild for man-in-the-middle attacks in that country. The certificate was issued on July 10th, and so could have been in use for several weeks prior to its discovery.

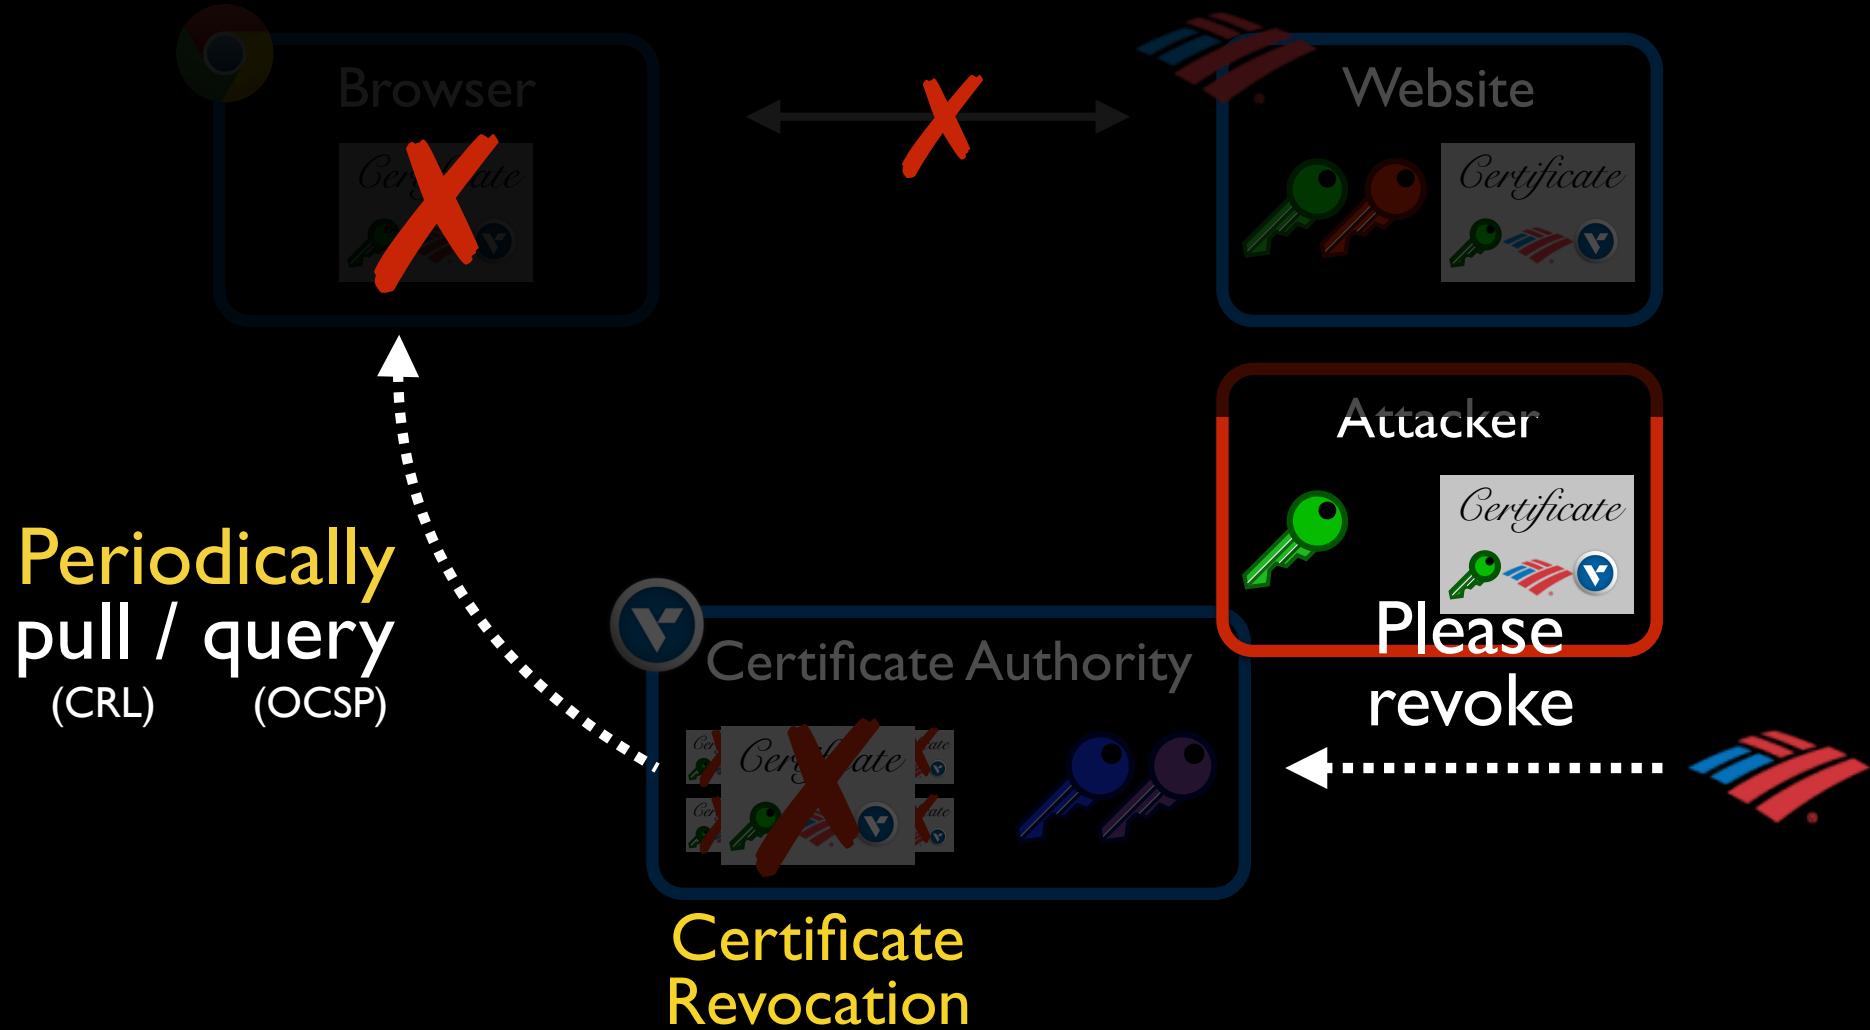
DigiNotar has revoked the certificate, which provides some protection to users (though many applications do not bother checking for revocations). However, the company has so far not disclosed how the certificate was issued in the first place, making it unclear that its integrity has been restored. As a result, Google and Mozilla have both made patches to [Chrome](#) and [Firefox](#) respectively that blacklist the entire certificate authority.

# How to handle those situations?

- A certificate has been mis-issued.
  - In the perspective of clients, the certificate seems legit
  - Still valid (not expired)
- Question:
  - How can we protect clients from accepting mis-issued certificates?
    - Revocation

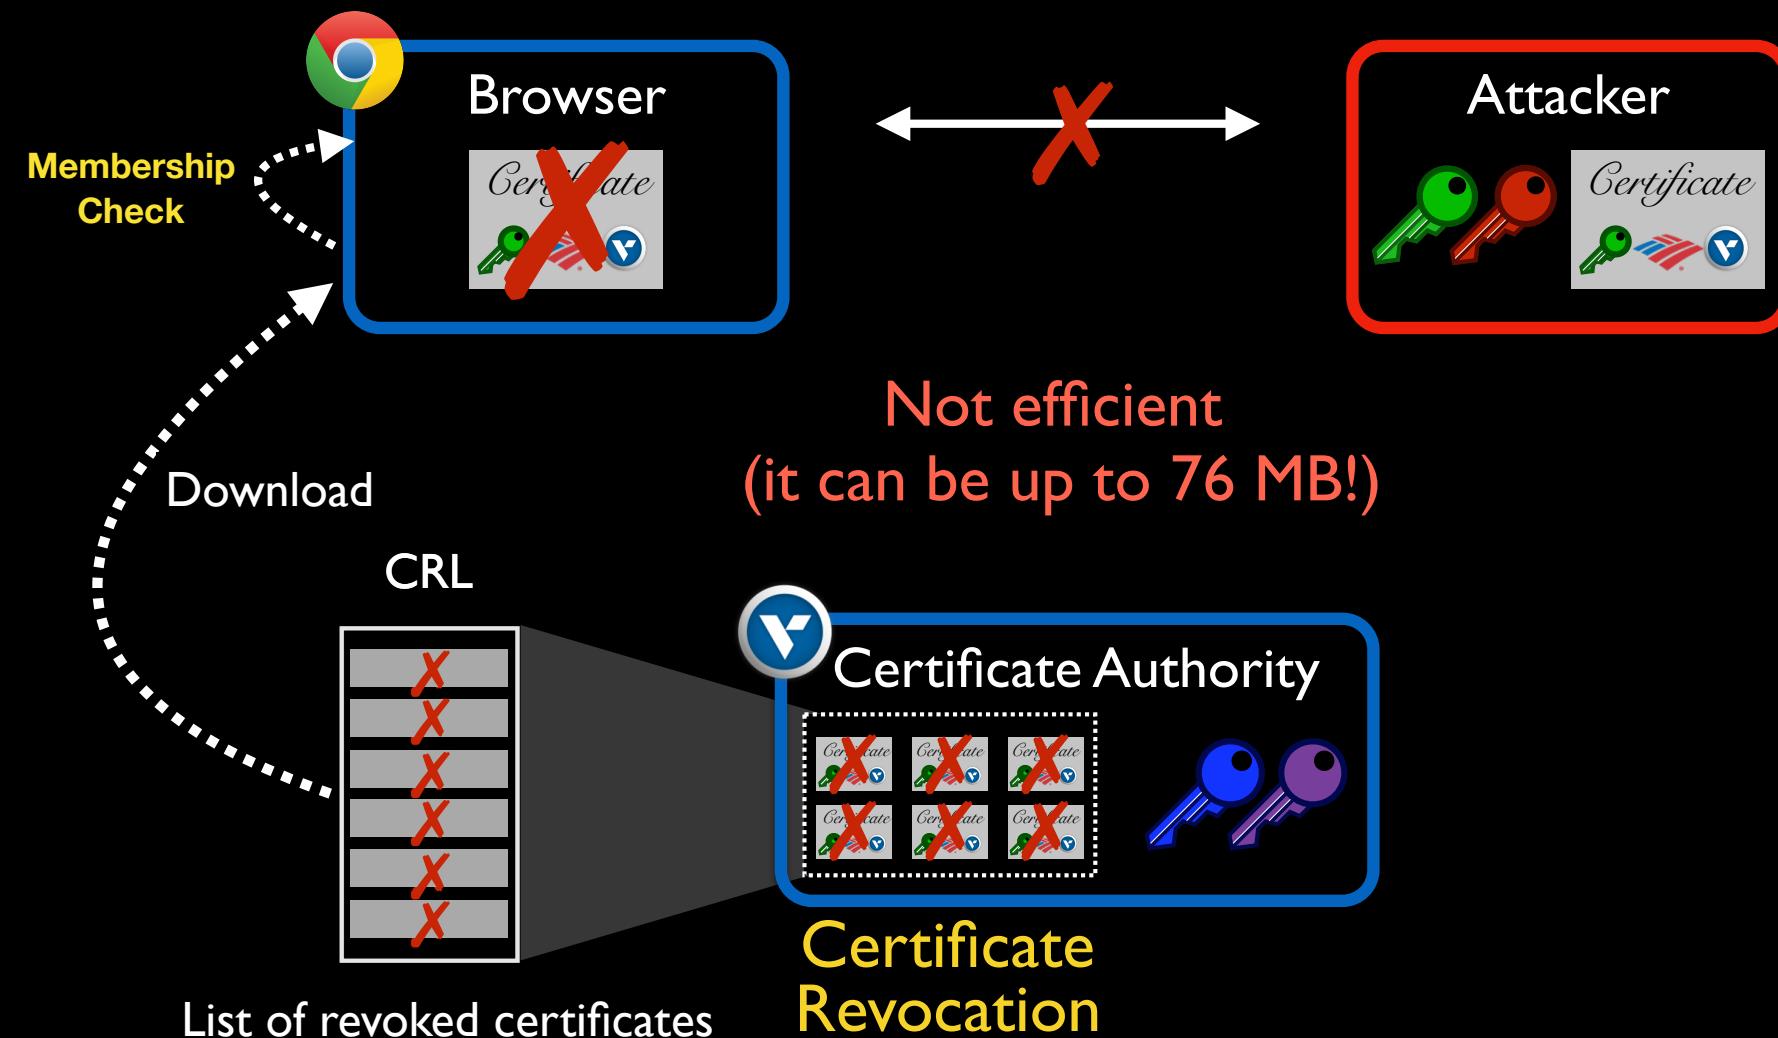
# Certificate revocation

What happens when a certificate is no longer valid?



# Revocation Check (I)

## Certificate Revocation List



# Revocation Check (I)

## Certificate Revocation List

https://www.rit.edu

USERTrust RSA Certification Authority

InCommon RSA Server CA

www.rit.edu

Extension Subject Alternative Name ( 2.5.29.17 )  
Critical NO  
DNS Name www.rit.edu  
DNS Name rit.edu

Extension Certificate Policies ( 2.5.29.32 )  
Critical NO  
Policy ID #1 ( 1.3.6.1.4.1.5923.1.4.3.1.1 )  
Qualifier ID #1 Certification Practice Statement ( 1.3.6.1.5.5.7.2.1 )  
CPS URI [https://www.incommon.org/cert/repository/cps\\_ssl.pdf](https://www.incommon.org/cert/repository/cps_ssl.pdf)  
Policy ID #2 ( 2.23.140.1.2.2 )

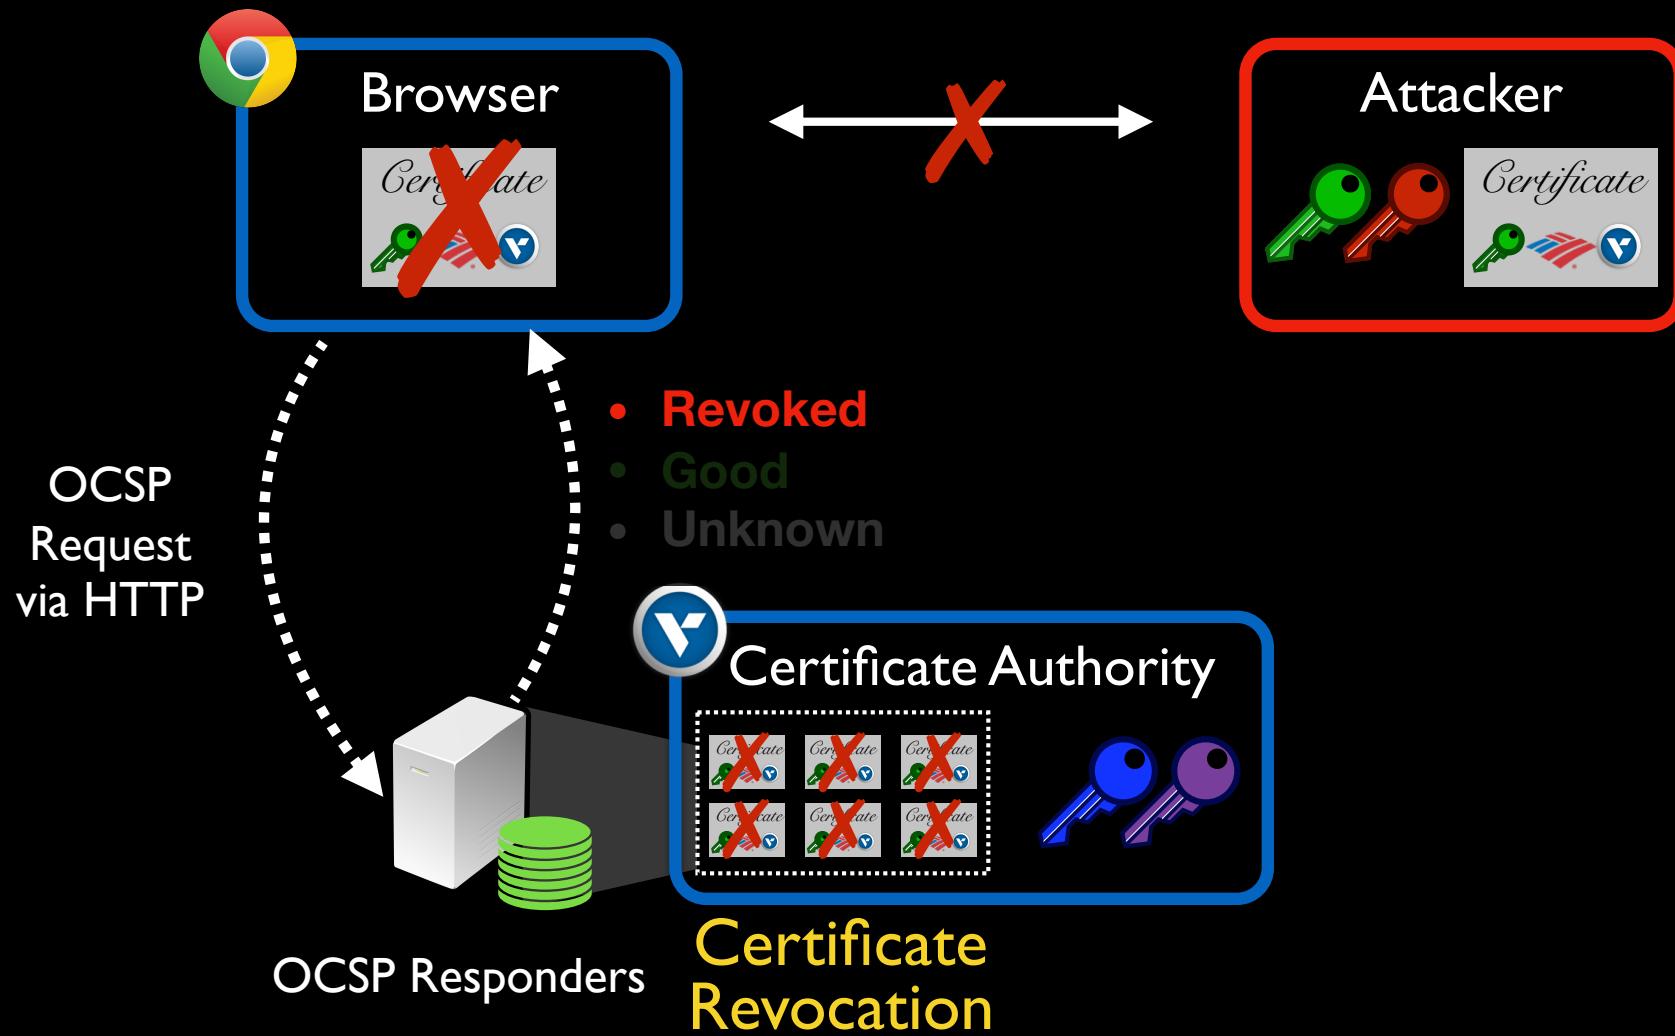
Extension CRL Distribution Points ( 2.5.29.31 )  
Critical NO  
URI <http://crl.incommon-rsa.org/InCommonRSAServerCA.crl>

Extension Embedded Signed Certificate Timestamp List ( 1.3.6.1.4.1.11129.2.4.2 )  
Critical NO  
SCT Version 1  
Log Key ID EE 4B BD B7 75 CE 60 BA E1 42 69 1F AB E1 9E 66 A3 0F 7E 5F B0 72 D8  
Timestamp Monday, August 20, 2018 at 4:35:53 PM Eastern Daylight Time  
Signature Algorithm SHA-256 ECDSA  
Signature 71 bytes : 30 45 02 21 00 F3 D6 BD ...

```
$ openssl crl -inform DER -text -noout -in InCommonRSAServerCA.crl
```

# Revocation Check (2)

## Online Certificate Status Protocol



# Revocation Check (2)

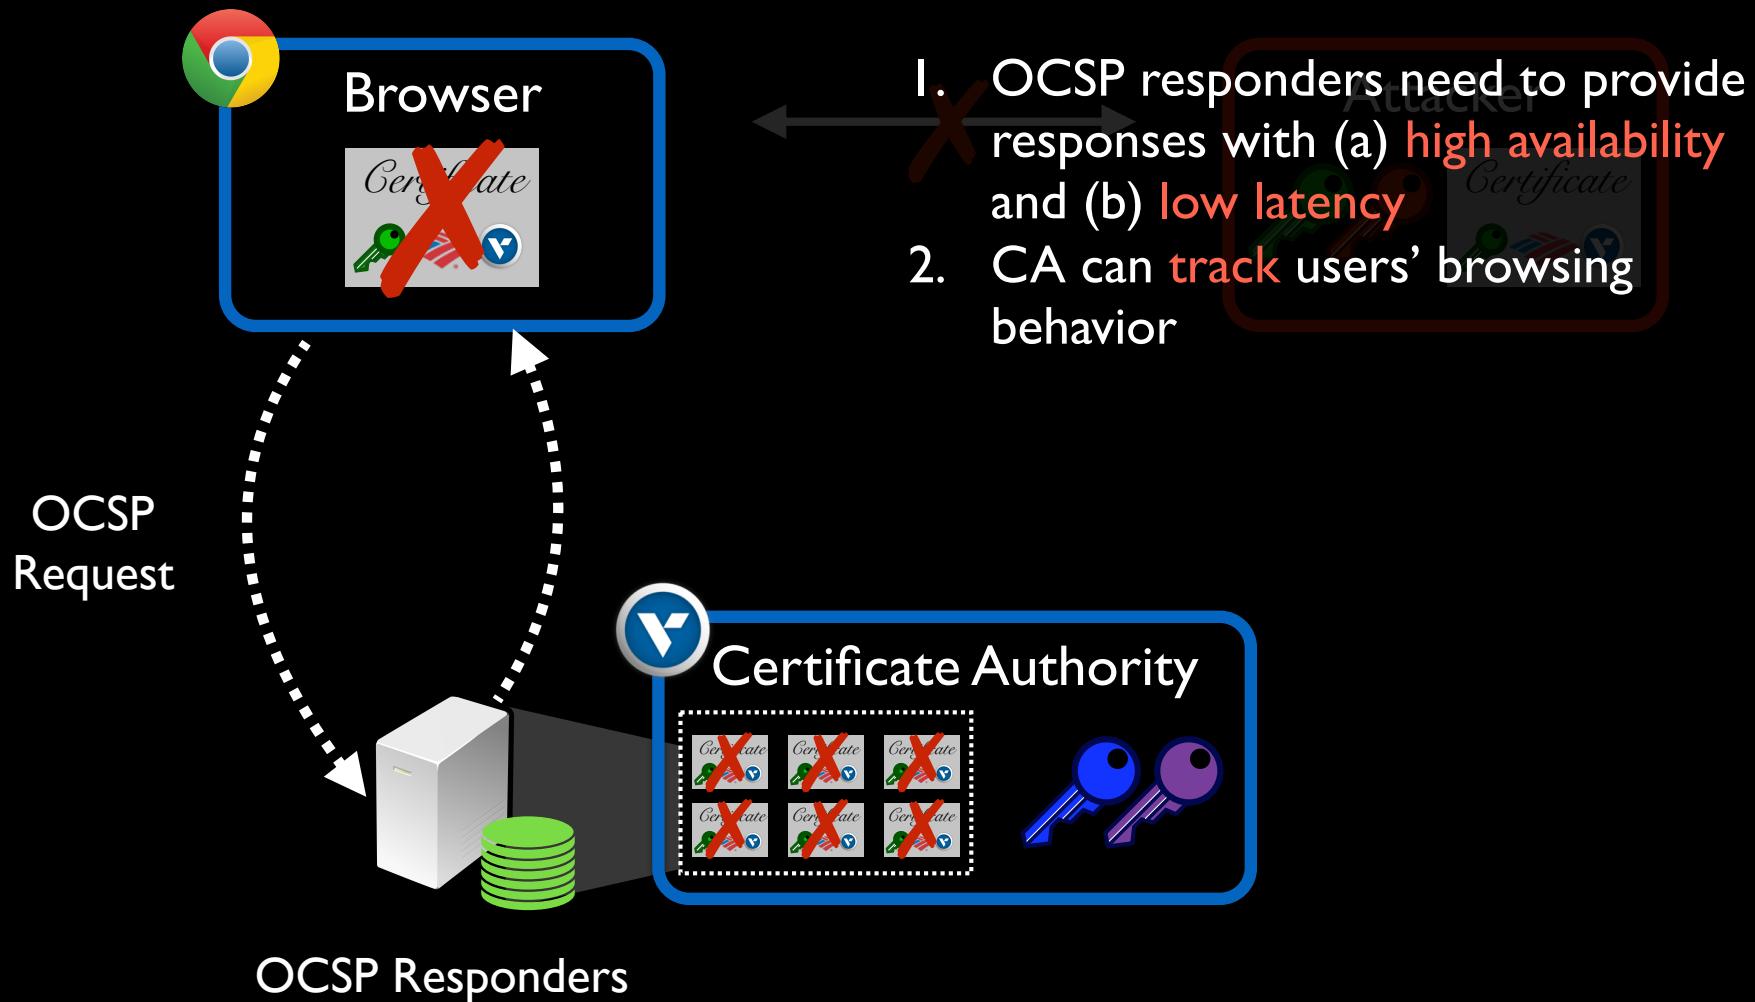
## Online Certificate Status Protocol

The screenshot shows a browser window with the URL <https://www.rit.edu>. The page displays the OCSP response for the certificate of [www.rit.edu](http://crl.incommon-rsa.org/InCommonRSAServerCA.crl). The response includes the following details:

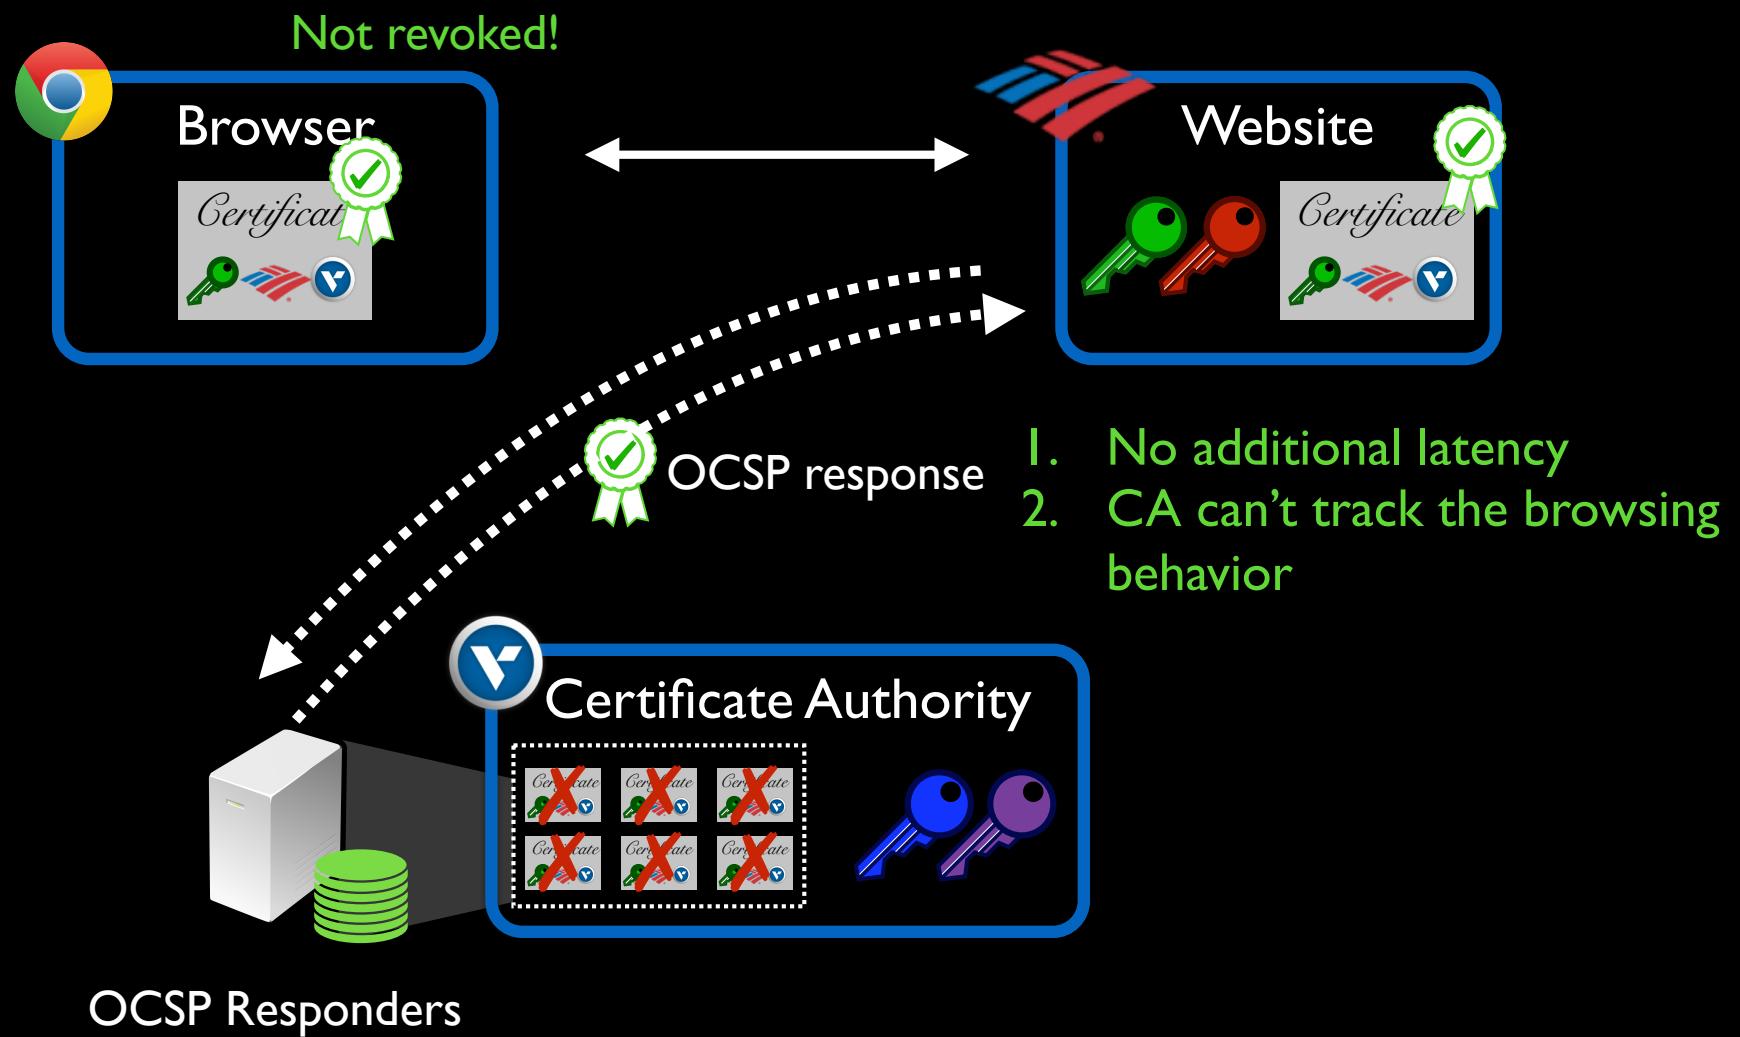
- CPS URI: [https://www.incommon.org/cert/repository/cps\\_ssl.pem](https://www.incommon.org/cert/repository/cps_ssl.pem)
- Policy ID #2 (2.23.140.1.2.2)
- Extension: CRL Distribution Points (2.5.29.31)
  - Critical: NO
  - URI: <http://crl.incommon-rsa.org/InCommonRSAServerCA.crl>
- Extension: Embedded Signed Certificate Timestamp List (1.3.6.1.4.1.11129.2.4.2)
  - Critical: NO
- SCT Version: 1
- Log Key ID: EE 4B BD B7 75 CE 60 BA E1 42 69 1F AB E1 9E 66 A3 0F 7E 5F B0 72 D8 83 00 C4 7B 89 7A A8 FD CB
- Timestamp: Monday, August 20, 2018 at 4:35:53 PM Eastern Daylight Time
- Signature Algorithm: SHA-256 ECDSA
- Signature: 71 bytes : 30 45 02 21 00 F3 D6 BD ...
- Extension: Certificate Authority Information Access (1.3.6.1.5.5.7.1.1)
  - Critical: NO
  - Method #1: CA Issuers (1.3.6.1.5.5.7.48.2)
    - URI: [http://crt.usertrust.com/InCommonRSAServerCA\\_2.crt](http://crt.usertrust.com/InCommonRSAServerCA_2.crt)
  - Method #2: Online Certificate Status Protocol (1.3.6.1.5.5.7.48.1)
    - URI: <http://ocsp.usertrust.com>

```
$ openssl ocsp -issuer cert.pem -serial  
5226810331521645508876562747113126991 -url http://ocsp.usertrust.com  
-header host ocsp.usertrust.com
```

# Challenges of Online Certificate Status Protocol

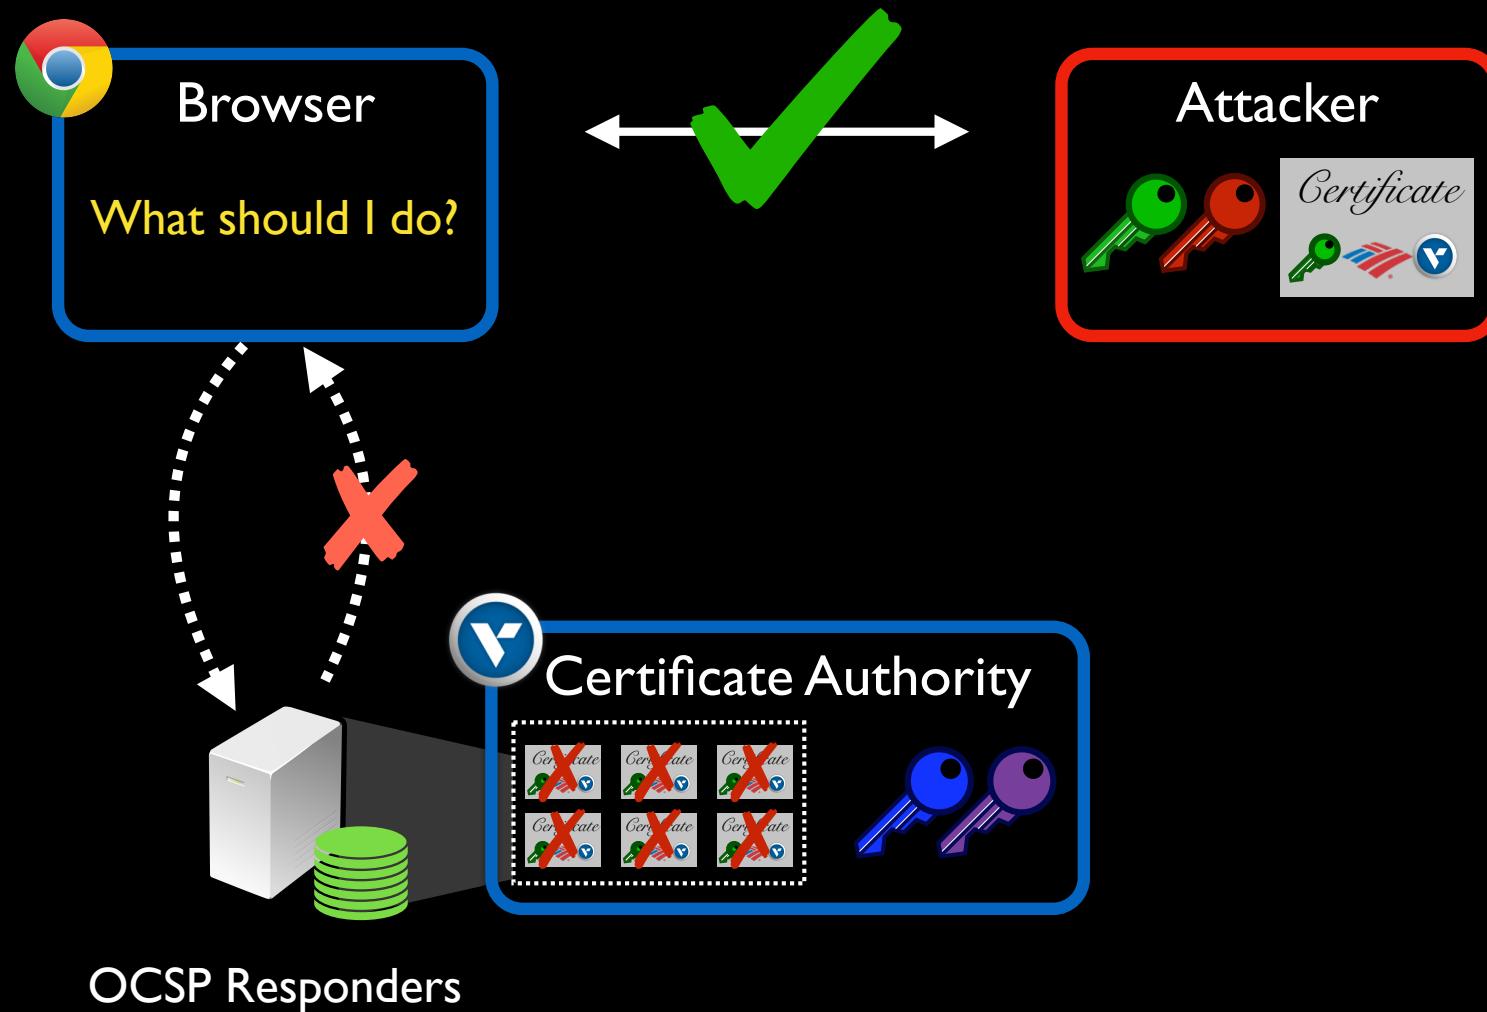


# OCSP Stapling



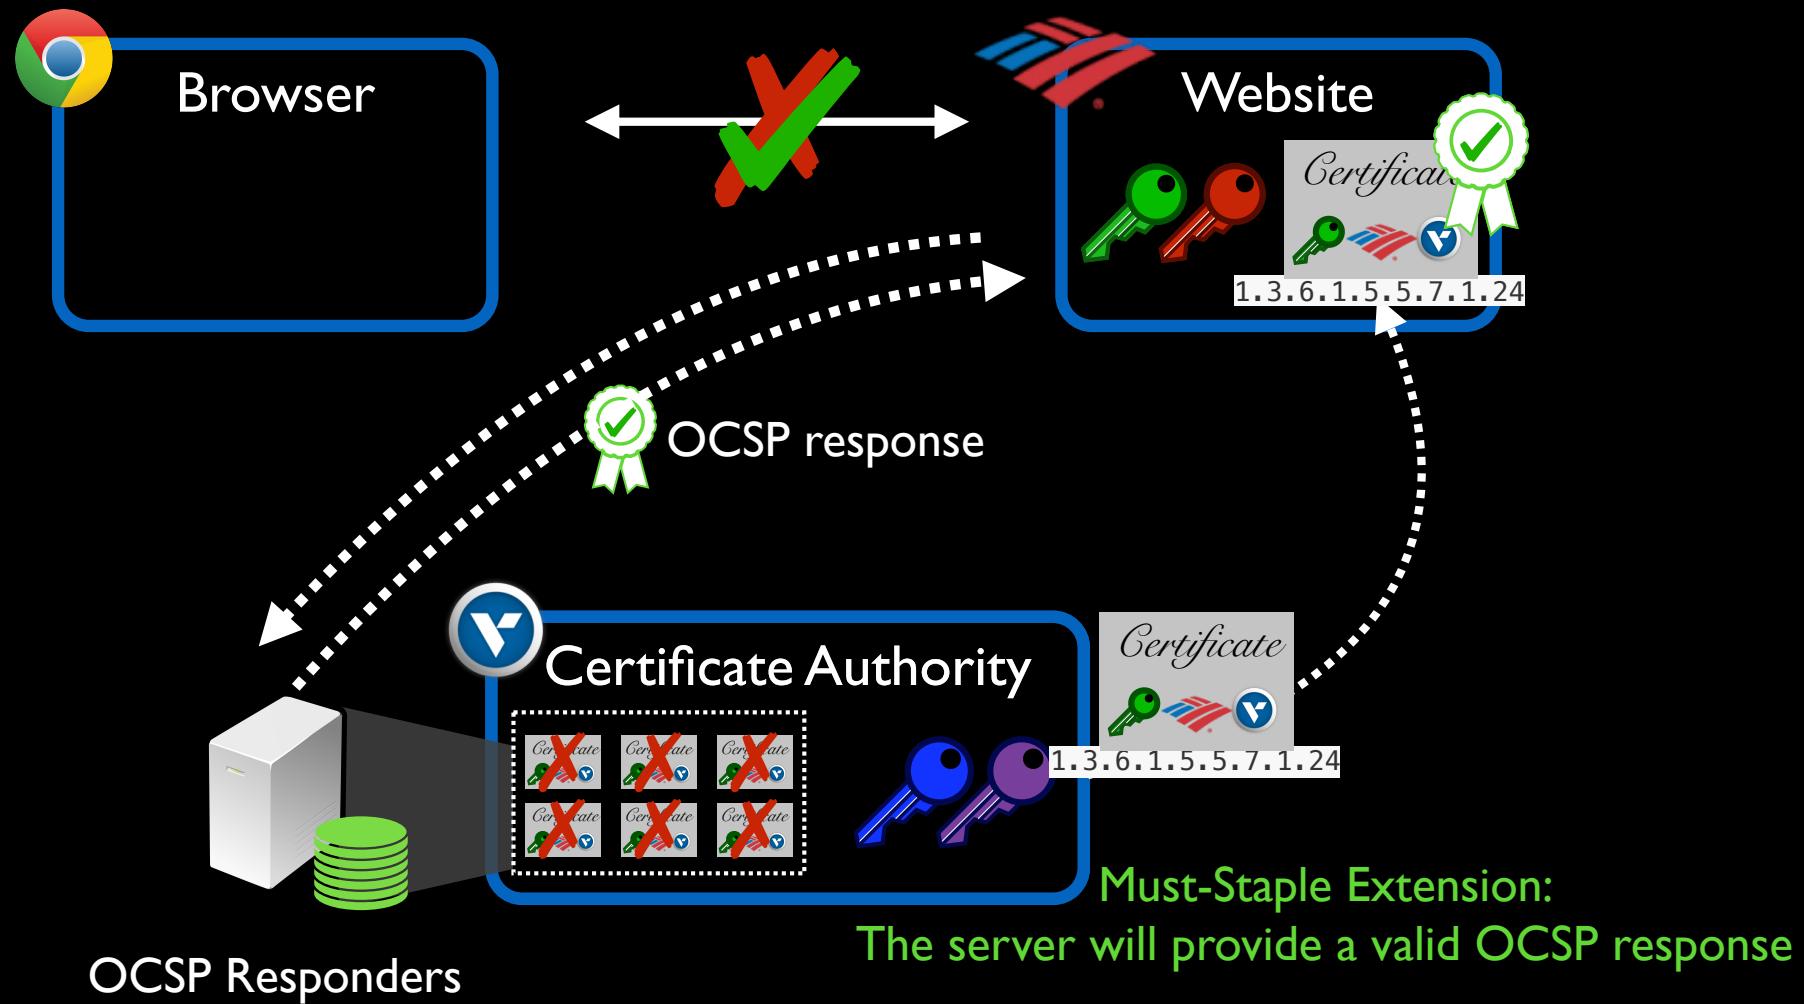
# Challenges still remain: Soft failure

Most clients will accept a certificate  
even if they are **unable** to obtain revocation information



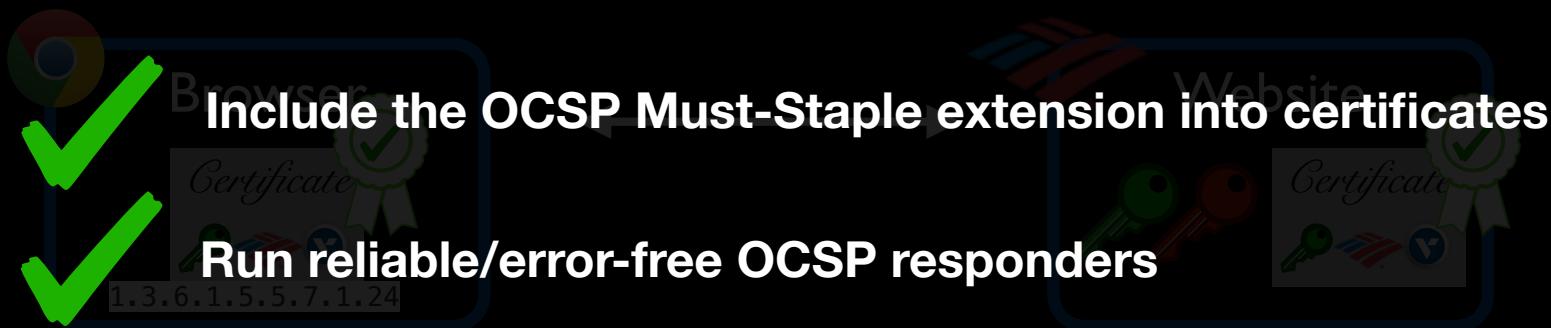
# OCSP Must-Staple

- ✓ No additional latency
- ✓ No privacy issues
- ✓ No soft failure

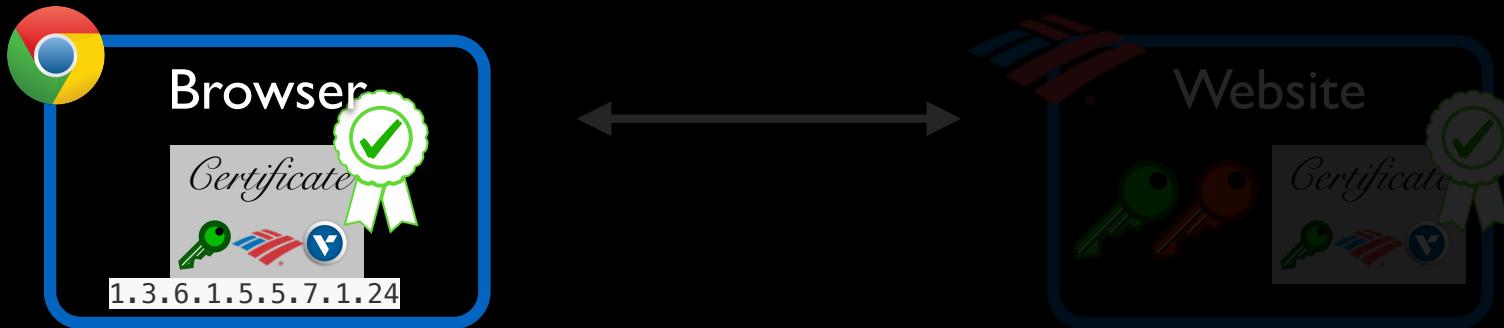


# To support OCSP Must Staple

## (I) CA



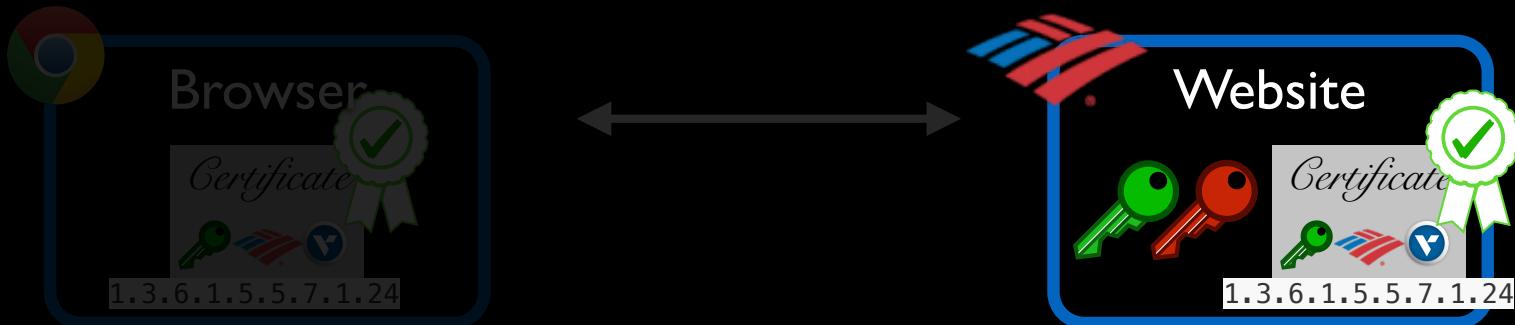
# To support OCSP Must Staple (2) Clients



- ✓ Understand the OCSP Must-Staple extension in the certificate
  - ✓ Present the Certificate Status Request (CSR) to the web servers
  - ✓ Reject the certificate if they do not receive OCSP responses
- OCSP Responders

# To support OCSP Must Staple

## (3) Web servers



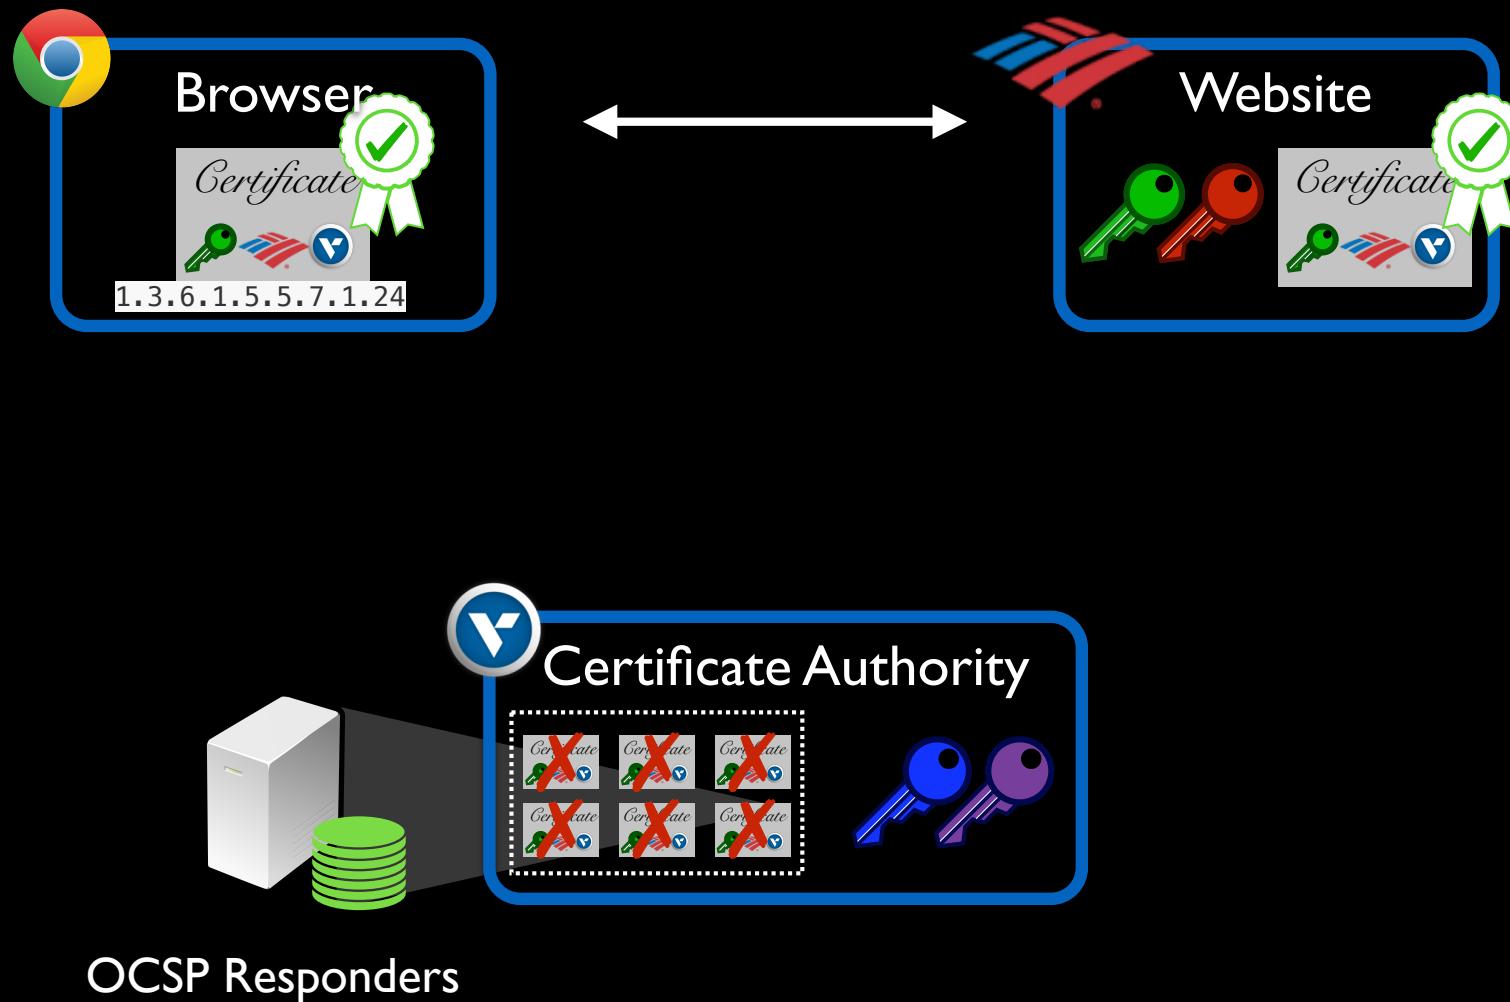
✓ **(Web server software) must fetch/cache OCSP responses**

✓ **(Web server administrators) must configure to use OCSP stapling**



OCSP Responders

# To support OCSP Must Staple



# Is the Web Ready for OCSP Must-Staple?



Certificate Authority  
(OCSP Responder)

- ✓ Availability
- ✓ Validity
- ✓ Consistency with CRL

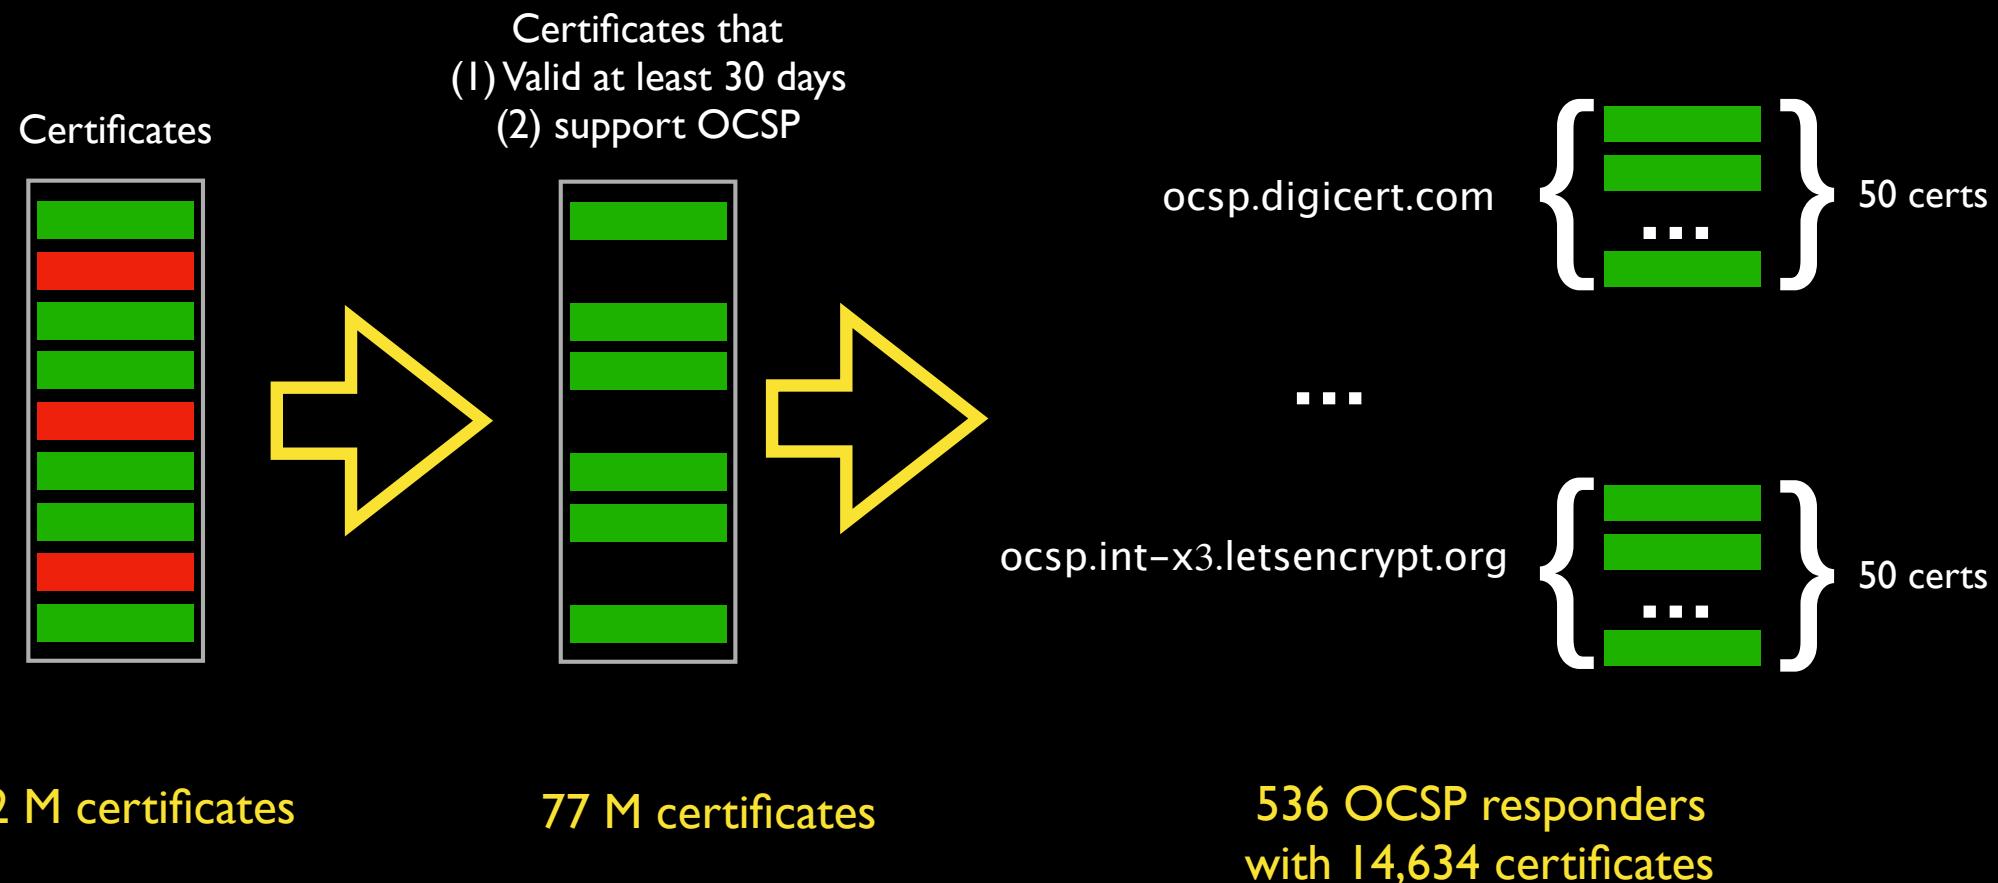


Website

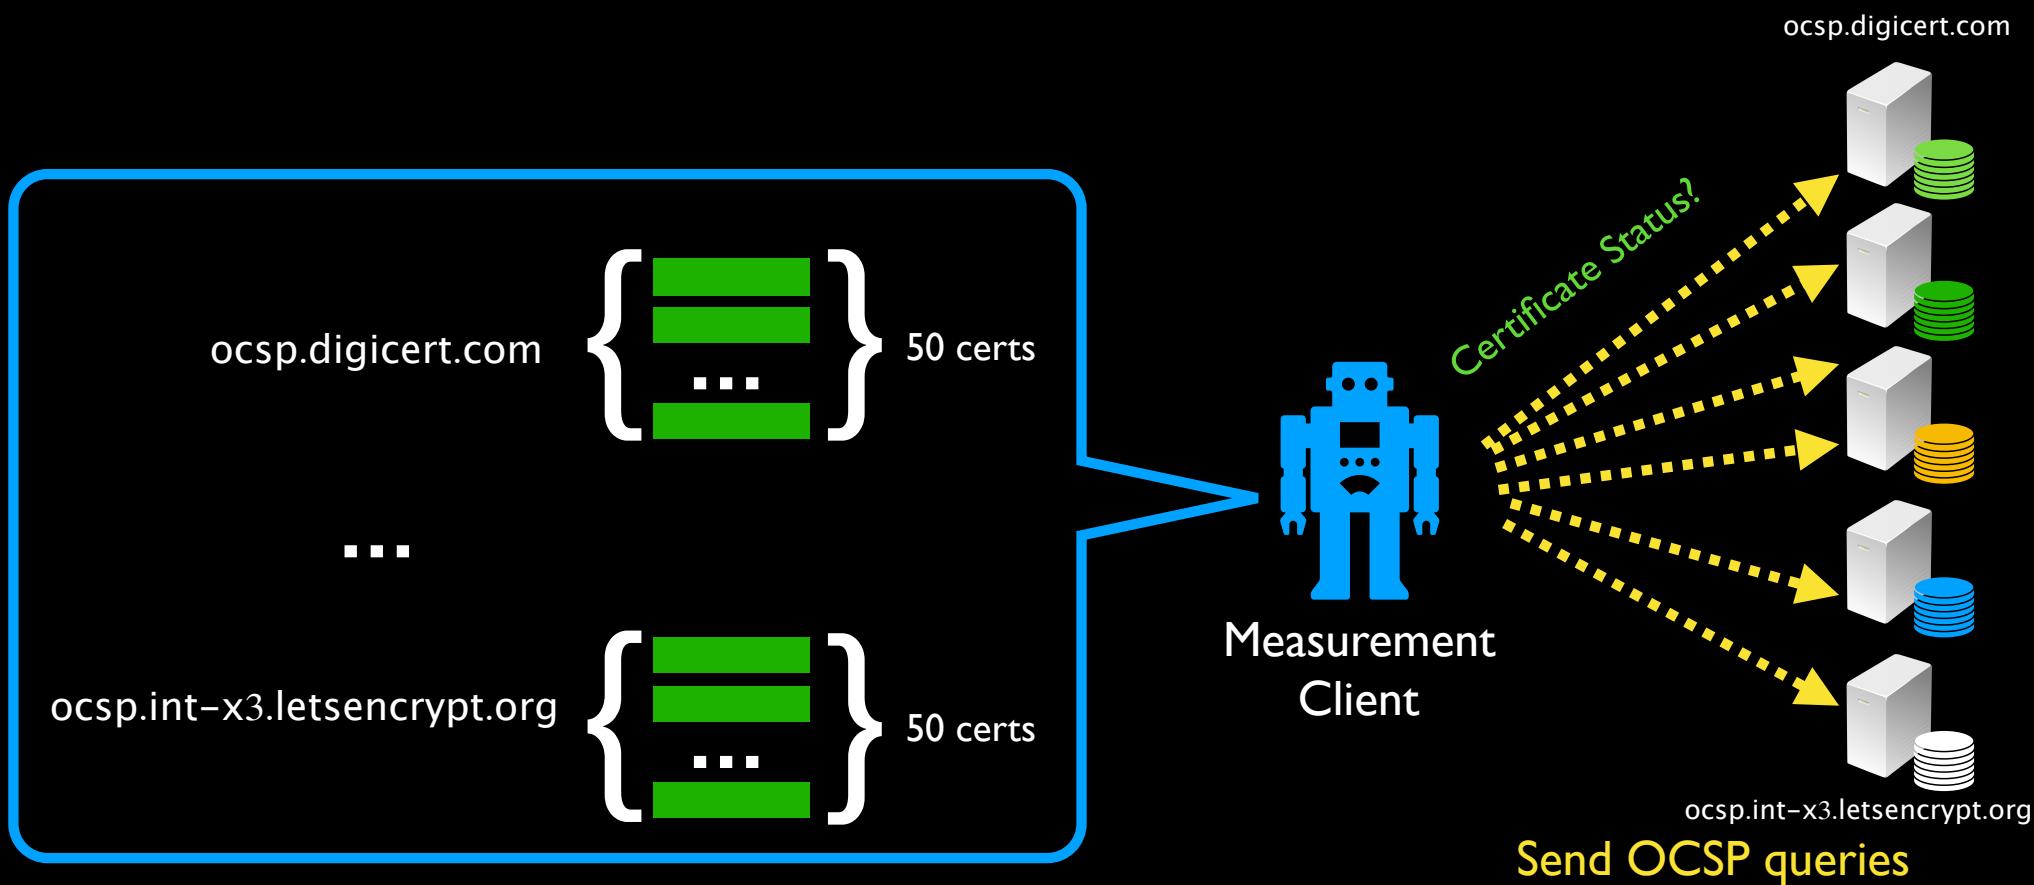


Browser

# Measuring OCSP Responders



# Measuring OCSP Responders



# Measurement



Oregon (US West)



Virginia (US East)



São Paulo (Brazil)



Paris (France)



Sydney (Australia)

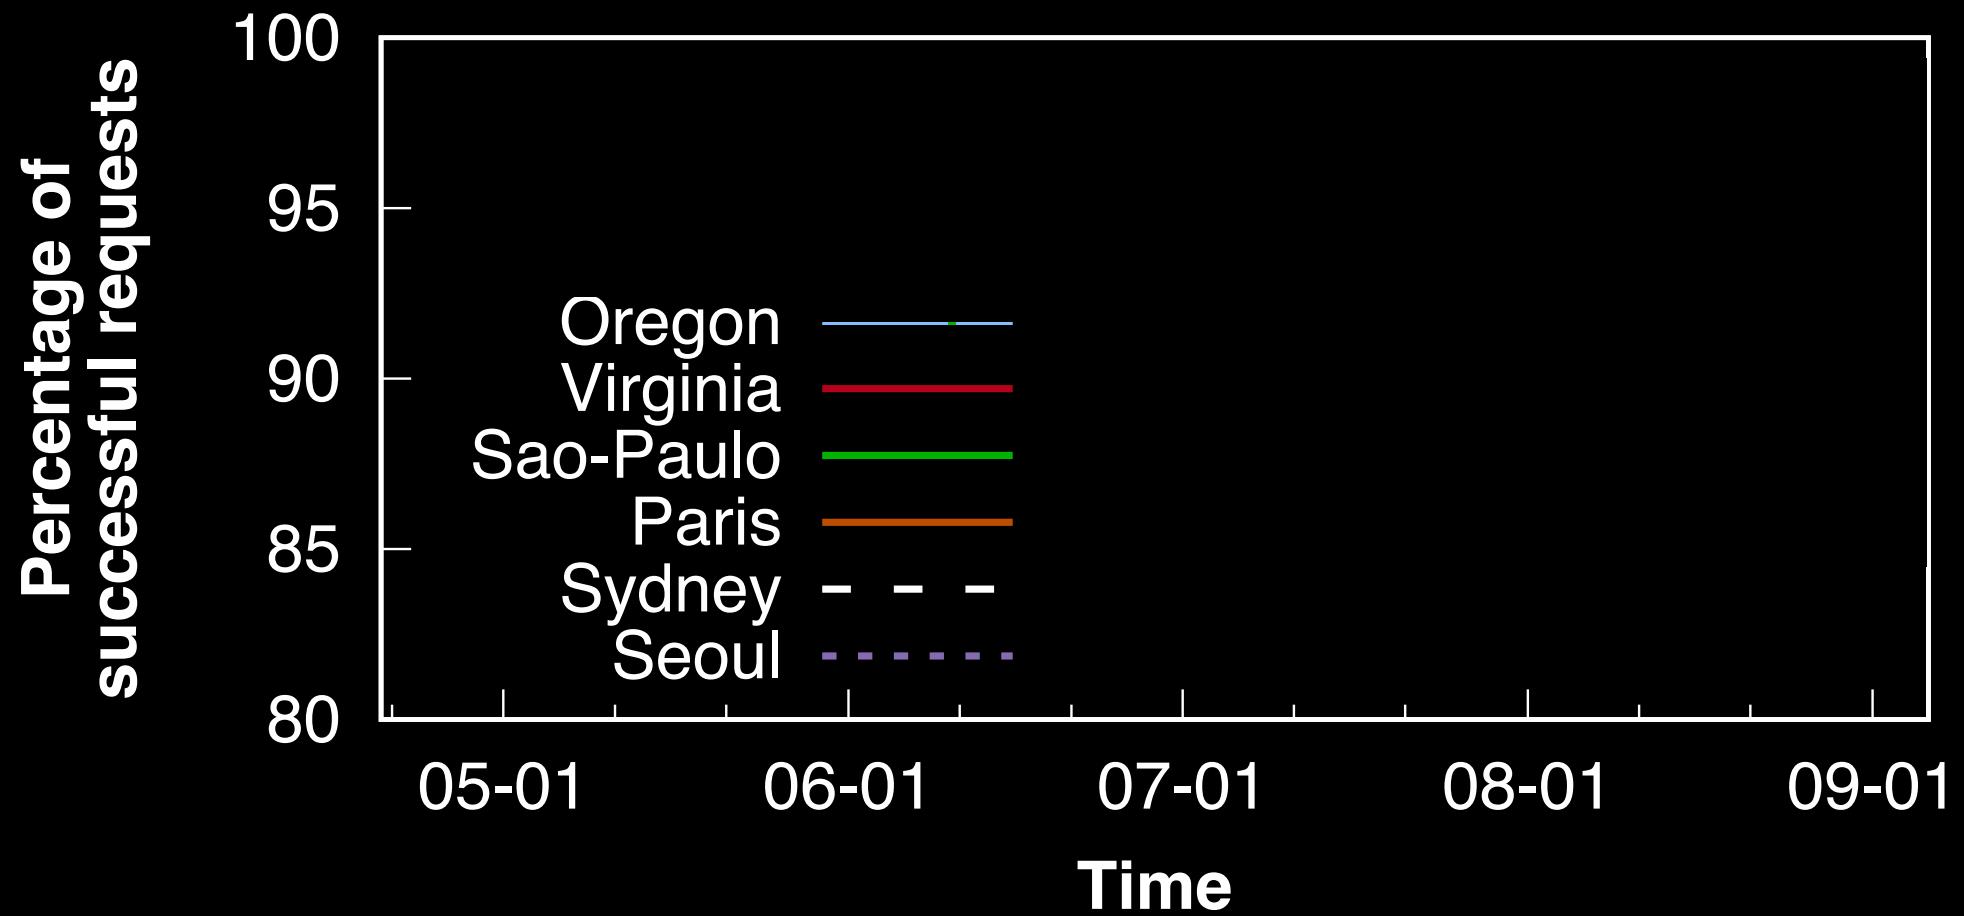


Seoul (Korea)

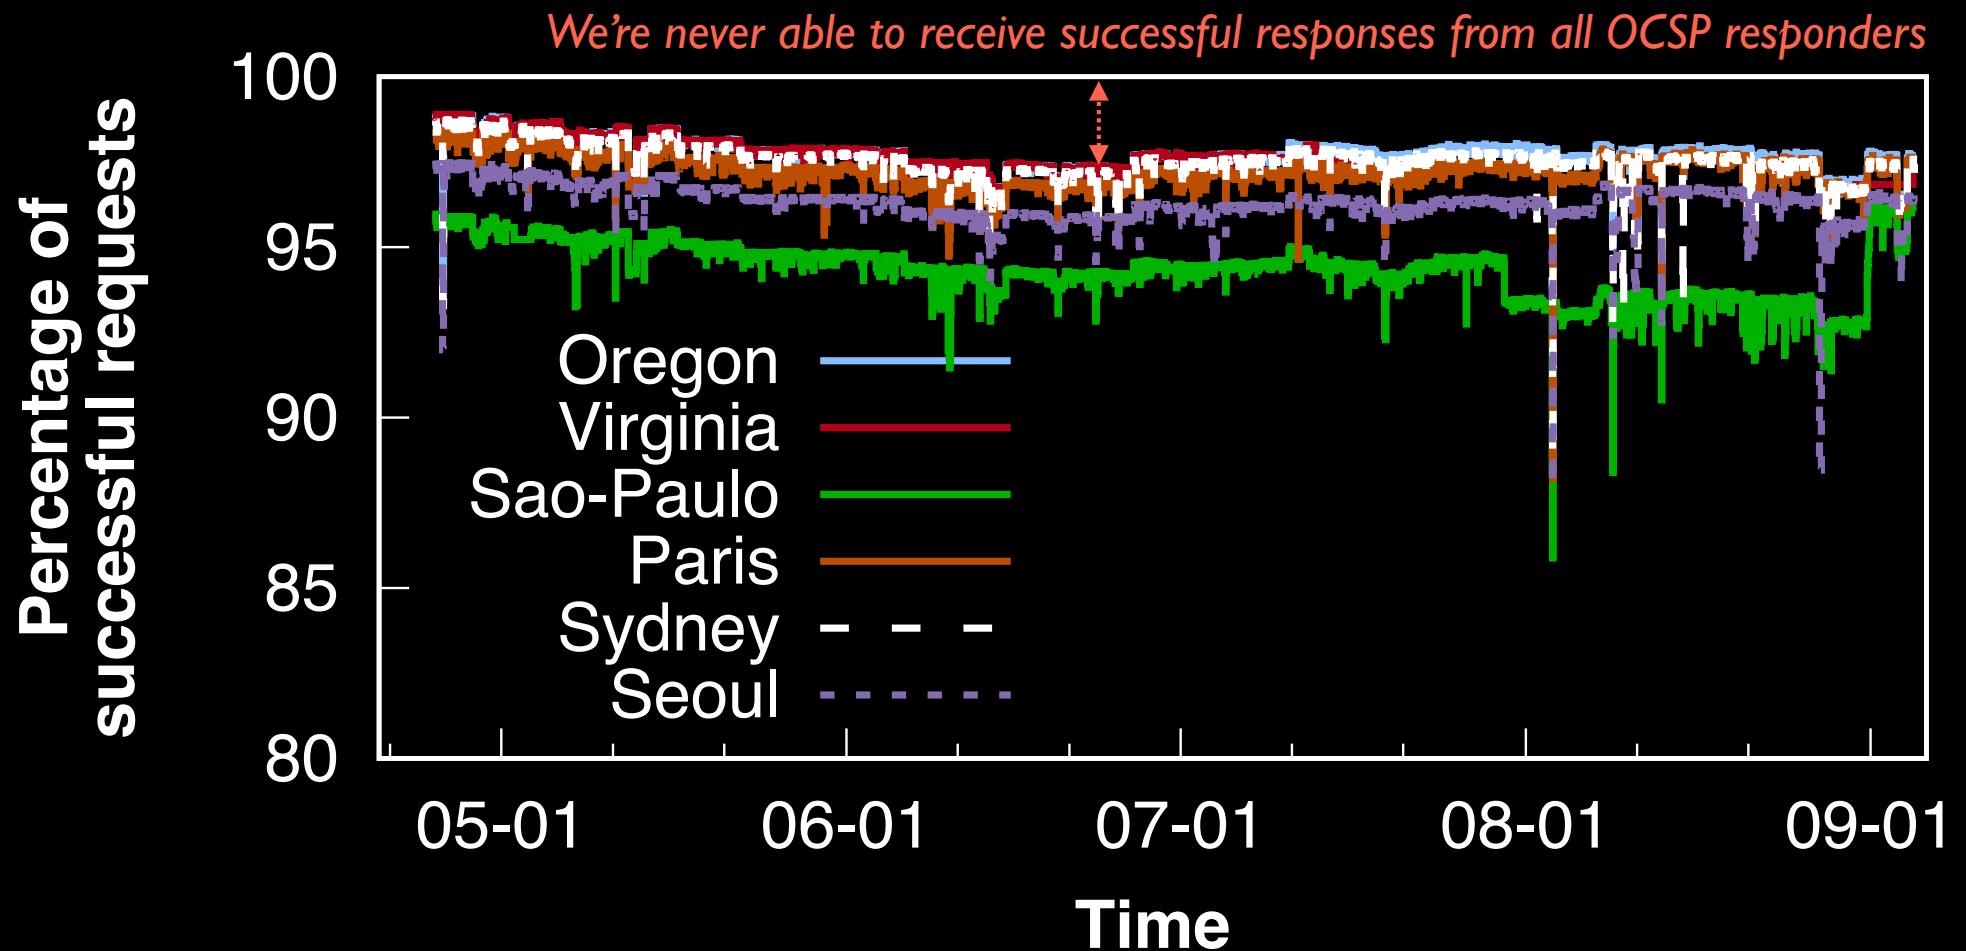
Scan them every hour  
April 25, 2018 ~ September 4, 2018

~ 46 M OCSP requests & responses

# (I) Availability



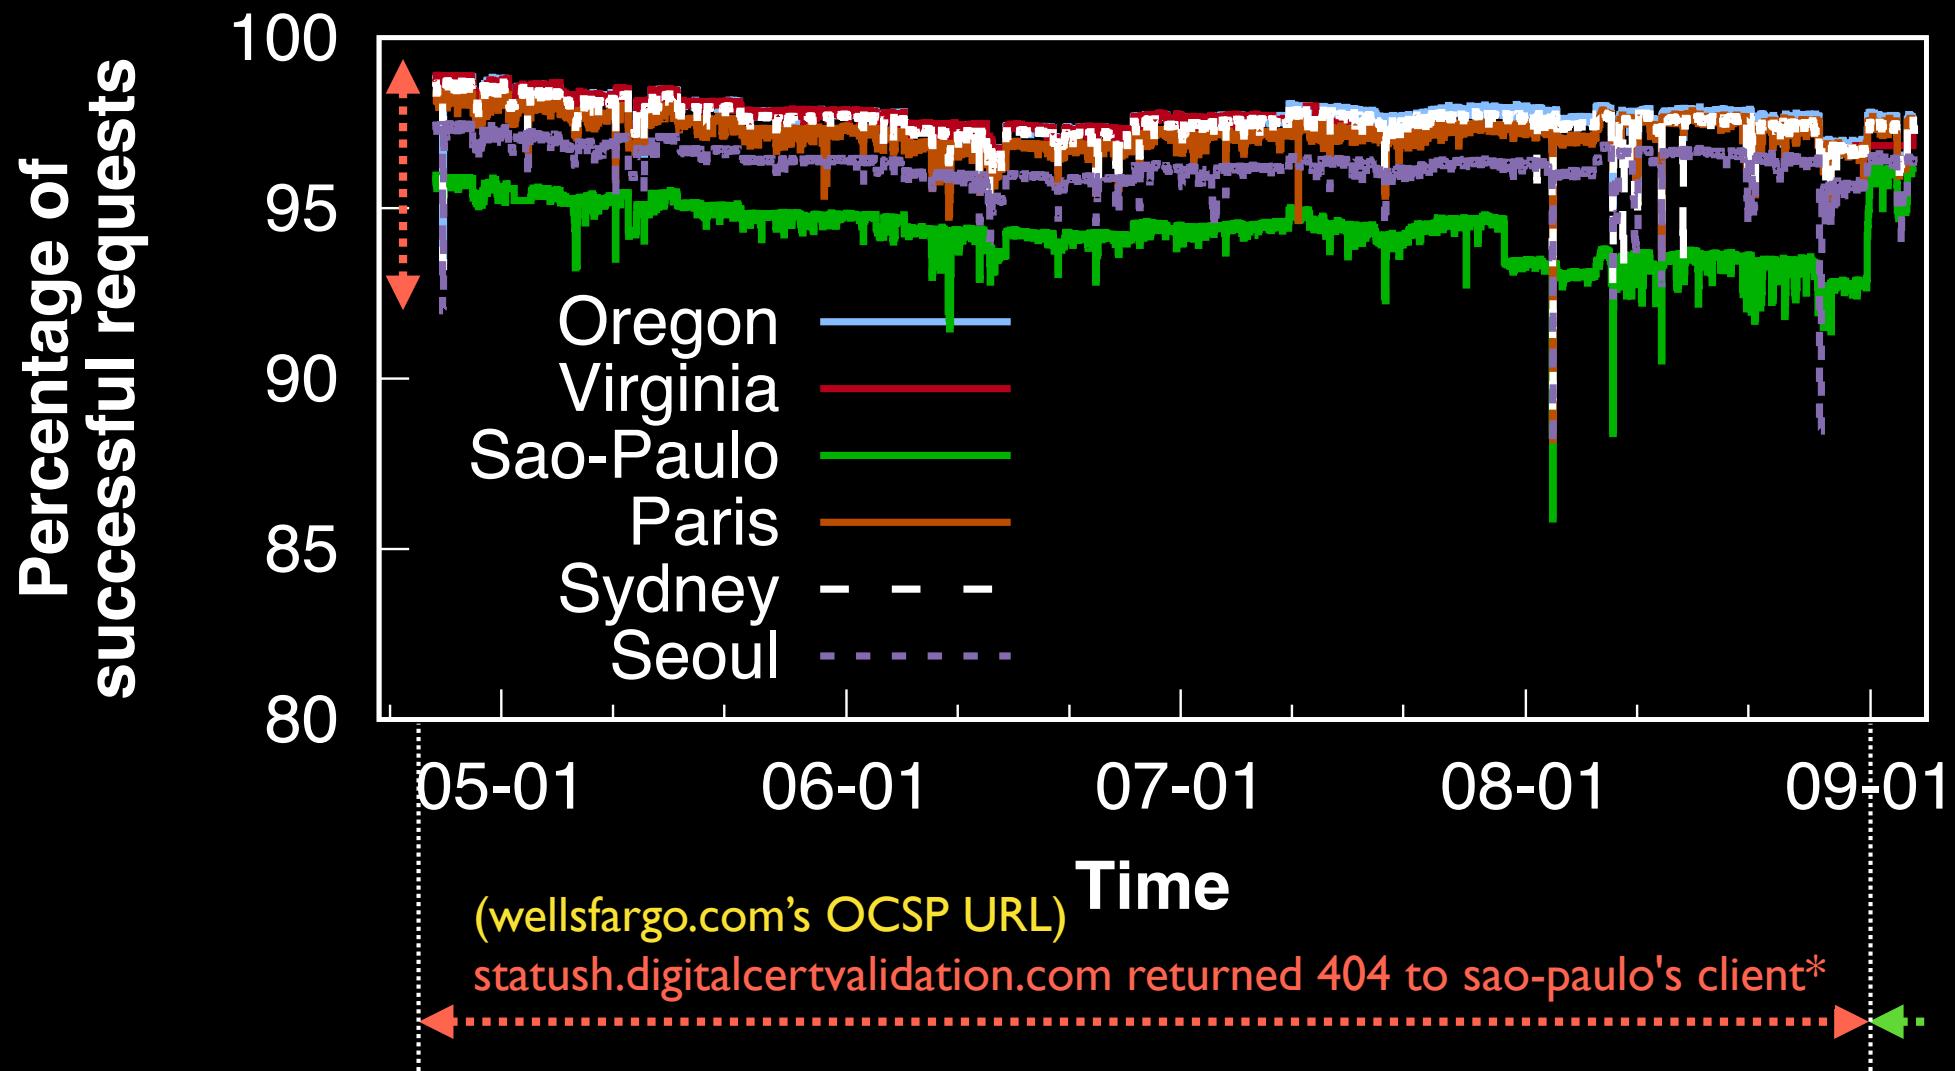
# (I) Availability Overview



*For 29 OCSP responders, there was at least one measurement client  
that was never able to make a successful request.*

*(16: DNS problem, 4: TCP connection errors, 8: HTTP problems, 1: HTTPS Error)*

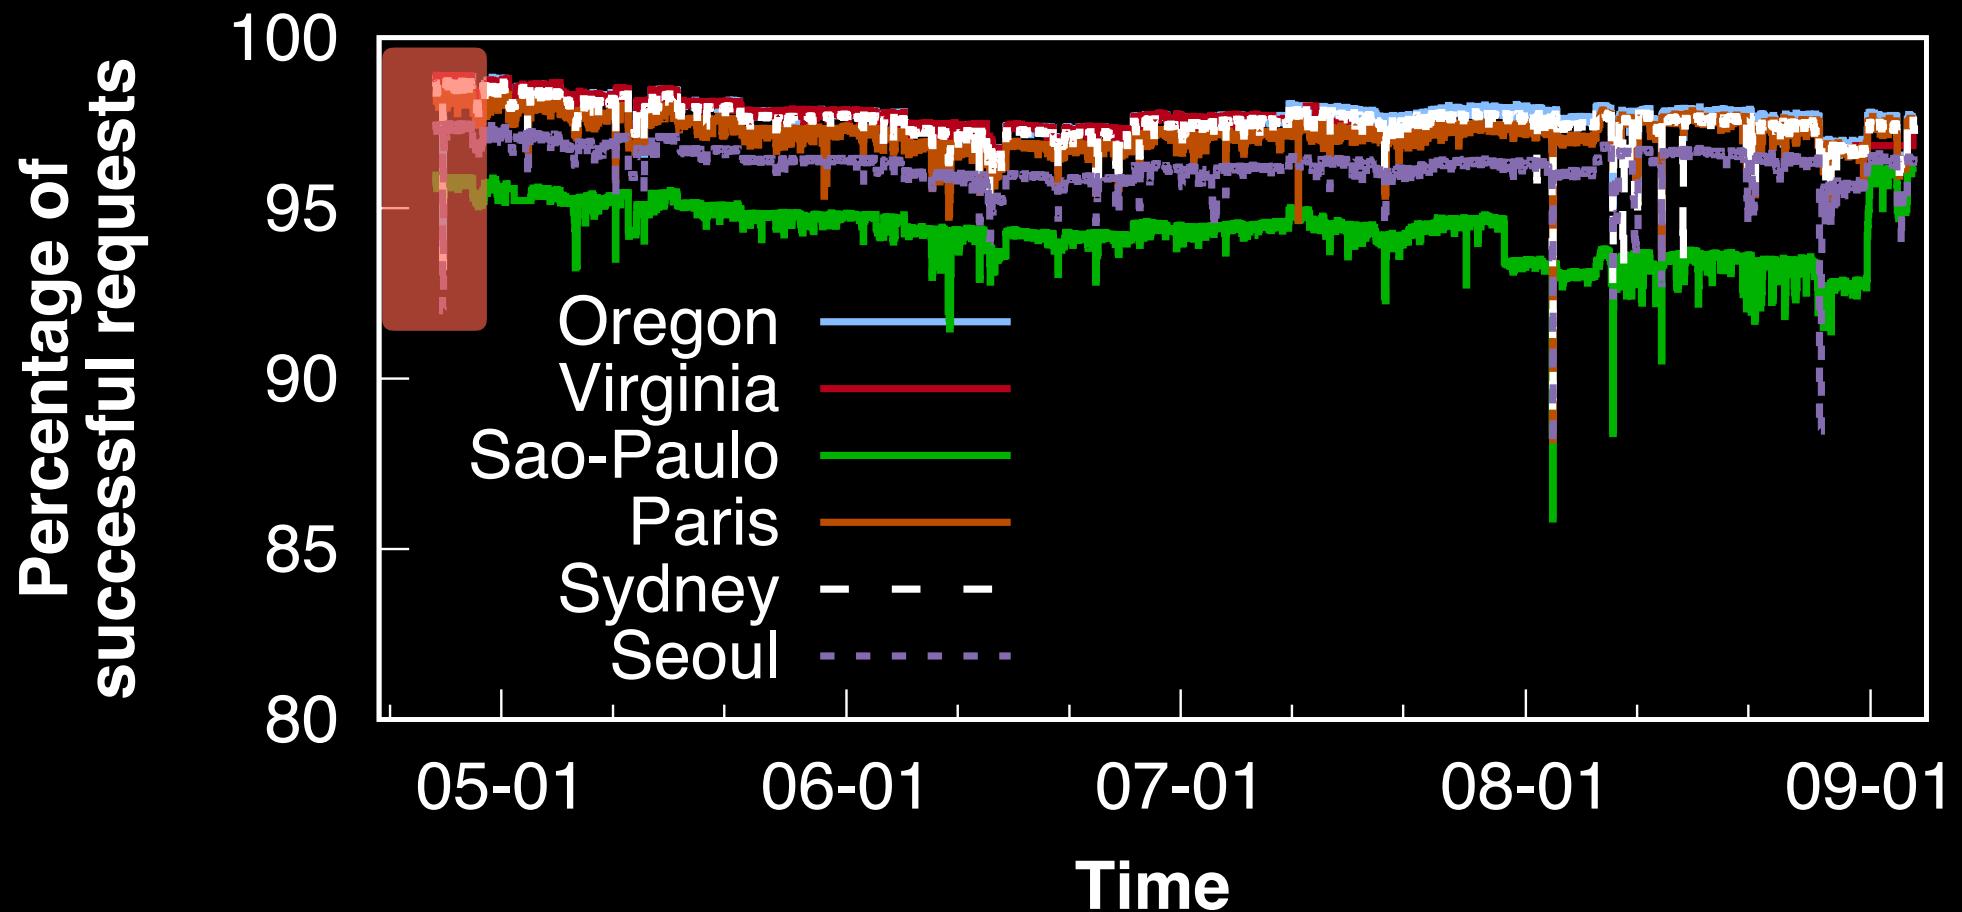
# (I) Availability: Geographical Differences



\*After we contacted them on August 29th, the issue was fixed at 11pm August 31st.

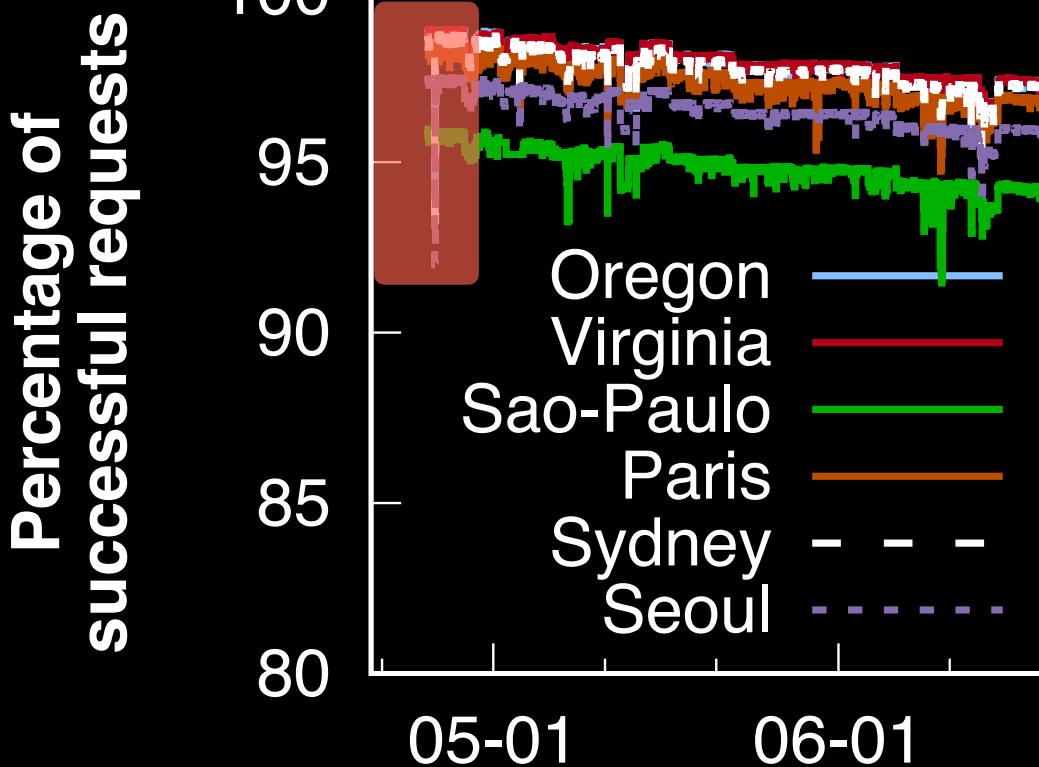
# (I) Availability: Transient Failure

Seoul, Sydney, and Oregon (Asia Pacific)



# (I) Availability: Transient Failure (Case-Study)

Seoul, Sydney, and Oregon (Asia Pacific)



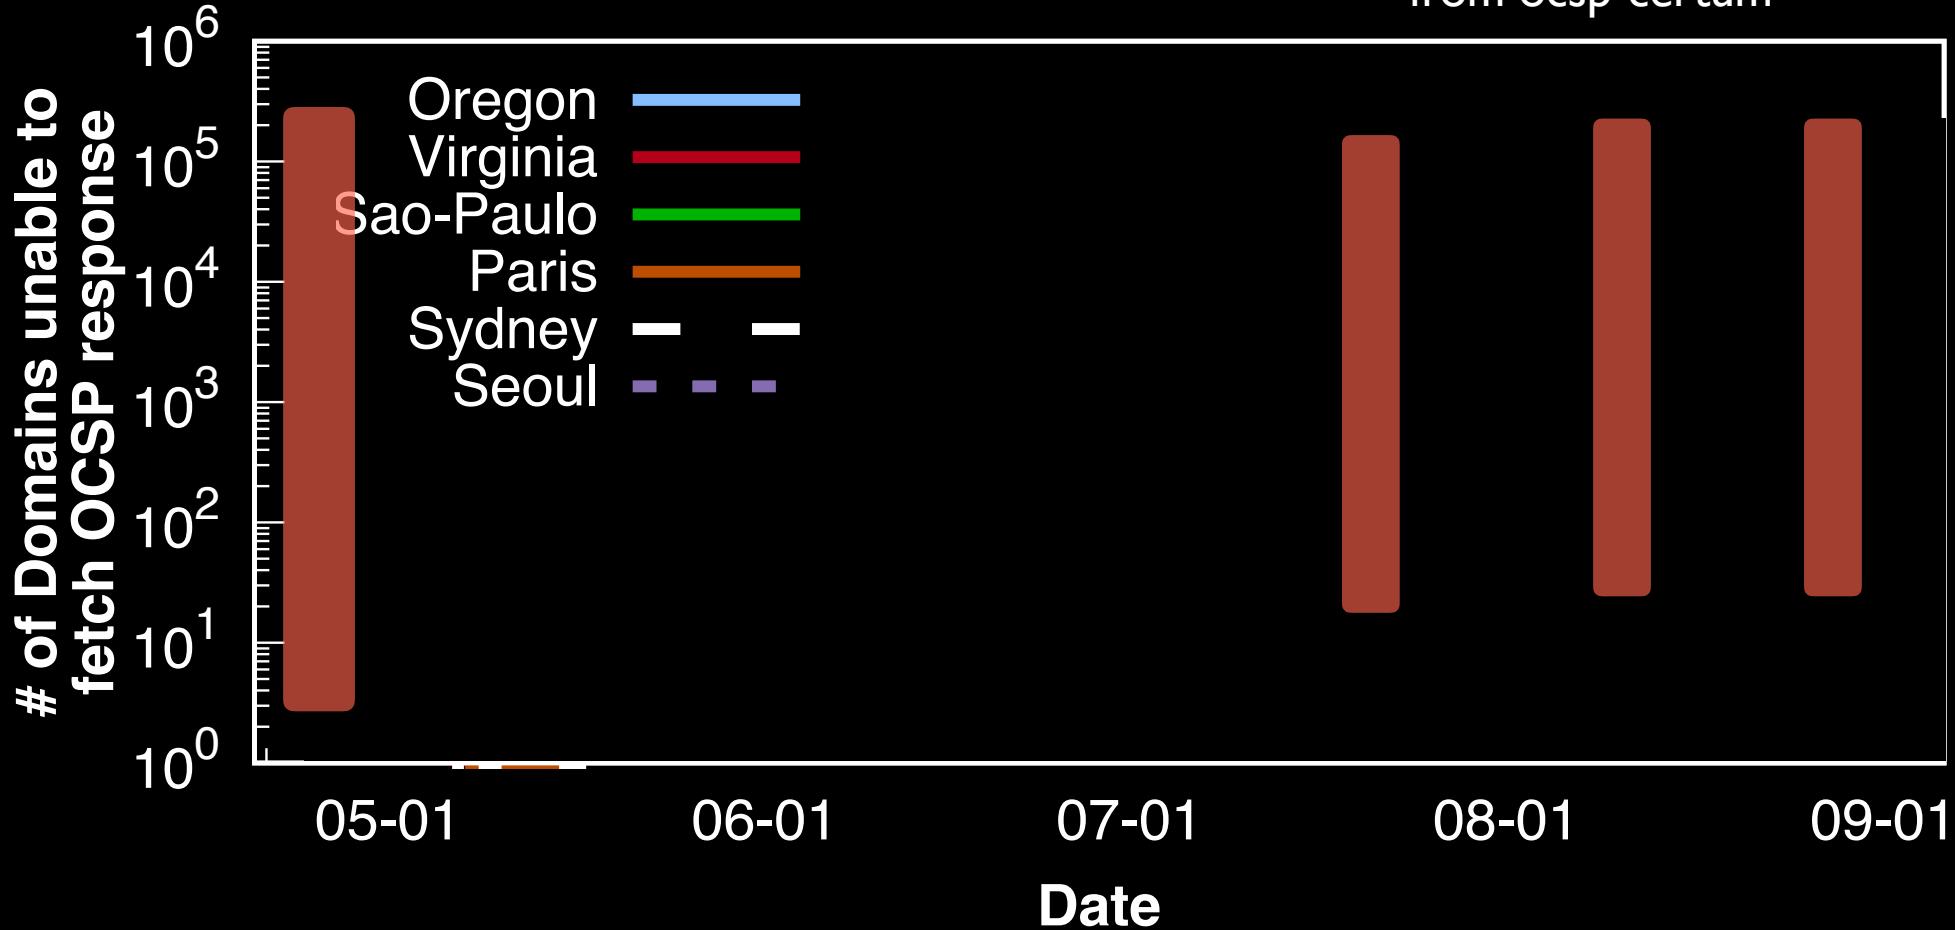
OCSP Server Name	DNS Records
ocsp.comodoca.com	
ocsp.comodoca4.com	
ocsp.gandi.net	CNAME: ocsp.comodoca.com
ocsp.globessl.com	CNAME: ocsp.comodoca.com
ocsp.incommon-ecc.org	CNAME: ocsp.comodoca.com
ocsp.incommon-igtf.org	NS: ns0.comododns.com.
ocsp.incommon-rsa.org	NS: ns0.comododns.com.
OCSP.intel.com	CNAME: ocsp.comodoca.com
ocsp.marketware.eu	CNAME: ocsp.comodoca.com
ocsp.netsolssl.com	CNAME: ocsp.comodoca.com
ocsp.register.com	CNAME: ocsp.comodoca.com
ocsp.securecore-ca.com	NS: ns0.comododns.com.
ocsp.sgssl.net.	NS: ns0.comododns.com.
ocsp.trustasiassl.com.	NS: ns0.comododns.com.
ocsp.trust-provider.com	CNAME: ocsp.comodoca.com
ocsp.usertrust.com	NS: ns0.comododns.com.

# (I) Availability: Impact on the Web

Comodo  
down for 2 hours

43 servers from wosign  
5 servers from startssl

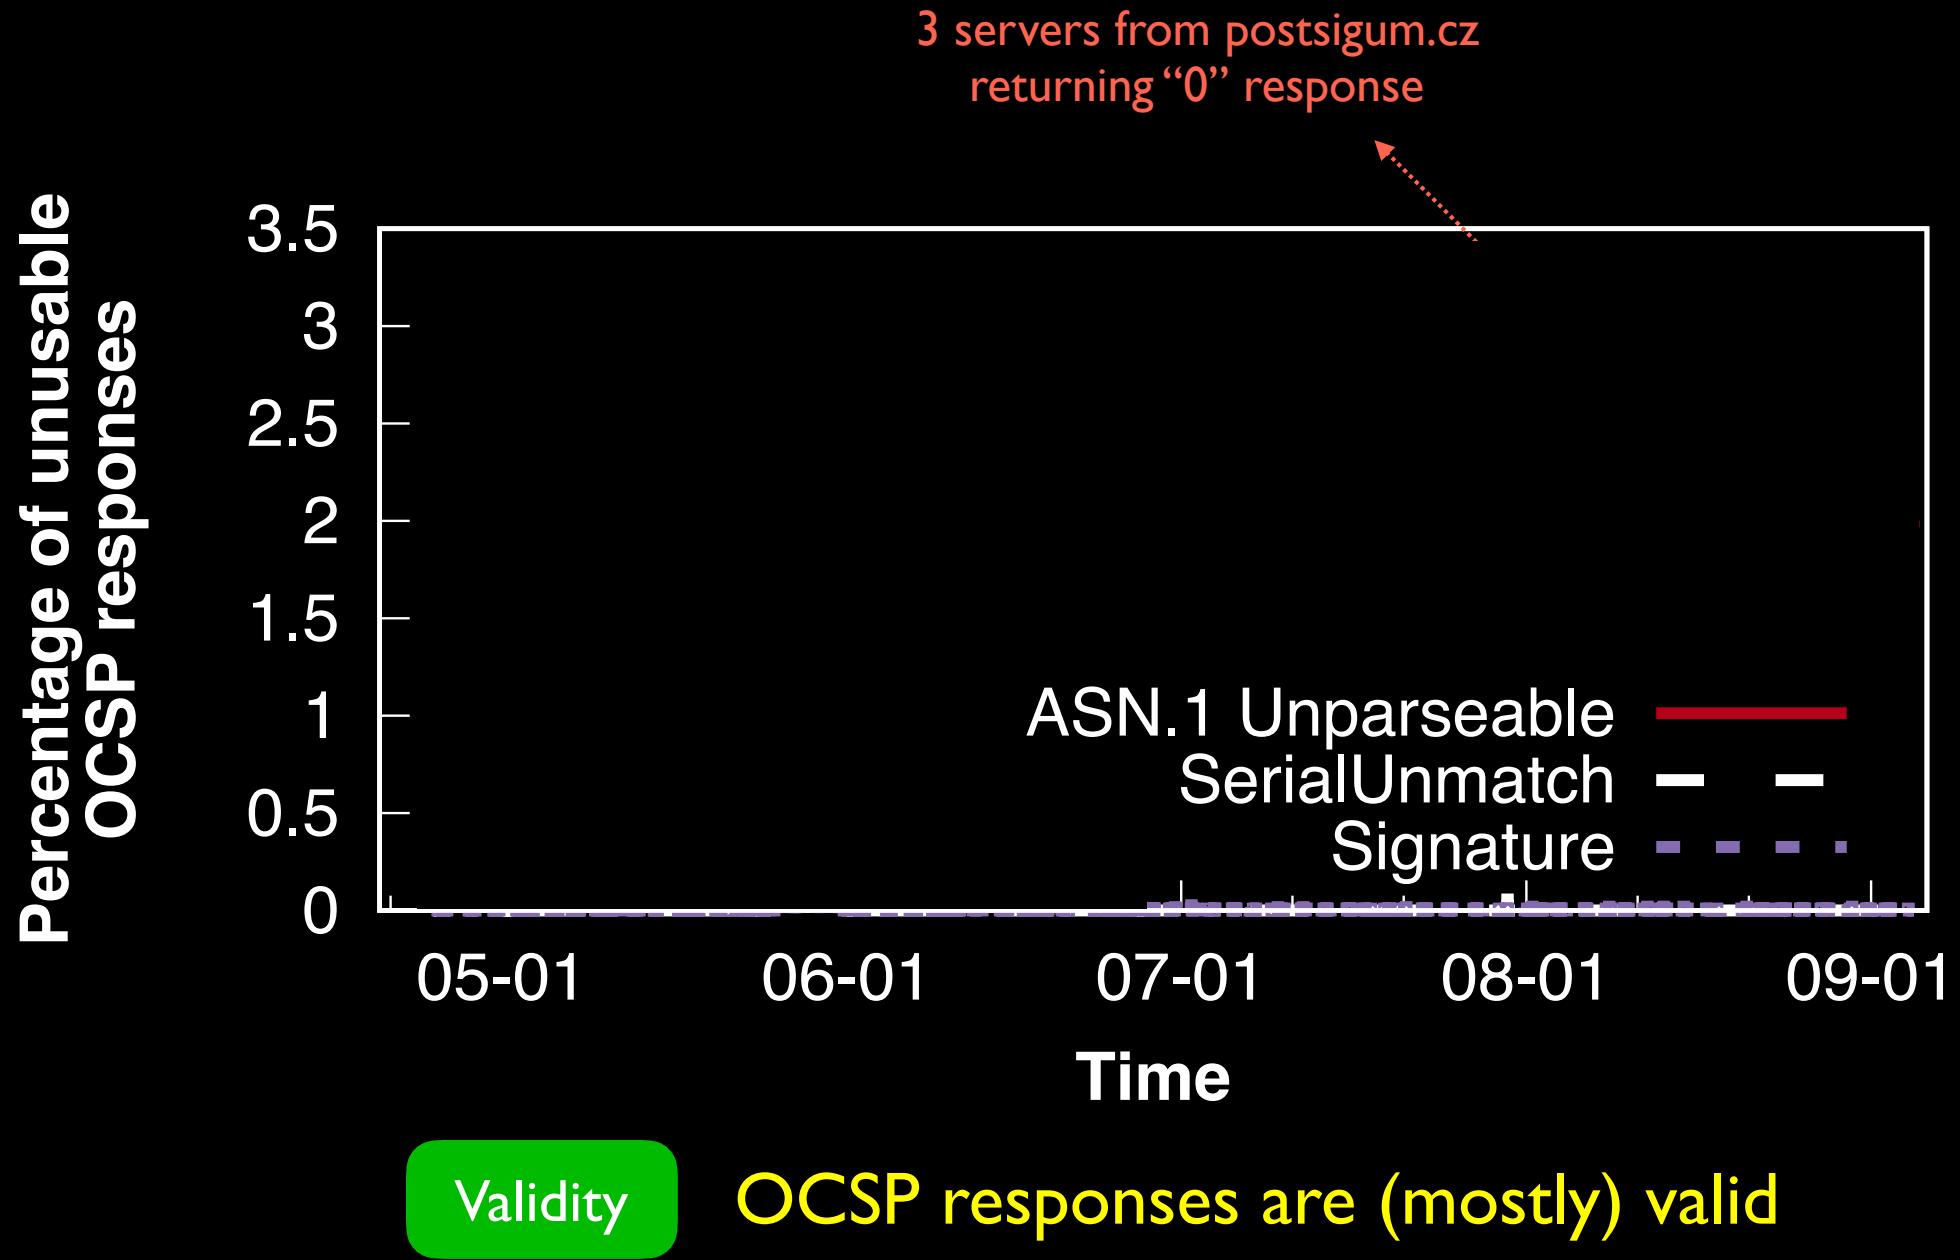
9 servers  
from digicert  
16 servers  
from ocsp-certum



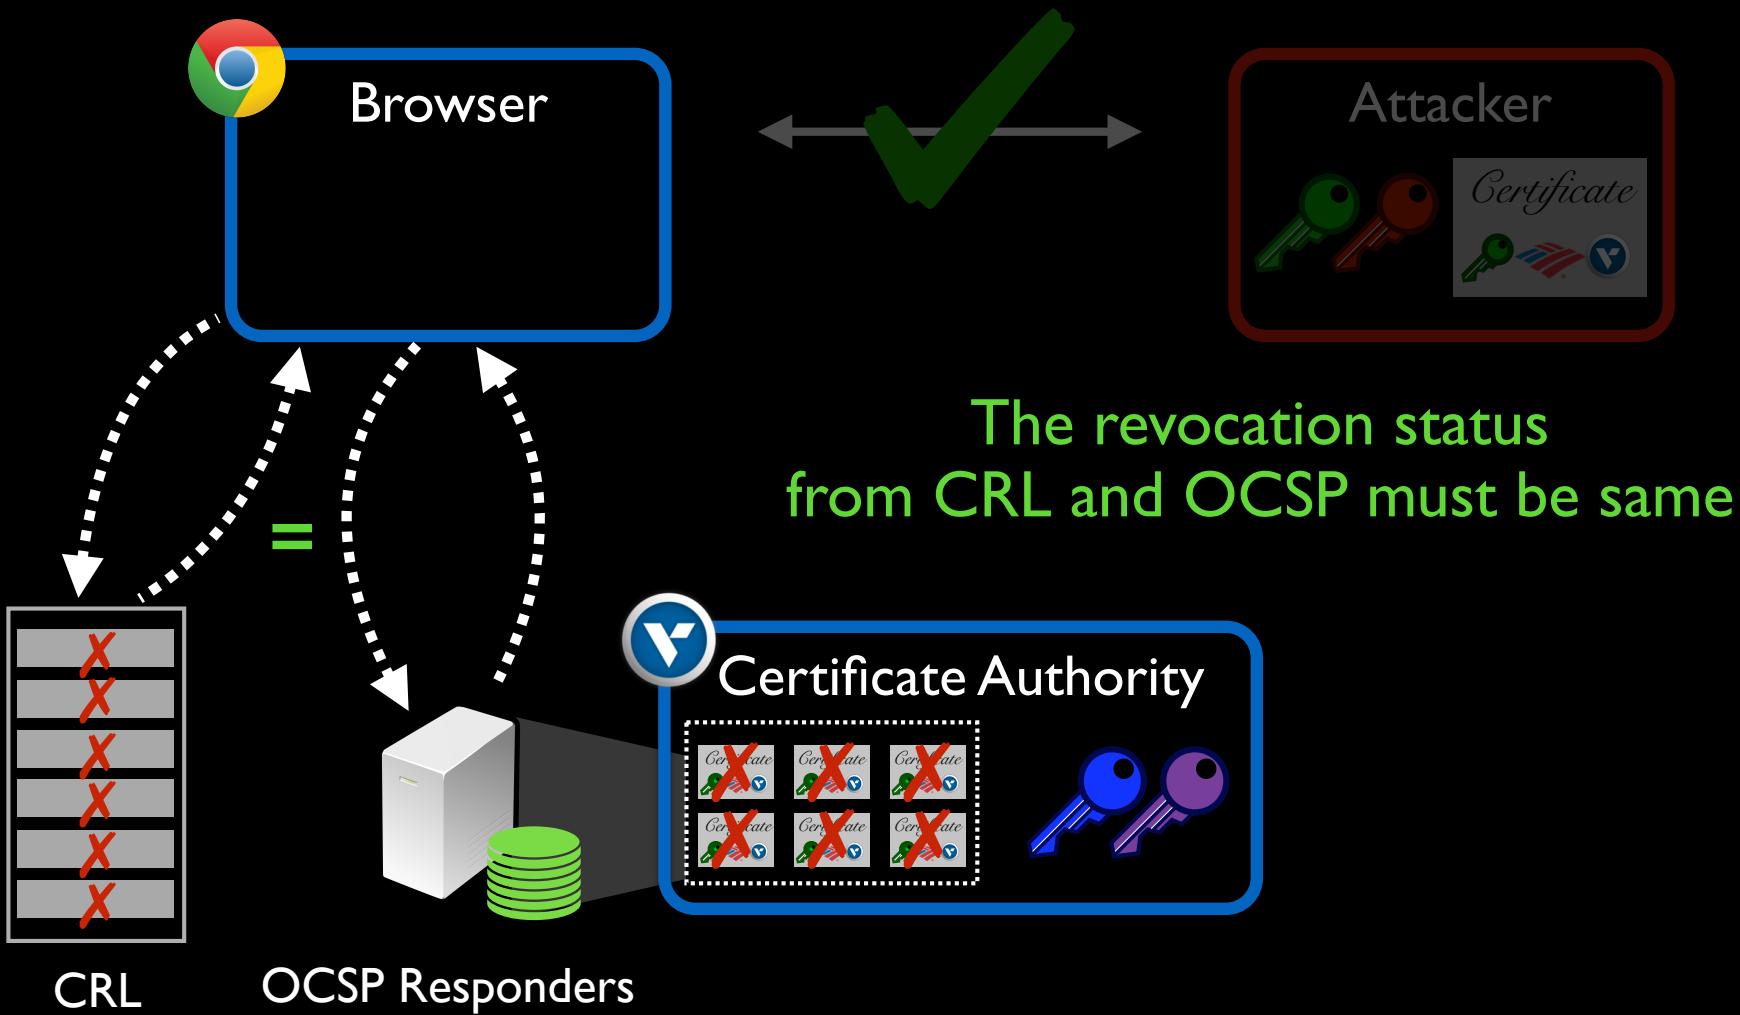
Availability

OCSP responders are not fully reliable

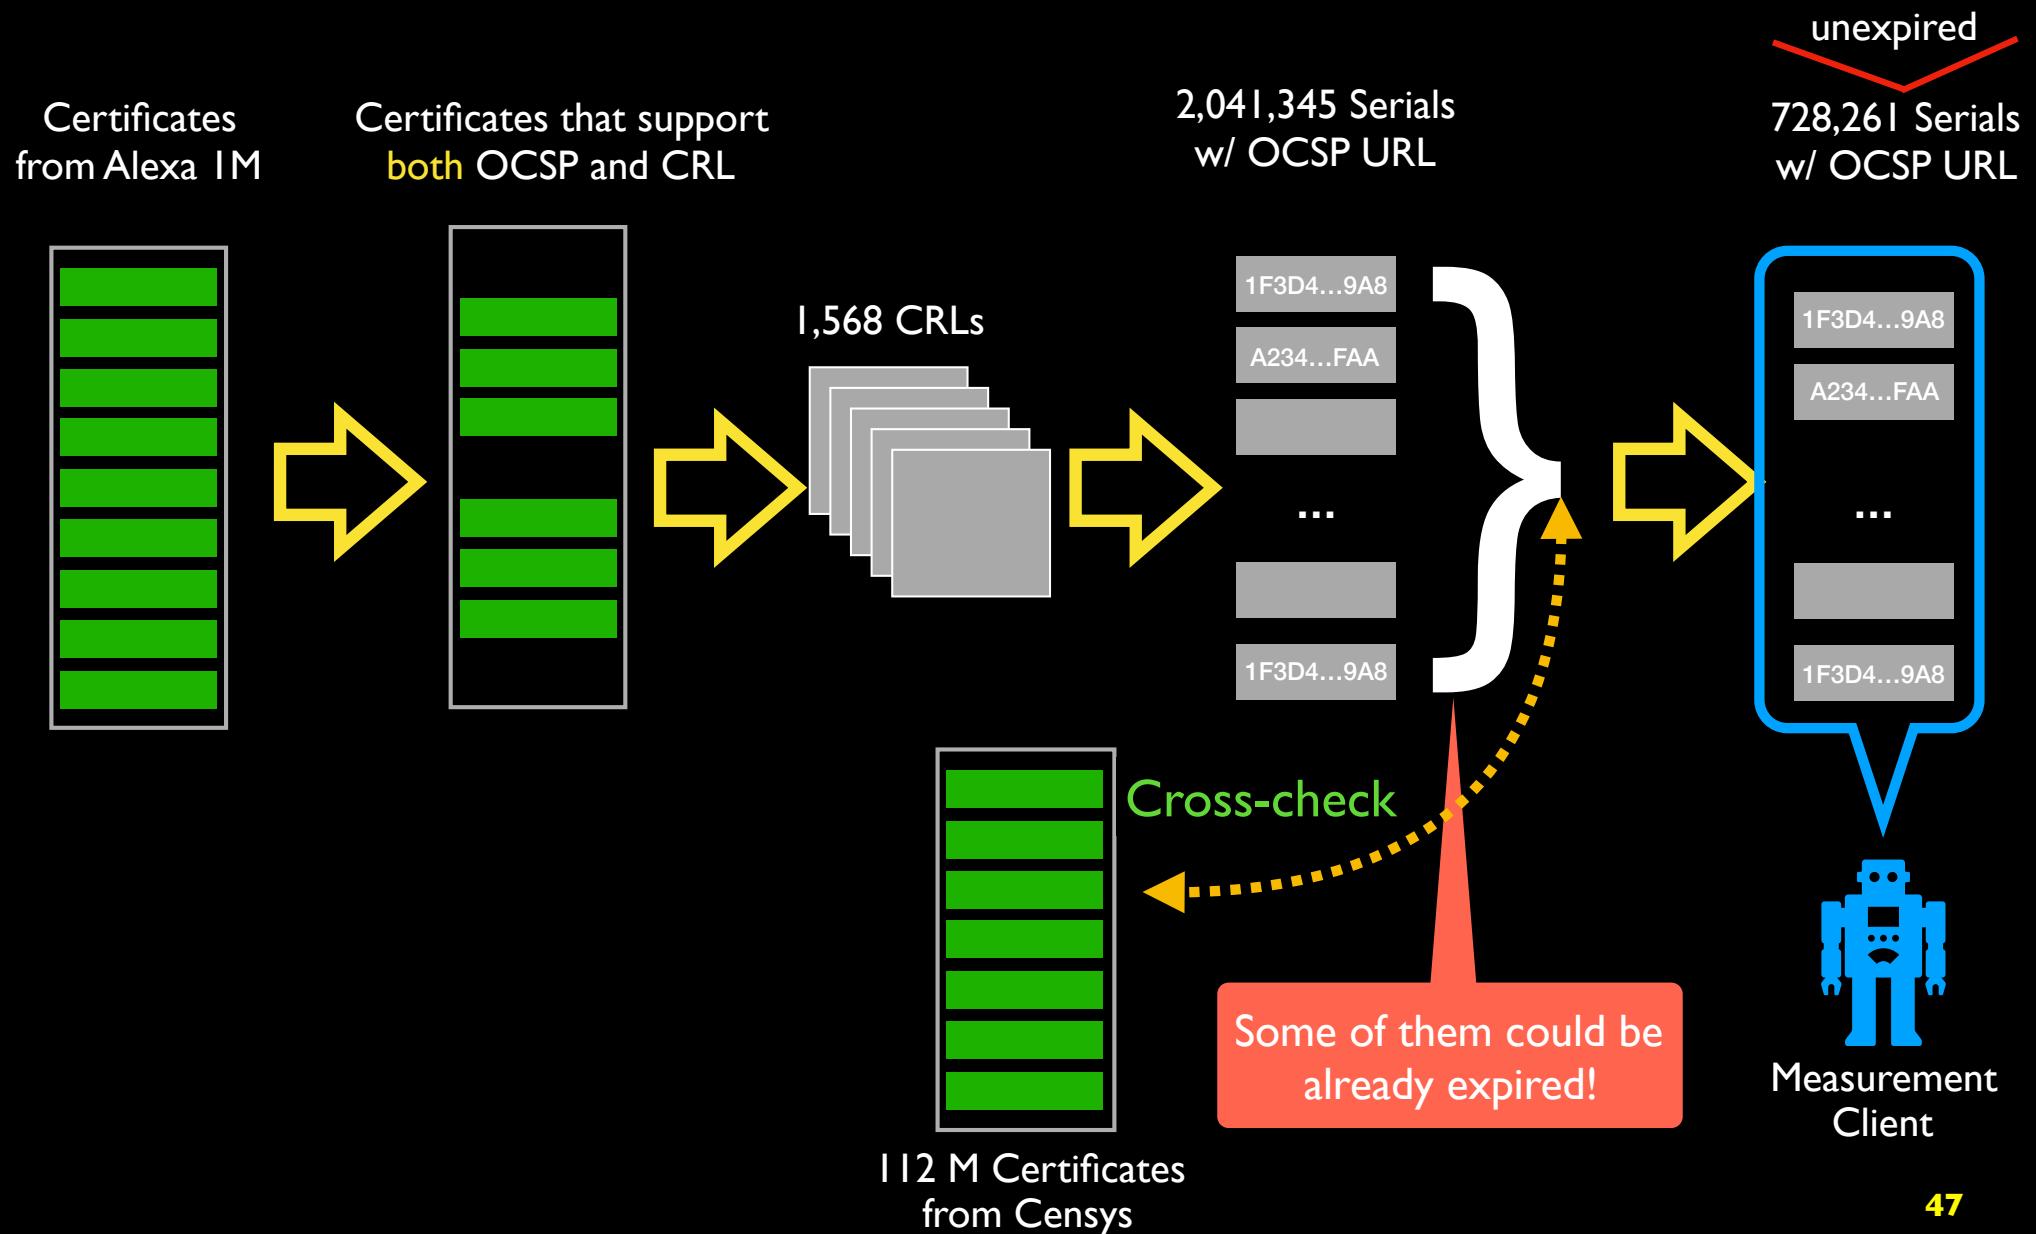
## (2) Validity of the Response



# (3) Consistency OCSP vs. CRL



# (3) Consistency OCSP vs. CRL



# (3) Consistency OCSP vs. CRL

OCSP URL	CRL	# of certificates where the OCSP response is		
		Unknown	Good	Revoked
ocsp.camerfirma.com	crl1.camerfirma.com/ camerfirma_cserverii-2015.crl			
ocsp.quovadisglobal.com	crl.quovadisglobal.com/qvsslg3.crl			
ocsp.startssl.com	crl.startssl.com/sca-server1.crl			
ss.symcd.com	ss.symcb.com/ss.crl			
twcasslocsp.twca.com.tw/	sslserver.twca.com.tw/sslserver/ securessl			
ocsp2.globalsign.com/gsalpha2g2	crl2.alphassl.com/gs/gsalpha2g2.crl			
ocsp.firmaprofesional.com	crl.firmaprofesional.com/ infraestructura.crl			
...	...			

# (3) Consistency OCSP vs. CRL

OCSP URL	CRL	# of certificates where the OCSP response is		
		Unknown	Good	Revoked
ocsp.camerfirma.com	crl1.camerfirma.com/ camerfirma_cserverii-2015.crl	0	7	369
ocsp.quovadisglobal.com	crl.quovadisglobal.com/qvssl3.crl	0	1	514
“OCSP and PKI Management are <i>two different platforms and are synchronized by means of some DDBB triggers</i> that are failing in some circumstances. Meanwhile CRL management is easier and simple, OCSP should give information about any certificate serial number issued by *** and the amount of information transmitted between them. That’s the source of this problem.”	sslserver.twca.com.tw/sslserver/ securessl	0	1	28,032
ocsp2.globalsign.com/ gsalphasha2g2	crl2.alphassl.com/gs/ gsalphasha2g2.crl	5,375	0	0
ocsp.firmaprofesional.com	crl.firmaprofesional.com/ infraestructura.crl	11	0	0
...	...	0	0	...

# Is the Web Ready for OCSP Must-Staple?



Certificate authority



Web server



Browser

- ✓ Fetch and cache OCSP responses
- ✓ Handling errors

# Web Server Methodology



## (1) Performance

- ? Prefetch OCSP response

## (2) Caching

- ? Cache OCSP response
- ? Respect nextUpdate\* in cache

## (3) Availability

- ? Retain OCSP response on error

# Web Server Administrator Result

	 THE APACHE™ SOFTWARE FOUNDATION	NGINX
Prefetch OCSP response	✗	✗
Cache OCSP response	✓	✓
Respect nextUpdate in cache	✗	✓
Retain OCSP response on error	✗	✓

\* Apache version 2.4.18 and Nginx version 1.13.12

# Is the Web Ready for OCSP Must-Staple?



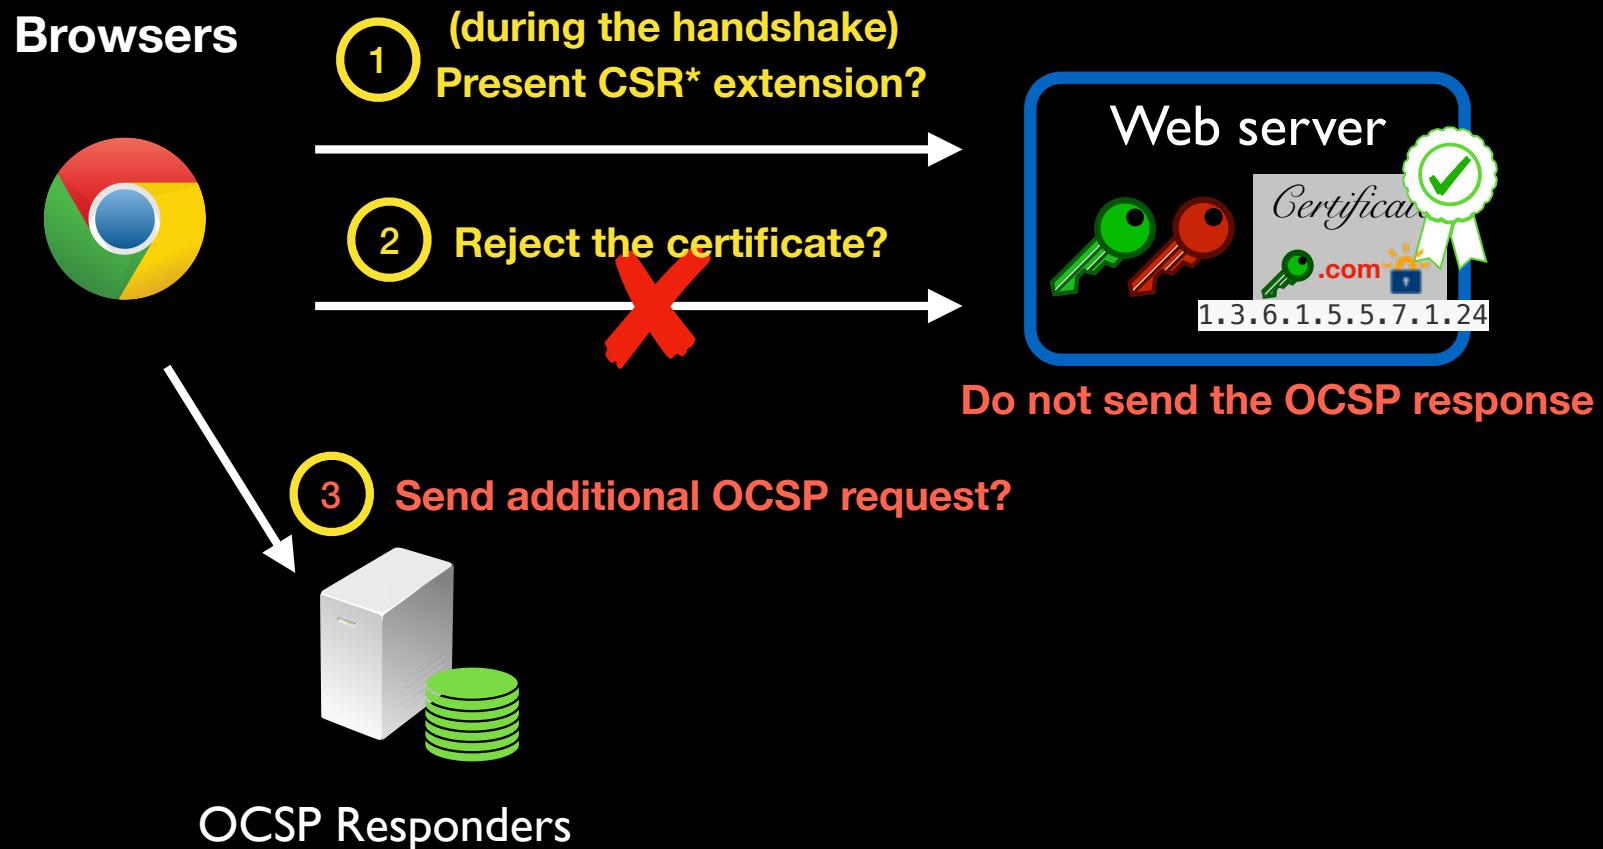
Certificate Authority



Browser

- ✓ Understand the extension
- ✓ Present Certificate Status Request extension
- ✓ Reject the certificate if the response is not provided

# Methodology



# Methodology and Result

	Desktop Browsers (OS X, Linux, Windows)						Mobile Browsers			
	Chrome 66	Firefox 60	Opera	Safari	IE	Edge	Safari	Chrome	Firefox/ iOS	Firefox/ Android
Request OCSP Response	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓
Respect OCSP Must-Staple	✗	✓	✗	✗	✗	✗	✗	✗	✗	✓
Send own OCSP Request	✗	-	✗	✗	✗	✗	✗	✗	✗	-

Clients

Clients are largely not yet ready for OCSP Must-Staple

(the additional coding work necessary to support OCSP Must-Staple is likely not too significant)

# Conclusion

- Considering OCSP Must-Staple can operate only if each of the principals in the PKI performs correctly.
  - OCSP servers: **not fully reliable**
  - Web server softwares: **not fully support**
  - Browsers: **not fully support**
- But the bright side is
  - **Only a few players** need to take action to make it possible for web server administrators to begin enabling OCSP Must-staple
  - Much wider deployment of OCSP Must-Staple is an **realistic** and **achievable goal**

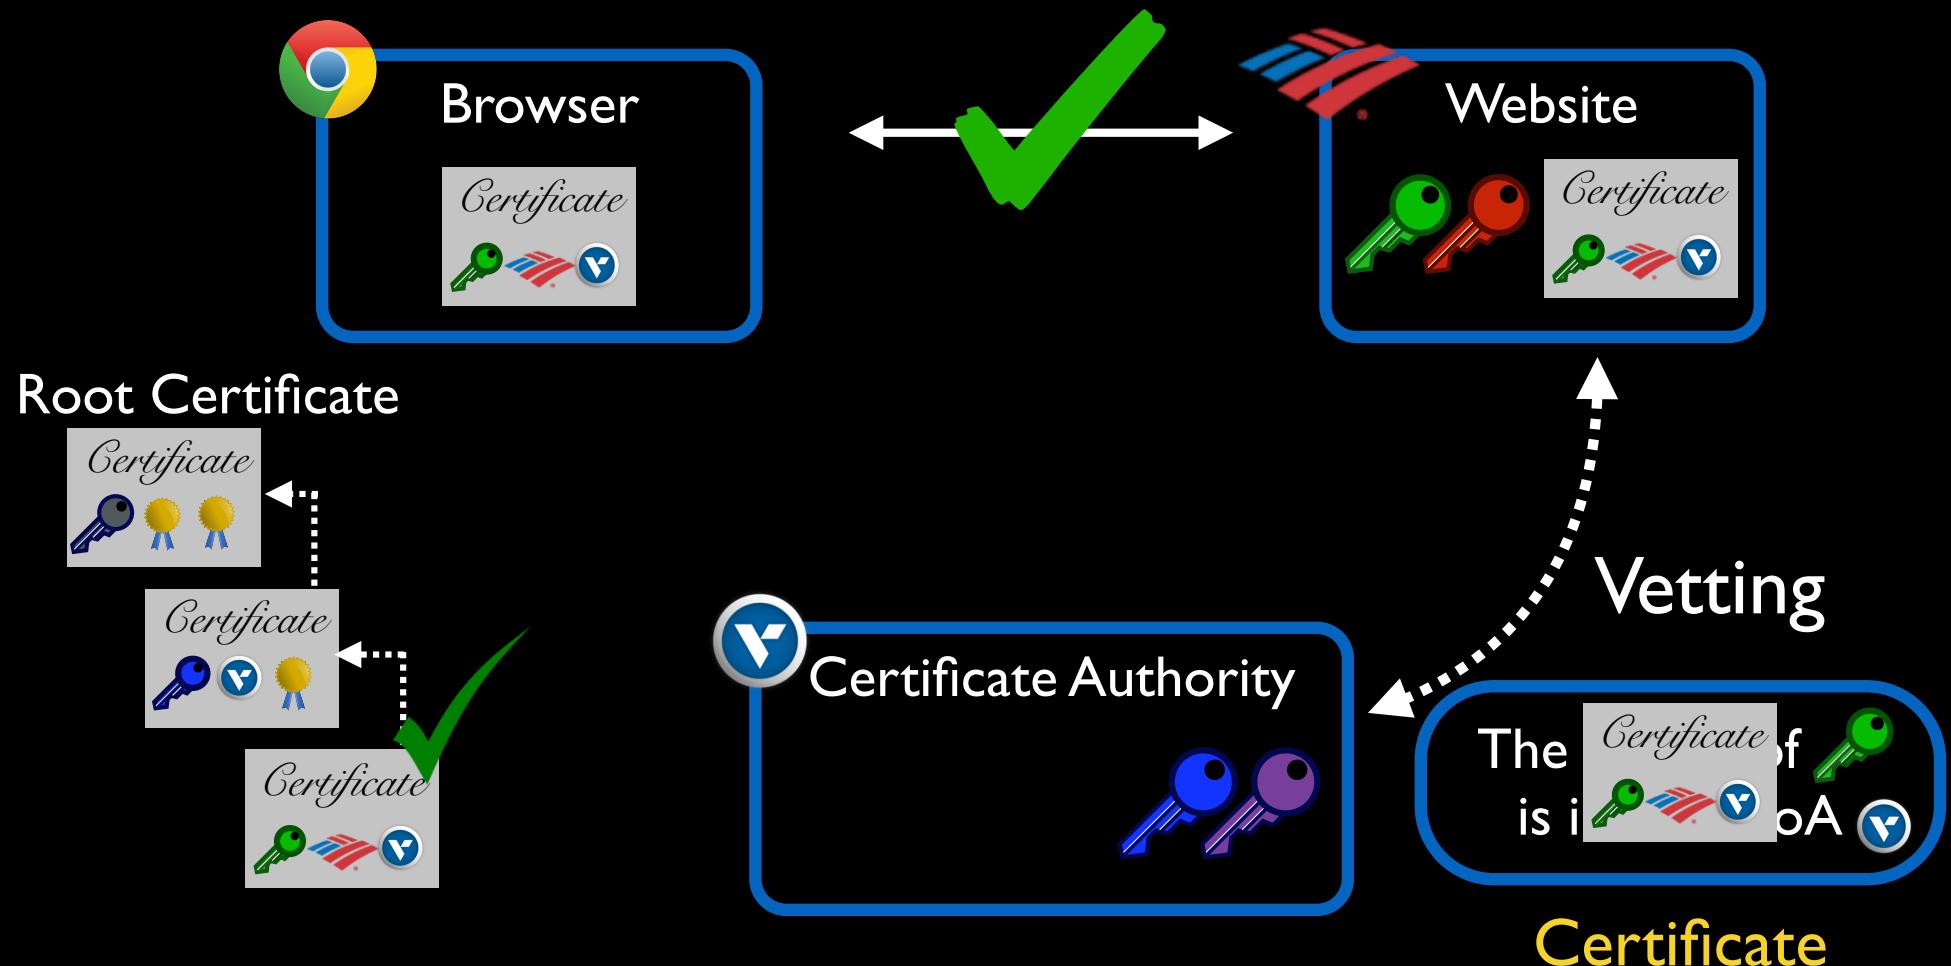
# Some protocols

- HSTS (HTTP-STRICT-TRANSPORT-SECURITY)
  - “Strict-Transport-Security” Header
- HSTS-preloaded list
- HPKP (HTTP Public Key Pinning)
- SNI (Server Name Indication)
- Certificate Transparency
-

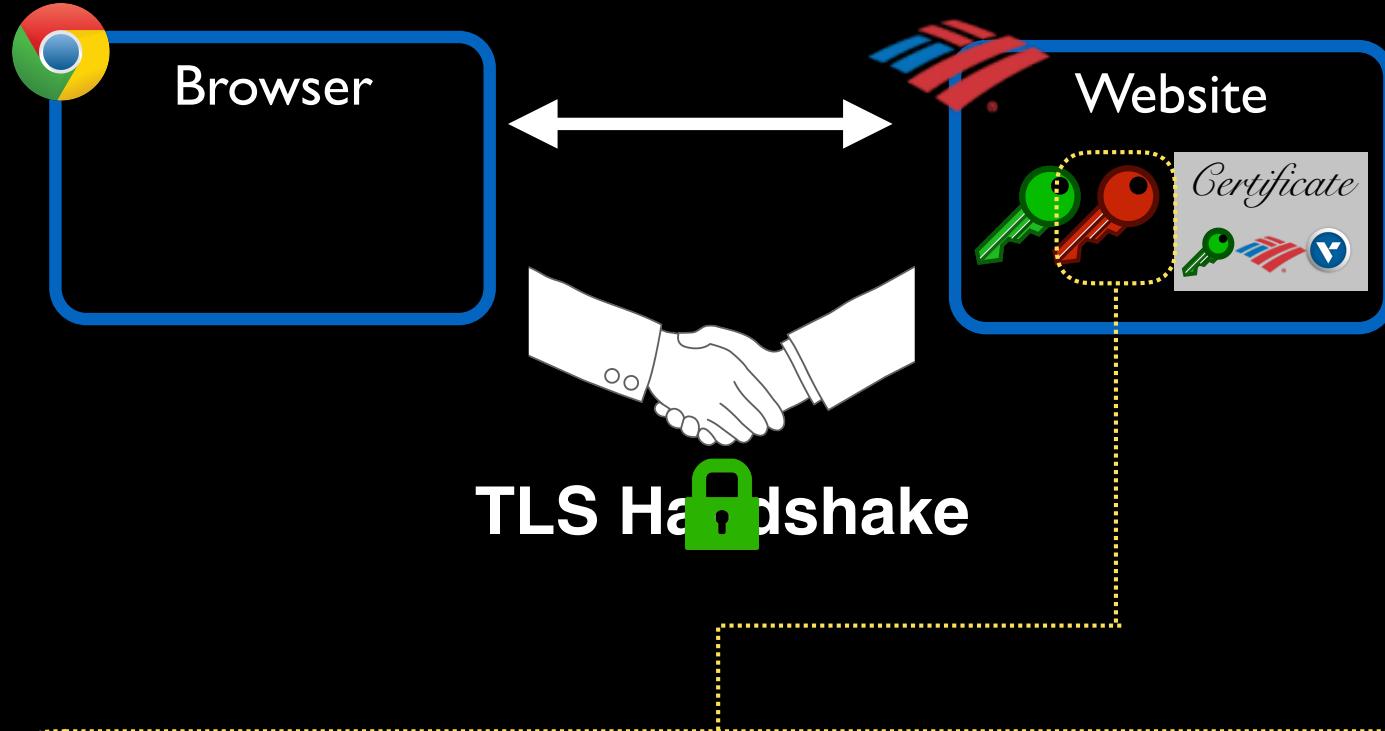
# Today's another problem of HTTPS

# How HTTPS Works

How can users truly know with whom they are communicating?



# Fundamental Assumption in HTTPS



Authentication **fundamentally** assumes:

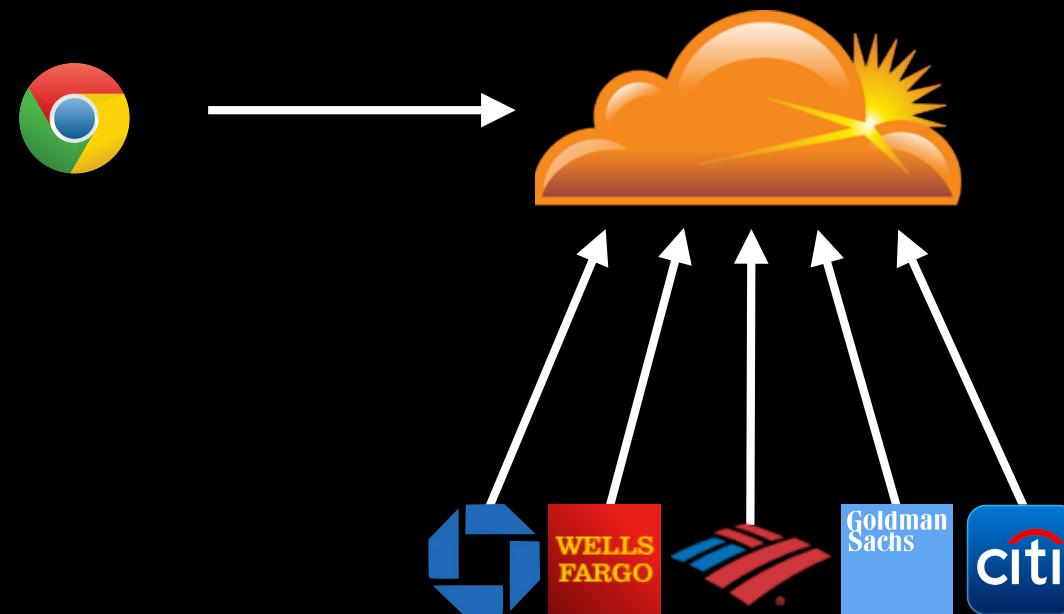
Only  knows 

# The PKI in Today's Web

Rare!



# The PKI in Today's Web



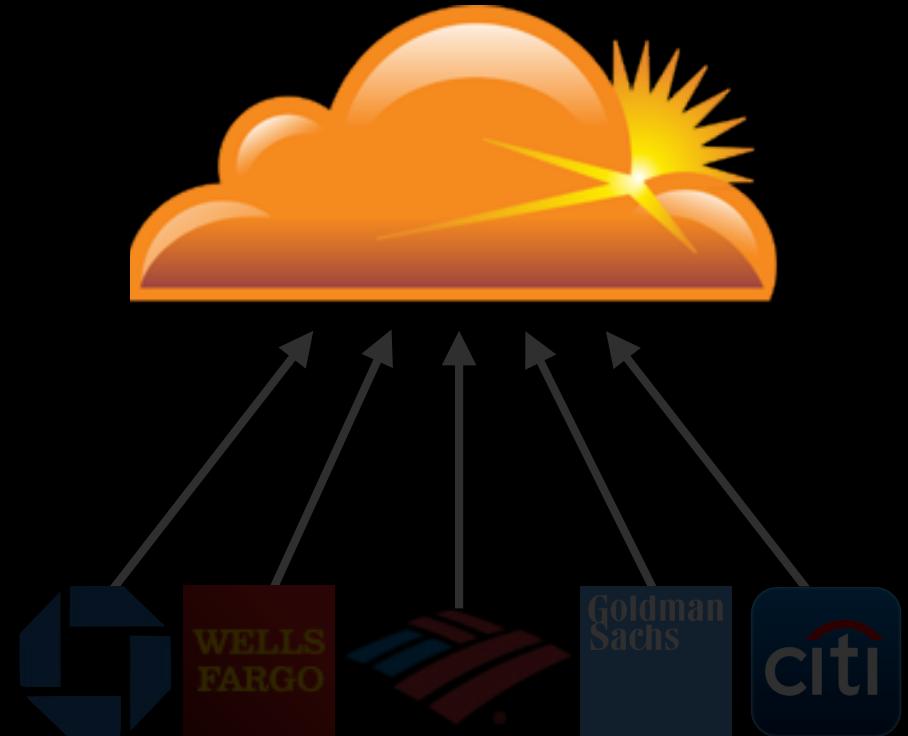
# The PKI in Today's Web

## Third-party Hosting Providers

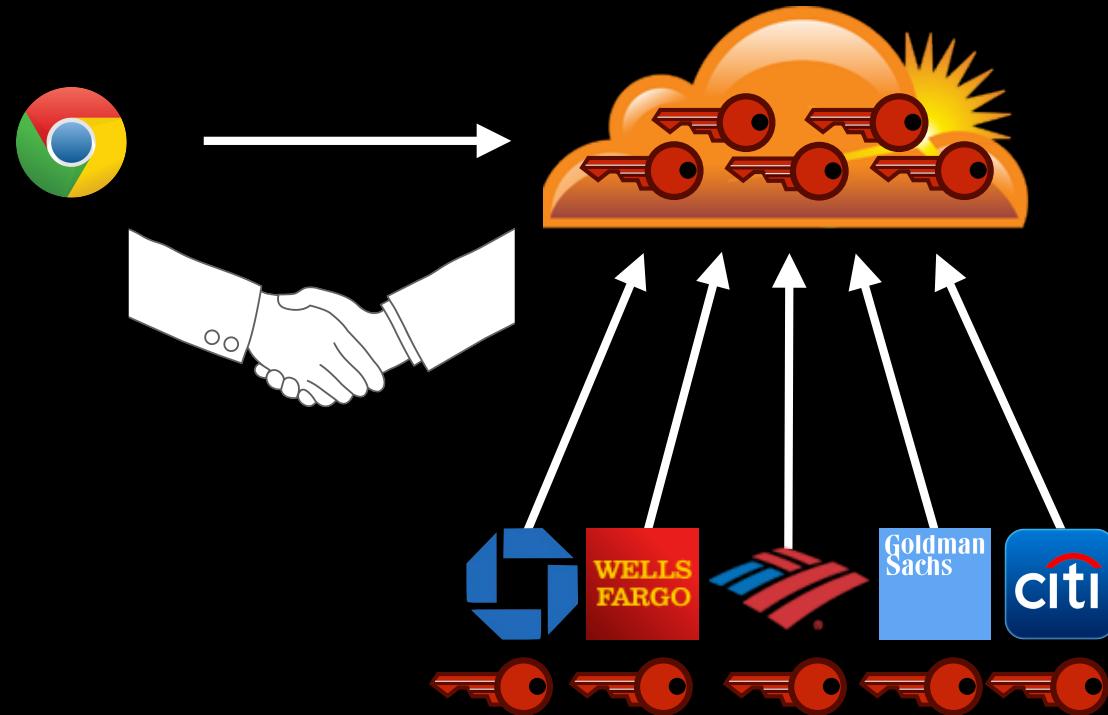
- Content delivery networks
- Web hosting services
- Cloud providers

Varying levels of involvement

But all trusted to deliver content

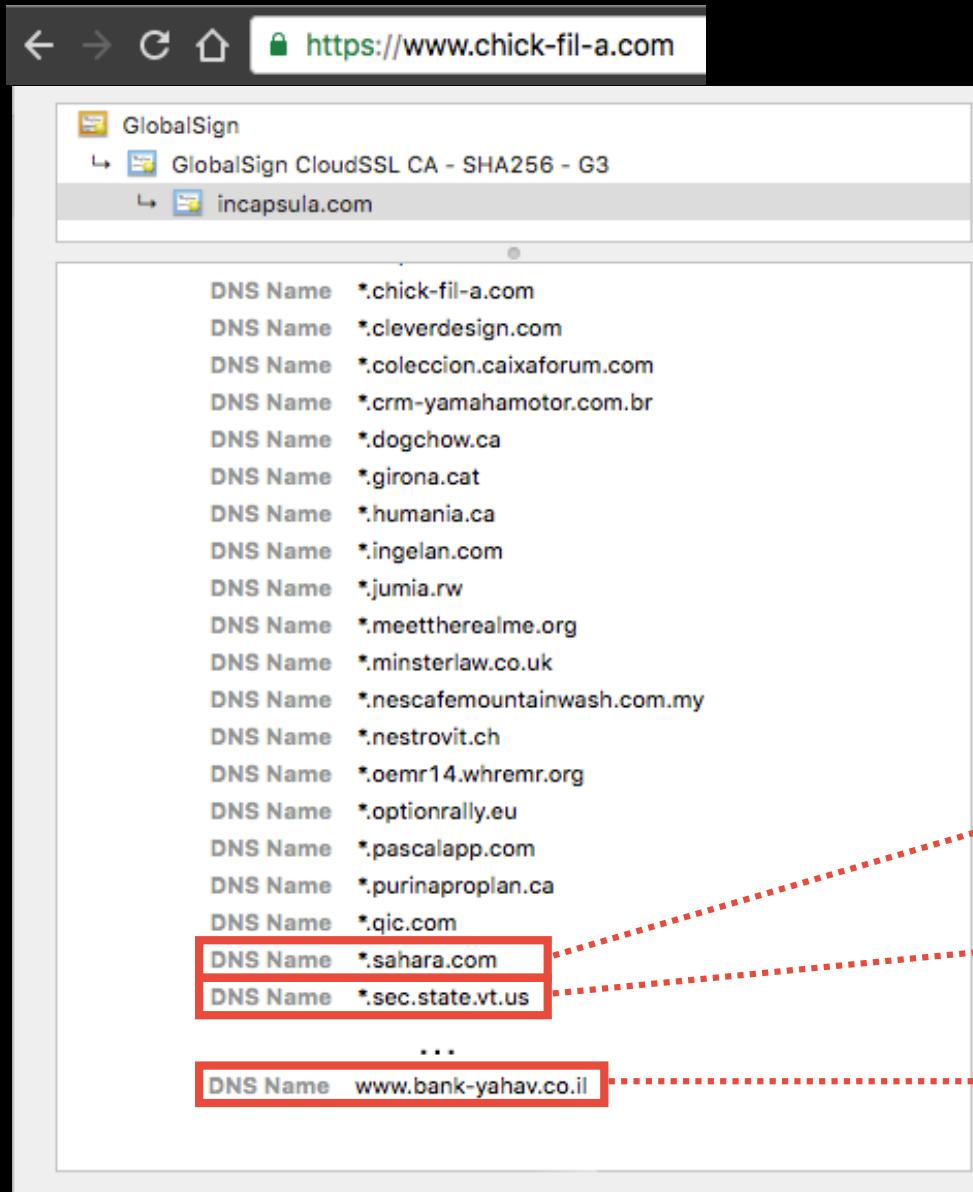


# The PKI in Today's Web



Third-party hosting providers  
know their customers' private keys

# Example of Key Sharing



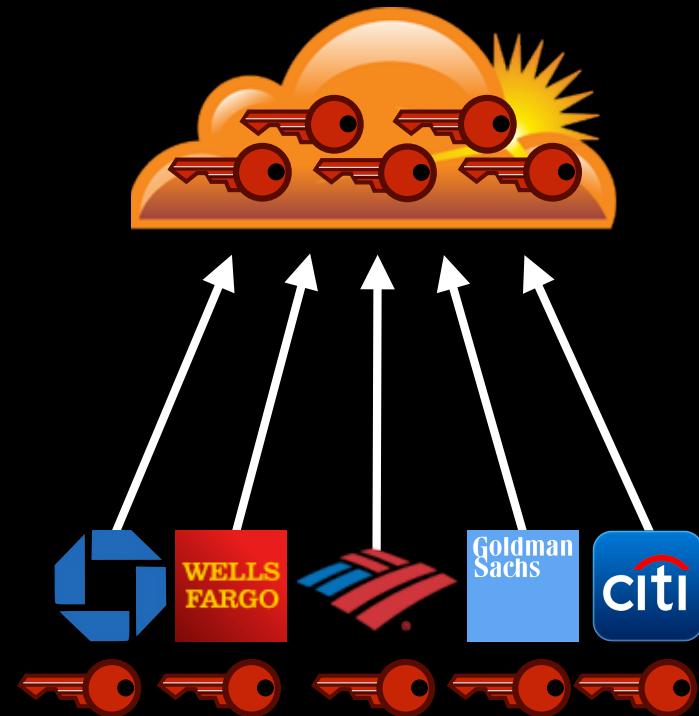
Large Saudi Arabian ISP

Secretary of State  
of Vermont

Israeli bank

# Problems of Key Sharing

1. Complicates the trust model,  
users don't know who they're  
really trusting
  
2. Potential to create  
centralization of trust
  
3. Potential to create  
single point of failure  
(in terms of management)



# Research Questions



How prevalent is the **key sharing**?



What's the **potential vulnerabilities**?

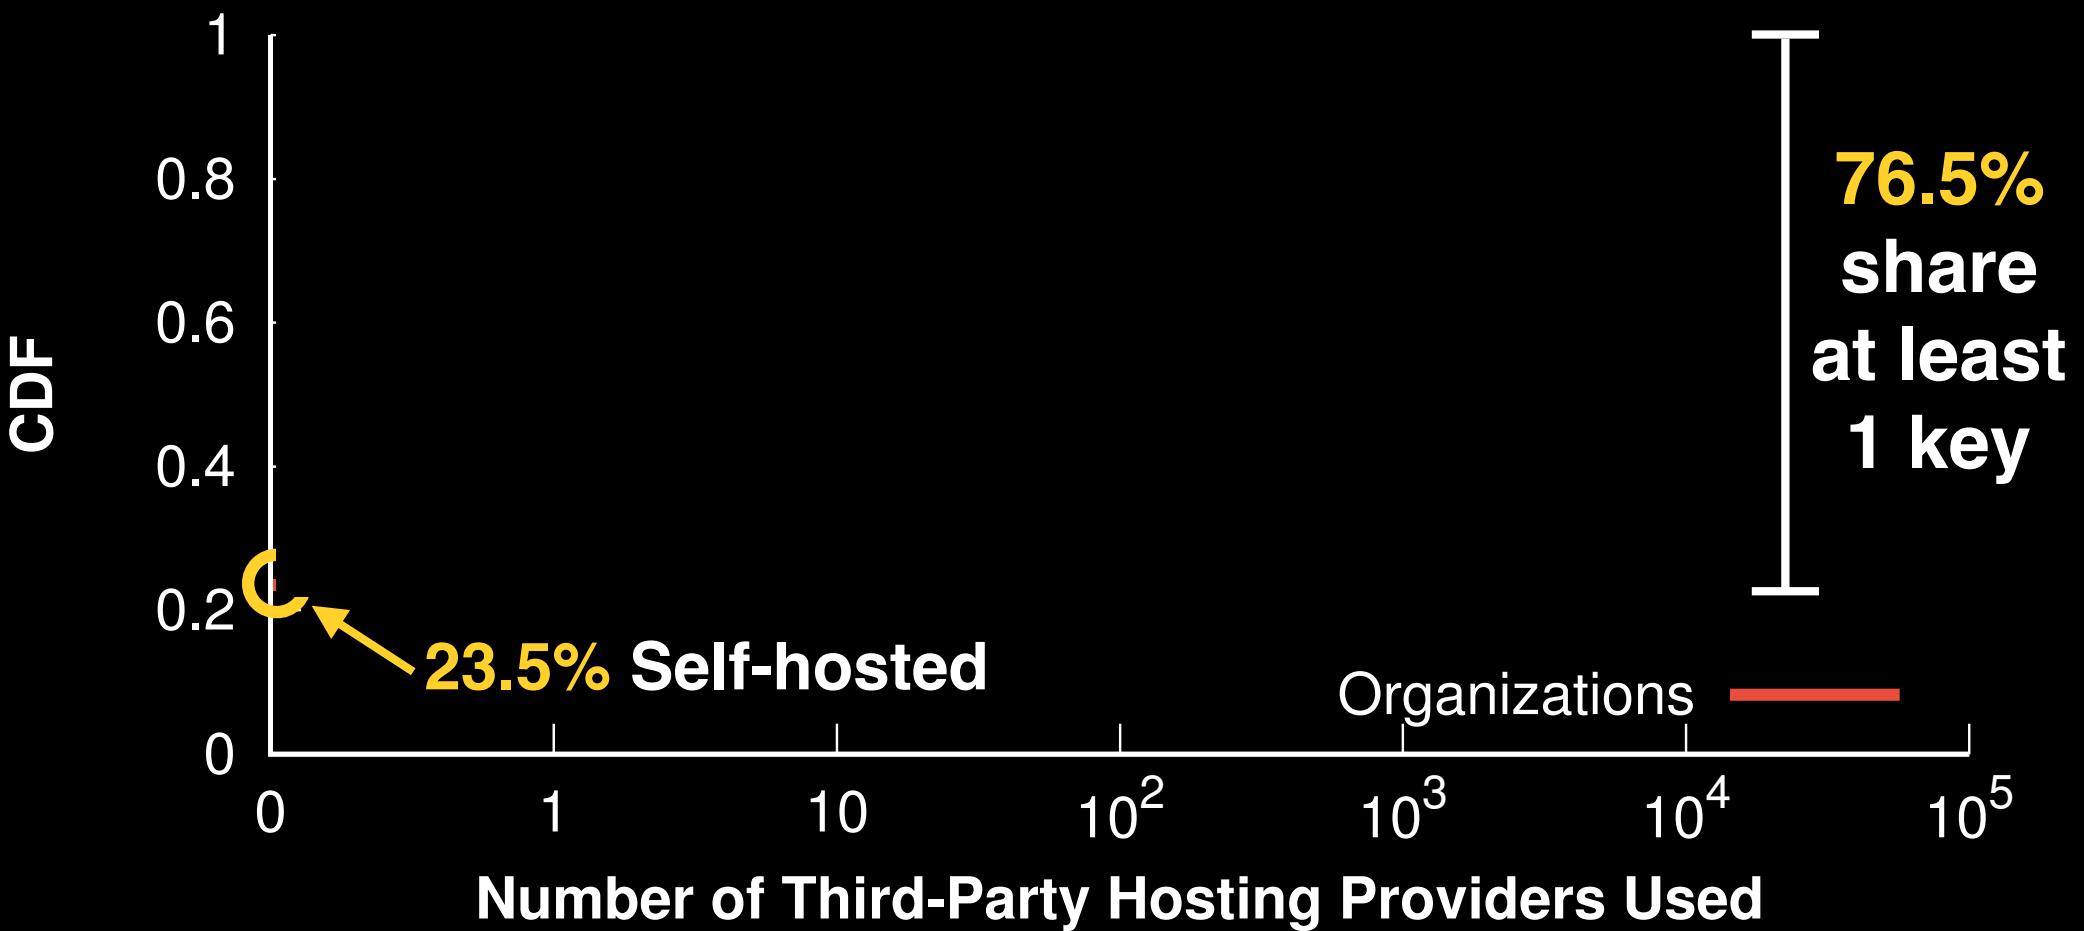


How can we **improve** it?

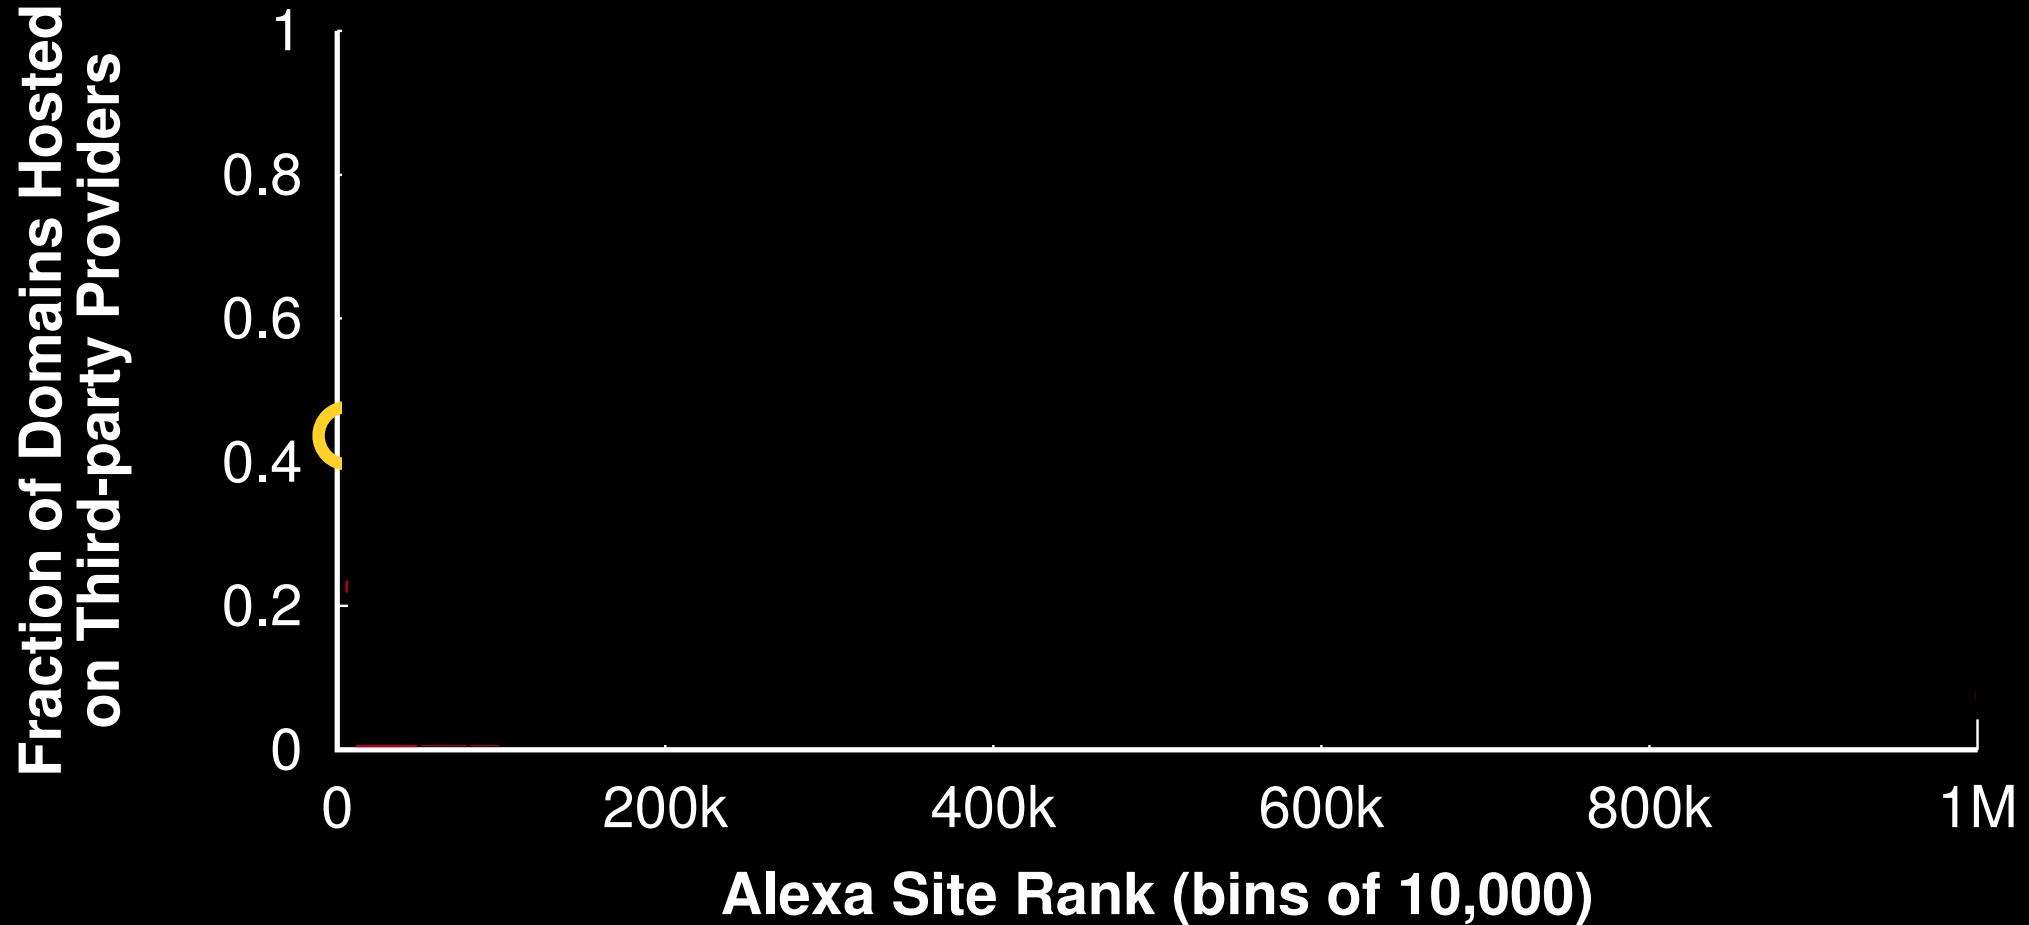
# Scanning All IPv4 Certificates

Dataset	
Period	2013/10/30 ~ 2016/04/30
# of IPs	101,306,358 (Full IPv4 scan)
Certificates	38,514,130
# of Domains	2,552,936

# How Prevalent of Key Sharing?

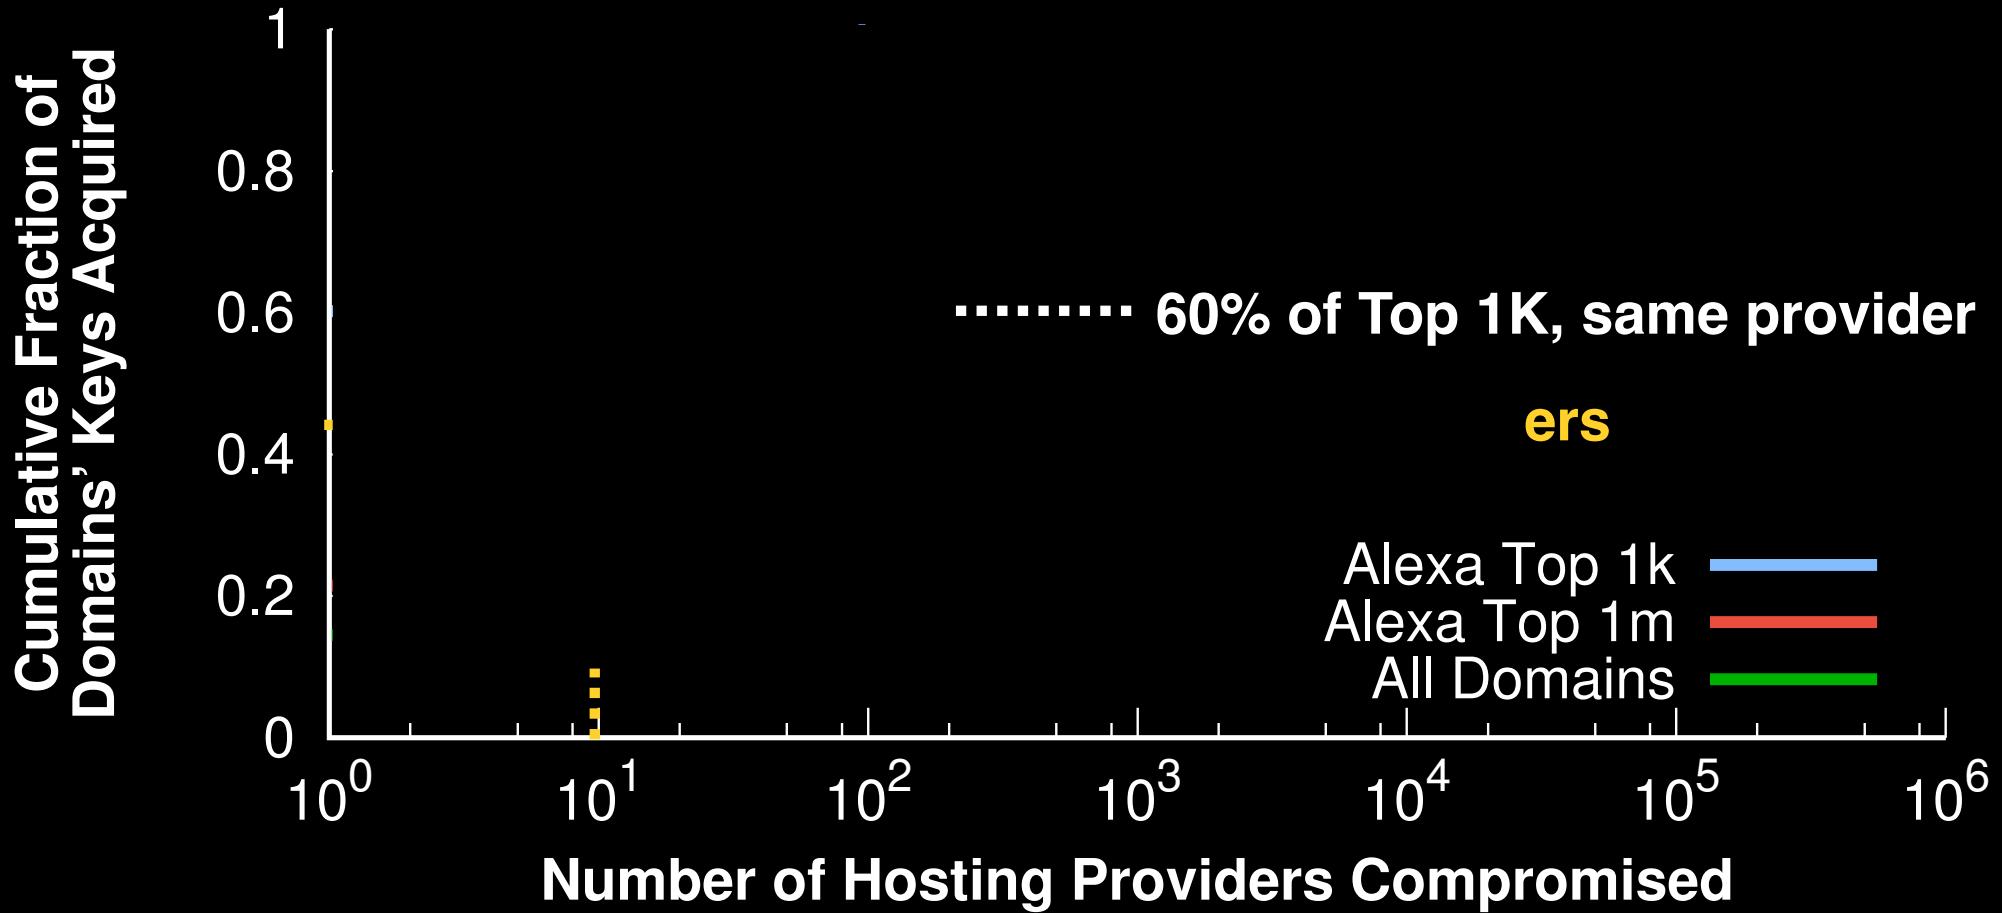


# Who Shares the Keys?



**Key sharing is common across the Internet**  
**Economic incentives drives key sharing**

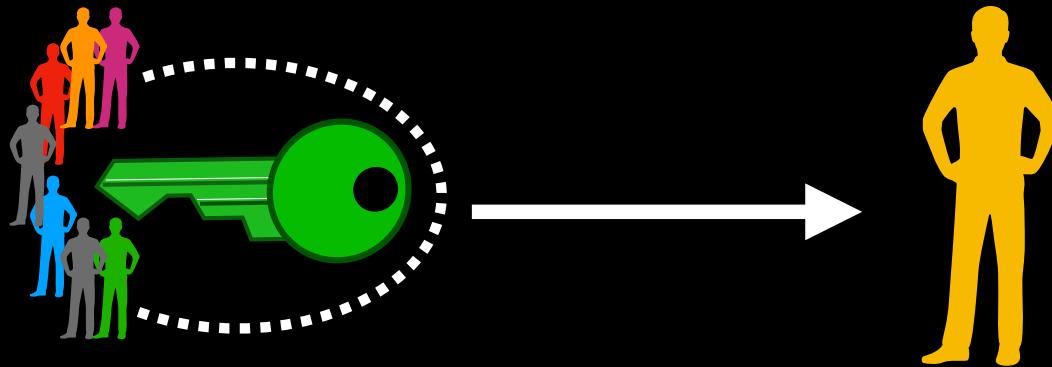
# New Attack Targets



Popular hosting services are prime targets for attack

# Summary

Whom am I talking with?



- Due to economic incentives, **key sharing** is prevalent in today's web
  - **76.5%** of keys are shared
  - **43.2%** (of top 10K webpages) share the private keys
  - Compromising a single hosting provider could put **60%** of top 1K webpages in danger

# Preparing the Final

- Okay to bring your cheat sheet
  - 1 page letter-size both sides
- Understanding concept is the most important thing.
- Please feel free to post a question on the discussion board; I've created another discussion board category "Questions for Final"



# Thanks again for taking my class

- Best wishes for the final